

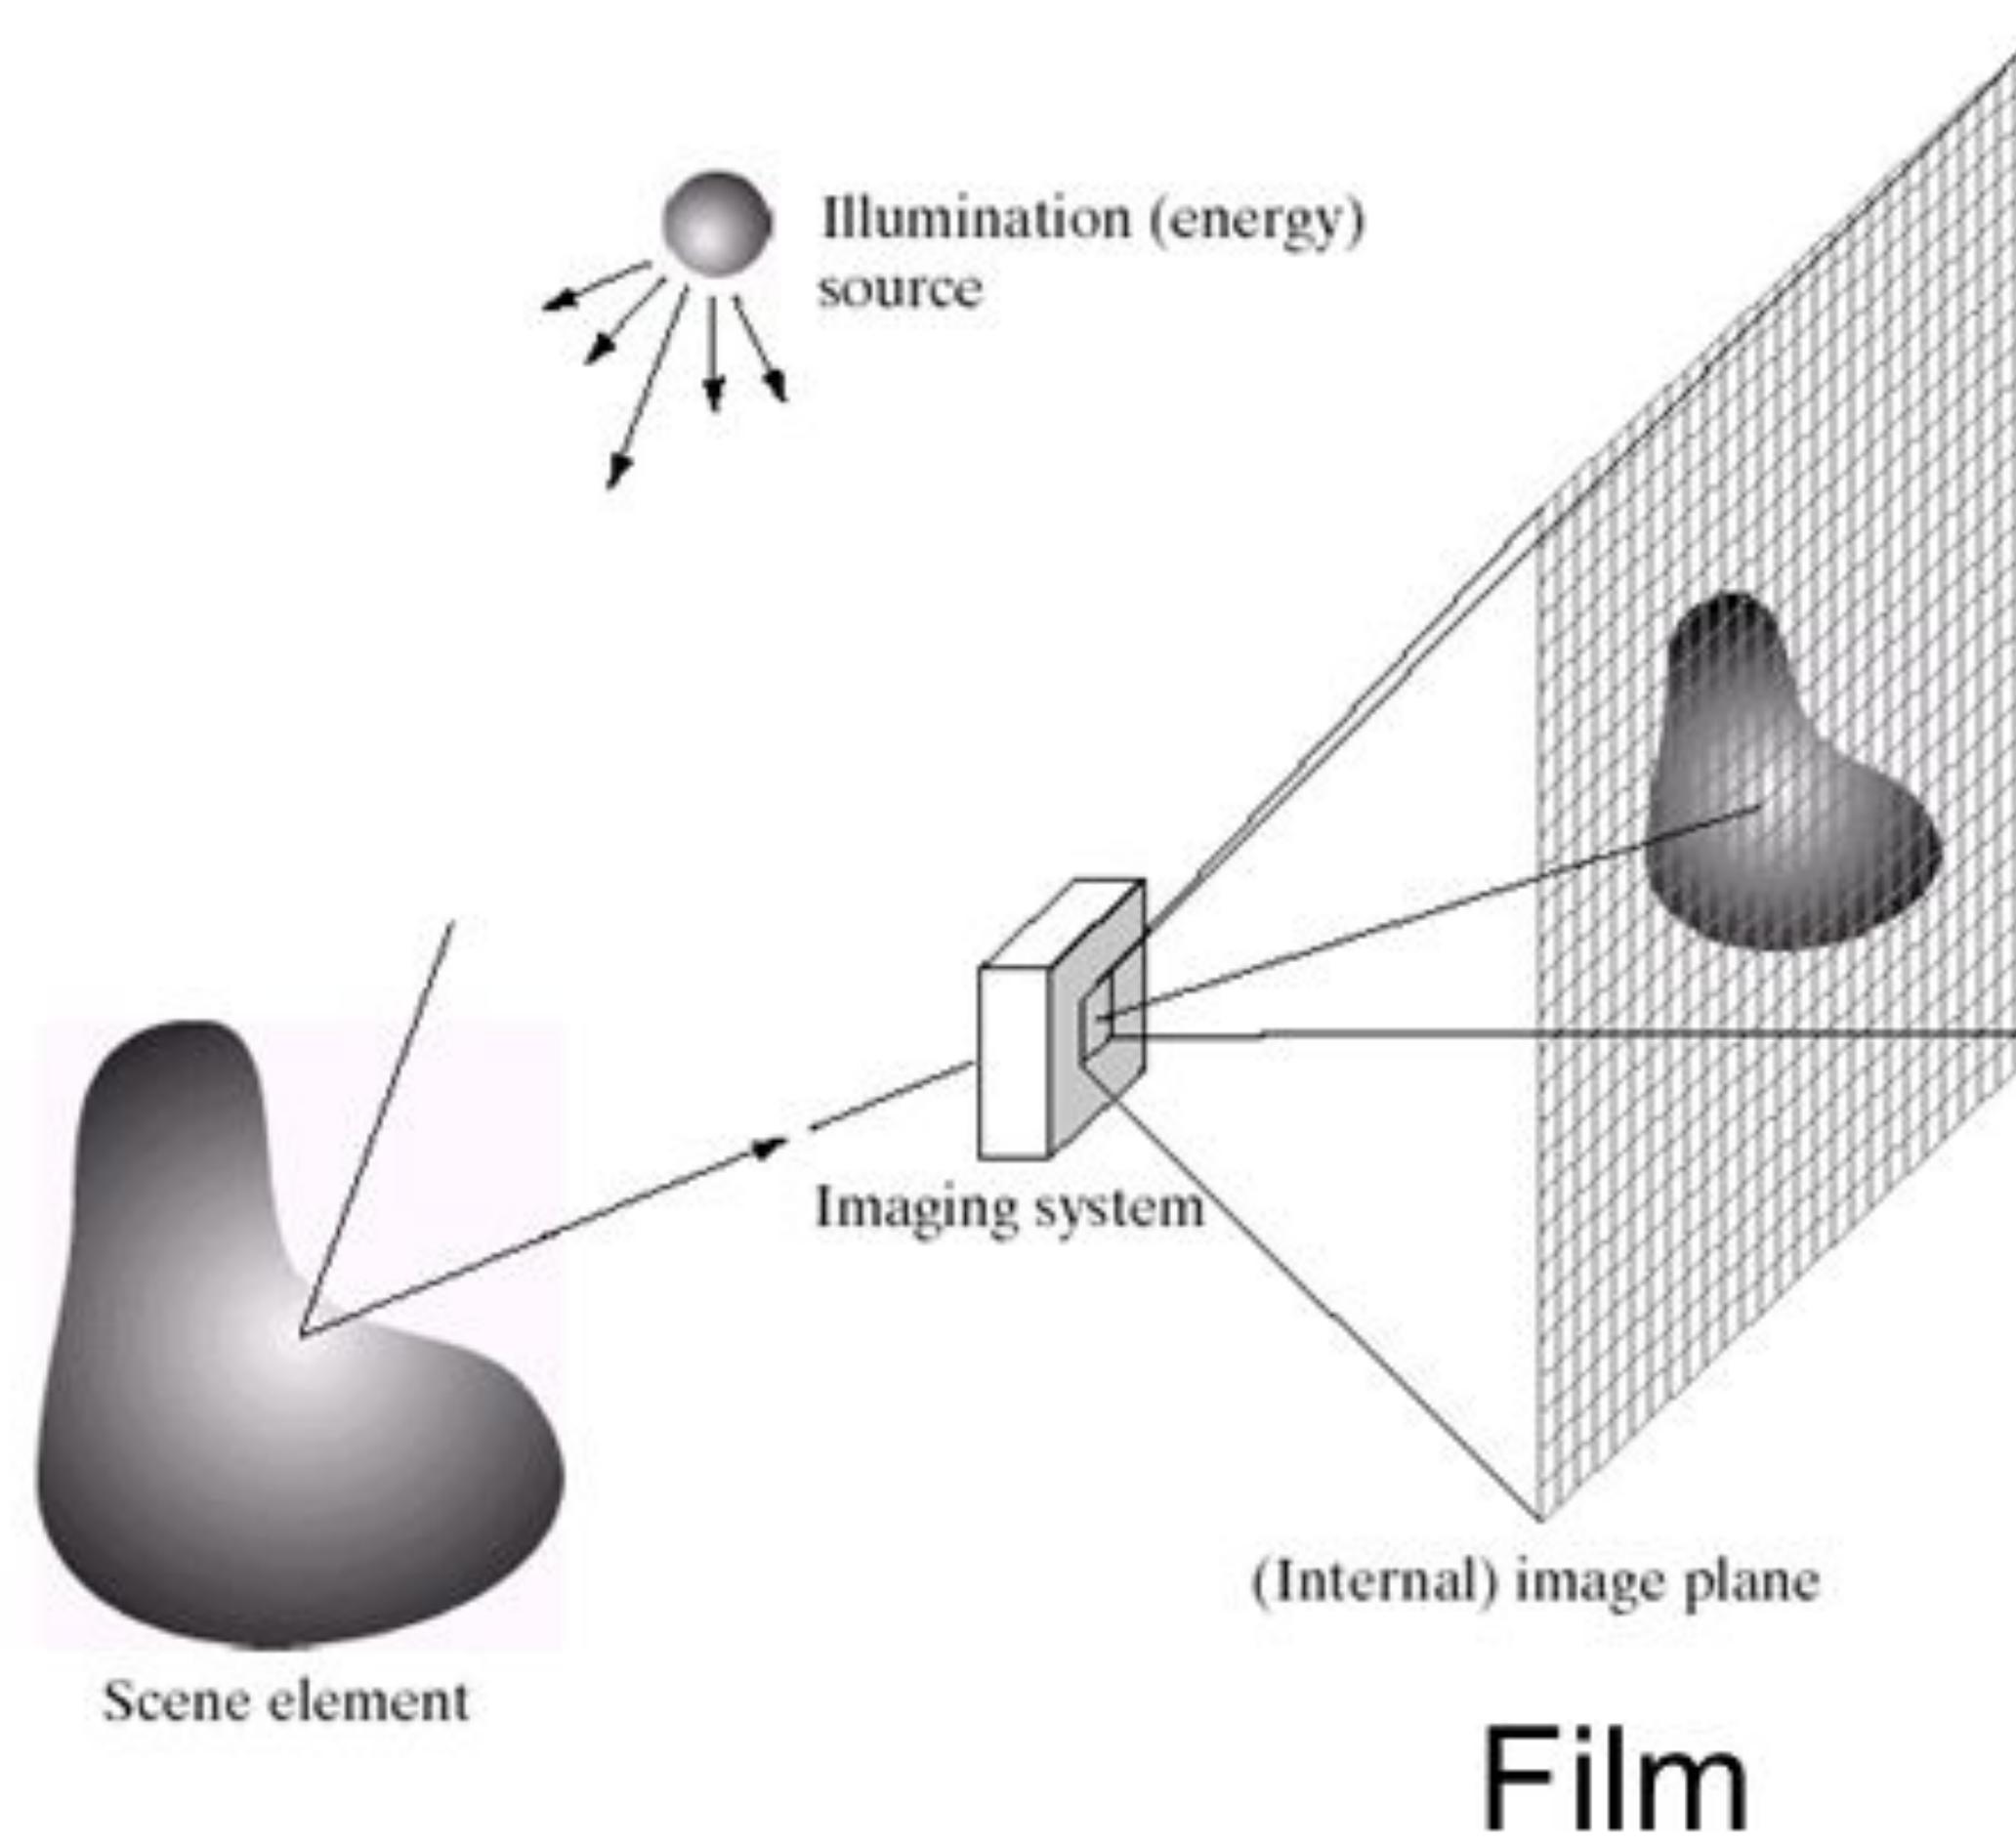
TDT4265:

Computer Vision & Deep Learning

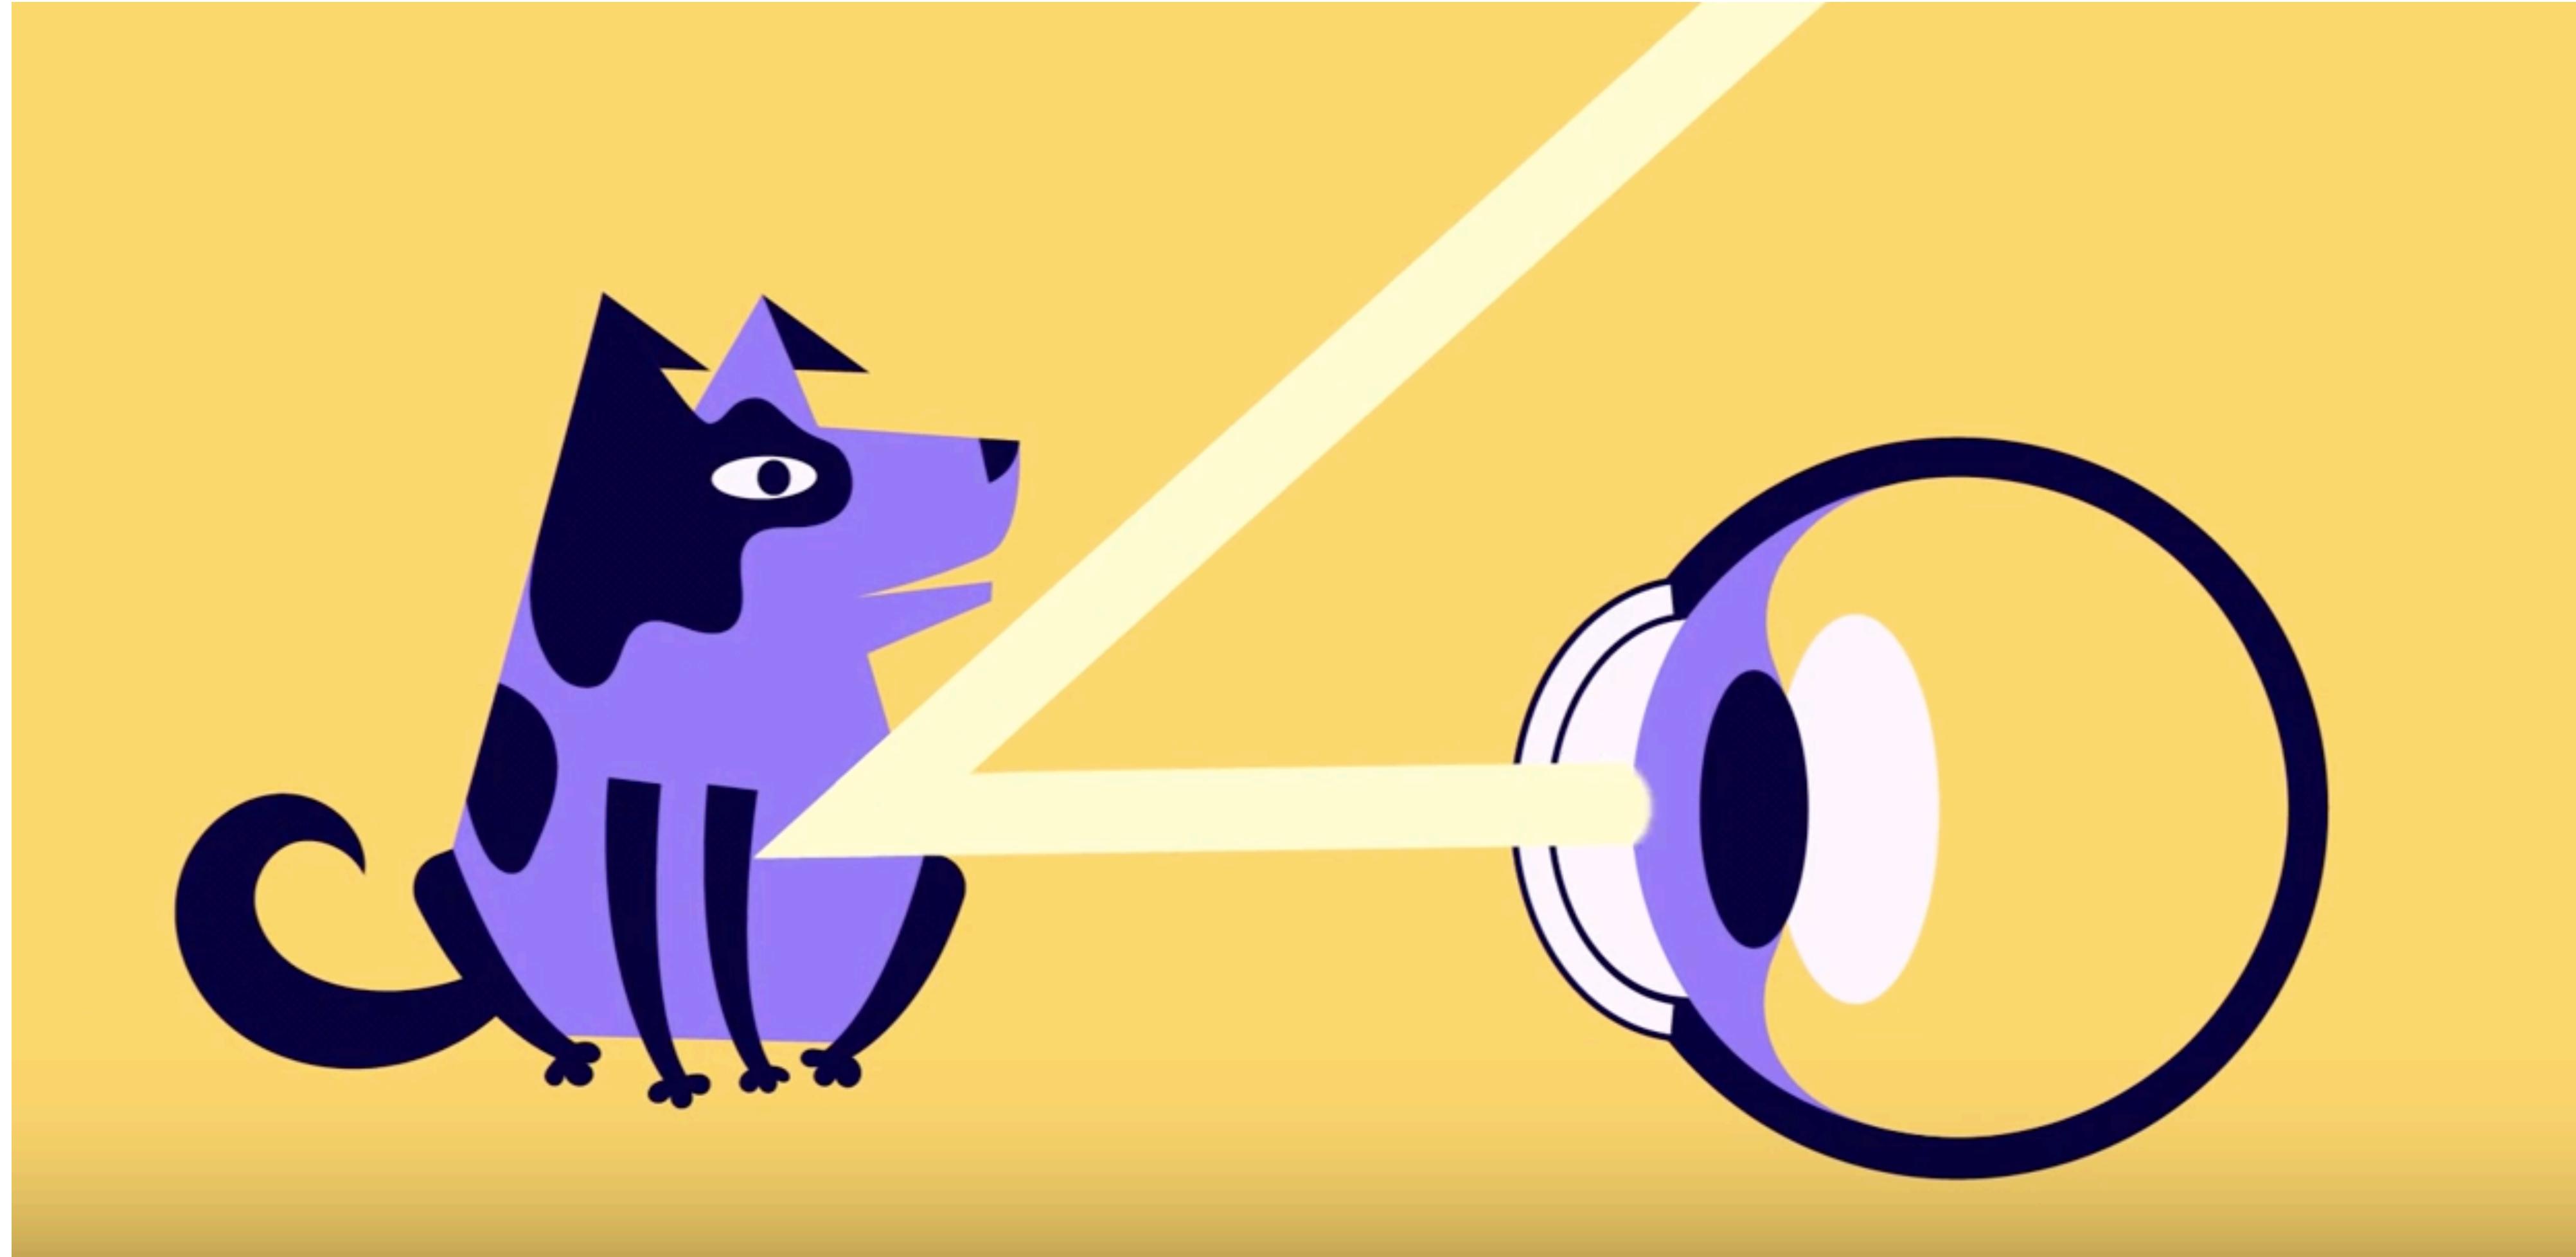
Intro

Frank Lindseth, IDI, IE, NTNU

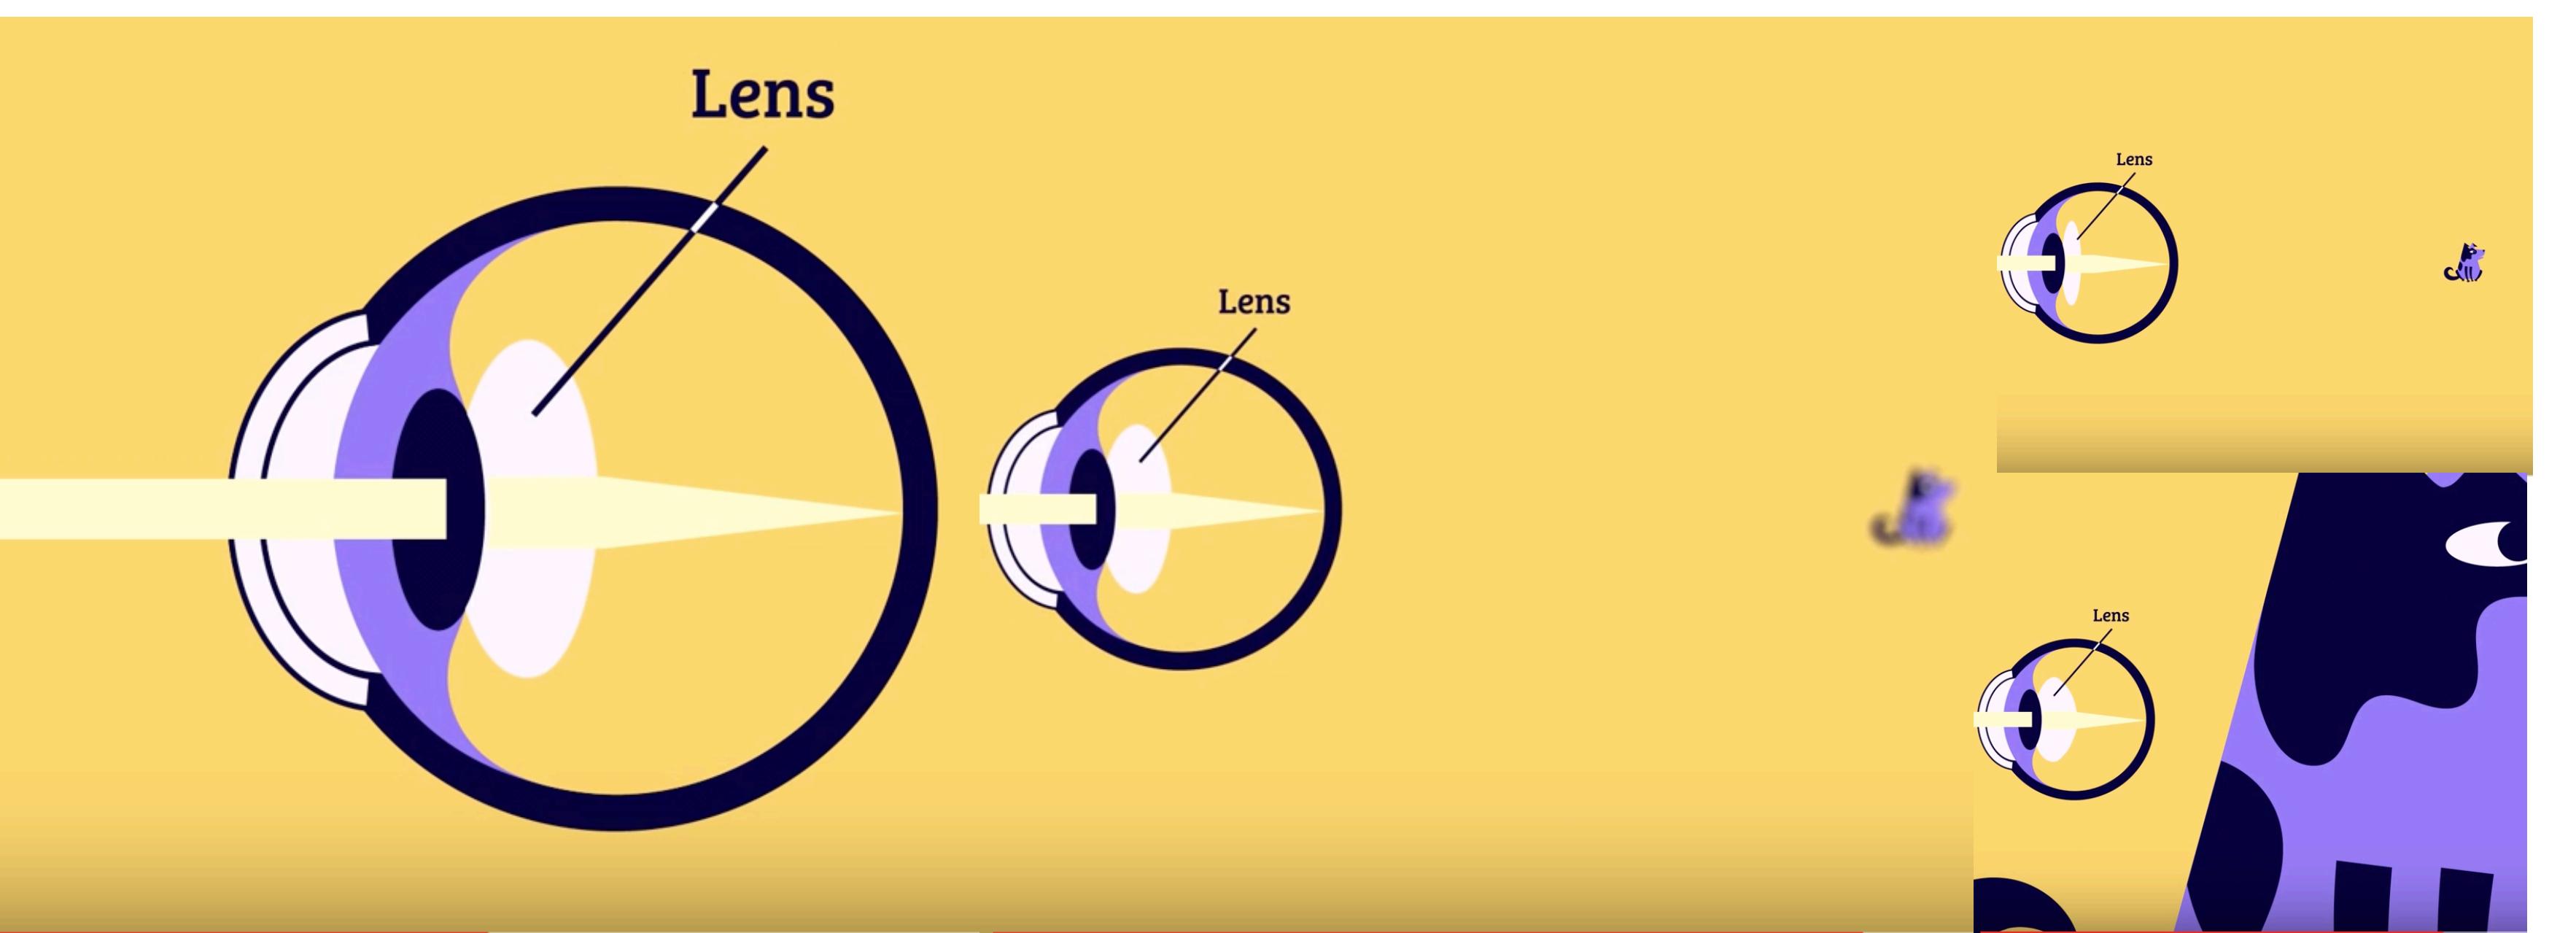
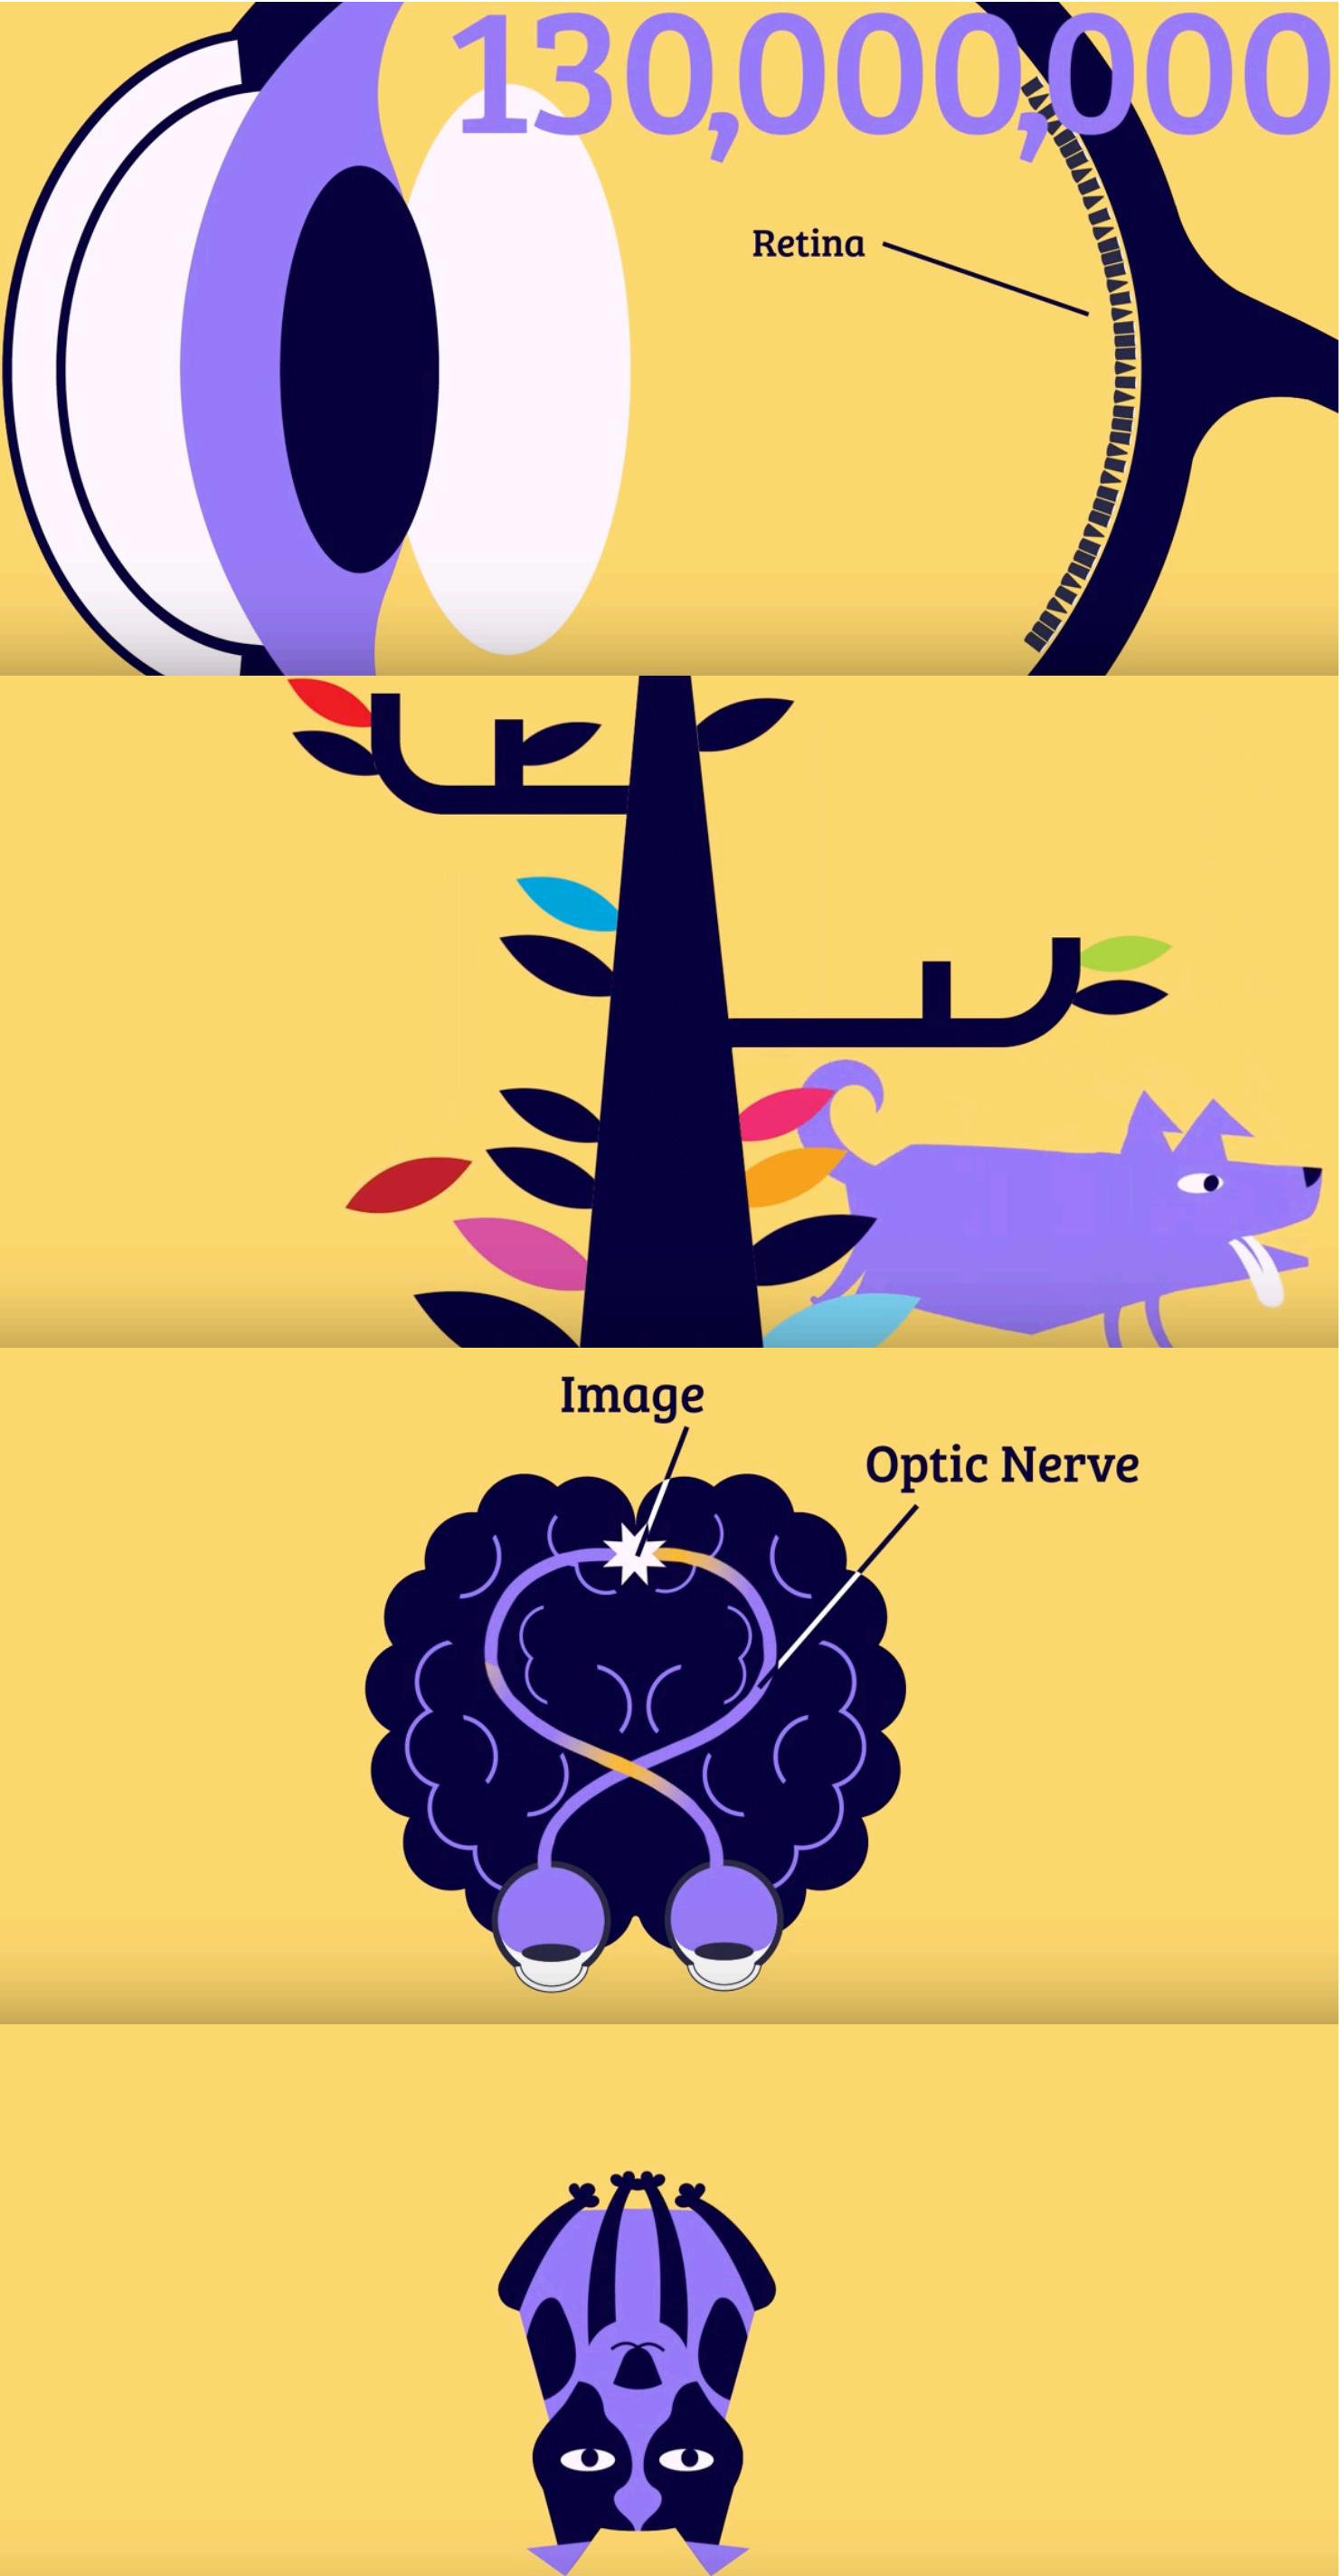
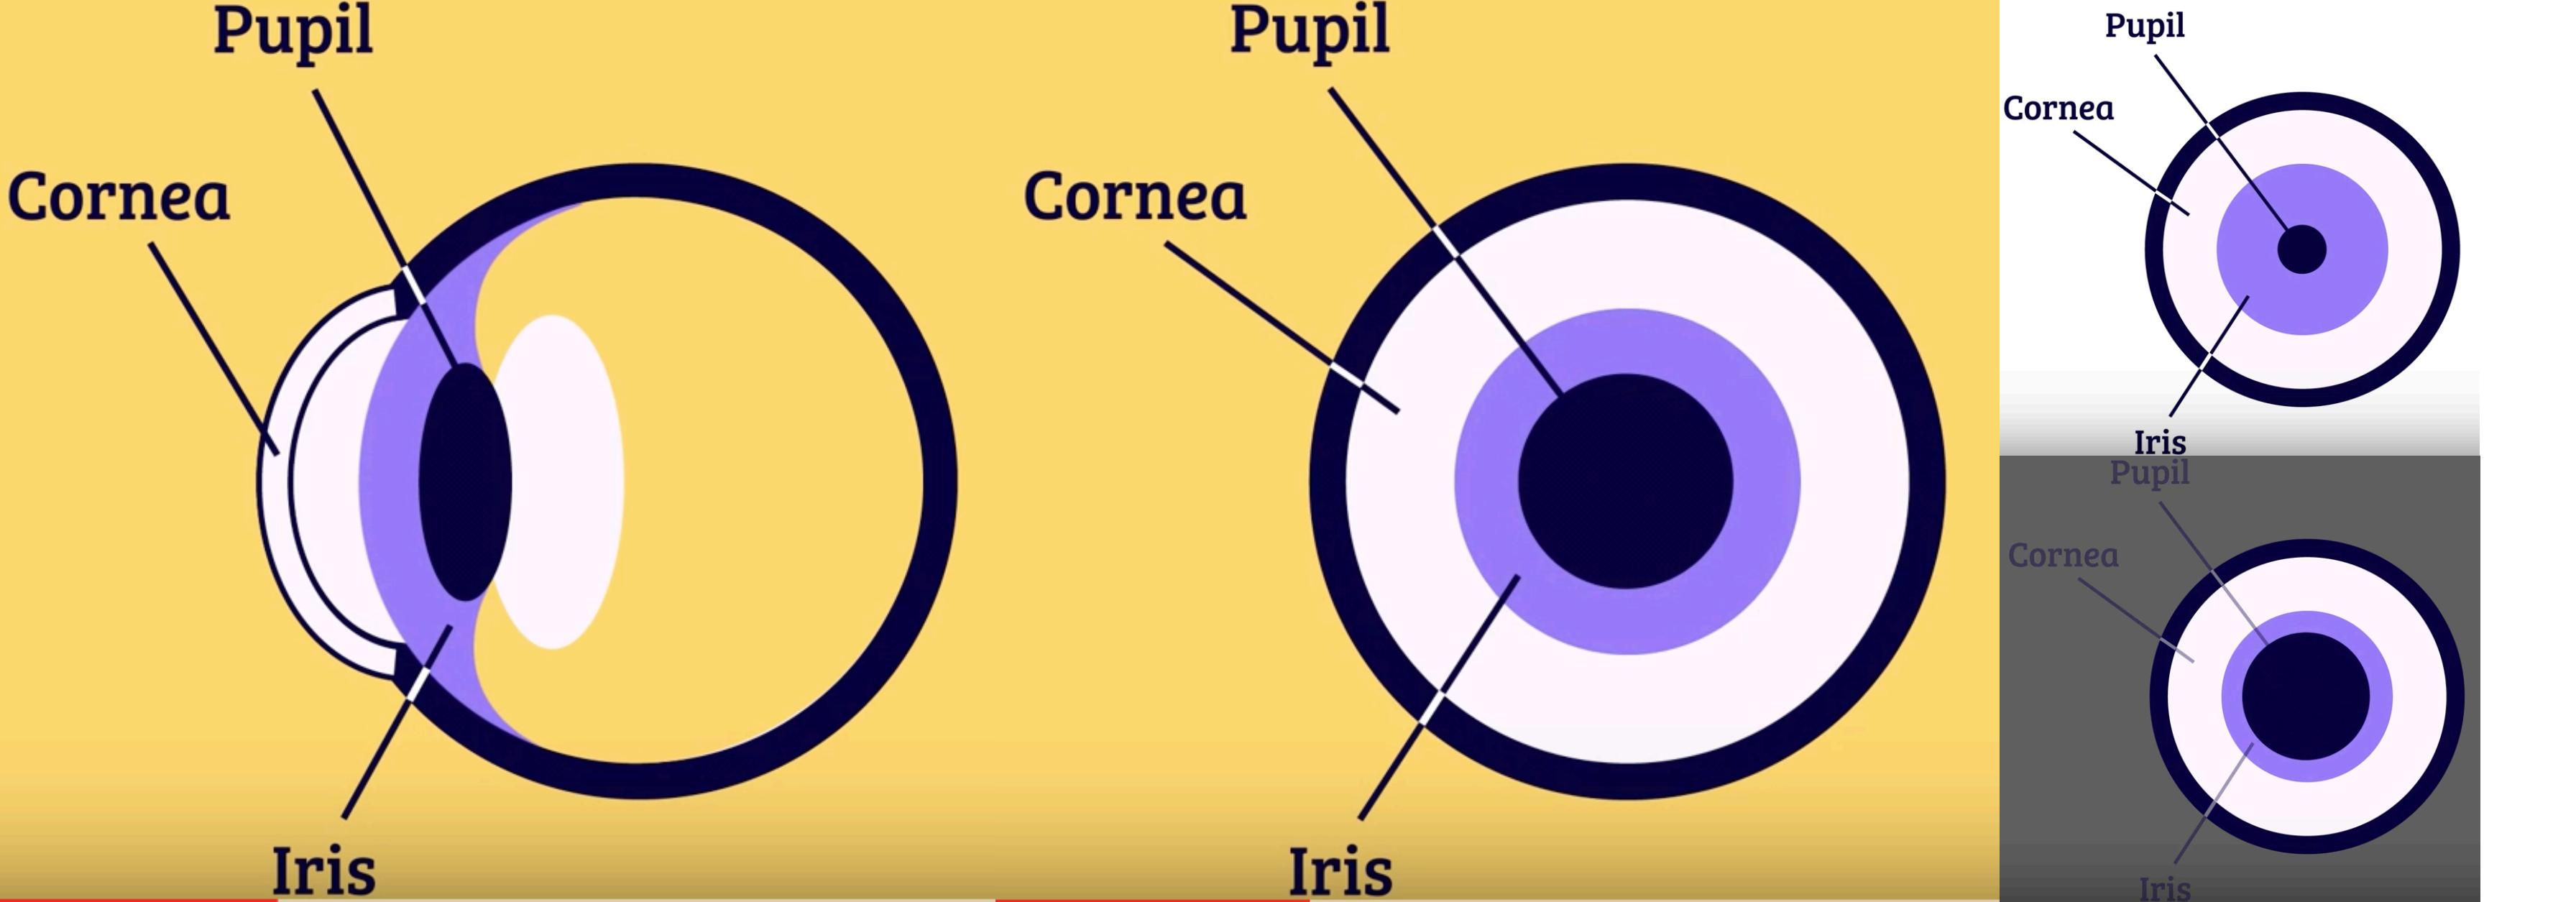
Image formation & Imaging systems



Human Eyes



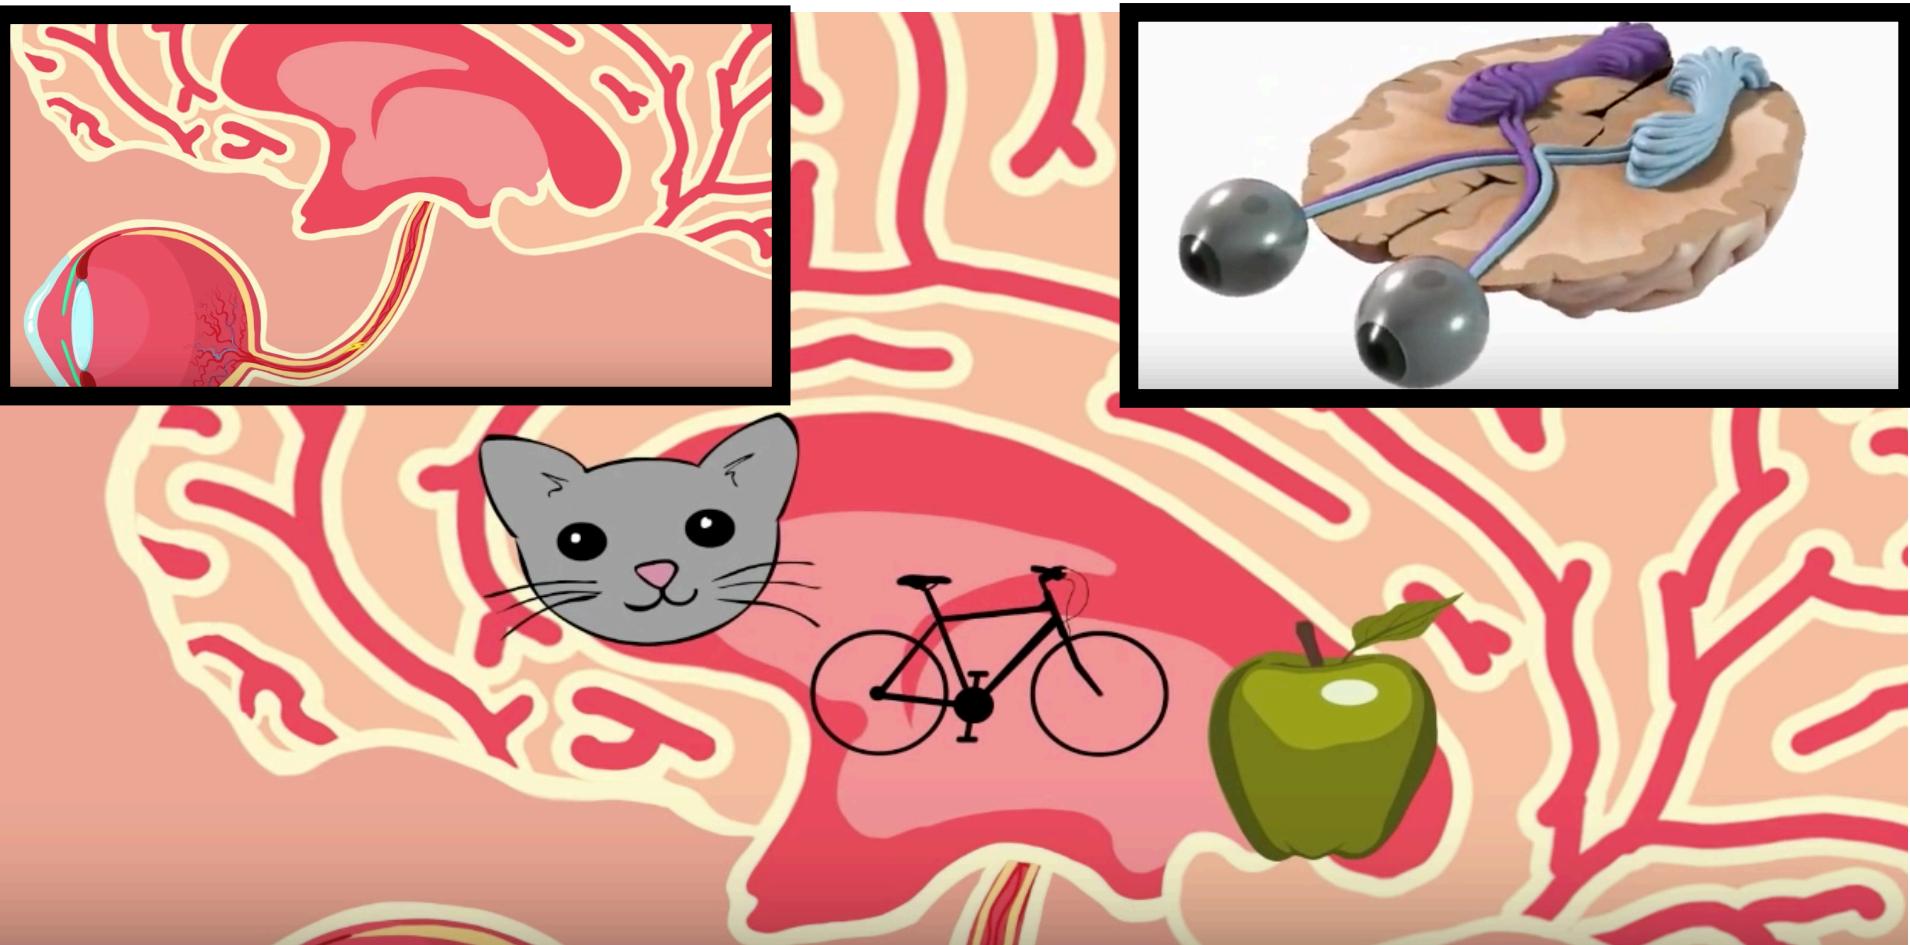
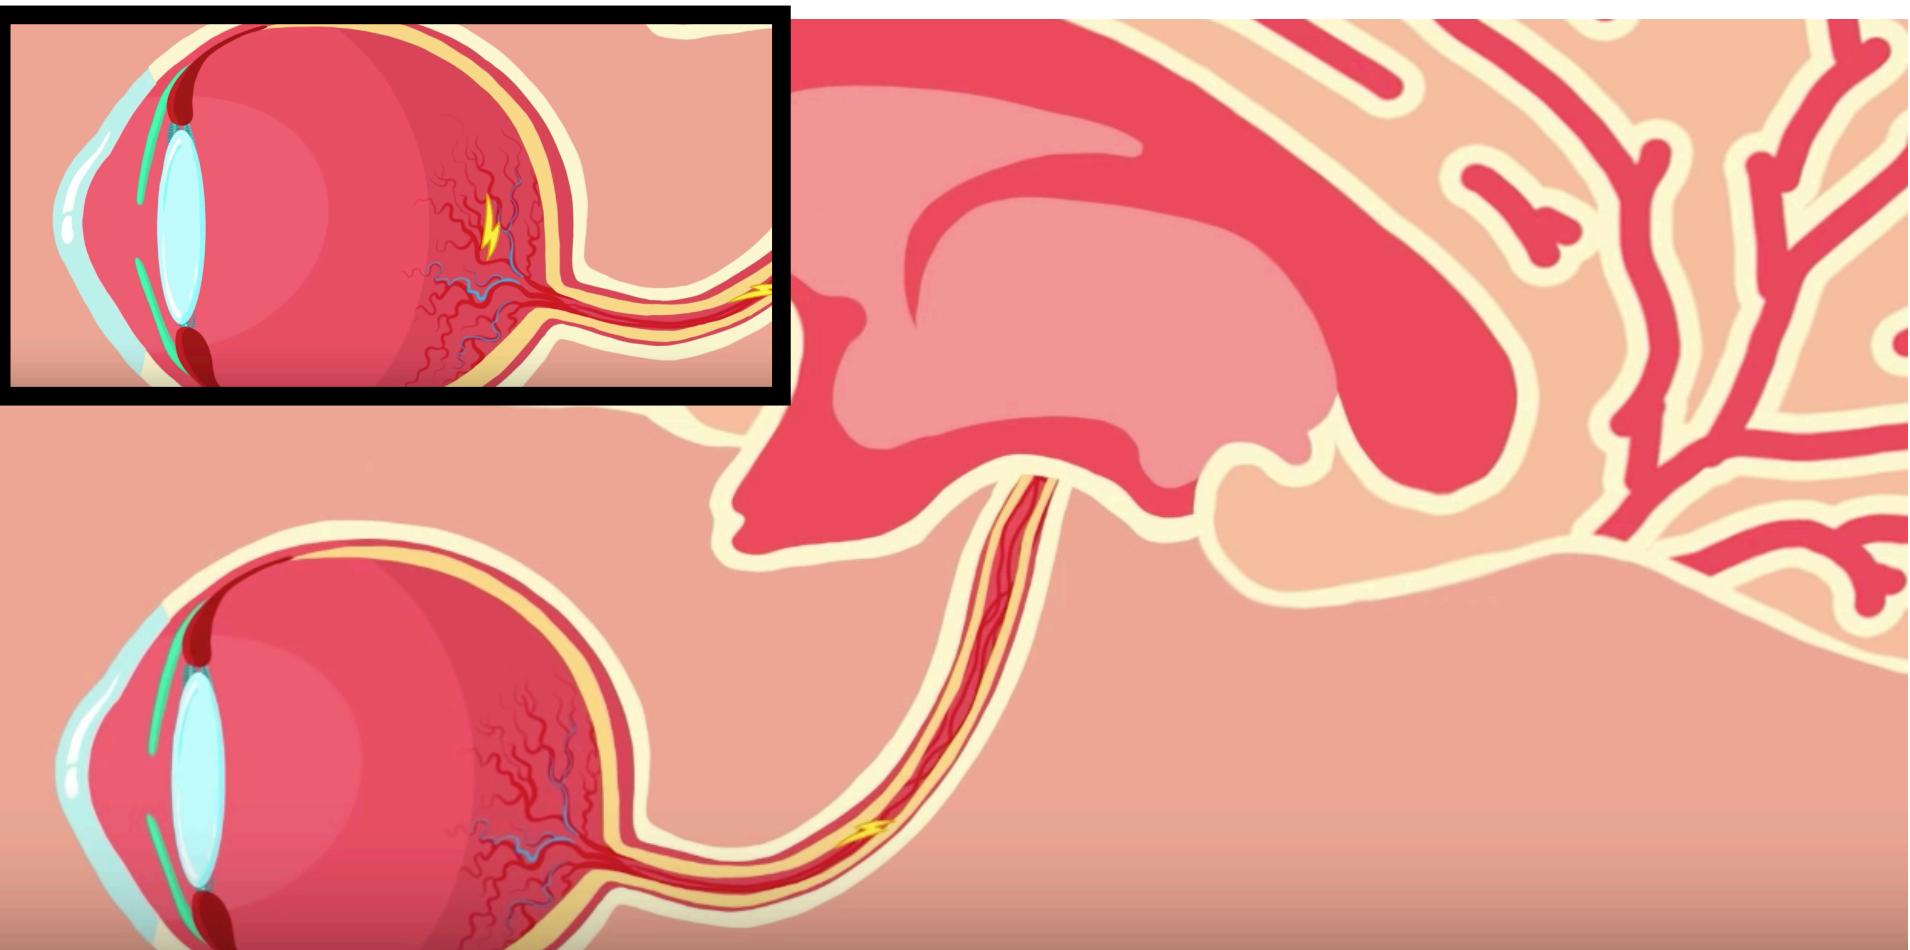
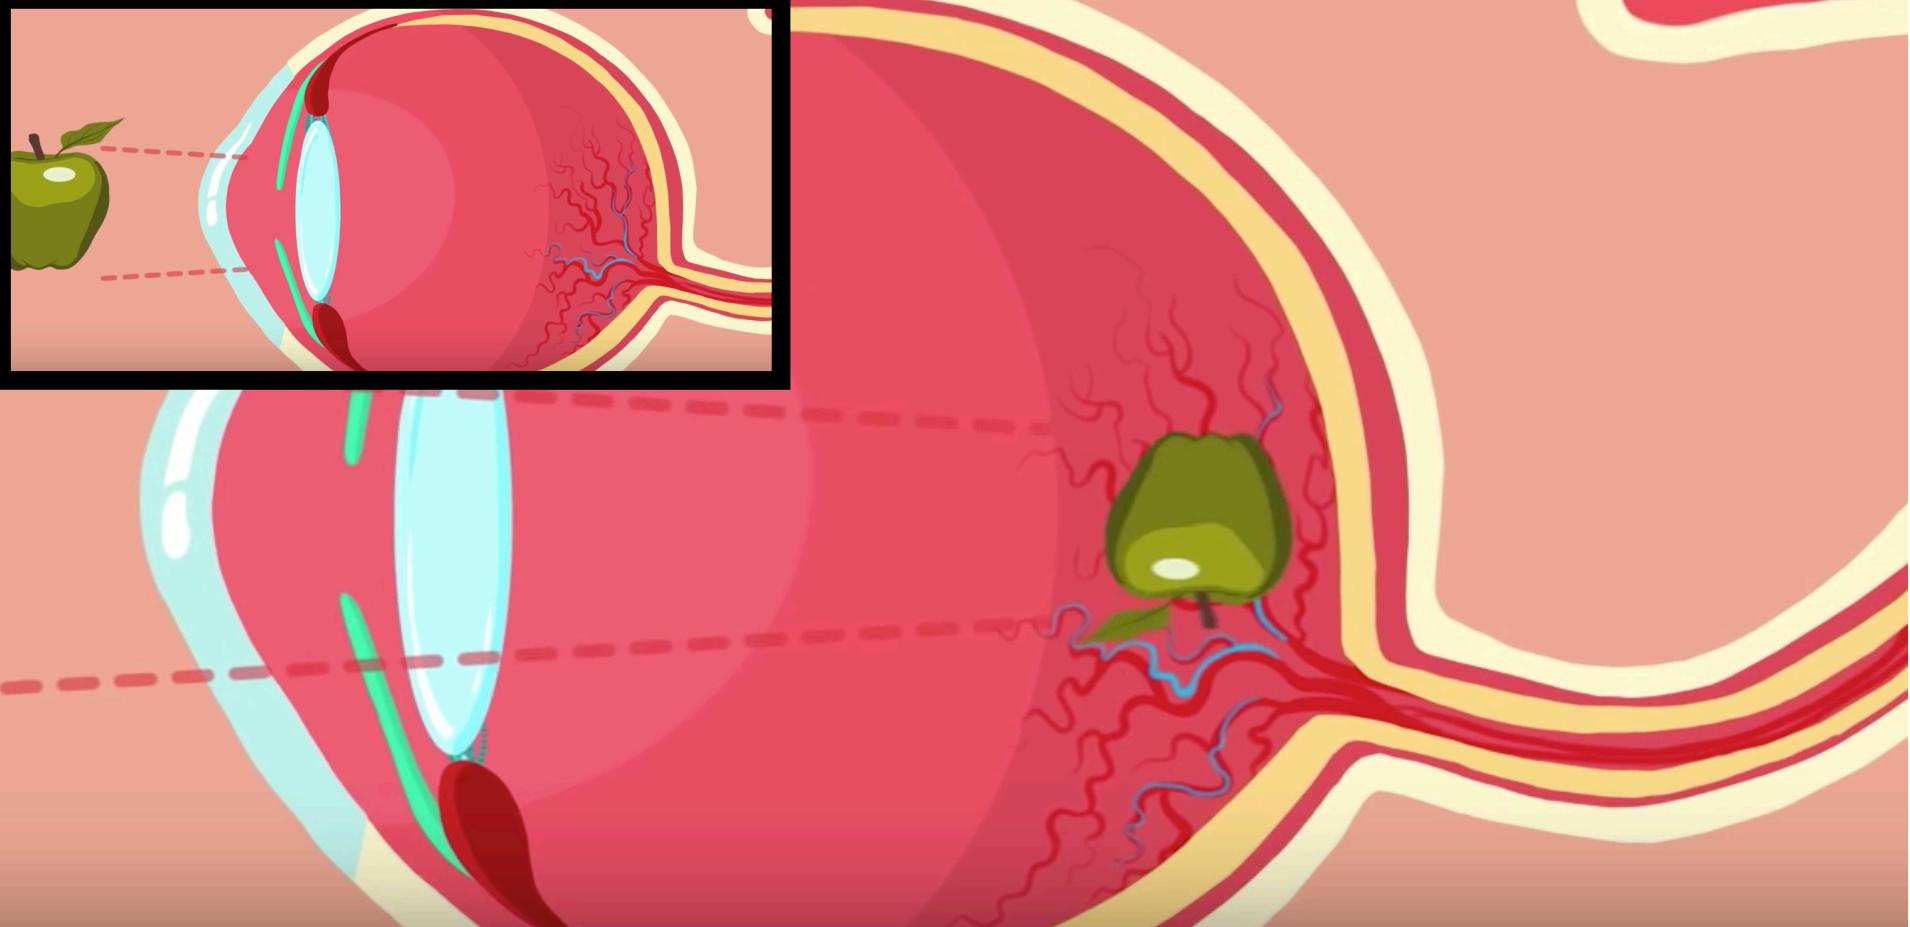
The Visual System: How Your Eyes Work



Human Eyes (2)



How Your Eyes Work (2)



Human Eyes (3)

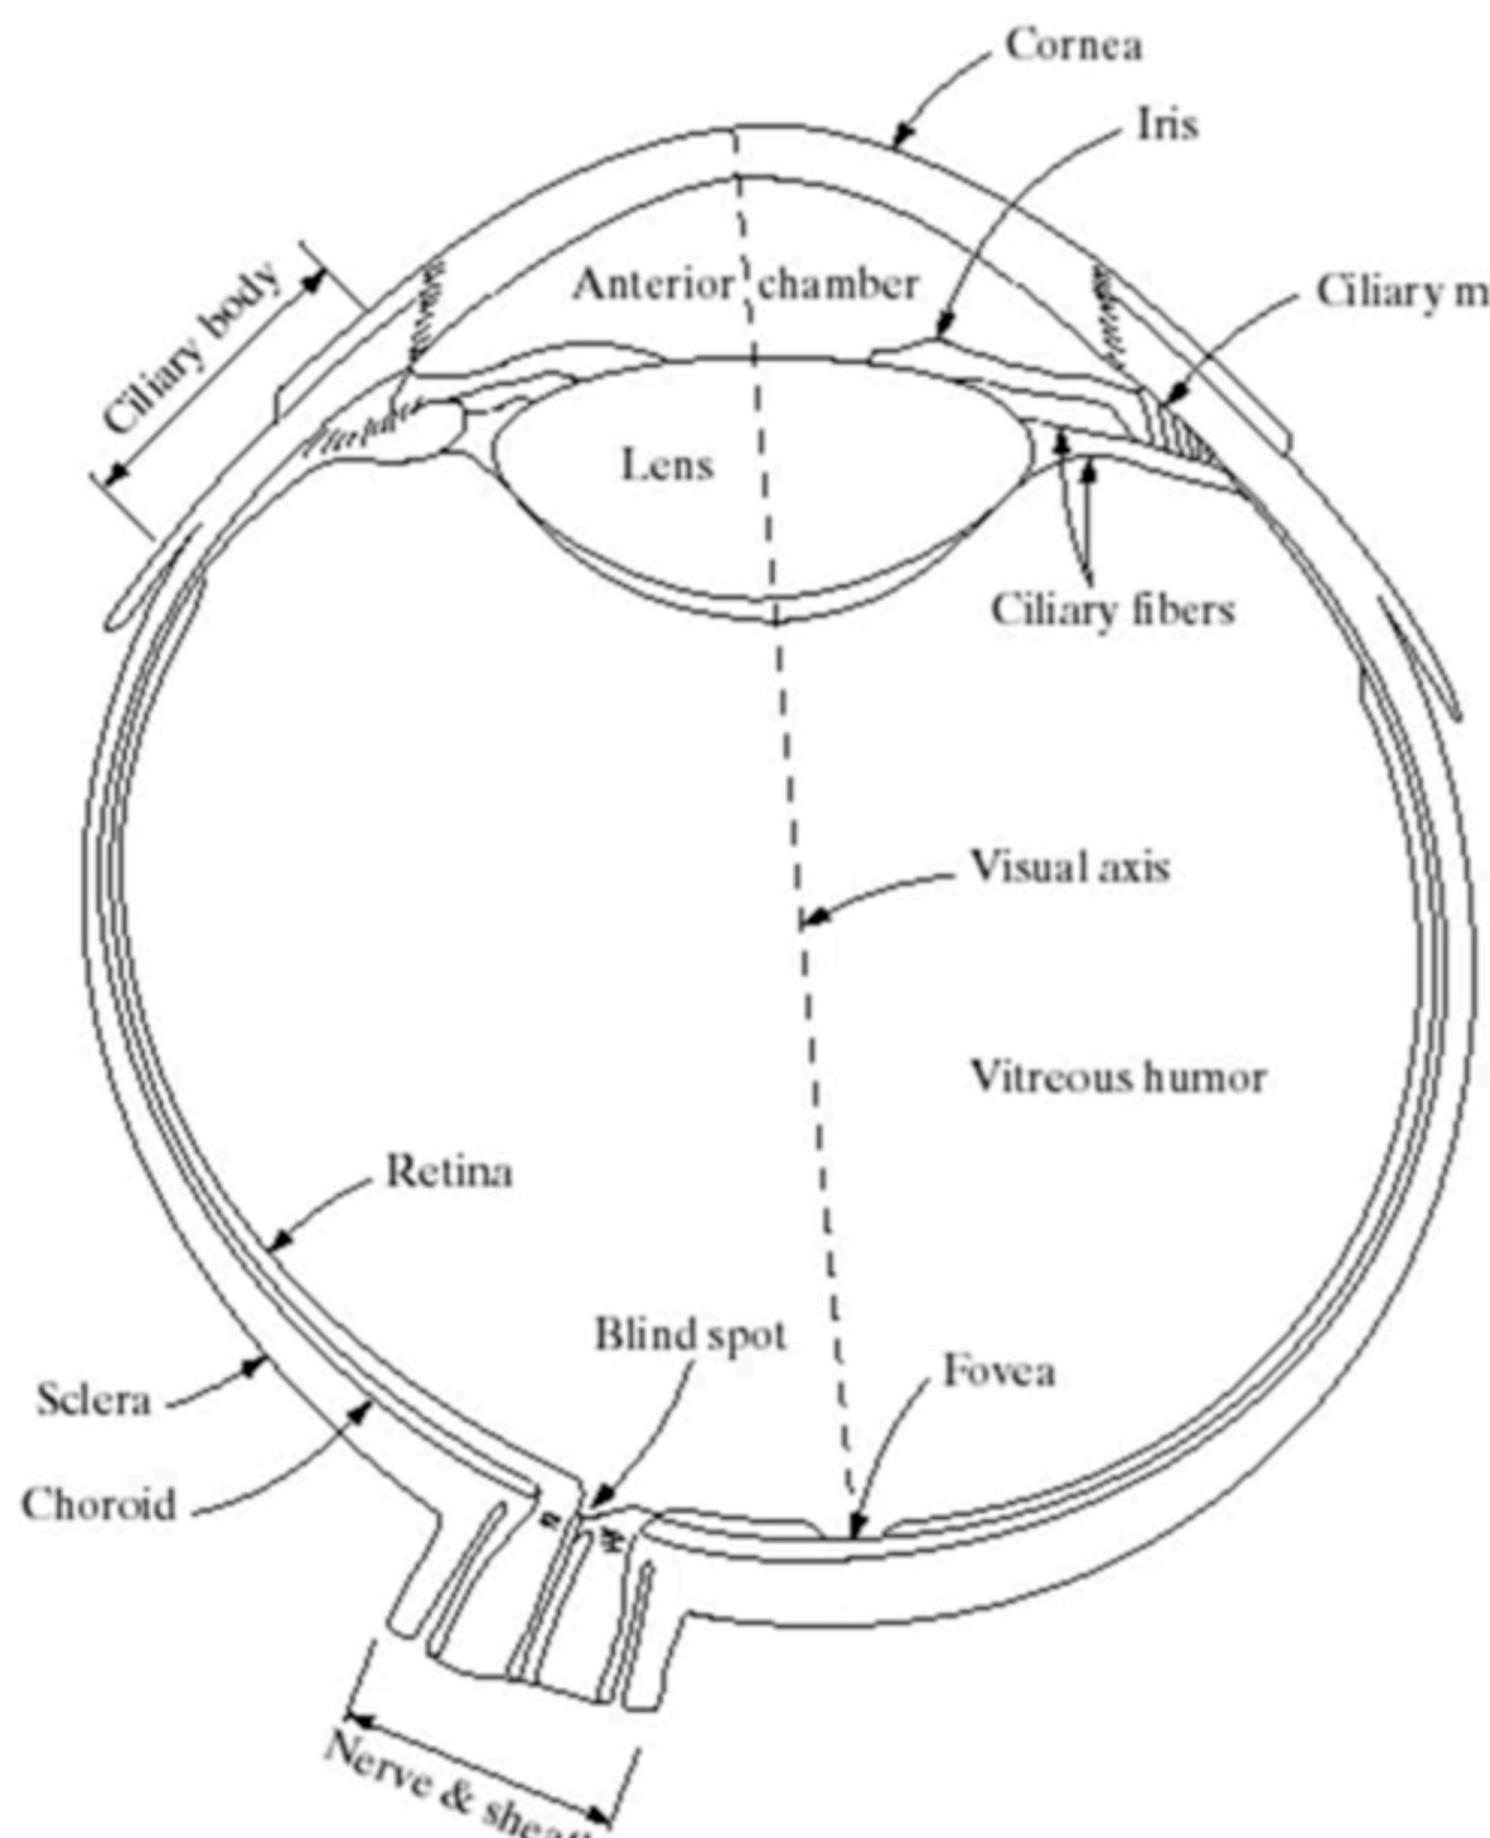


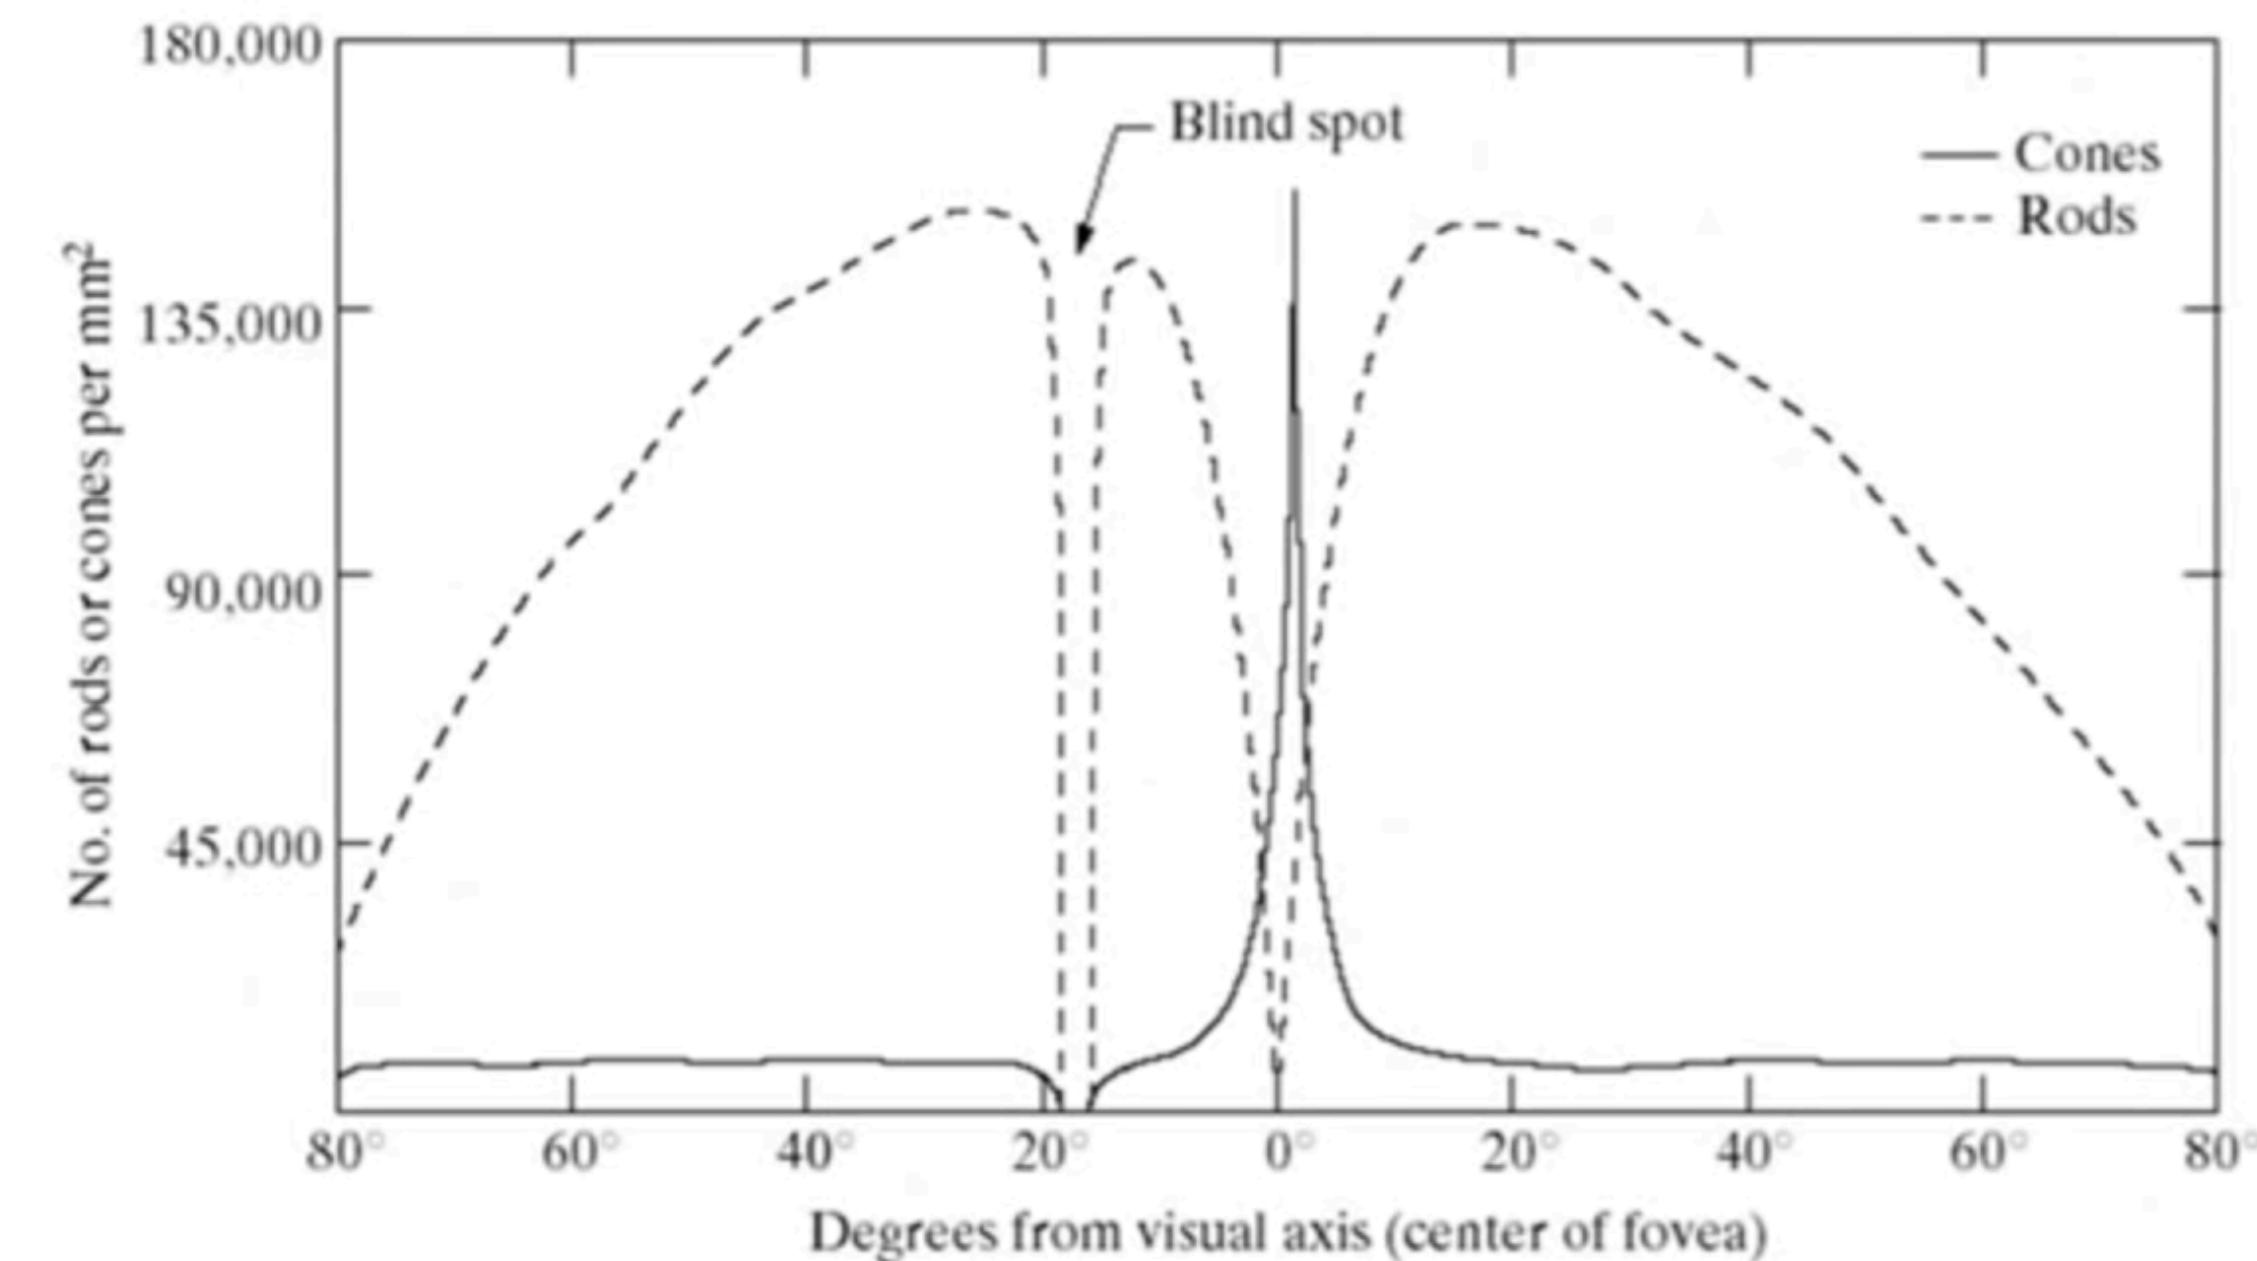
FIGURE 2.1
Simplified
diagram of a cross
section of the
human eye.

Notes:

- Eyeball is about 20 mm in diameter
- Retina contains both rods and cones
- Fovea is about 1.5 mm in width, contains about 337,000 cones

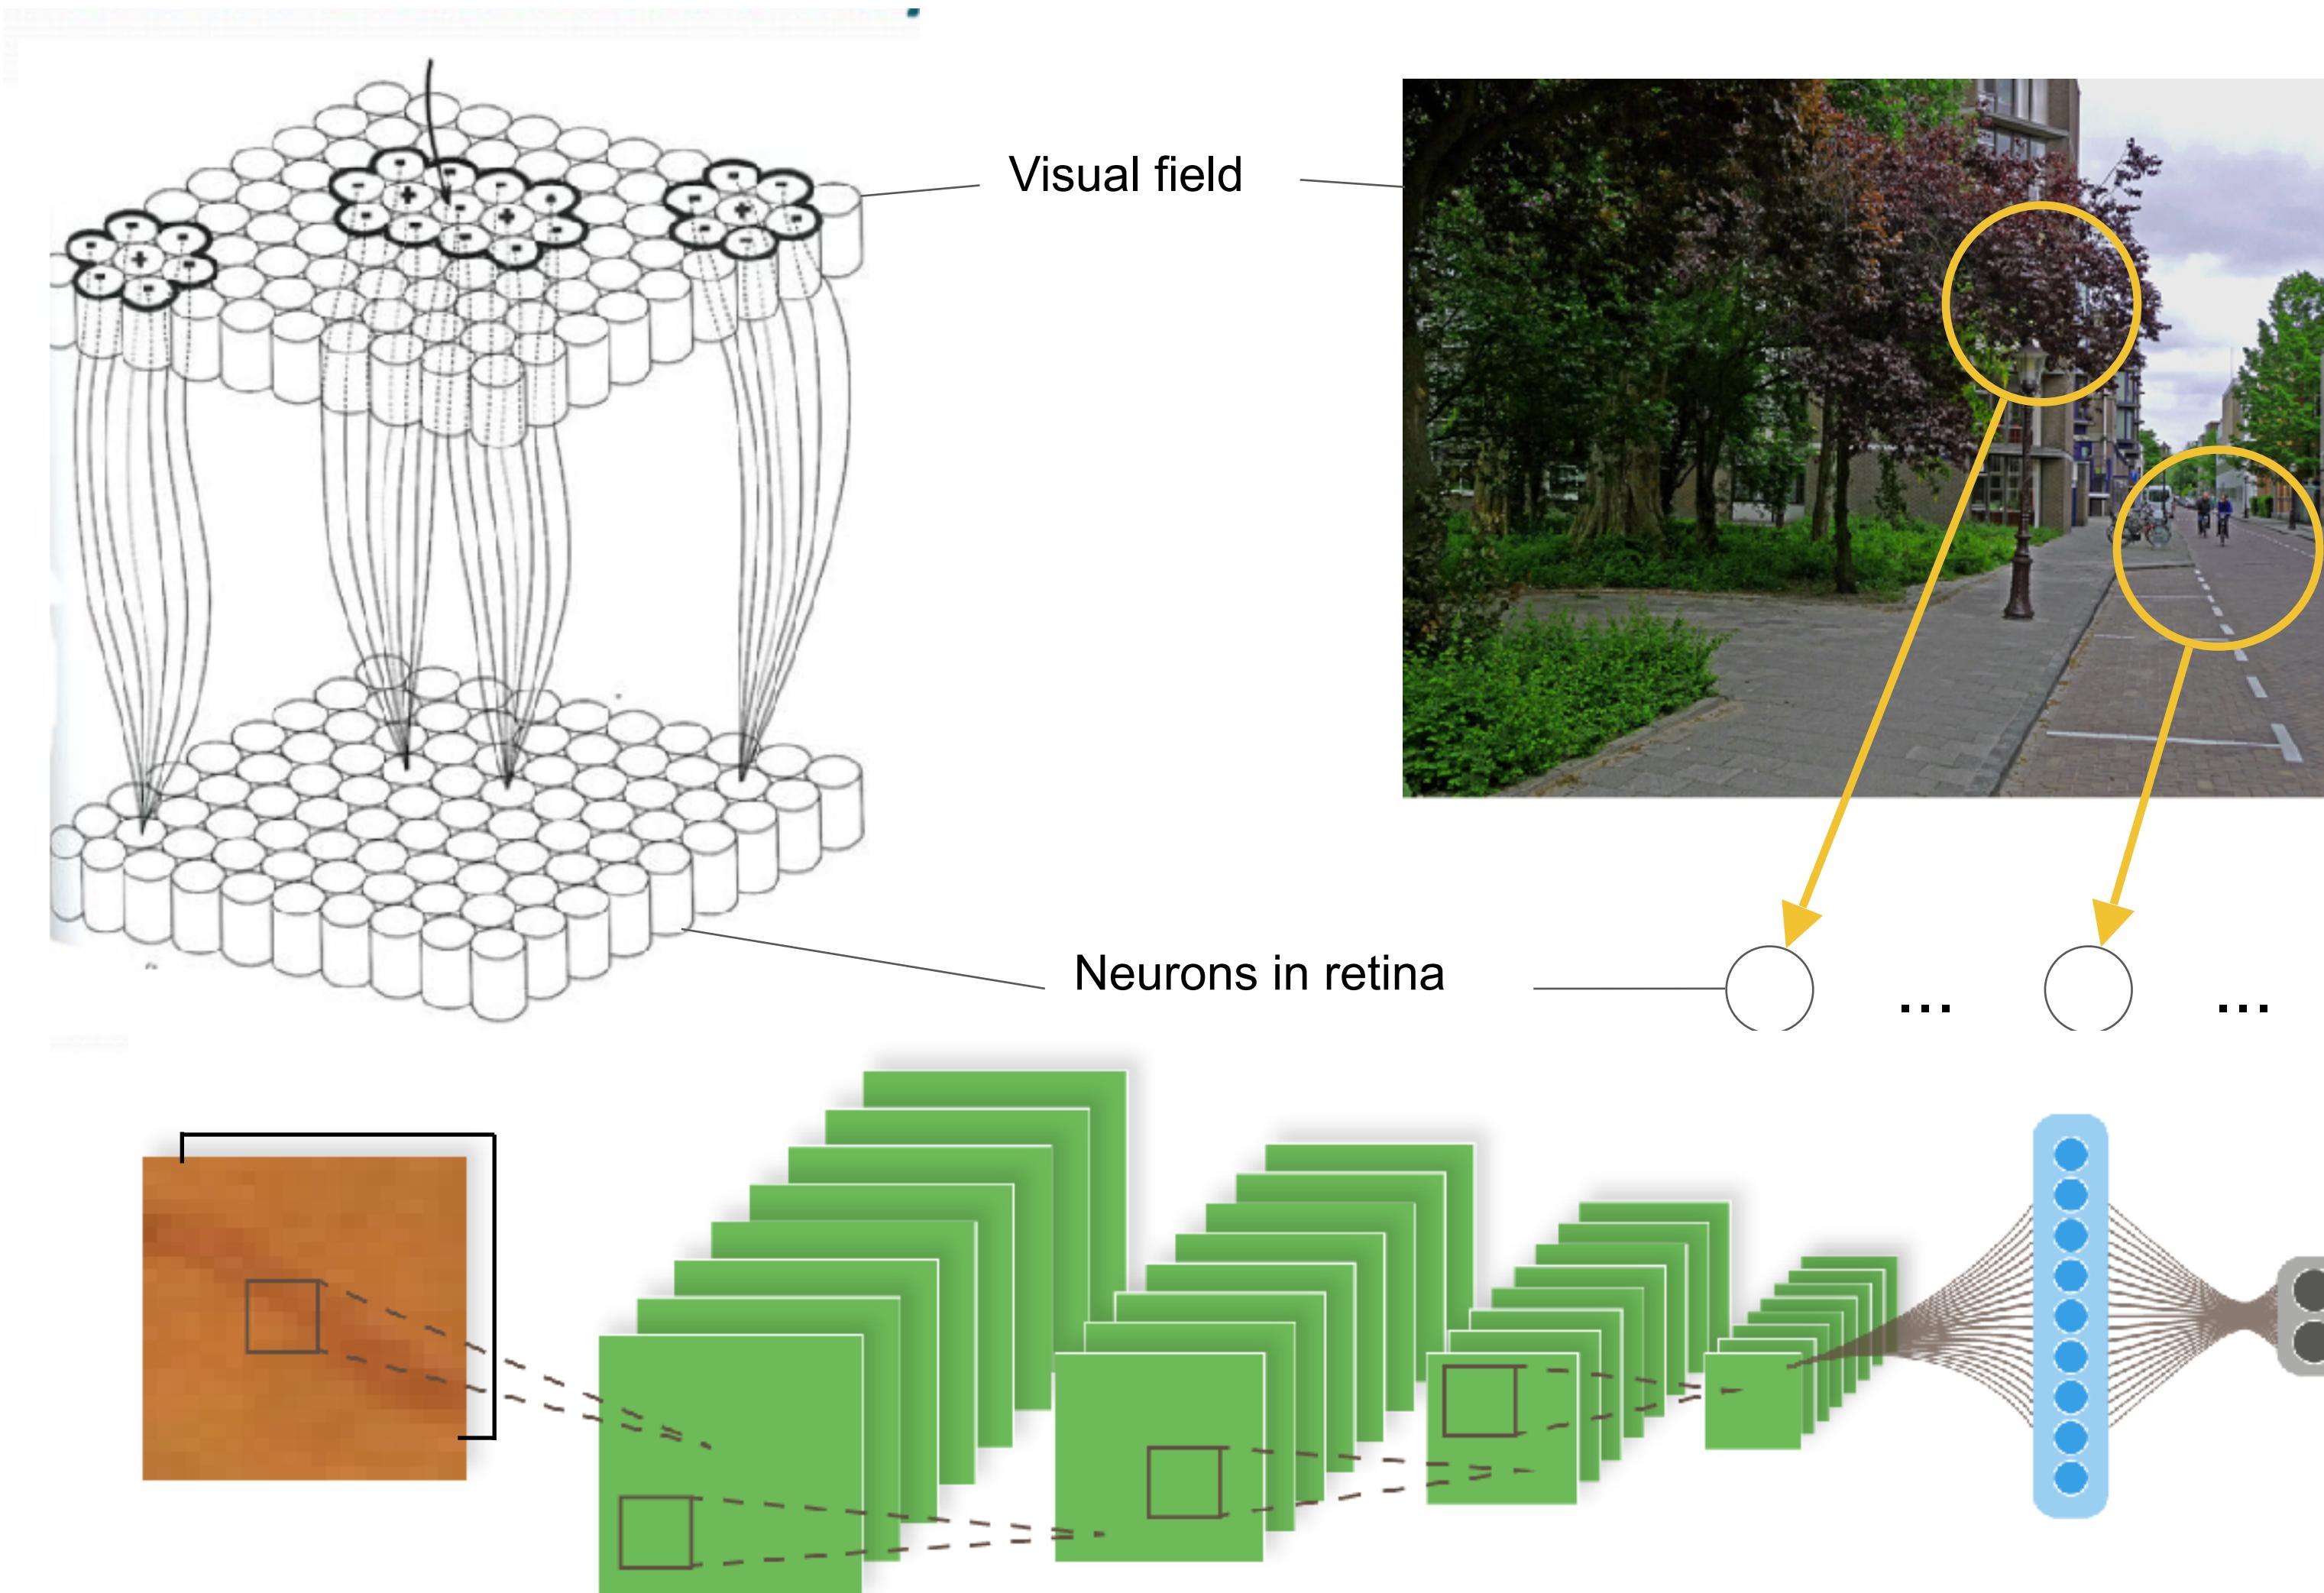
Focal length about 17 mm

- Rods sensitive to low light (scotopic vision)
- Cones detect color, work in bright light (photopic vision)



Biological receptive fields

- Specific neurons in the retina cares about only a small part of the visual field



Biological vs. Artificial Neuron and Network

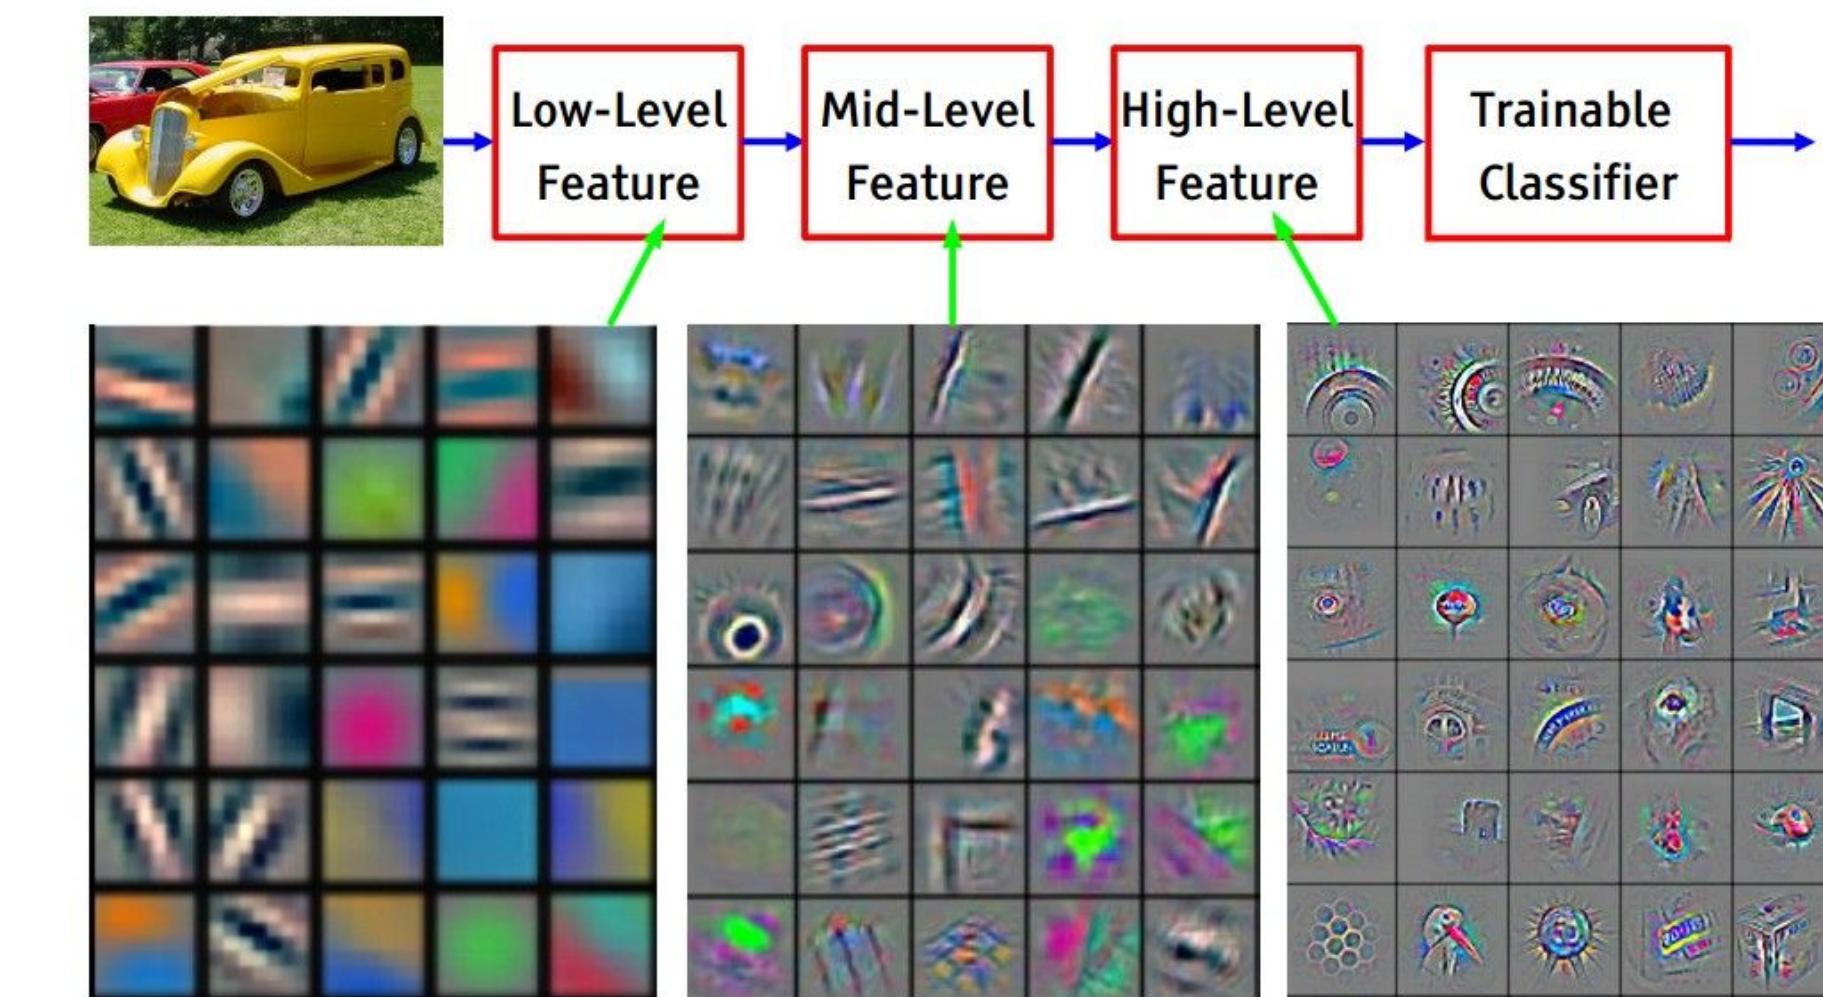
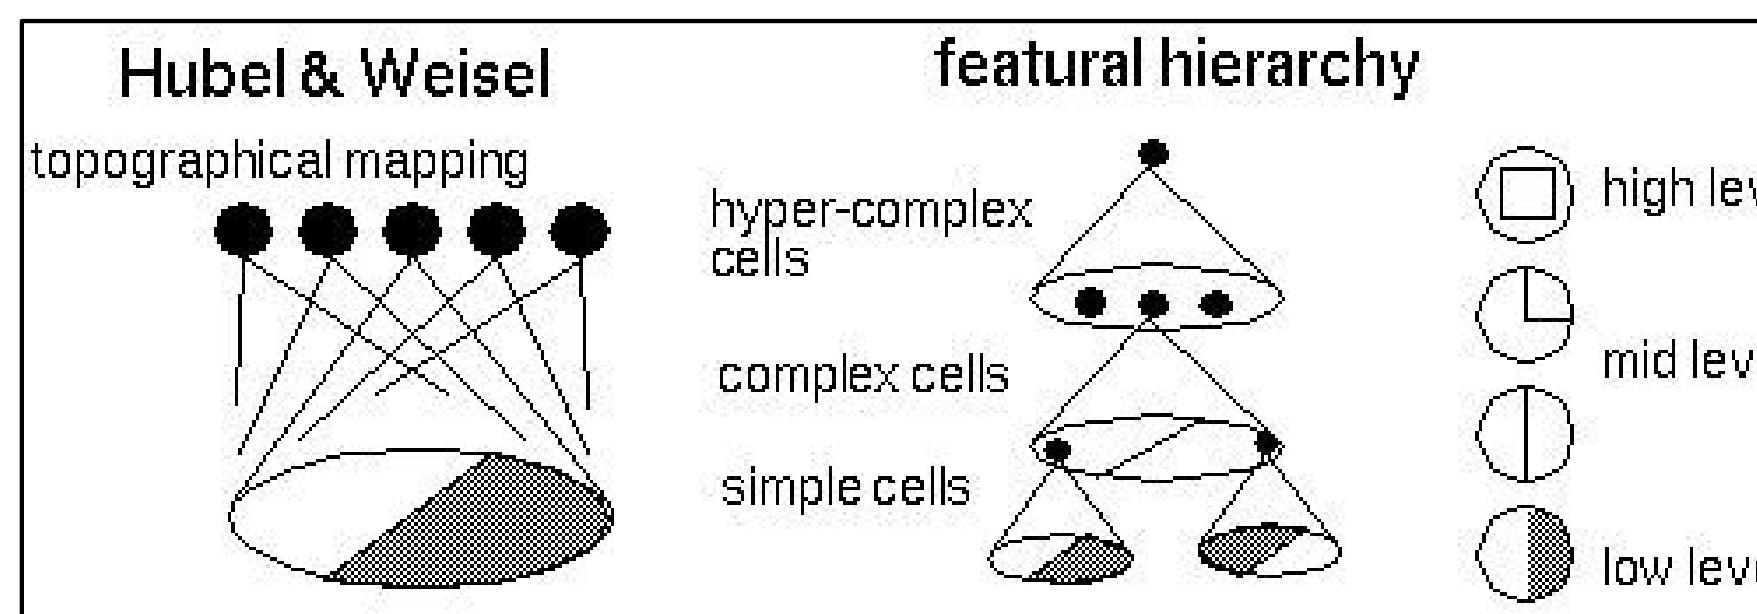
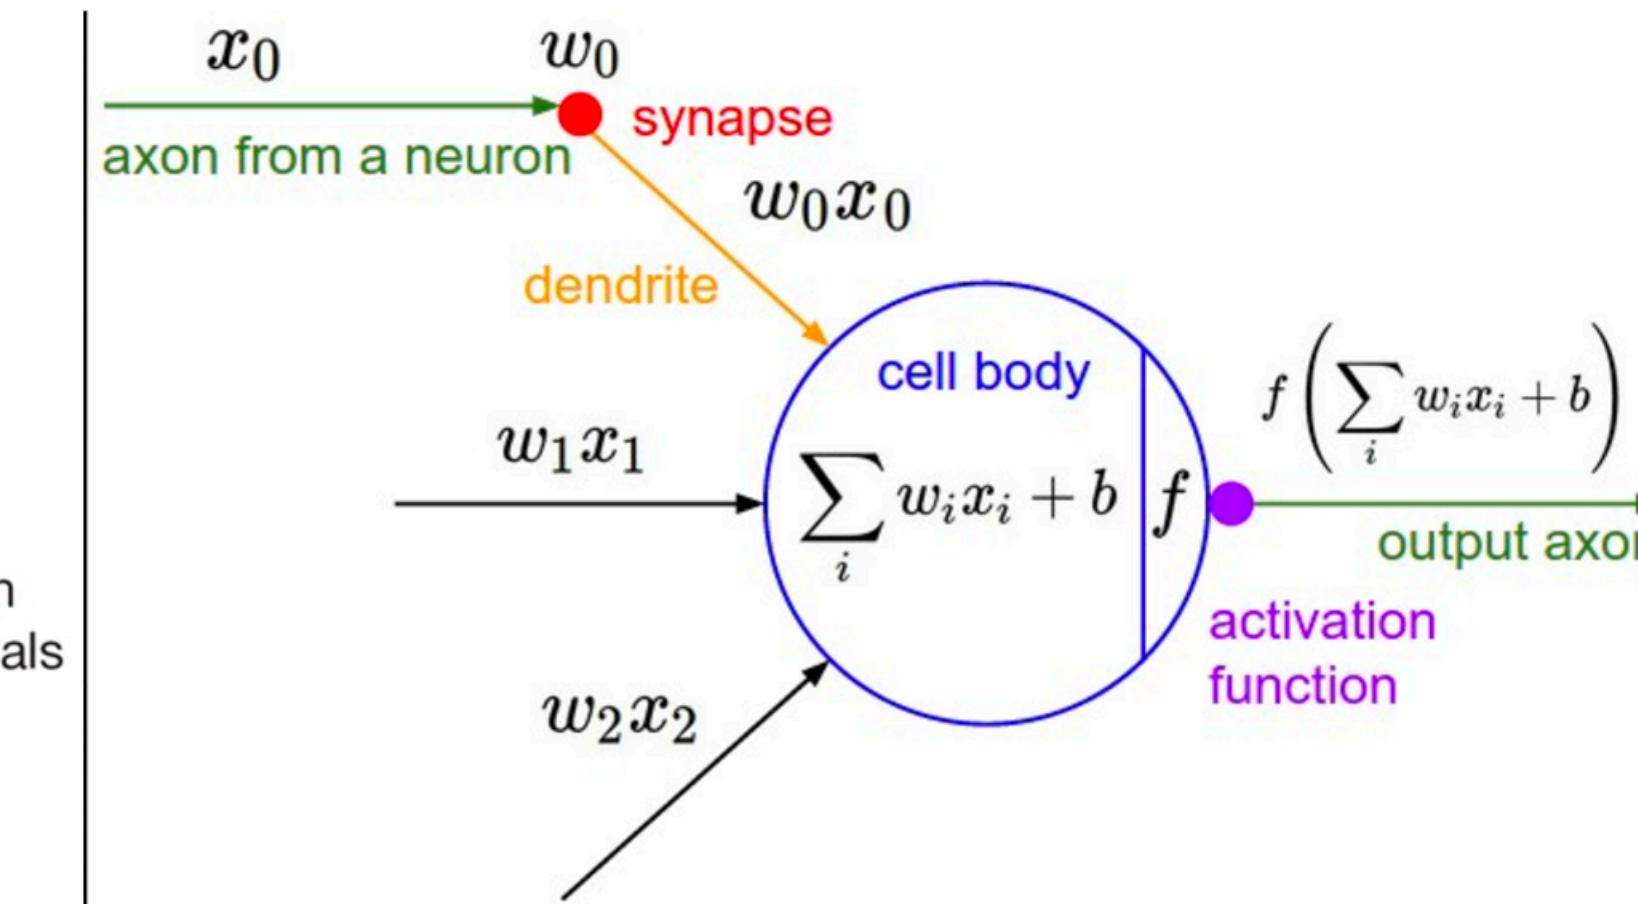
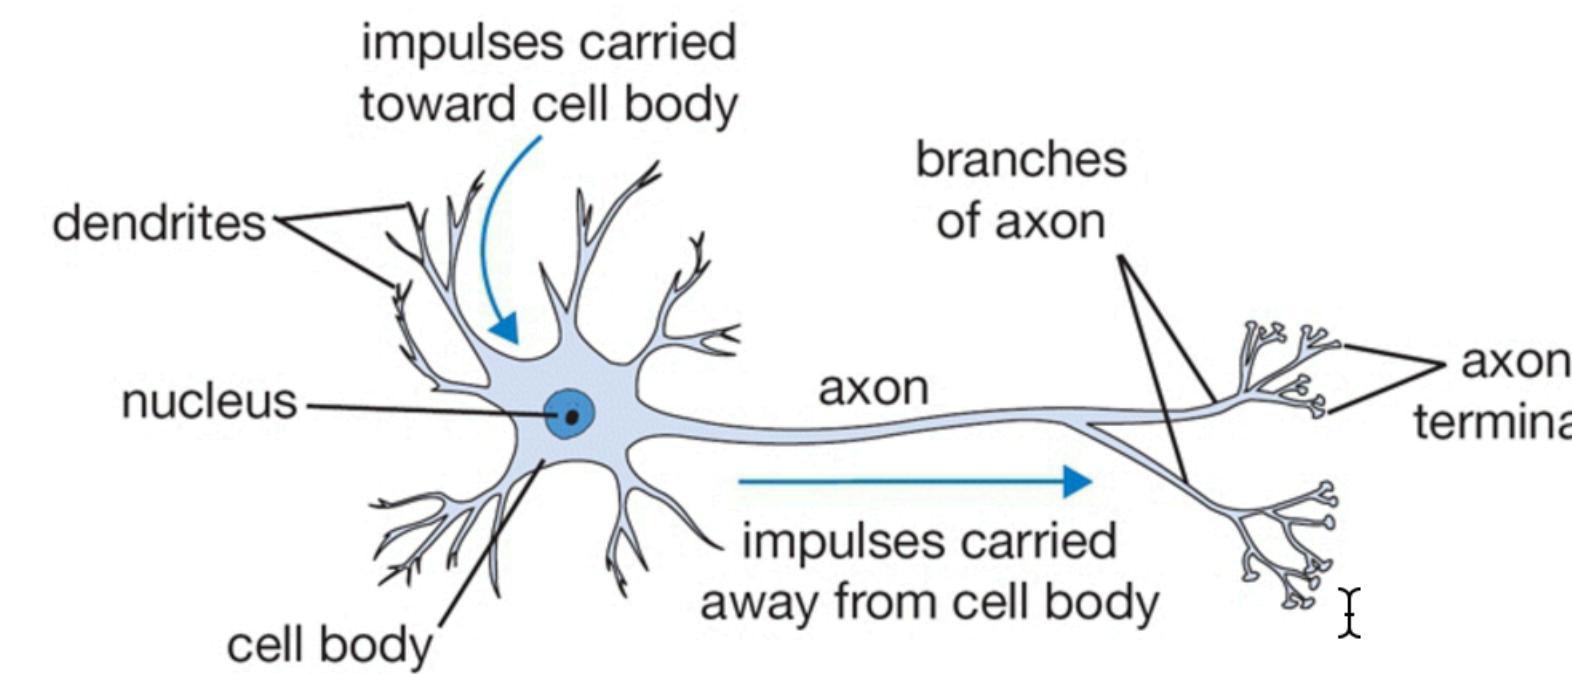
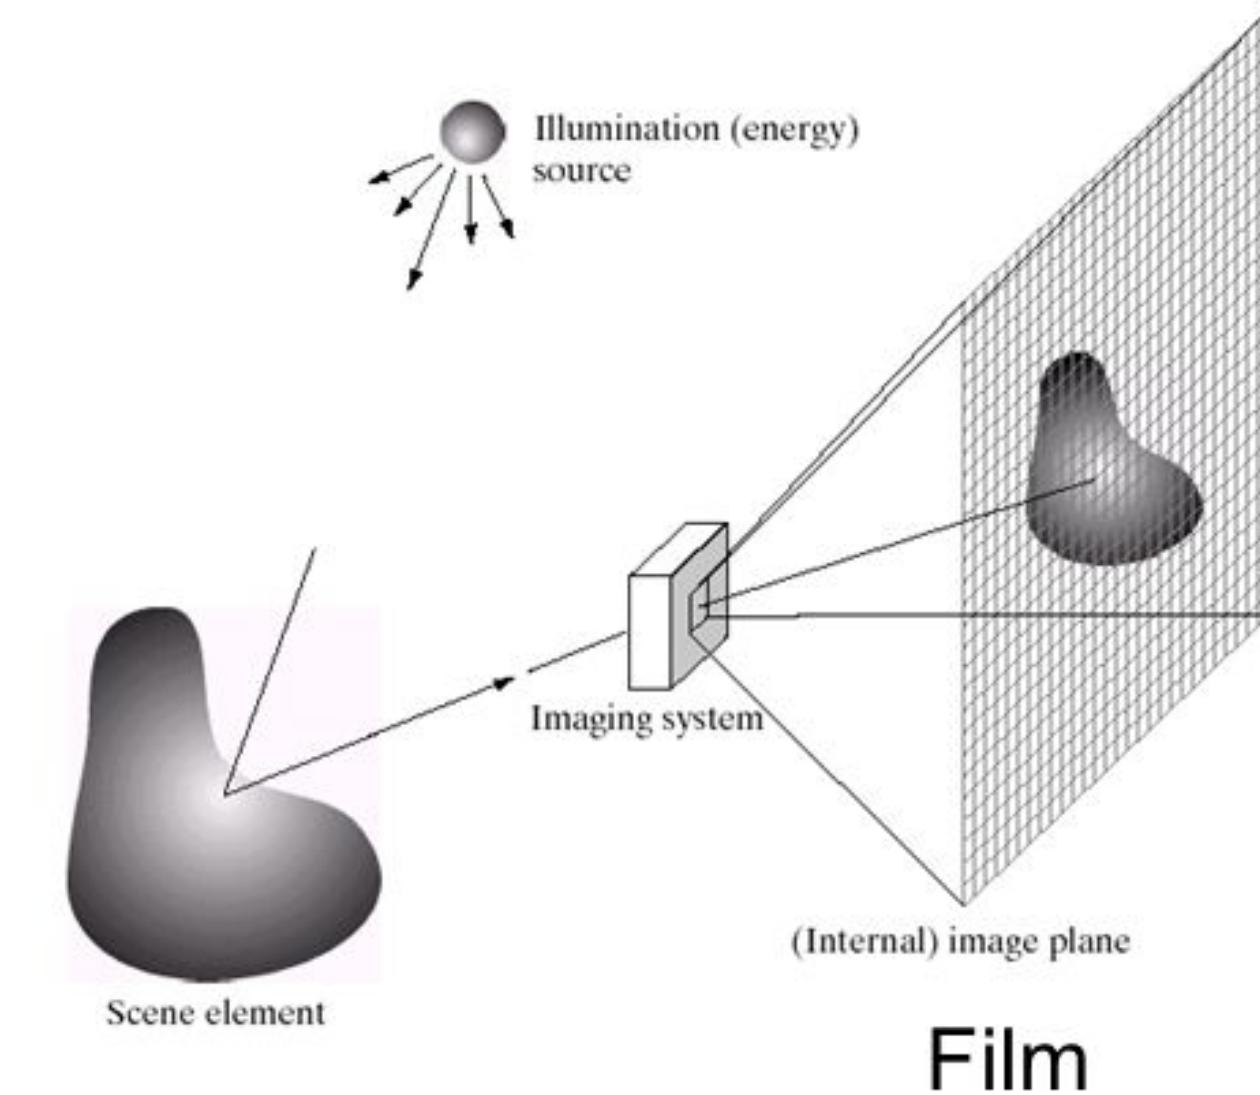
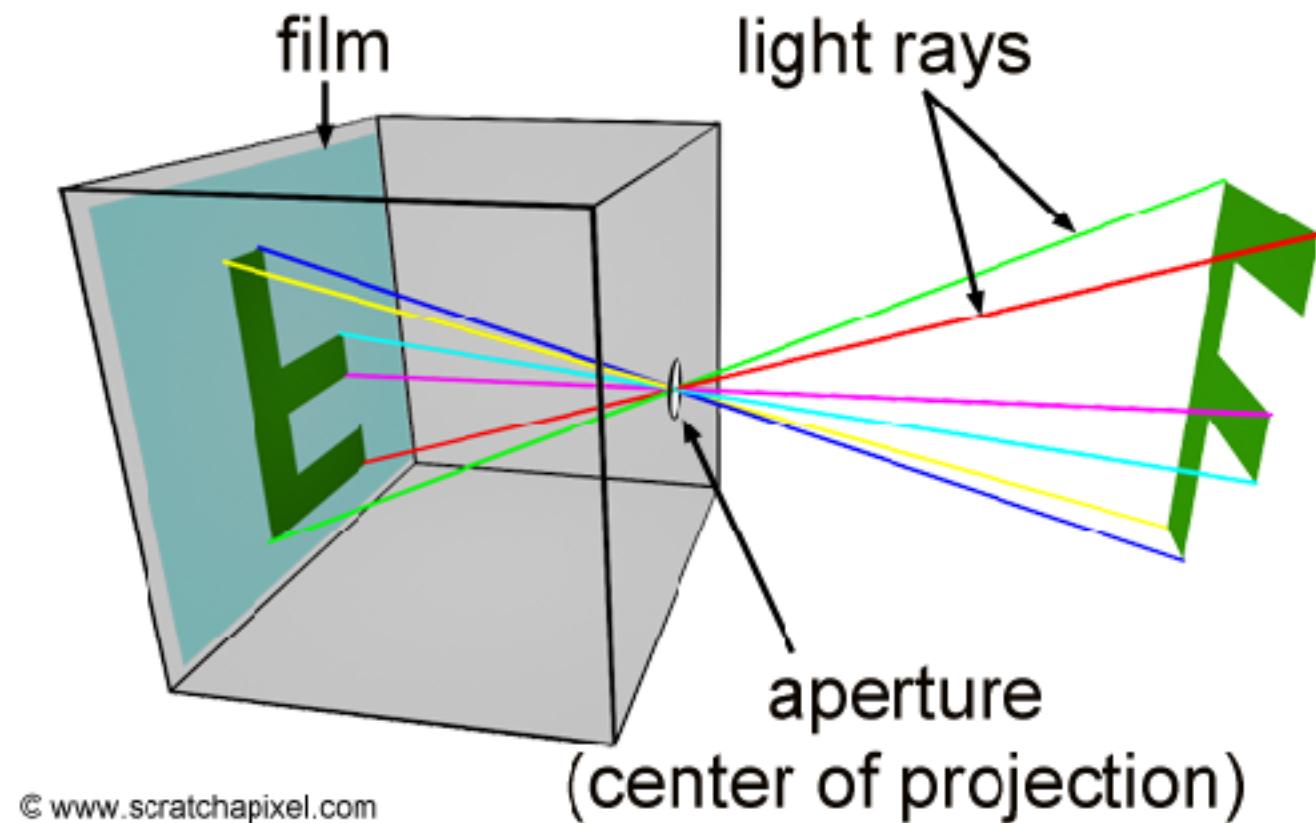


Image formation & the Pinhole camera

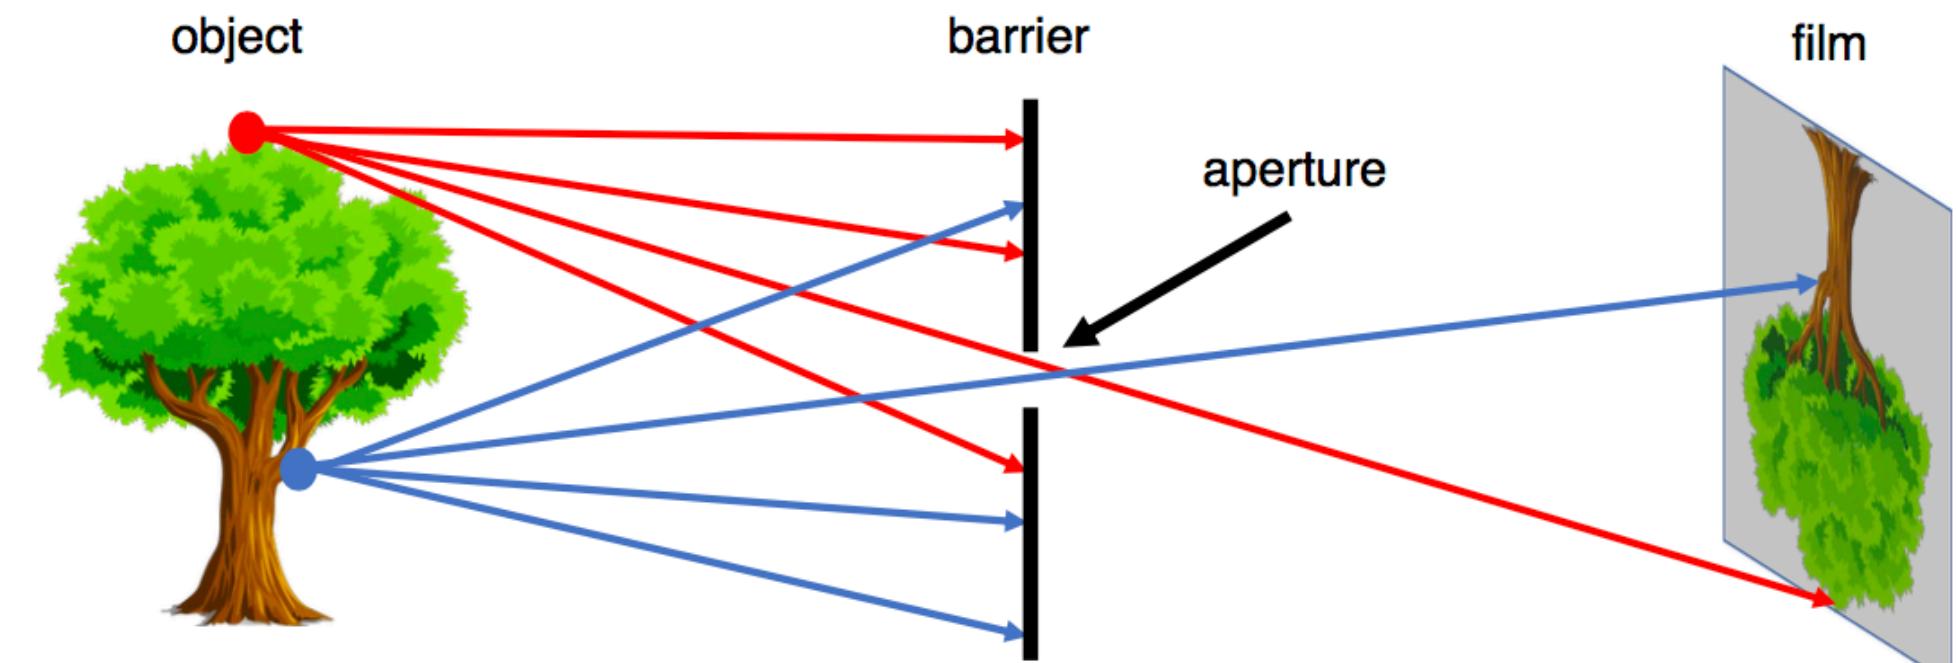
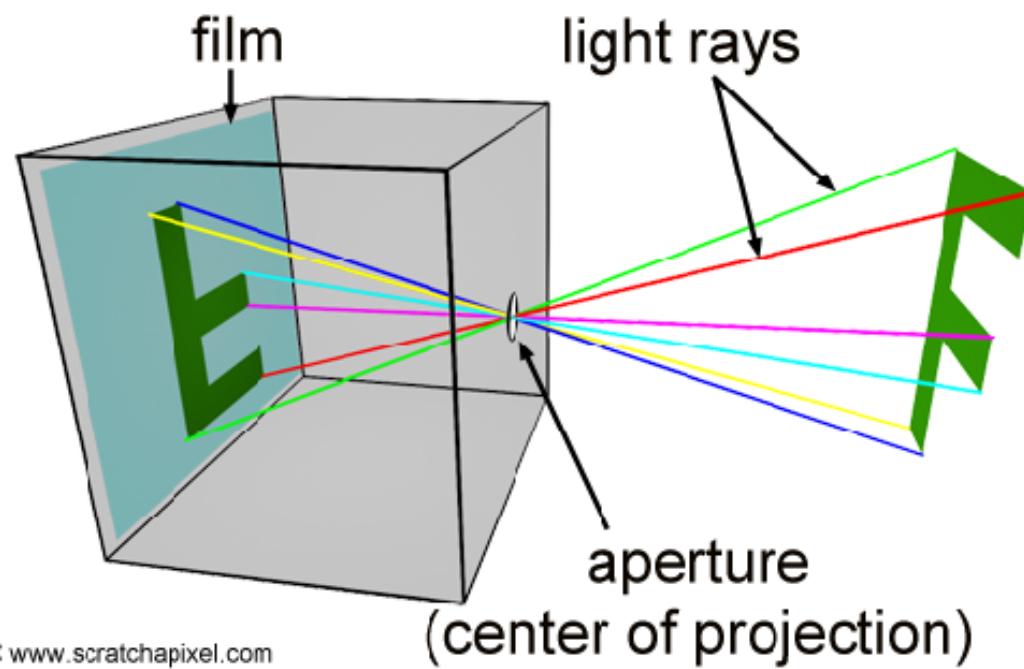
- The **image** $I(x,y)$ measures how much light that is captured at pixel (x,y)



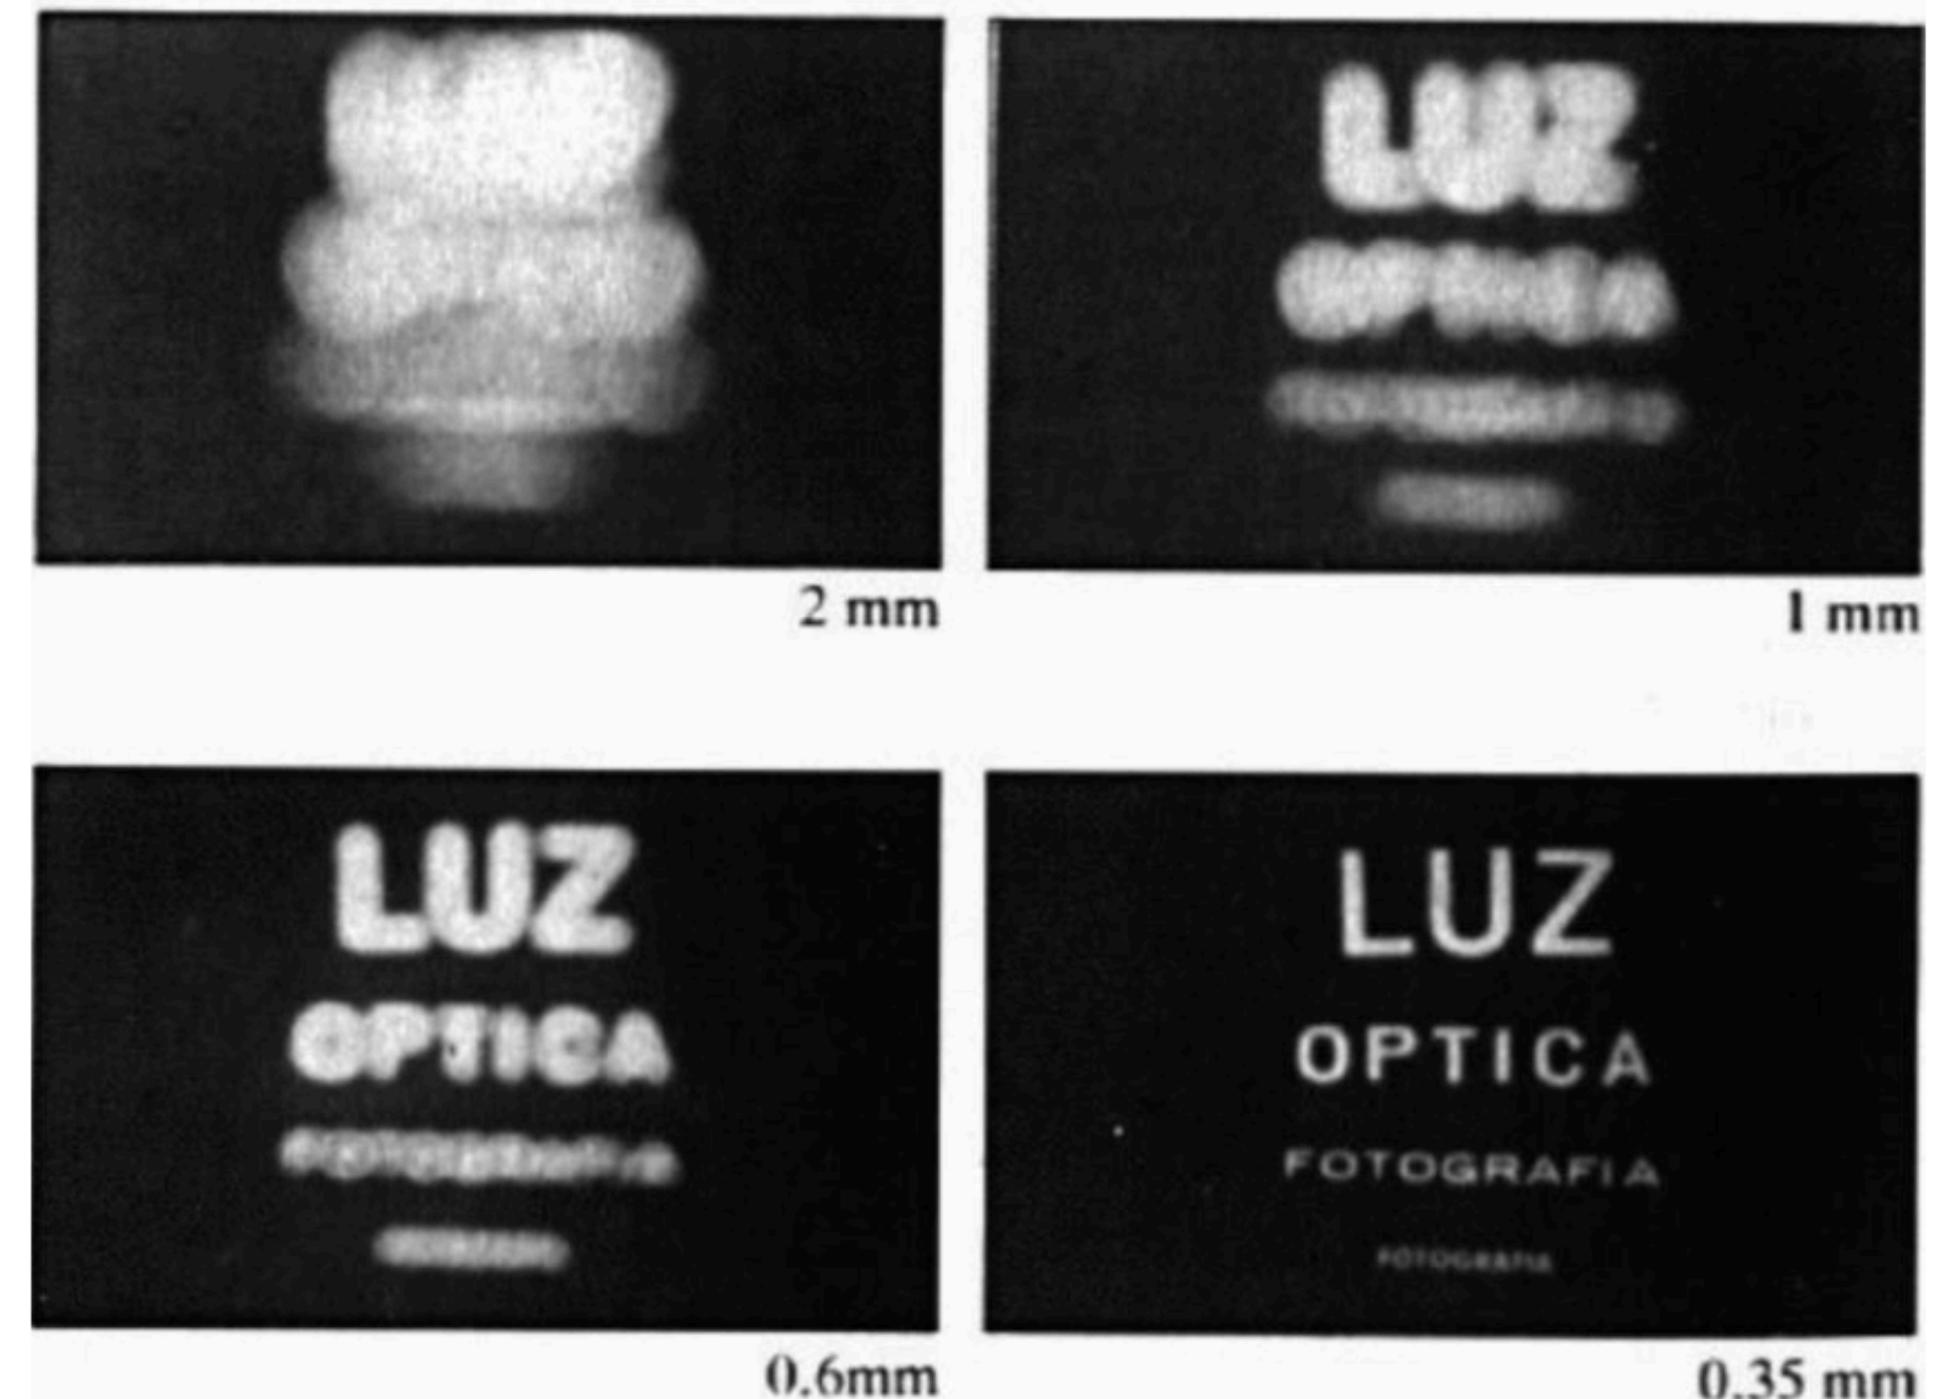
- We want to know:
 - Where does a point $P(X,Y, Z)$ in the world get imaged?
 - What is the brightness at the resulting point $p(x,y)$?



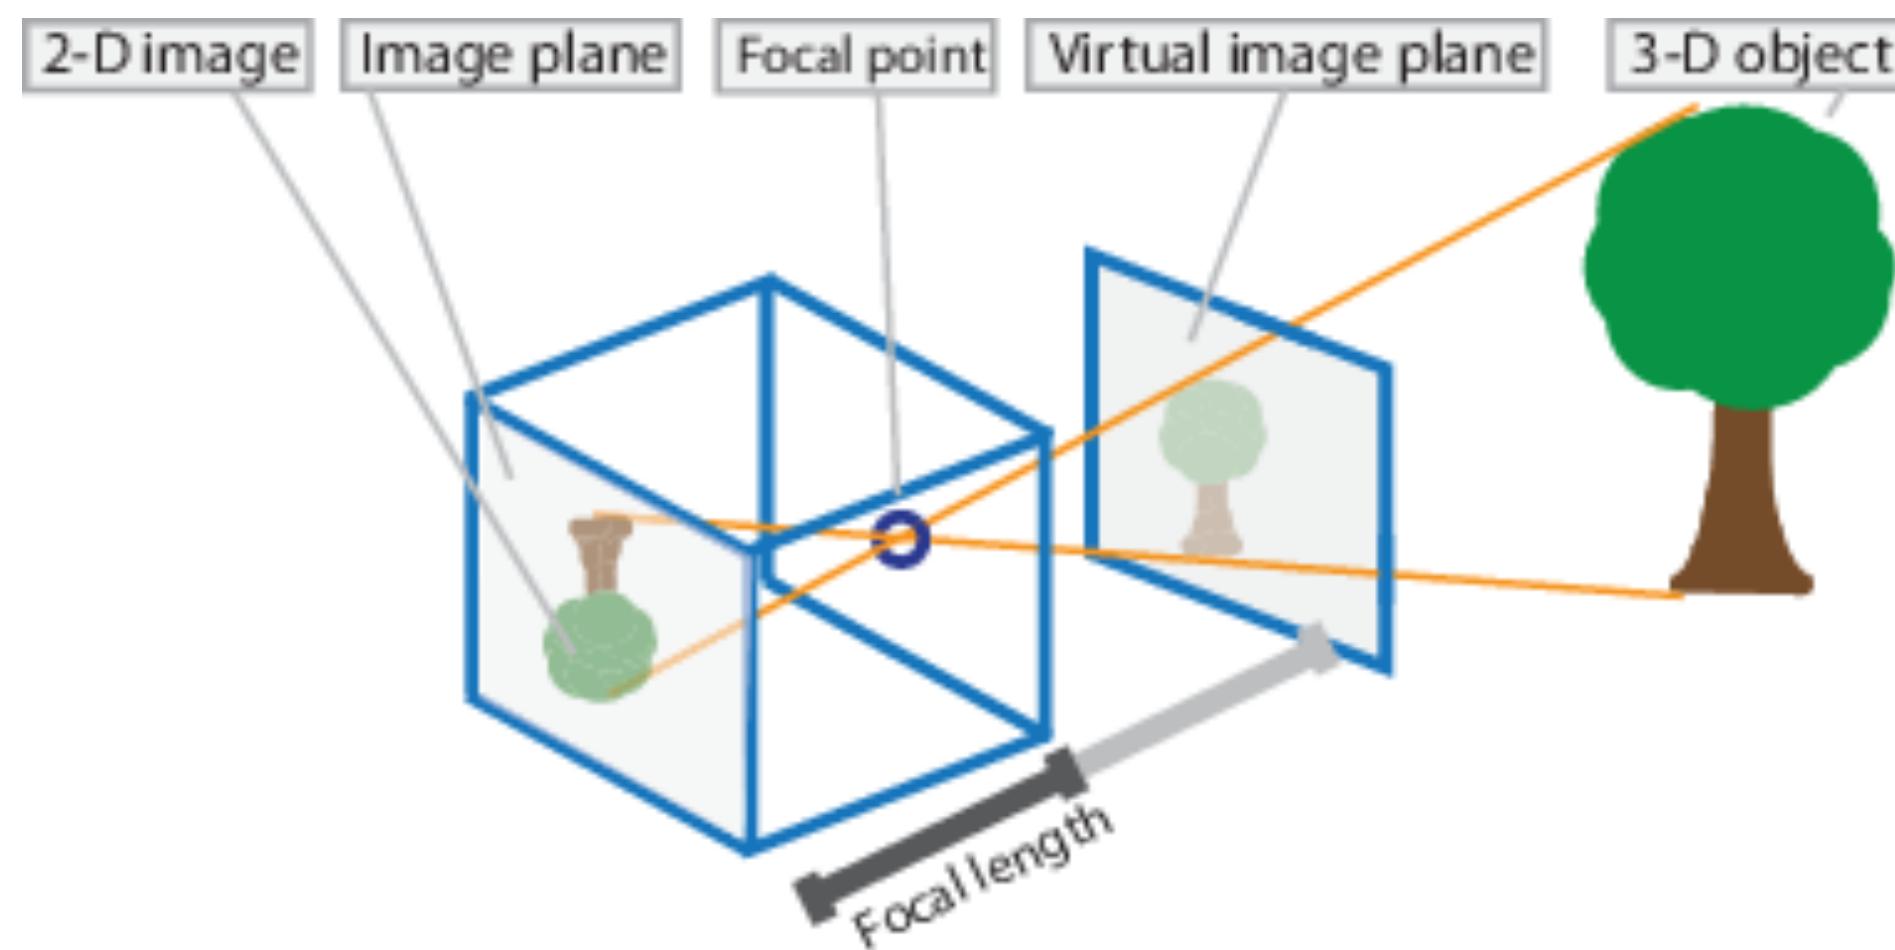
The Pinhole camera model (2)



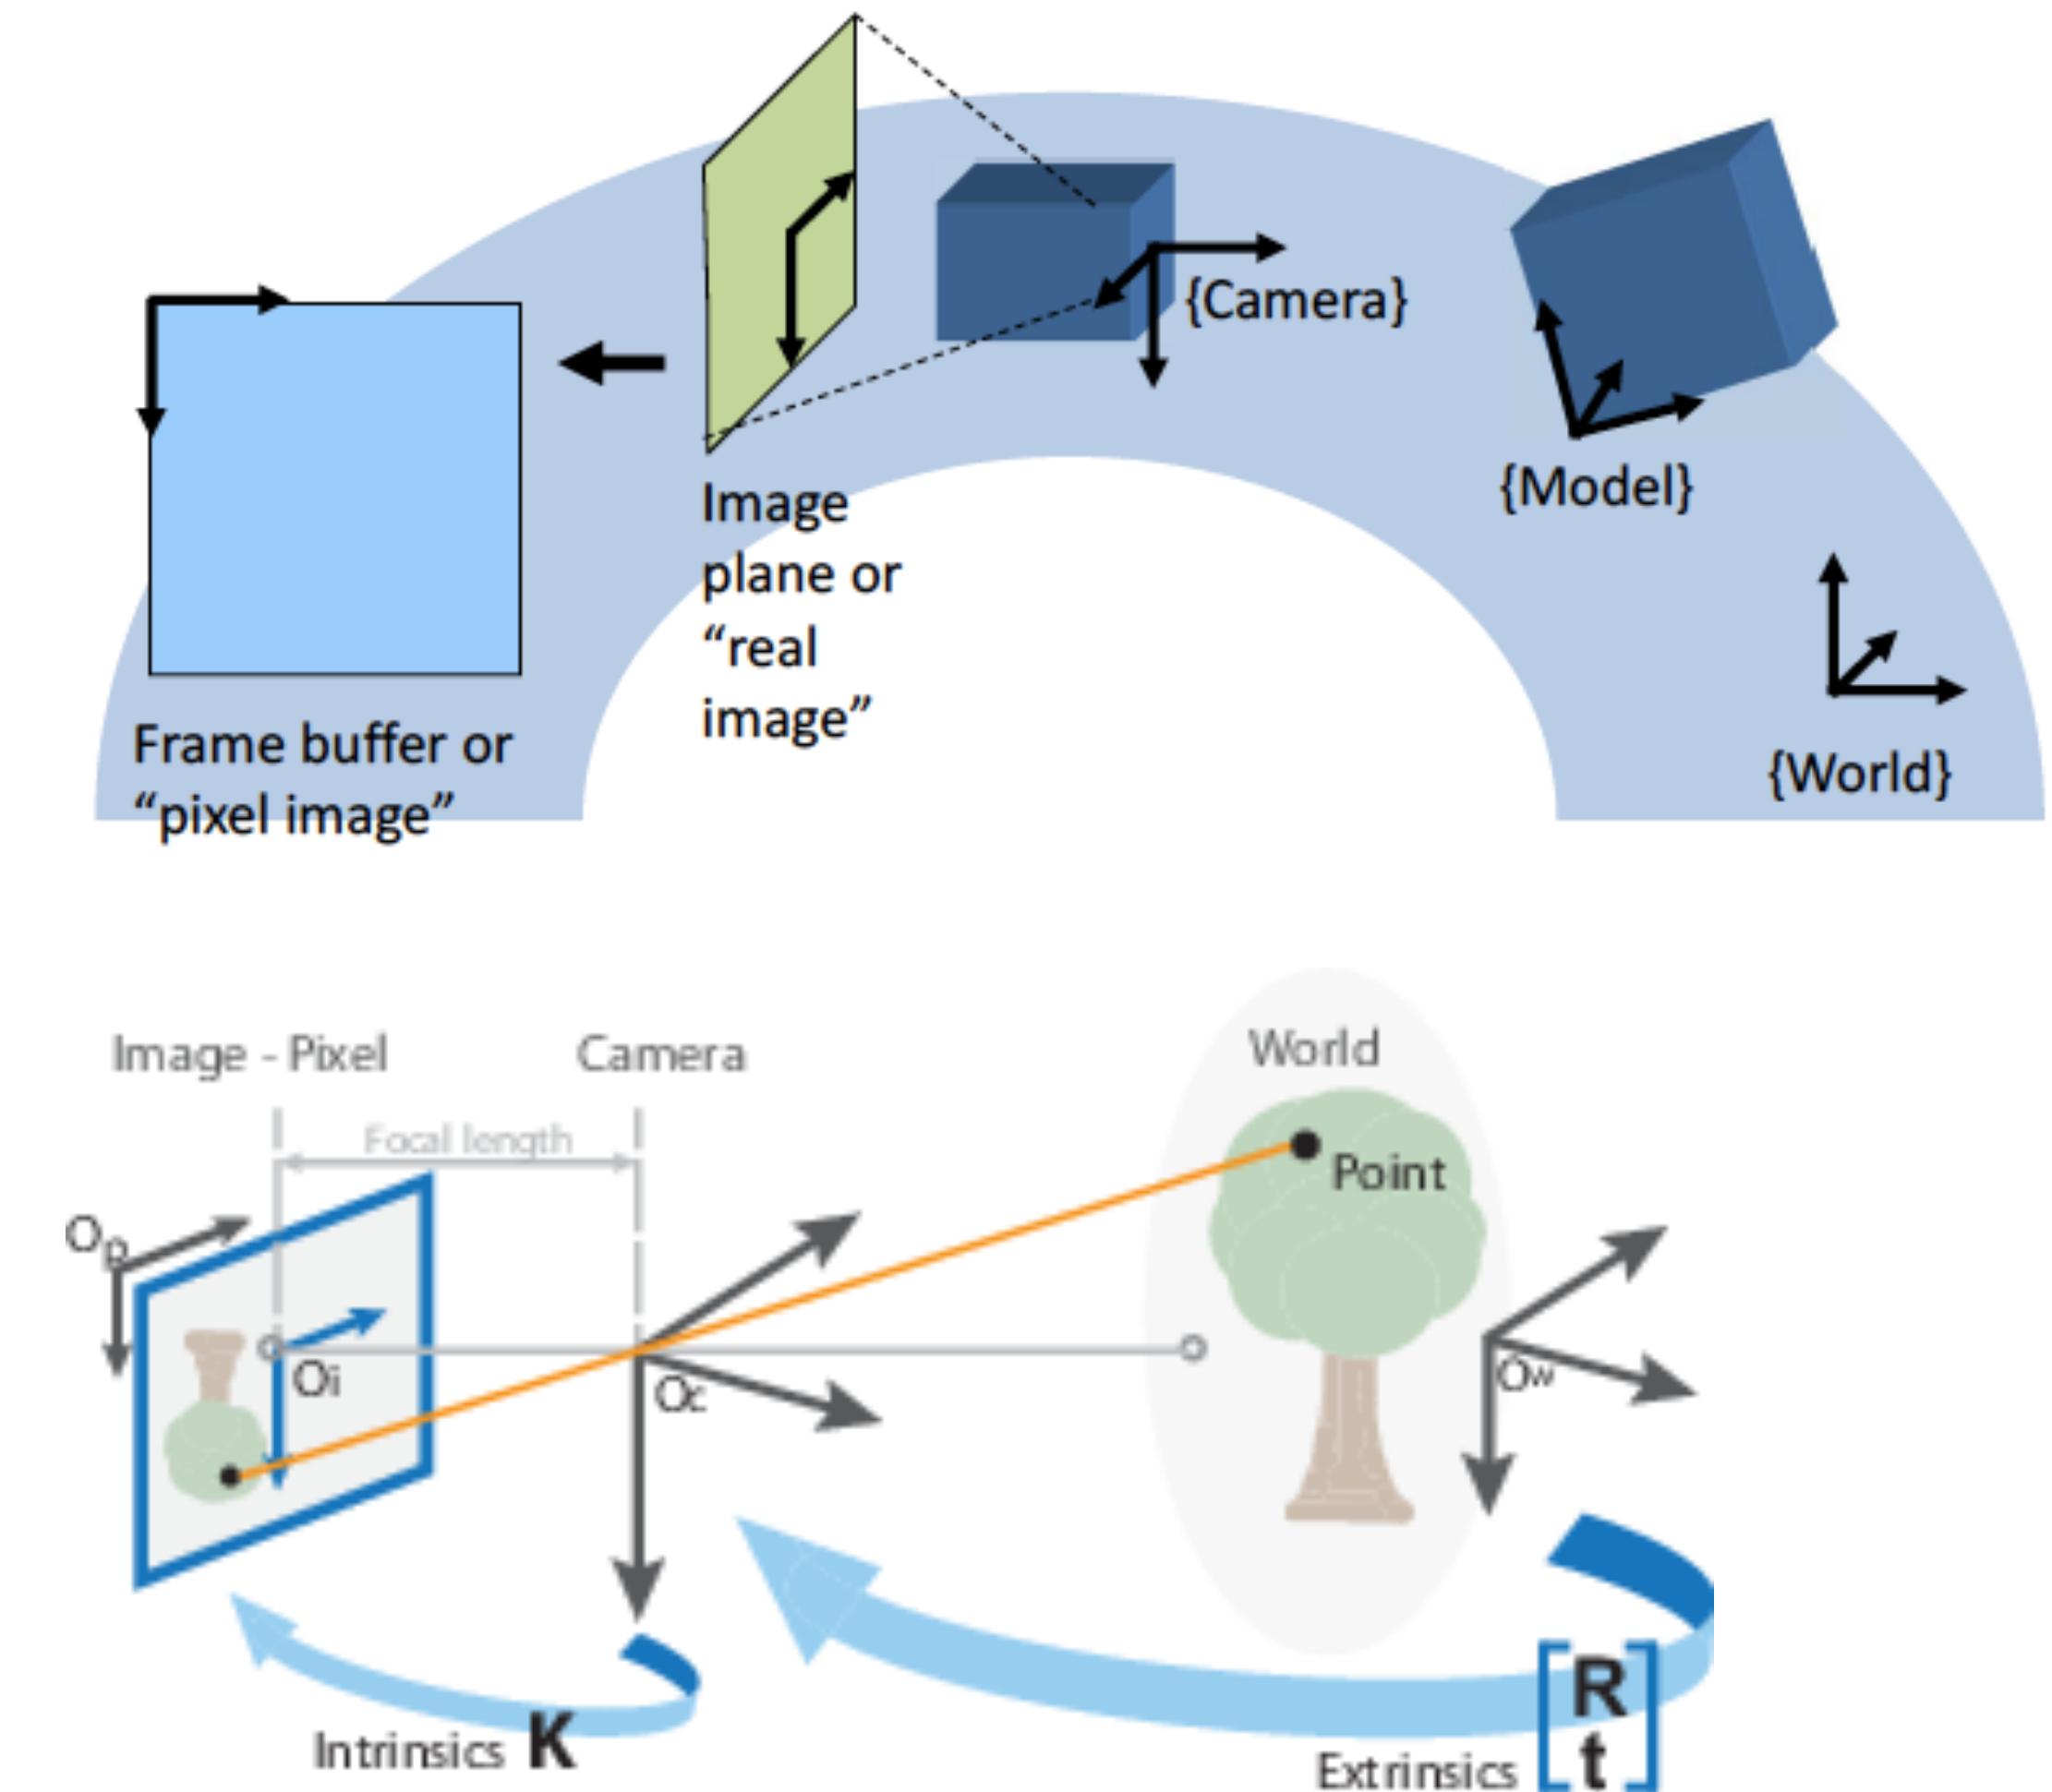
- **Without** a barrier in place, every point on the film will be influenced by light rays reflected from *every* point on the 3D object. Due to the **barrier**, only *one* (or a few) of these rays of light passes through the aperture, hits the film and generates an inverted image
- As the **aperture** size **decreases**, the image gets **sharper, but darker**.



The Pinhole camera model (3)



Simple camera without a lens and with a single small aperture. Light rays pass through the aperture and project an inverted image.



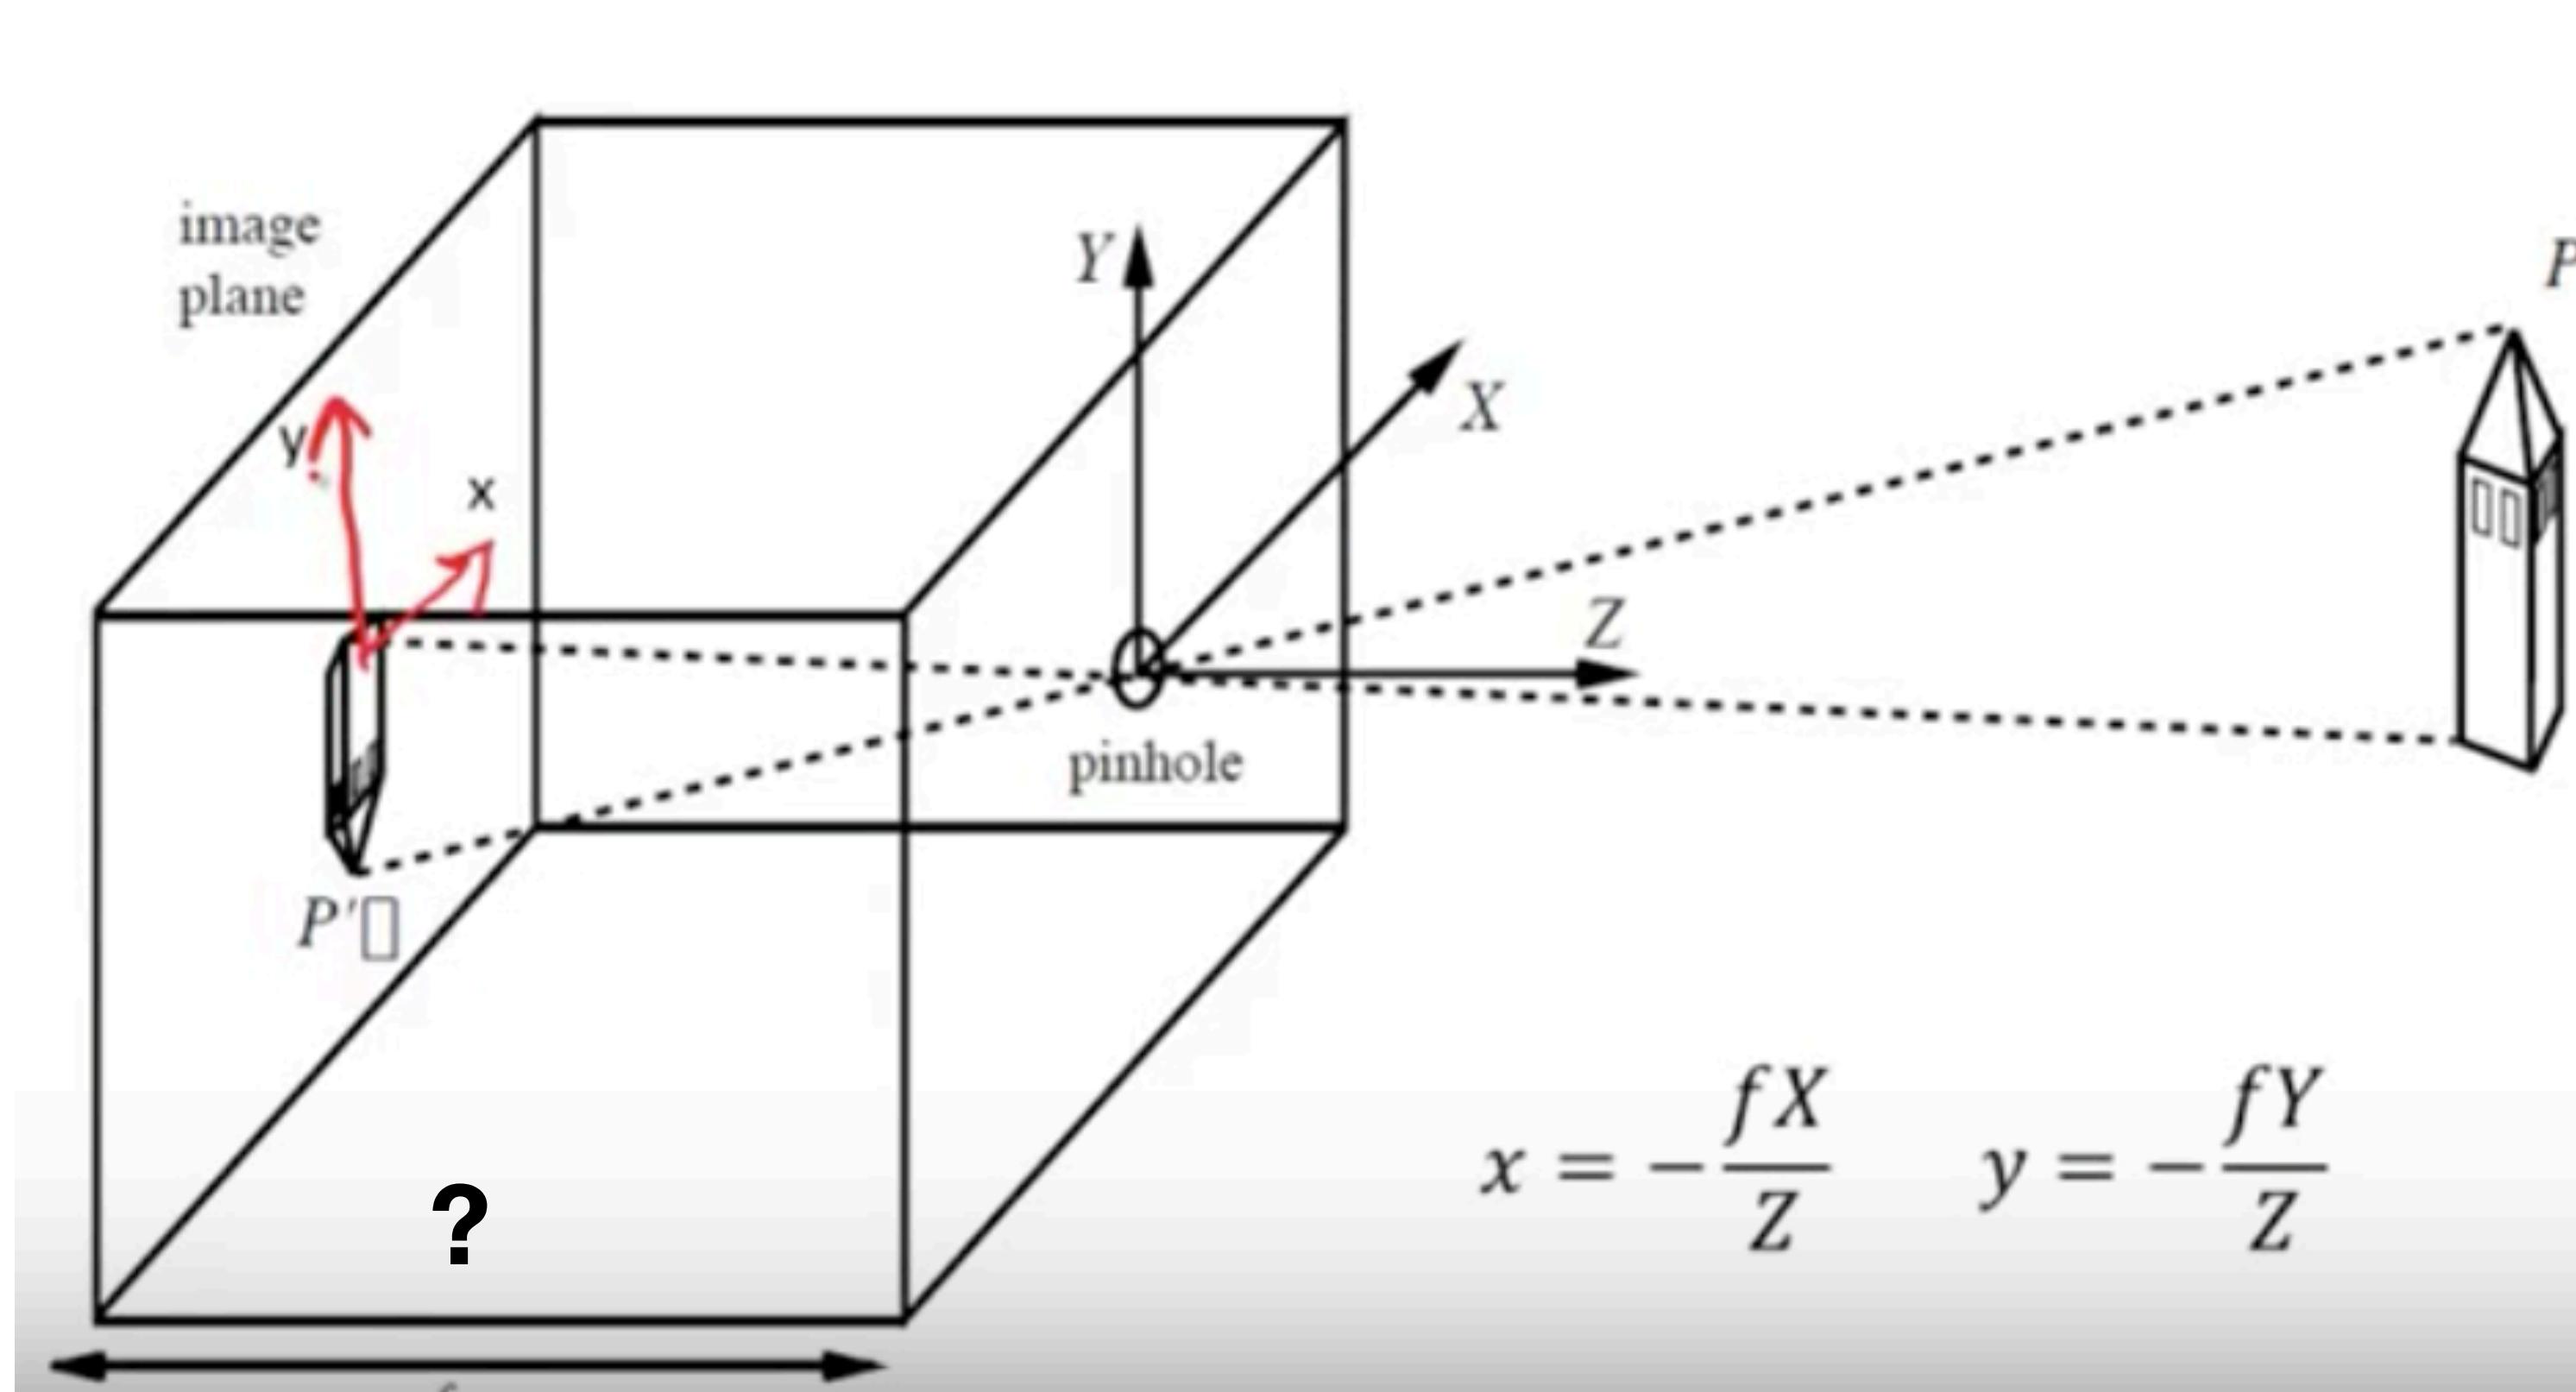
Extrinsic parameters:

world coordinates \rightarrow camera coordinates

Intrinsics parameters:

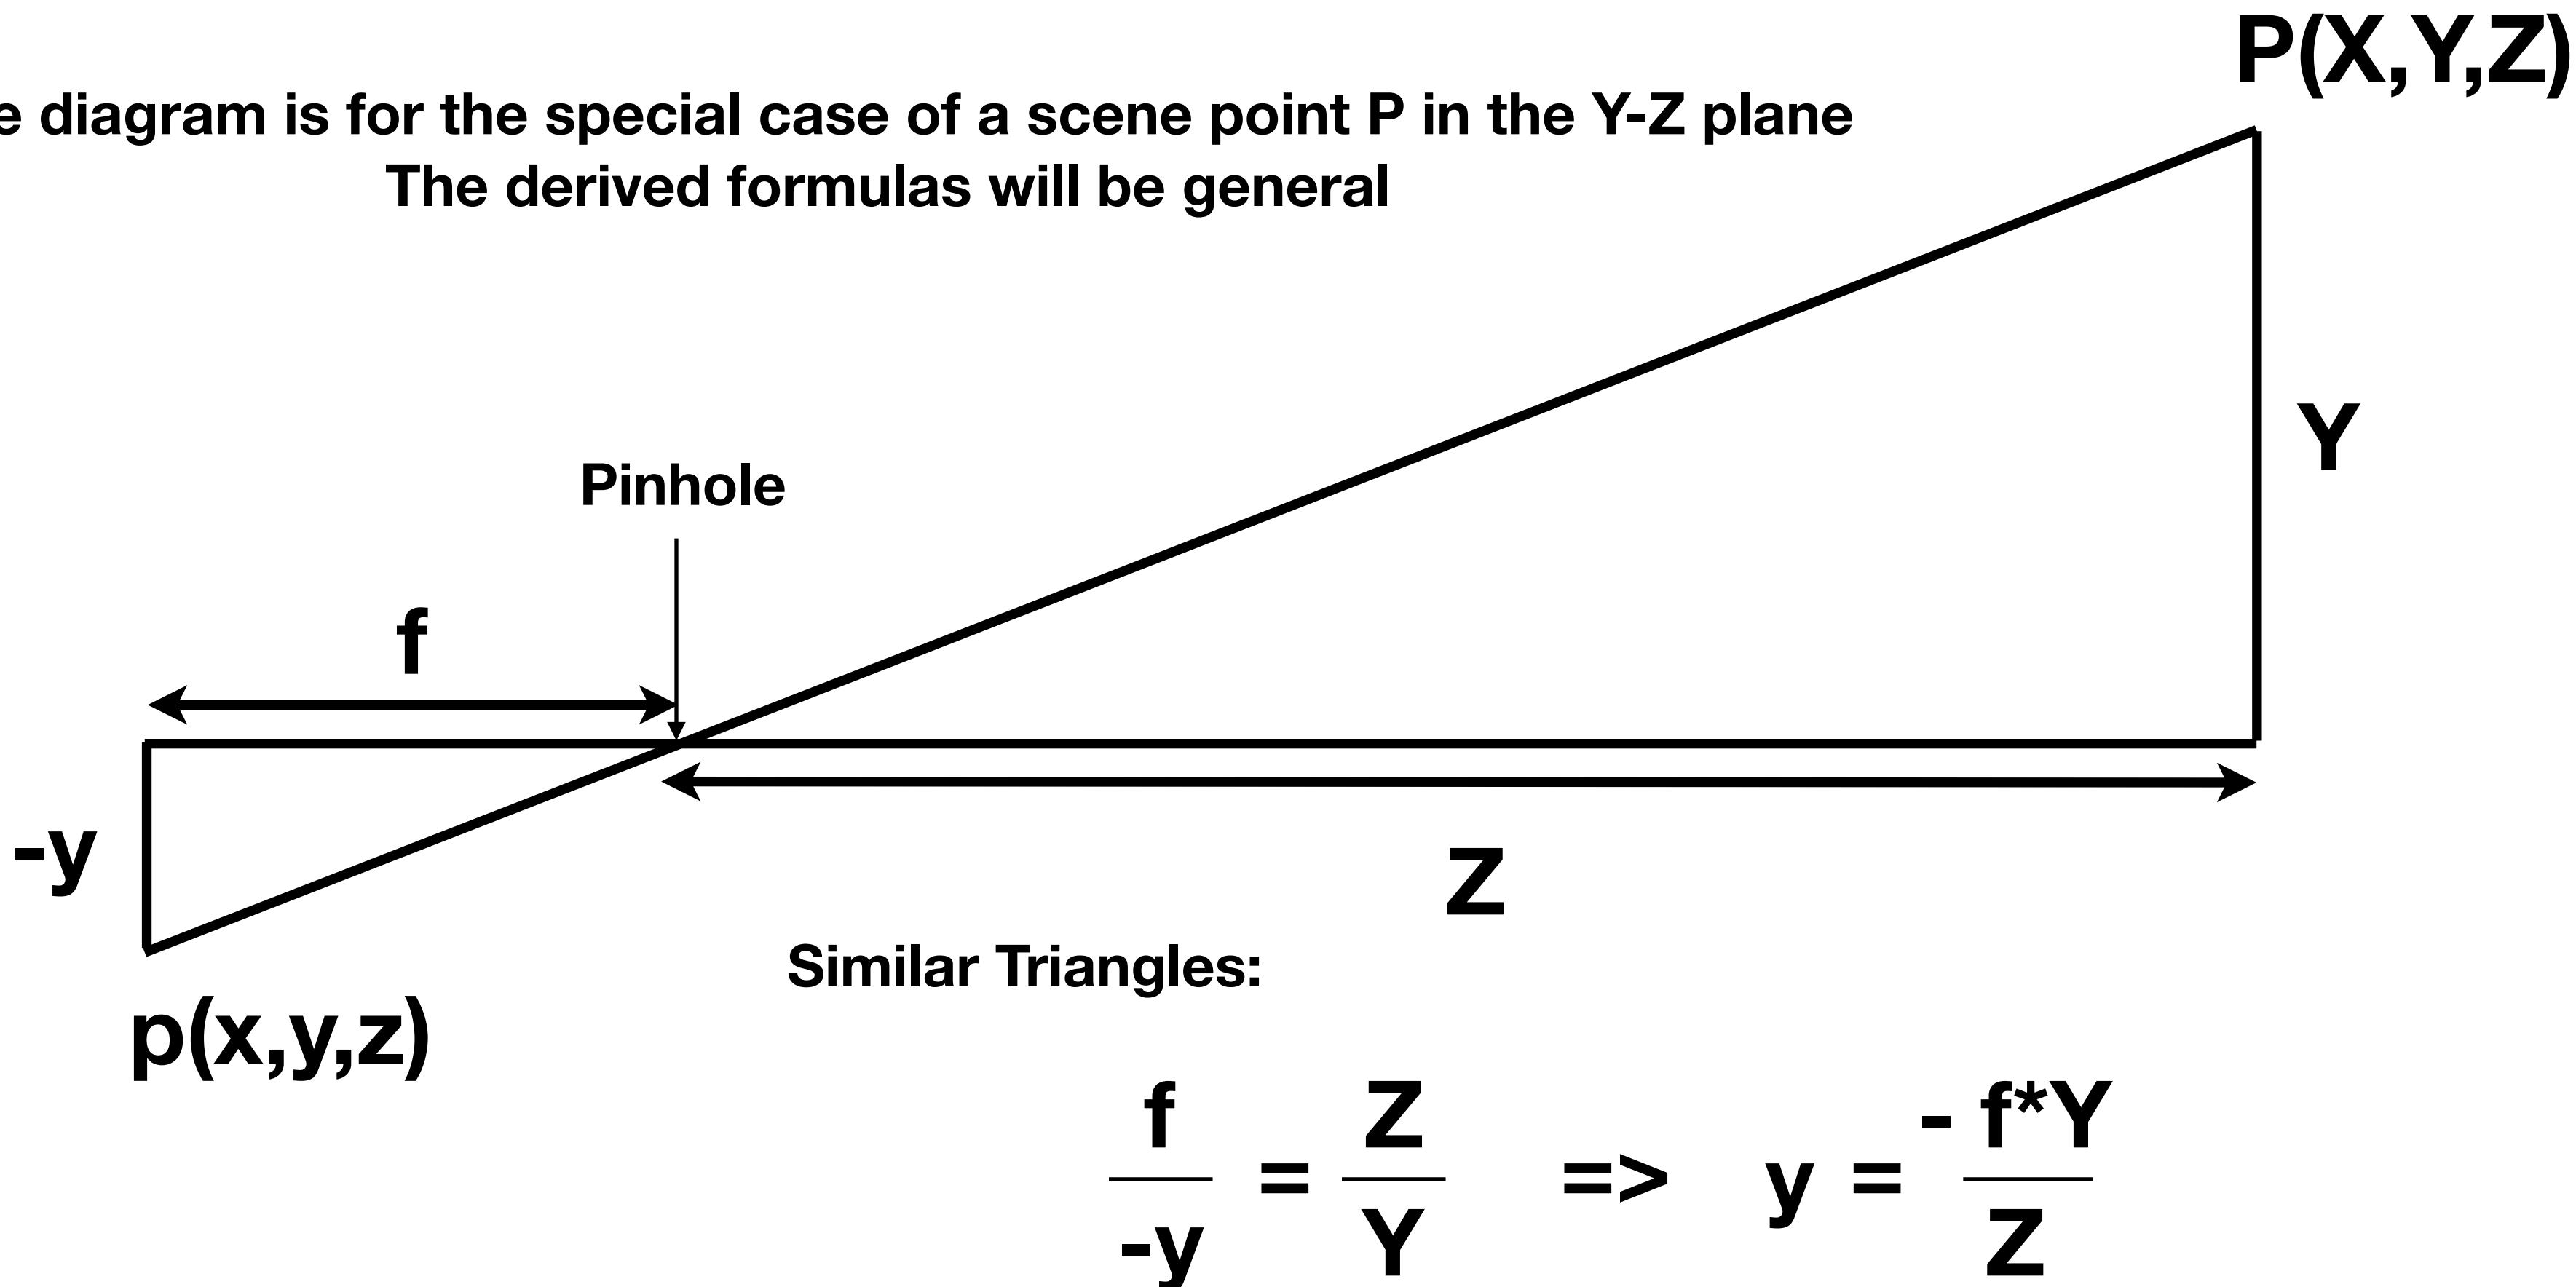
camera coordinates \rightarrow image plane coordinates

The Pinhole camera model (4)



The Pinhole camera model (5)

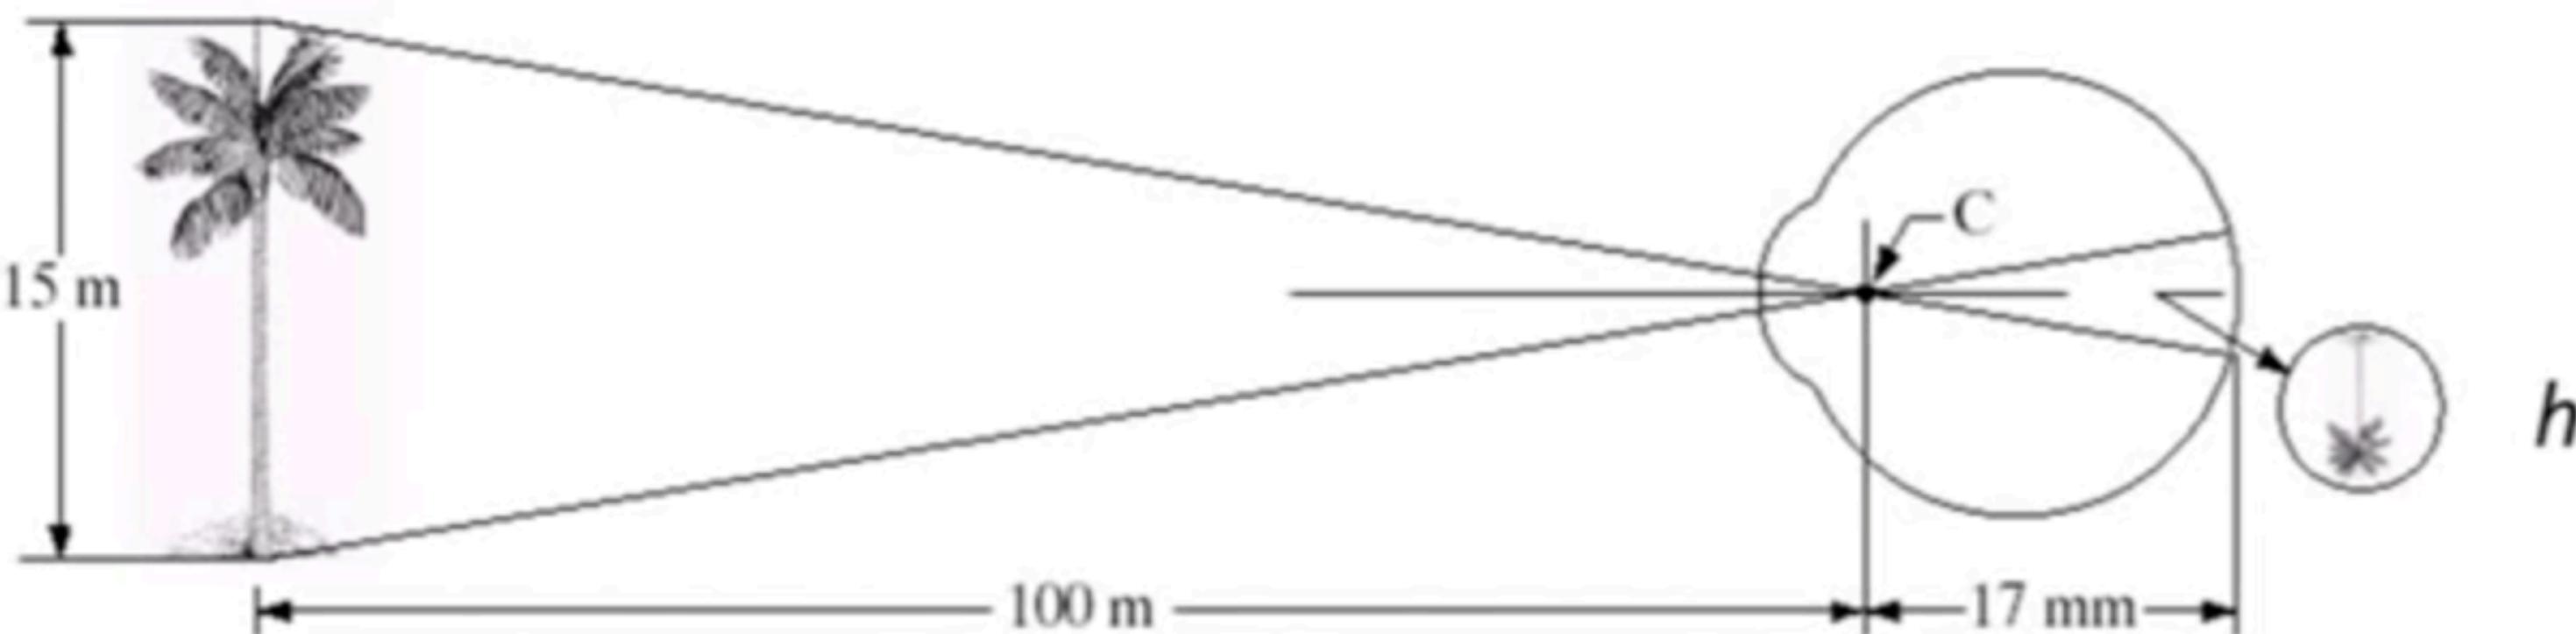
The diagram is for the special case of a scene point P in the Y-Z plane
The derived formulas will be general



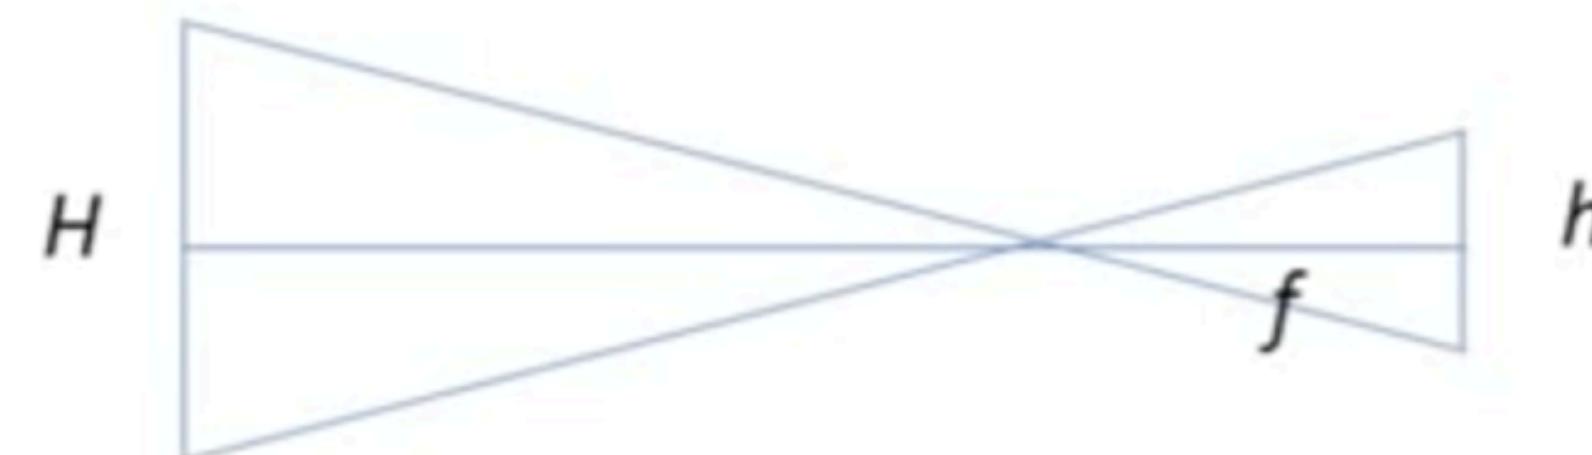
Example

FIGURE 2.3

Graphical representation of the eye looking at a palm tree. Point C is the optical center of the lens.



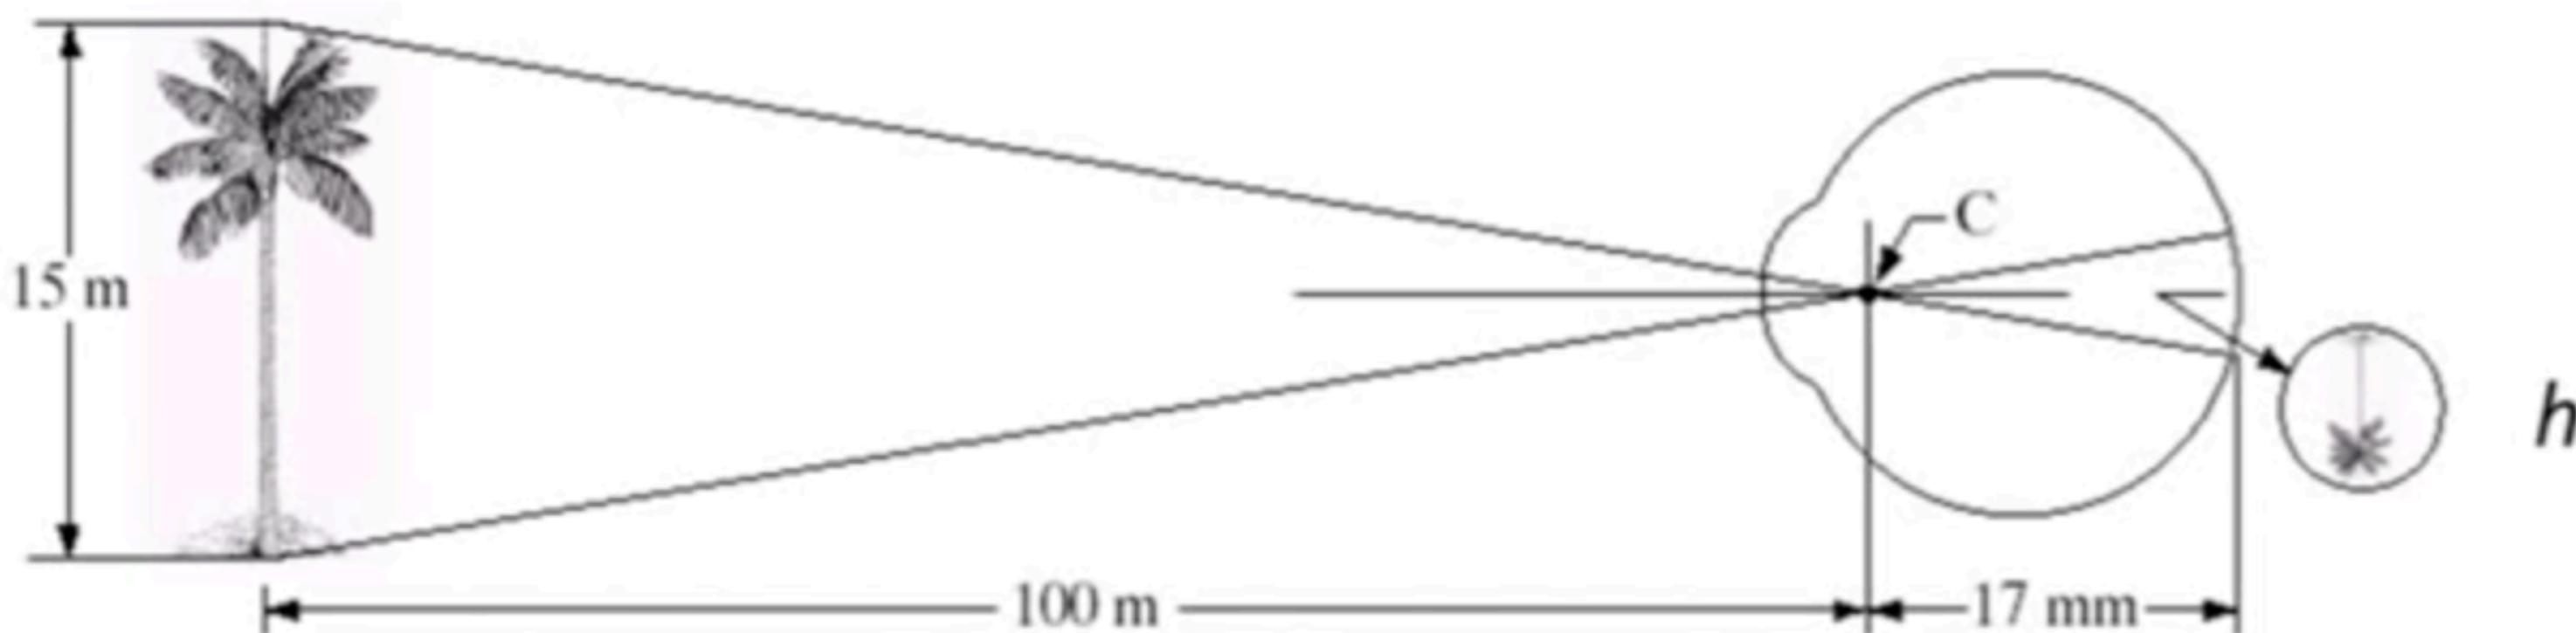
- Treat as pinhole camera – good approximation
 - Similar triangles
- What is the size of the tree on the retina?



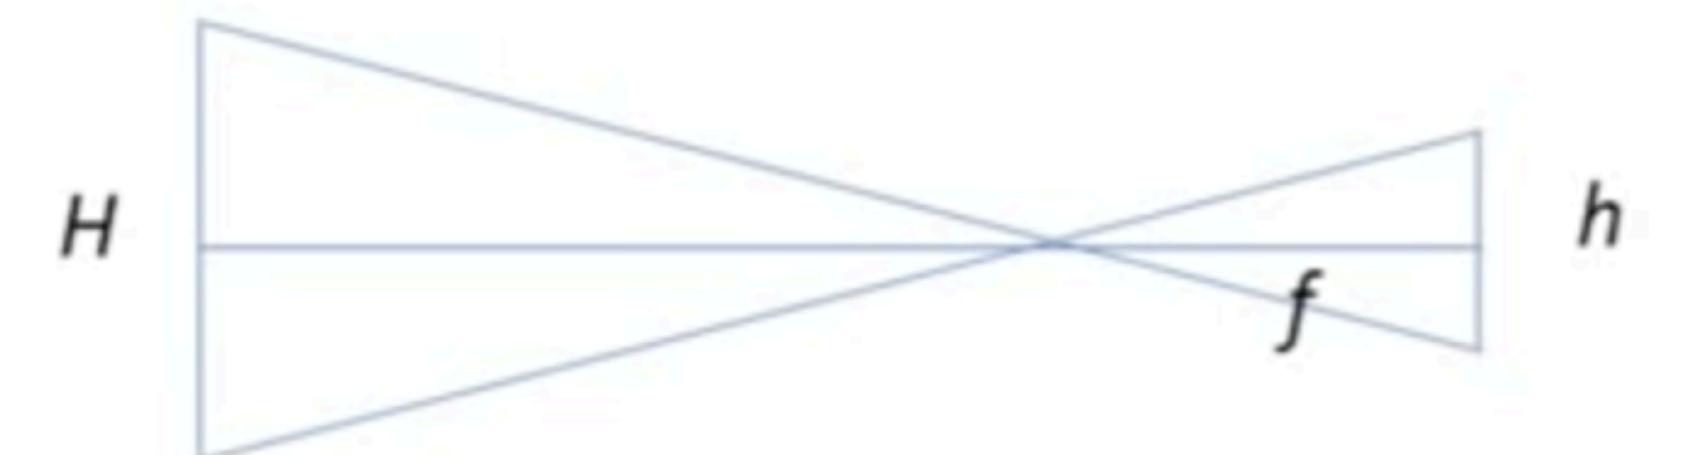
Example (2)

FIGURE 2.3

Graphical representation of the eye looking at a palm tree. Point C is the optical center of the lens.



- Treat as pinhole camera – good approximation
 - Similar triangles
- What is the size of the tree on the retina?
- What if the tree were twice as far away?



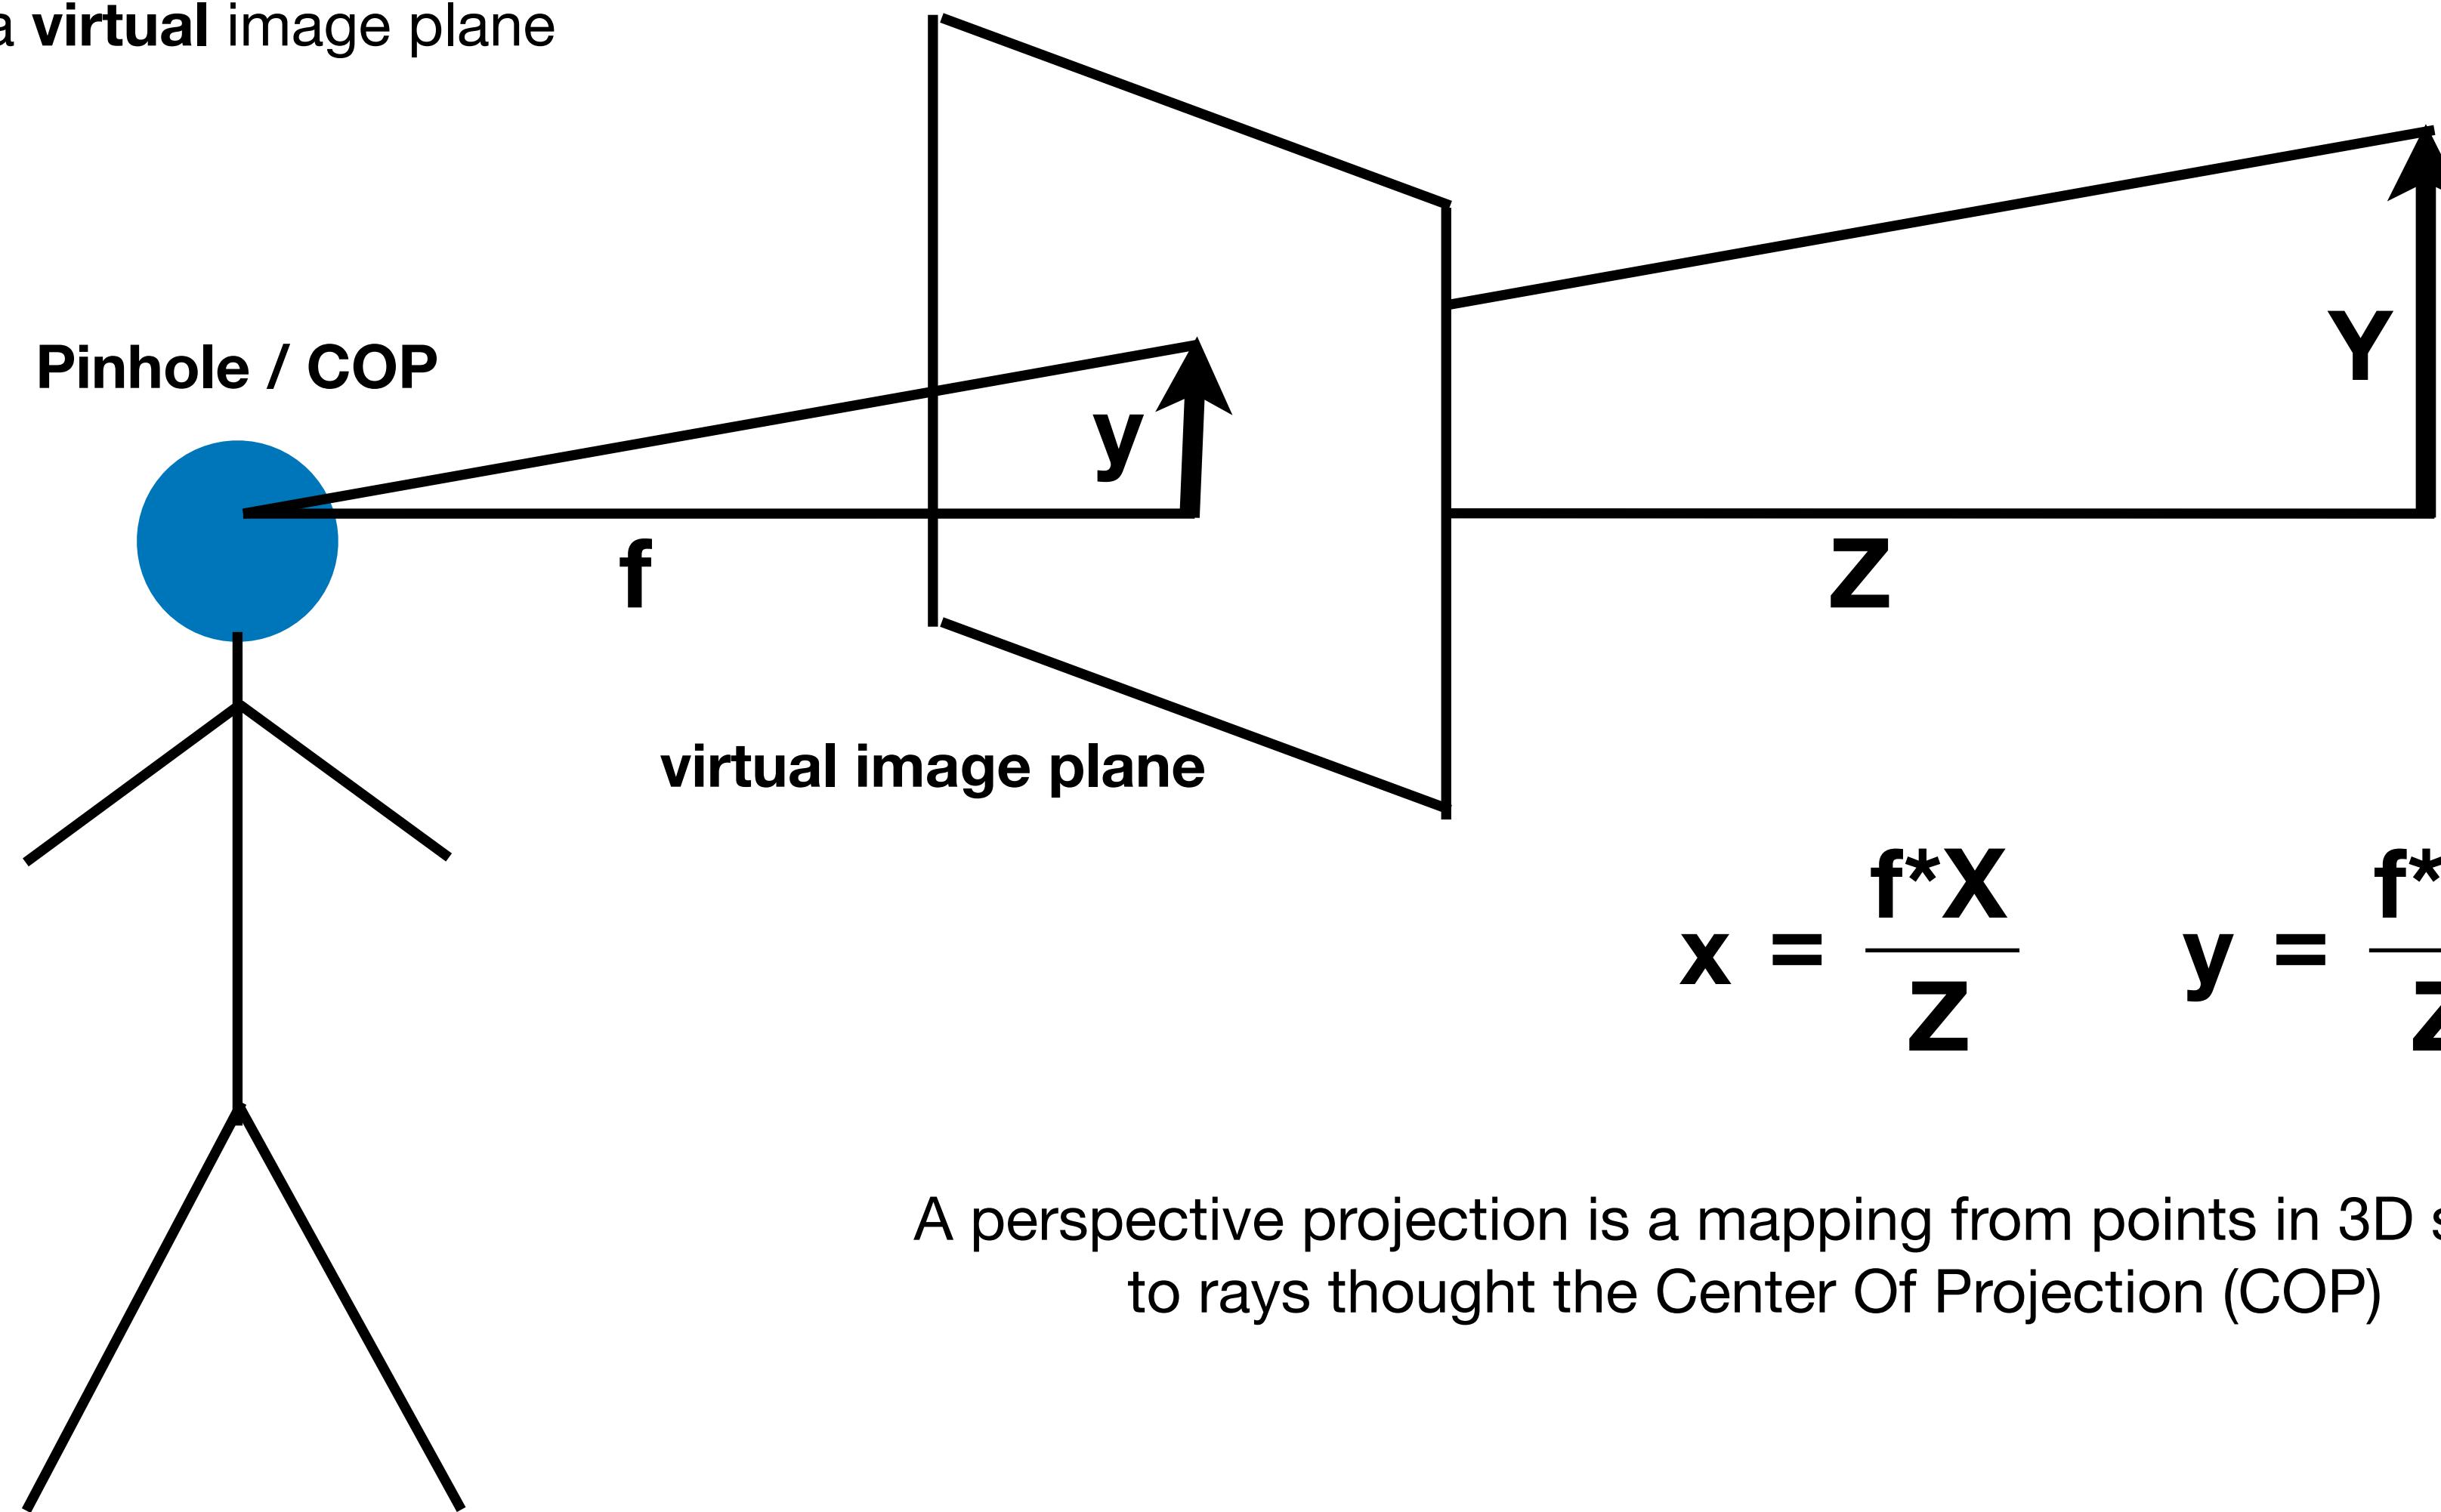
$$\frac{H}{h} = \frac{Z}{f}$$
$$15/100 = h/17 \text{ or}$$
$$h = 2.55 \text{ mm}$$

$$Y/Z = y/f$$

Pinhole camera: trick (no inversion)

A projection model that avoids **inversion**

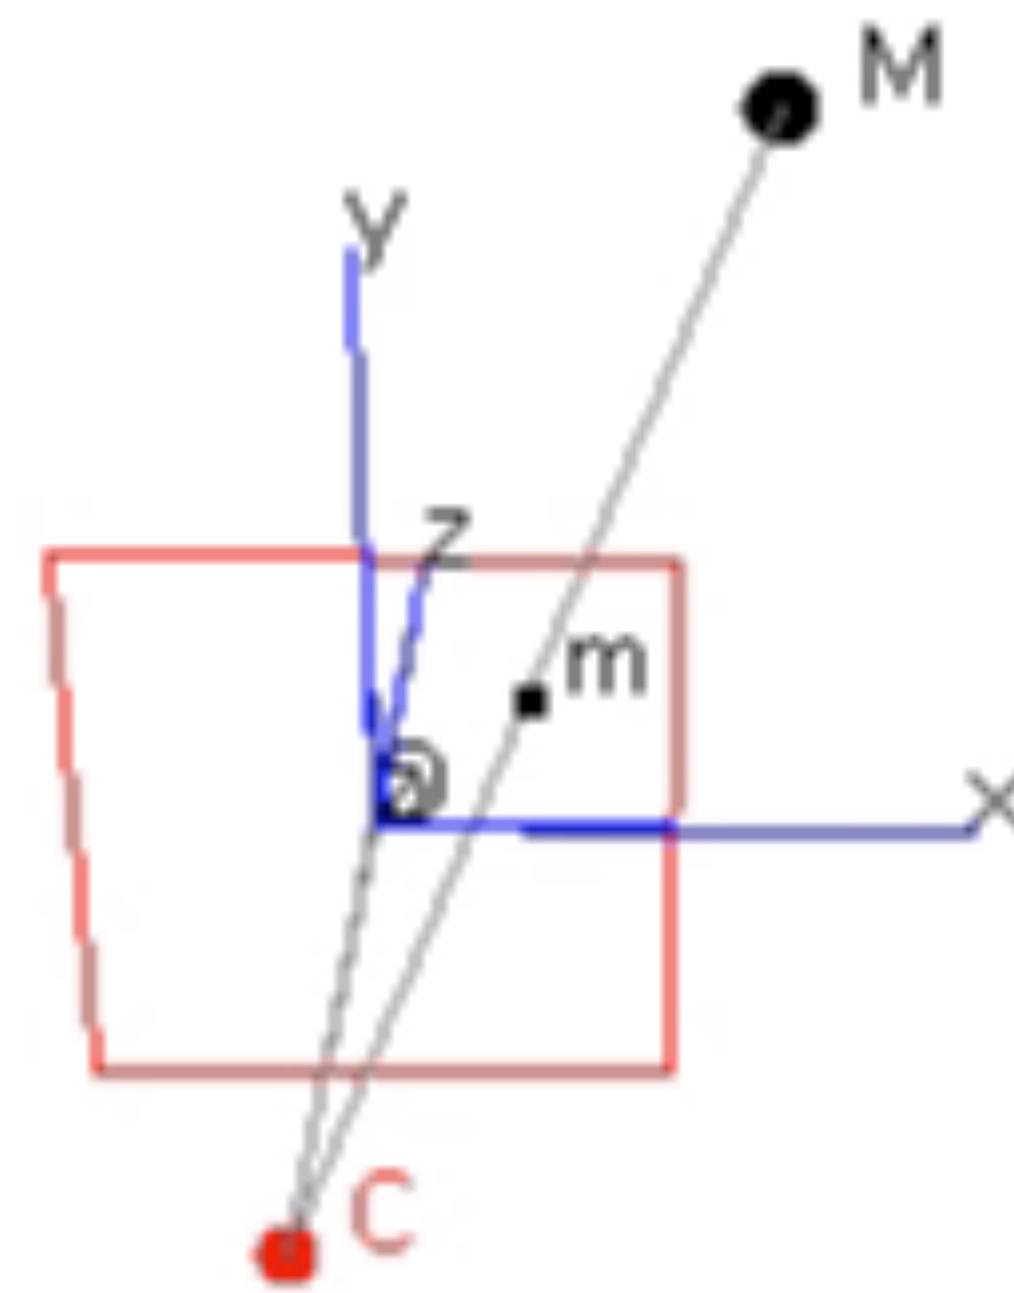
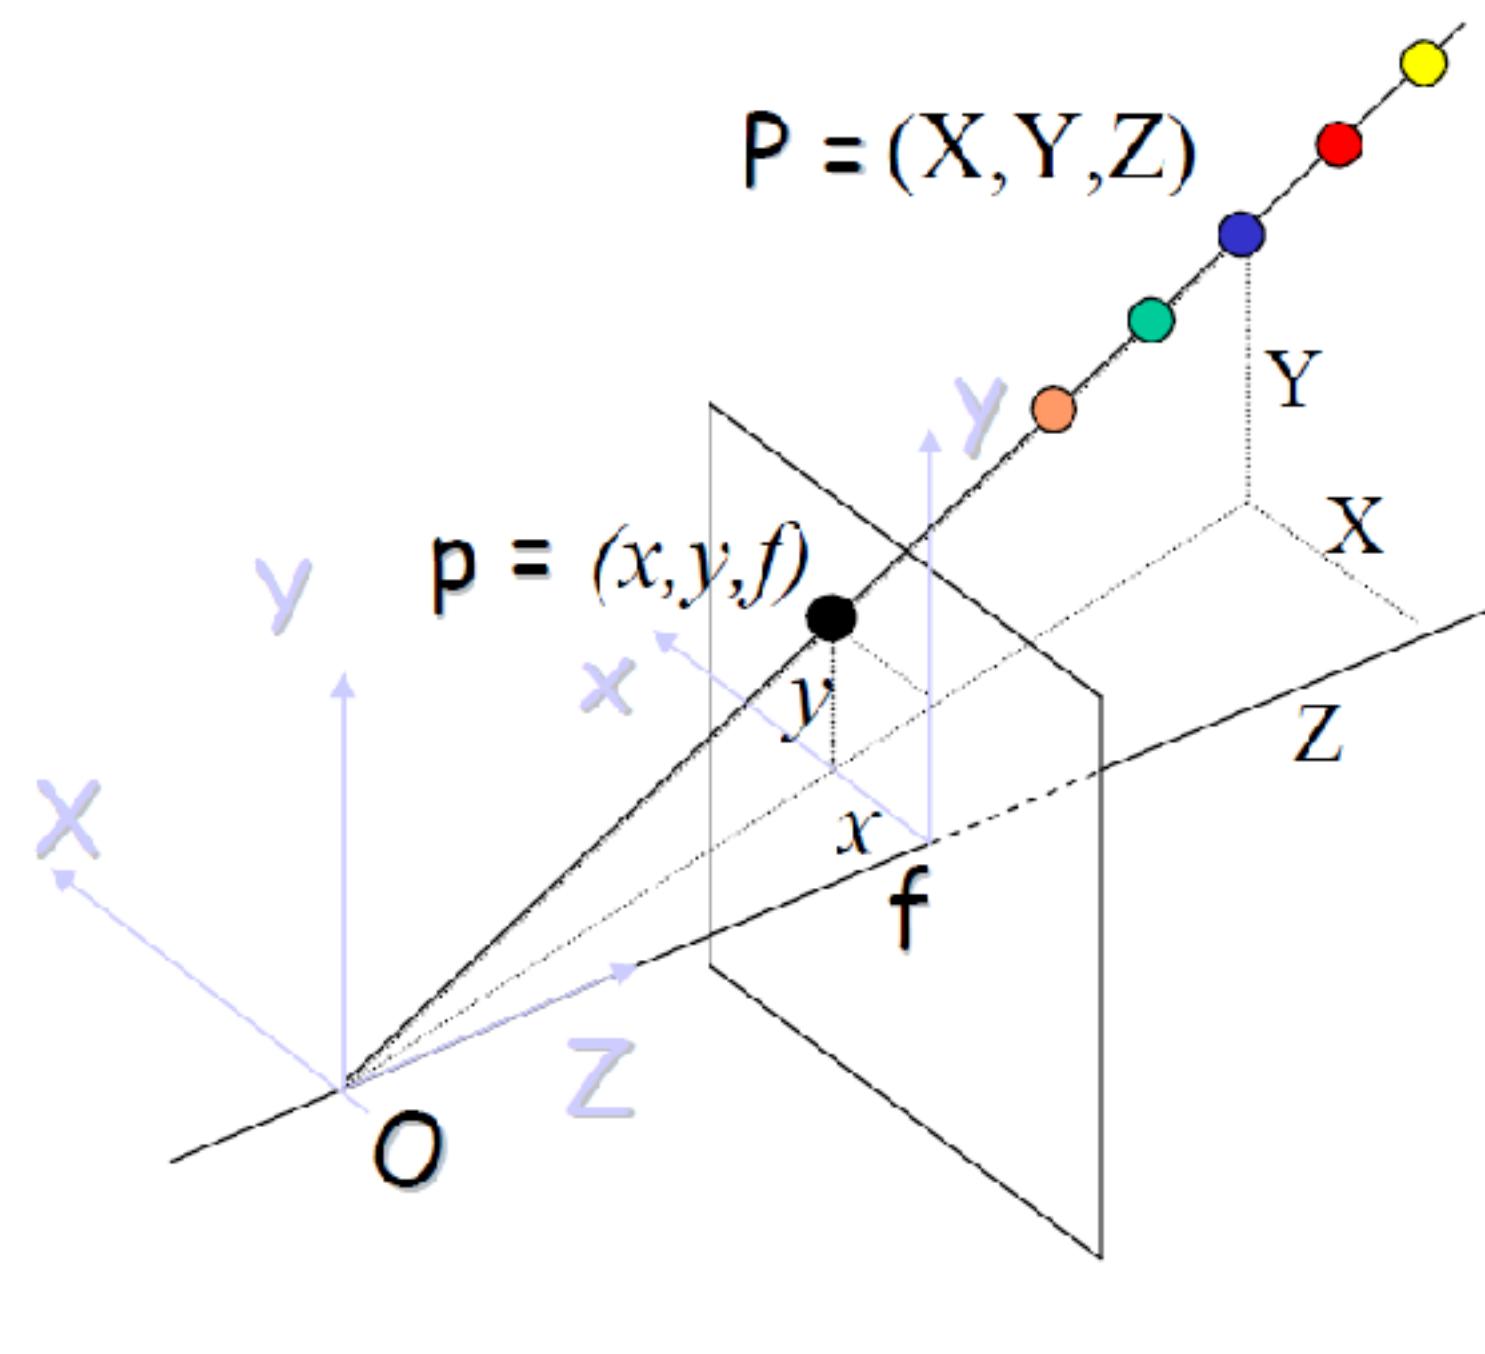
Trick: use a **virtual** image plane



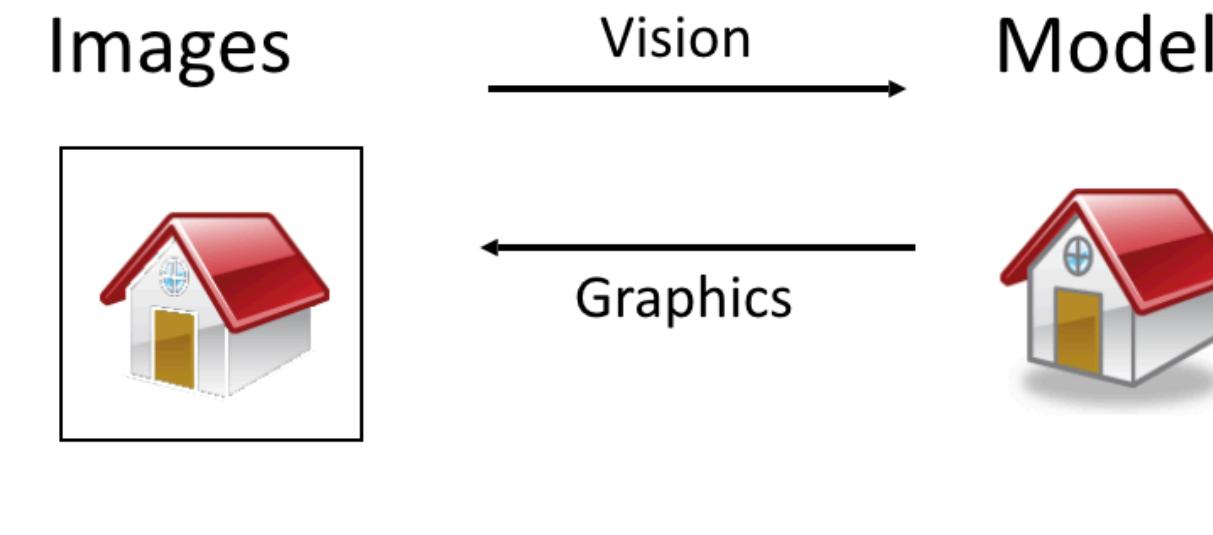
$$x = \frac{f^*X}{z} \quad y = \frac{f^*Y}{z}$$

A perspective projection is a mapping from points in 3D space
to rays thought the Center Of Projection (COP)

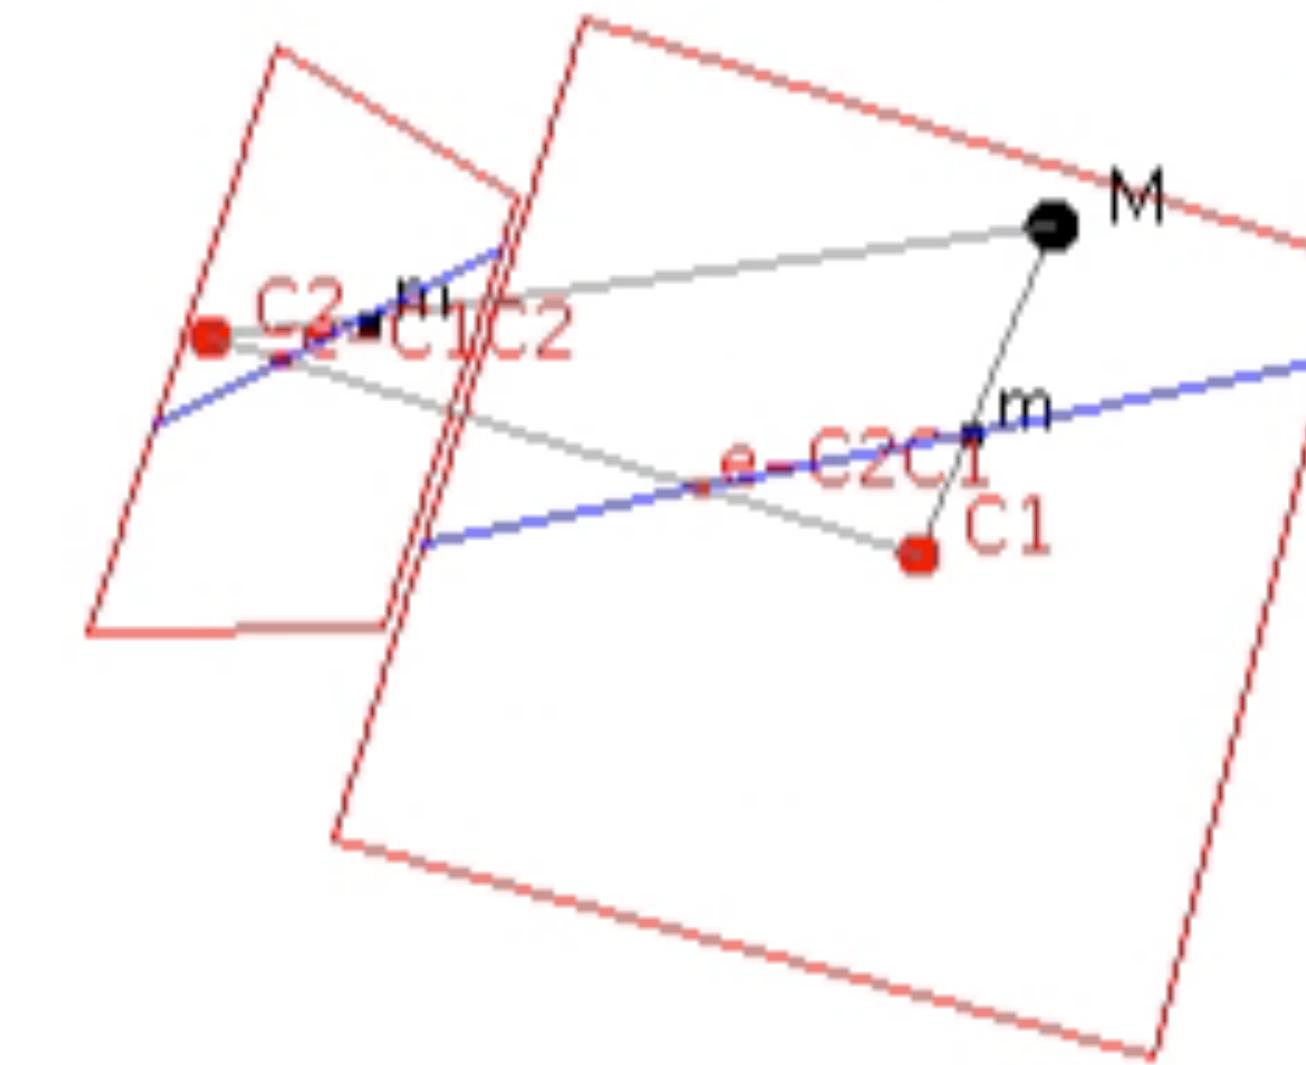
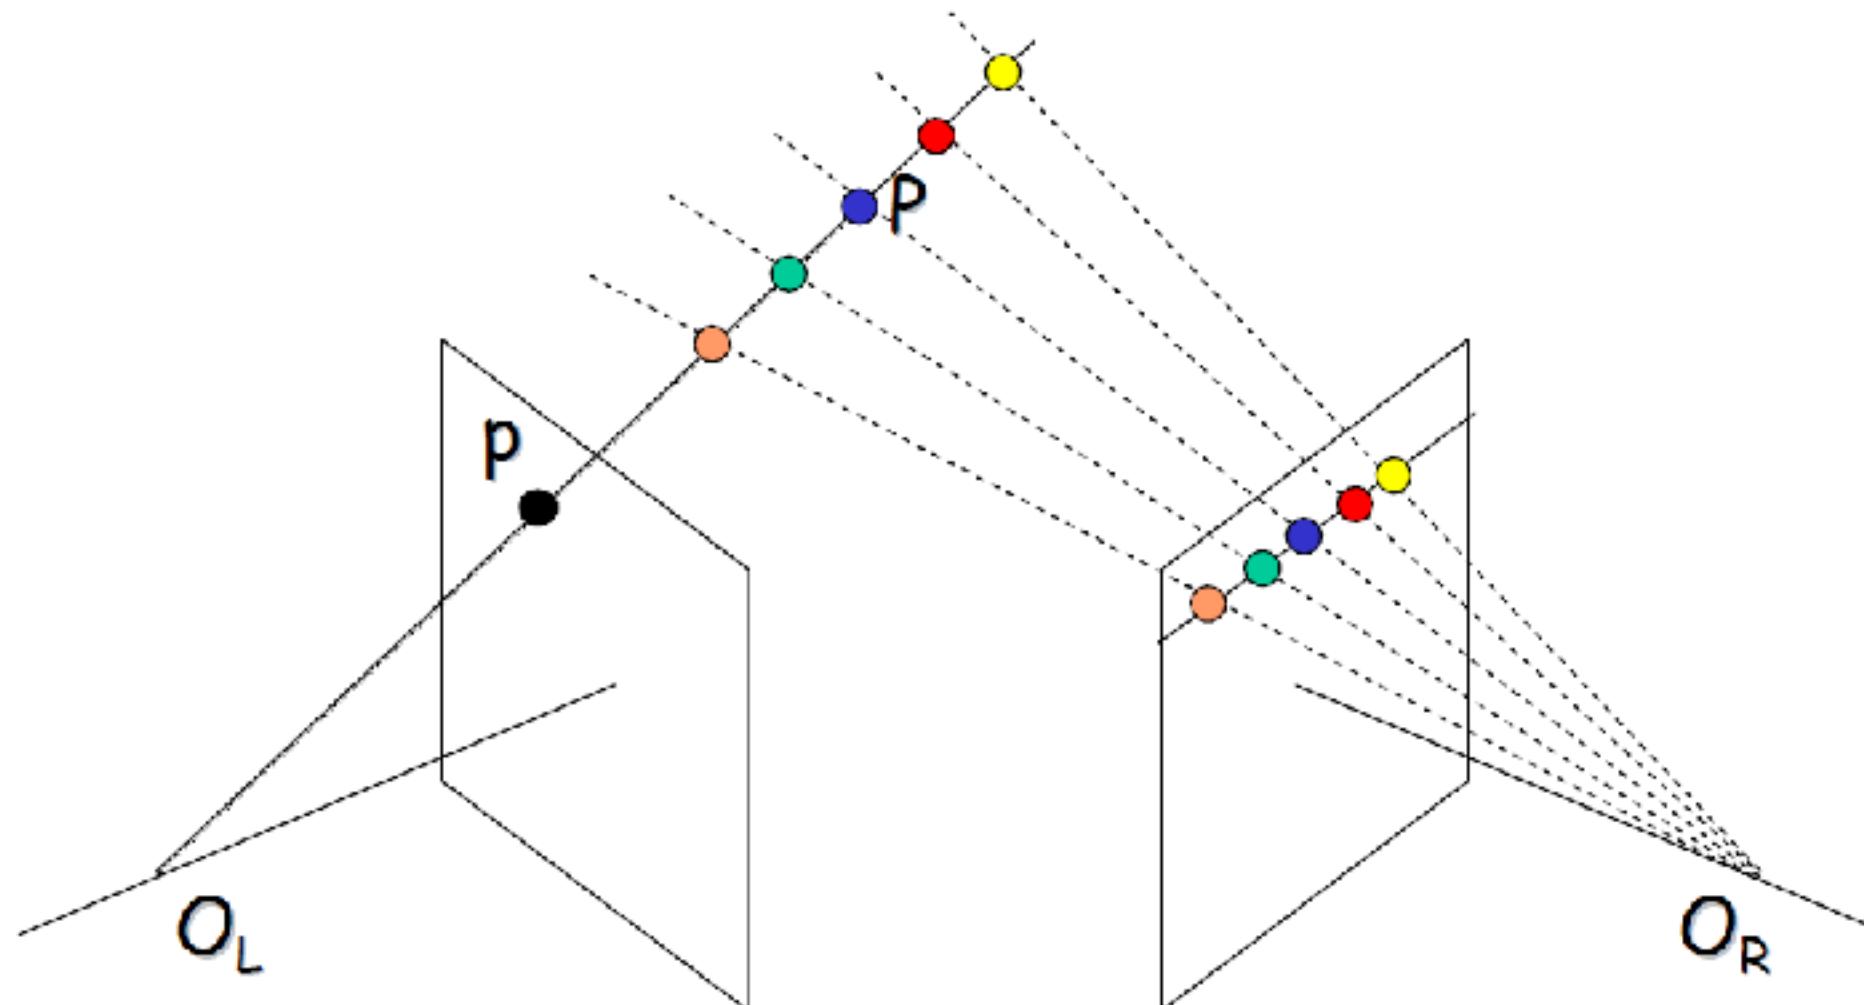
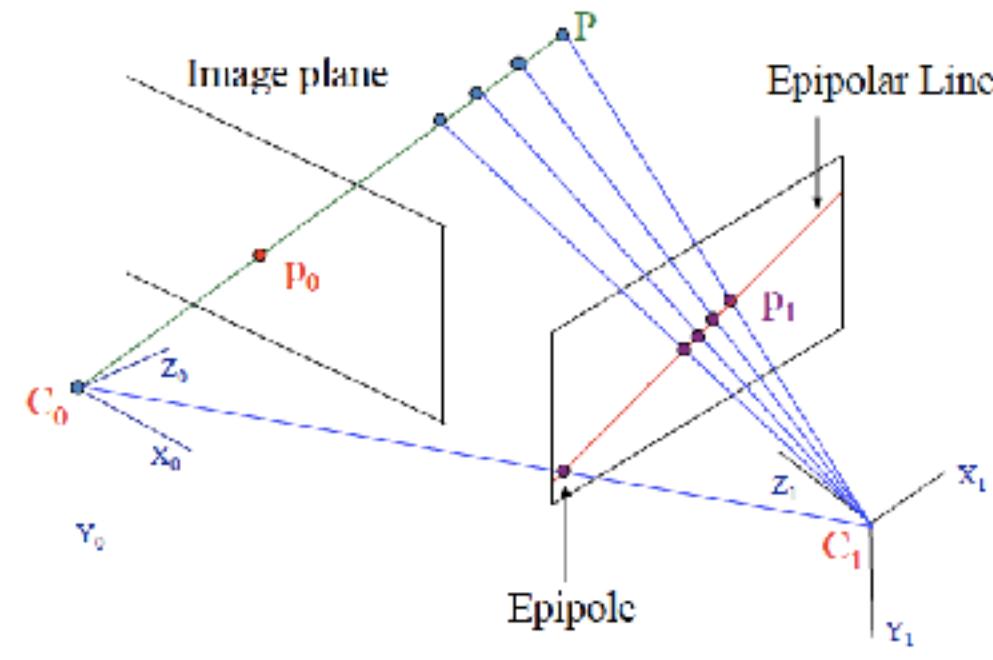
Single Camera / View



Fundamental ambiguity: Any scene point on the ray OP (CM) has image point p (m)

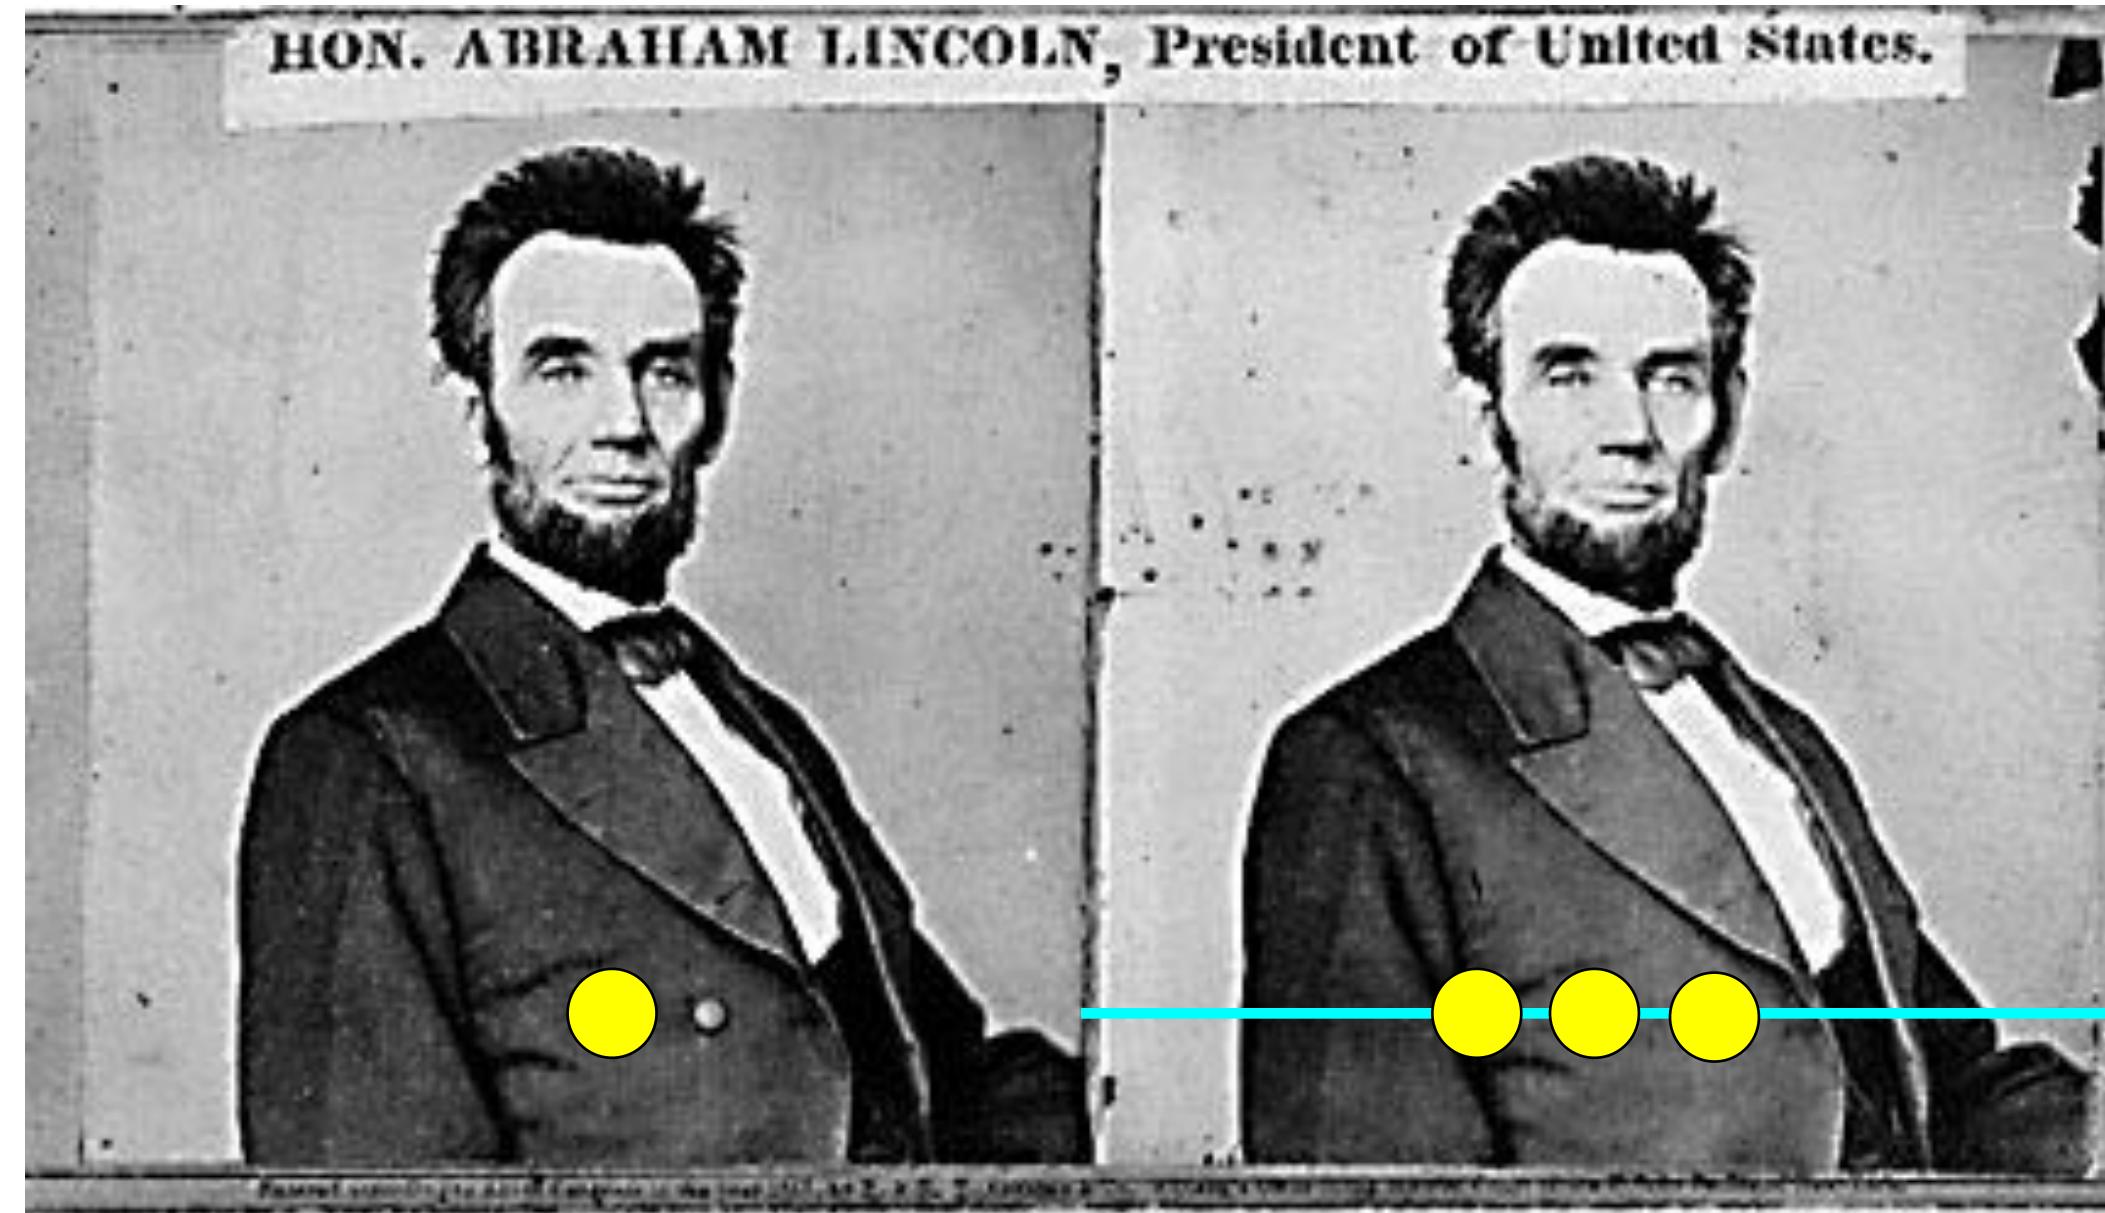


Stereo (multi-view)



A second camera / view can **resolve the ambiguity**,
enabling measurements of **depth** via triangulation

Stereo, correspondence, disparity and depth



- **For each pixel x in the first image**
 - Find corresponding epipolar **scanline** in the right image
 - Examine all pixels on the scanline and pick the best **match x'**
 - Compute **disparity $x-x'$** and set **depth(x) = $b*f/(x-x')$**

Digital Images

- Components
 - Illumination $i(x,y)$
 - Reflection $r(x,y)$

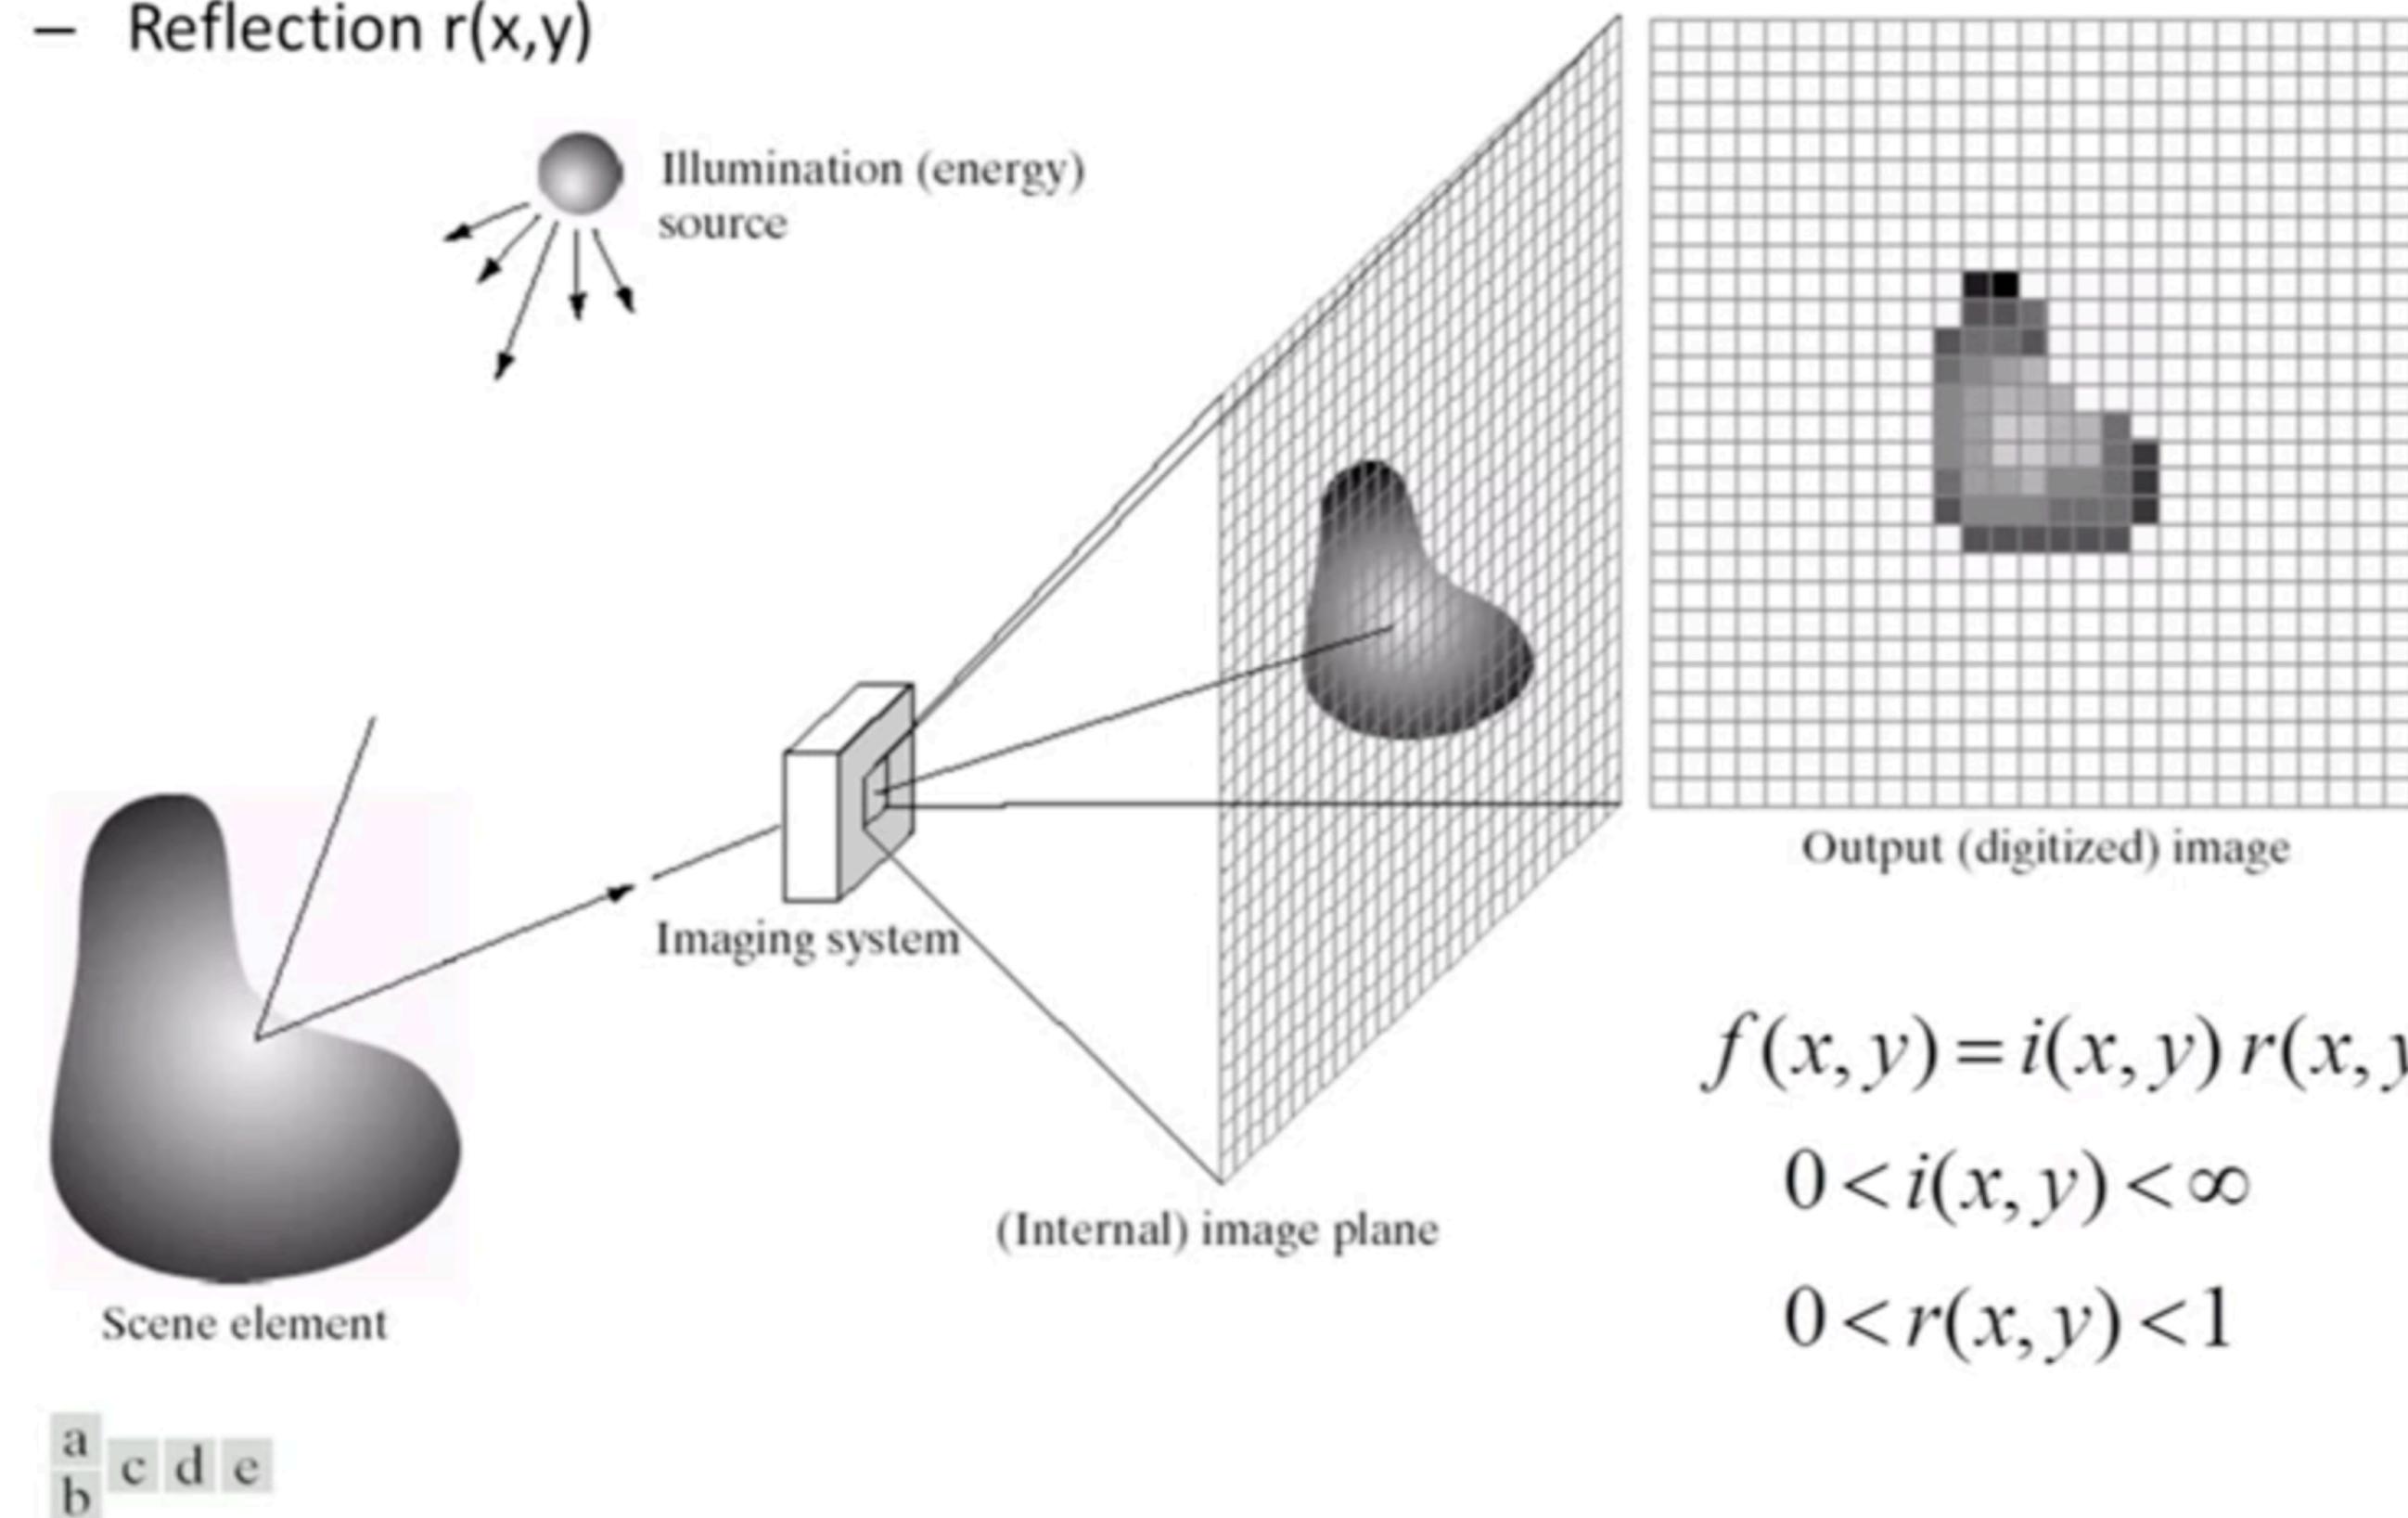


FIGURE 2.15 An example of the digital image acquisition process. (a) Energy (“illumination”) source. (b) An element of a scene. (c) Imaging system. (d) Projection of the scene onto the image plane. (e) Digitized image.

Sampling & Quantization

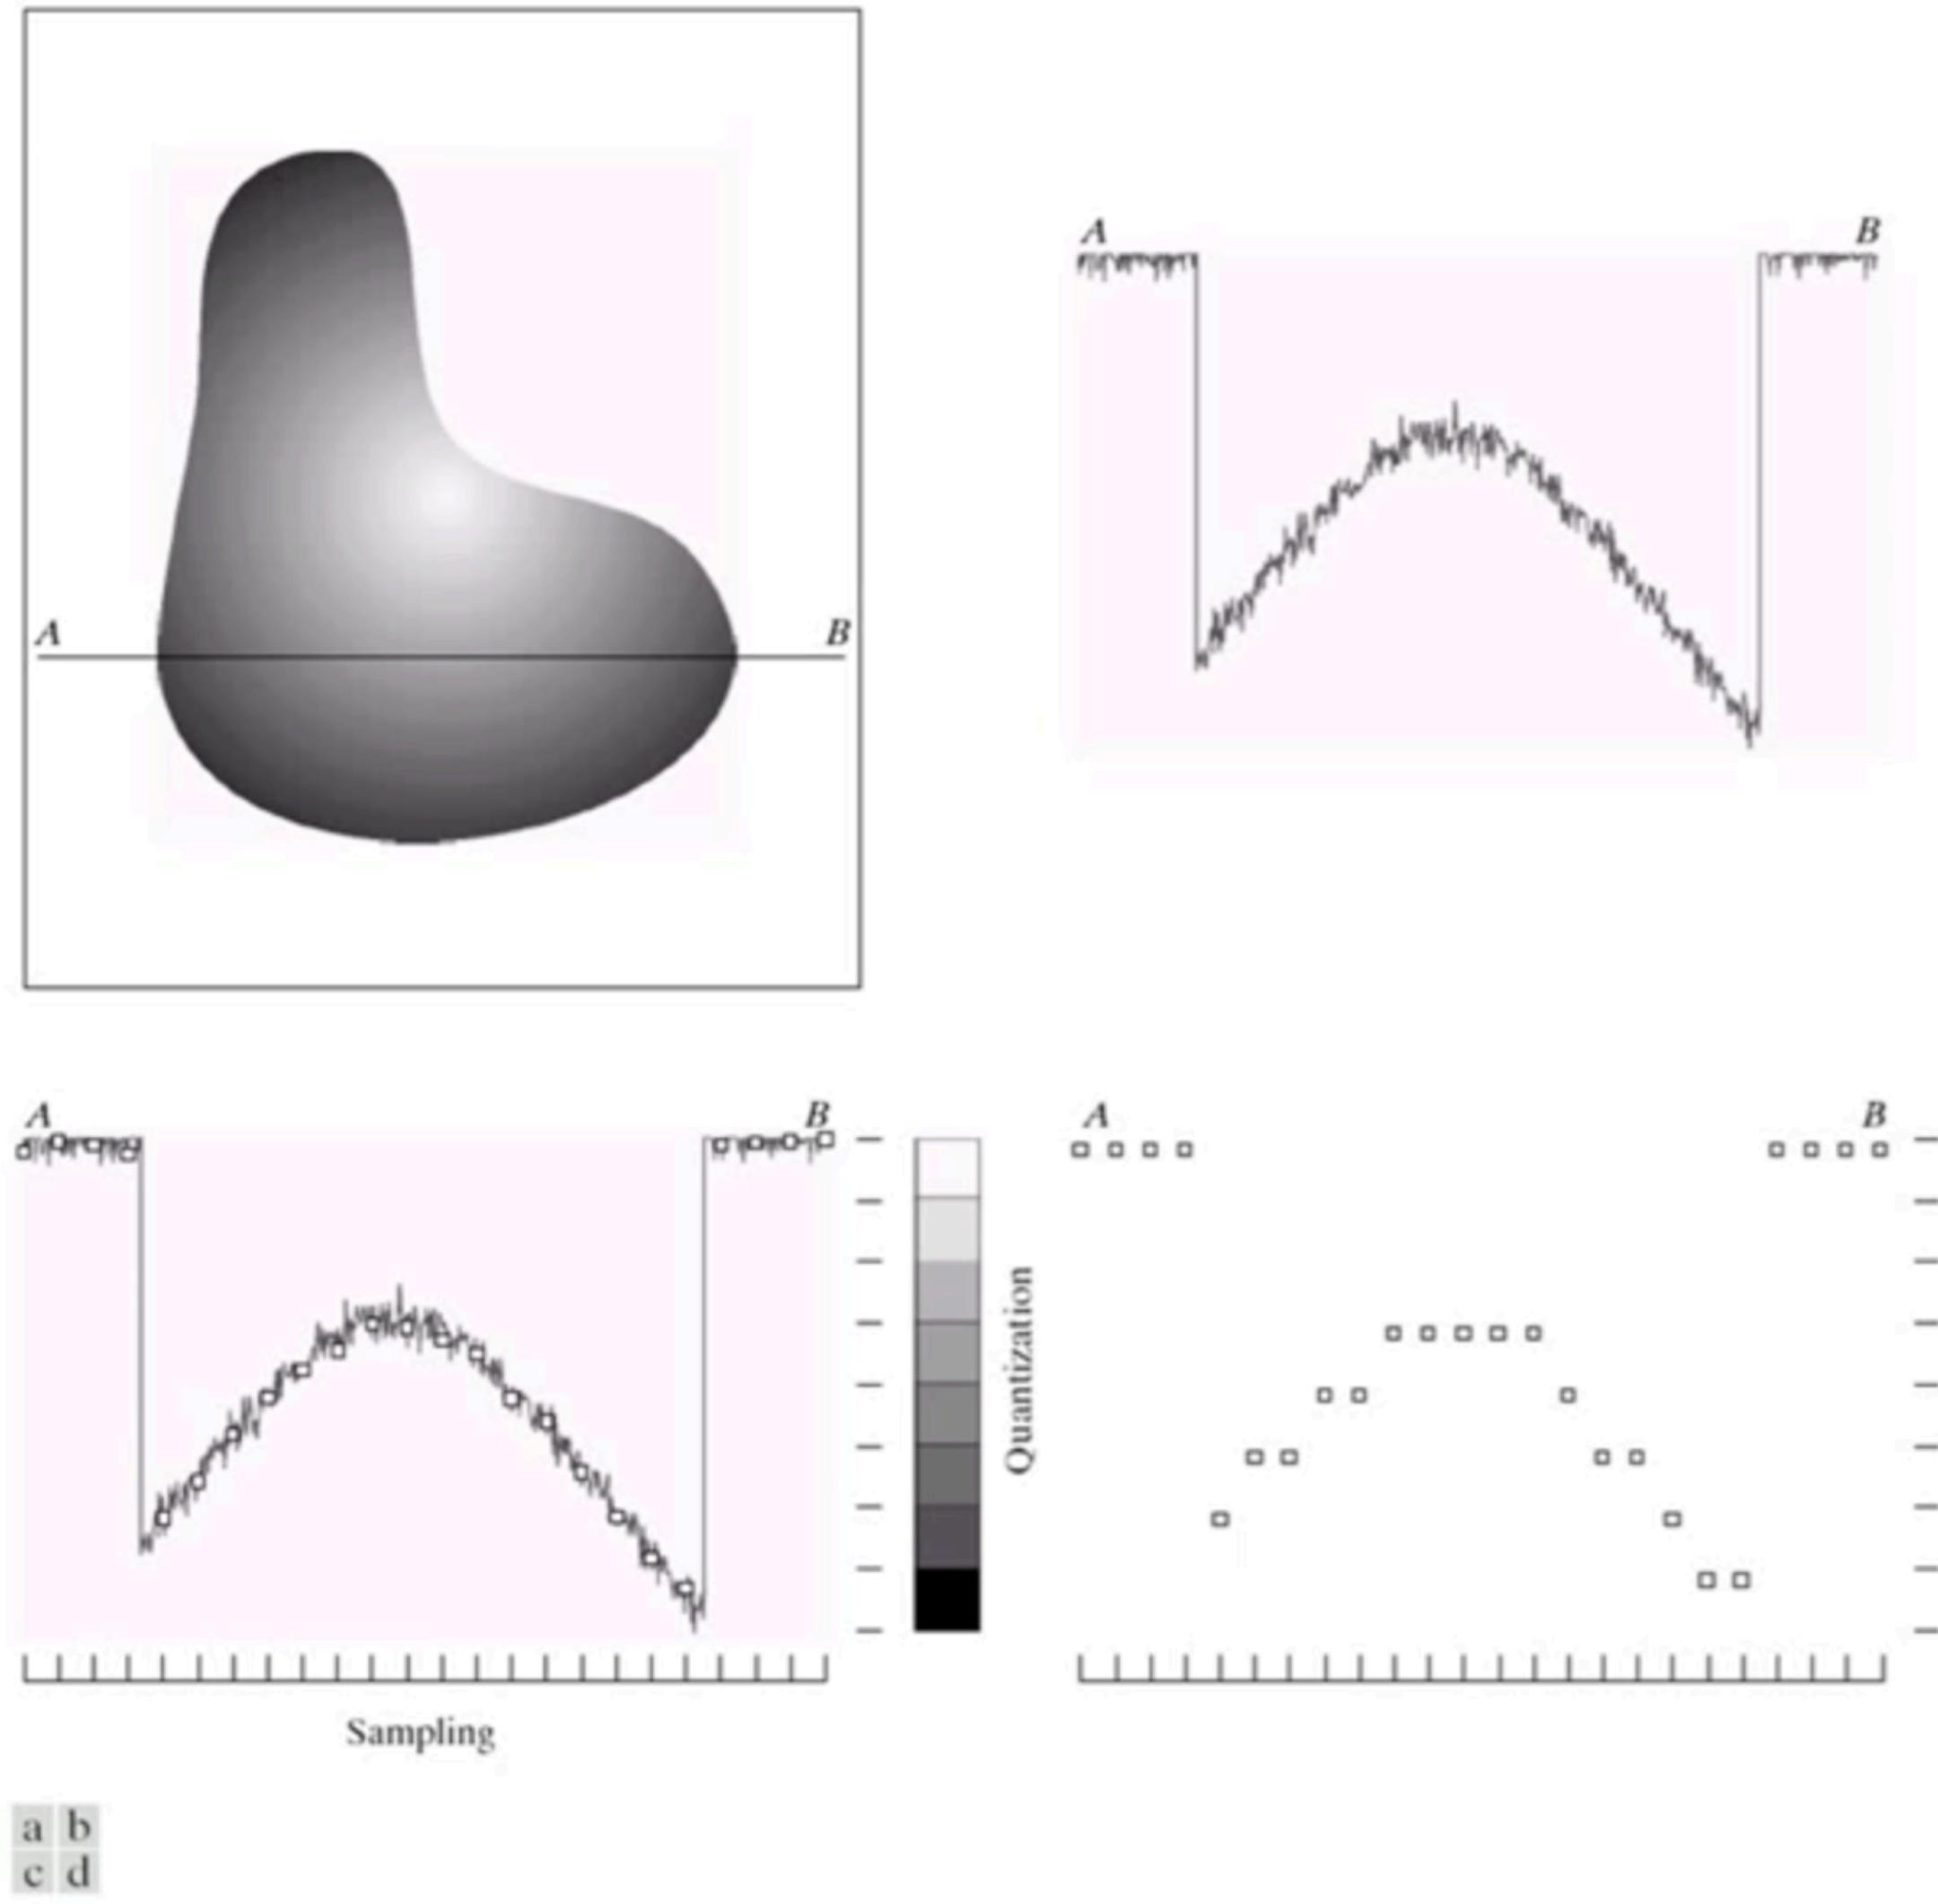
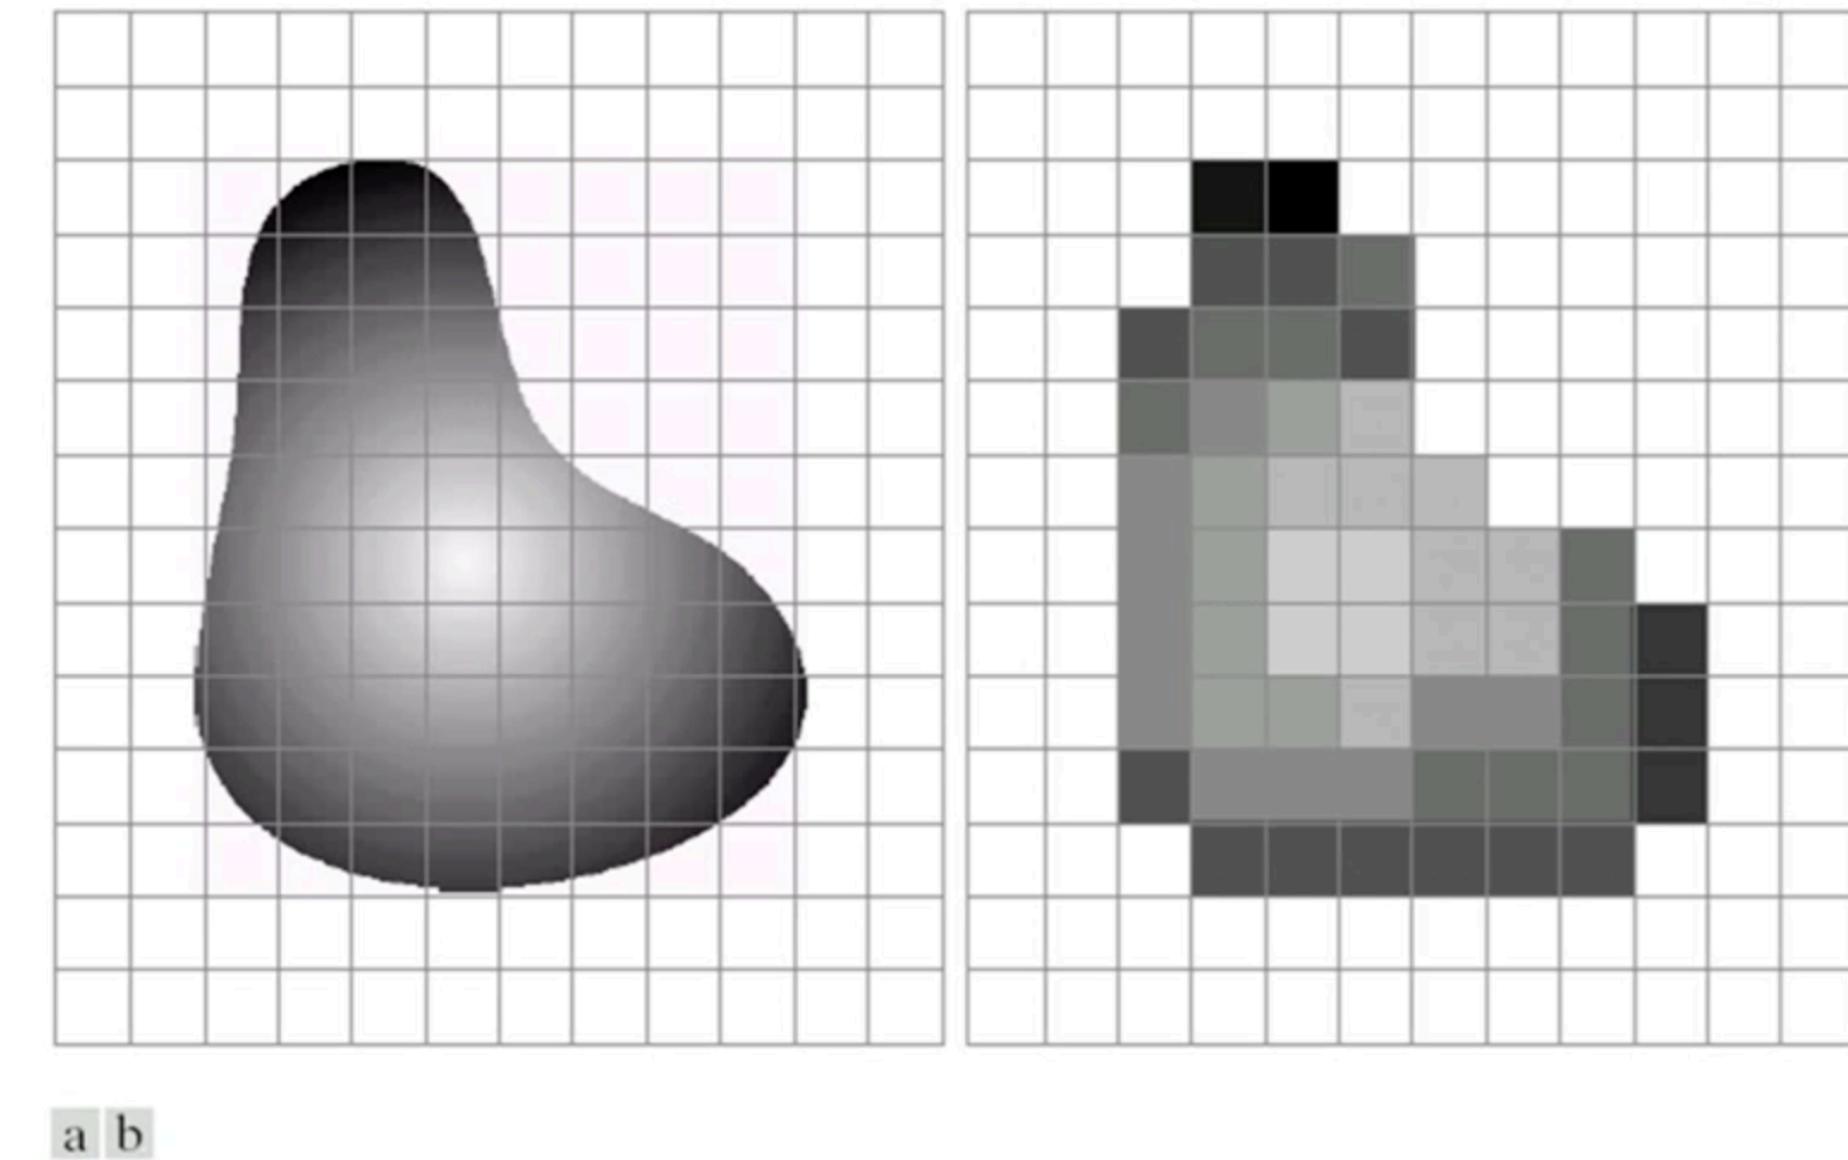


FIGURE 2.16 Generating a digital image. (a) Continuous image. (b) A scan line from *A* to *B* in the continuous image, used to illustrate the concepts of sampling and quantization. (c) Sampling and quantization. (d) Digital scan line.



a b

FIGURE 2.17 (a) Continuous image projected onto a sensor array. (b) Result of image sampling and quantization.

Images as functions

An image can be thought of as:

A 2-dimensional array of numbers ranging from some minimum to some maximum

A function I of x and y : $I(x, y)$

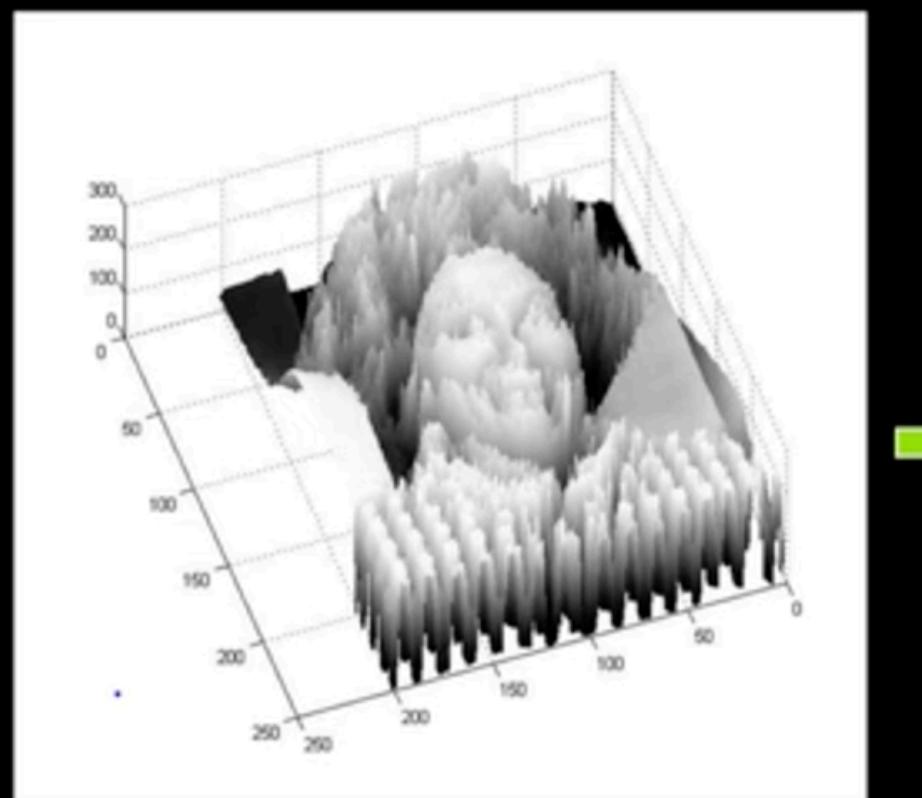
Something generated by a camera.

We think of an image as a **function**, f or I , from R^2 to R :

$f(x, y)$ gives the intensity or value at position (x, y)

Practically define the image over a rectangle, with a finite range:

$$f: [a, b] \times [c, d] \rightarrow [min, max]$$



j

i

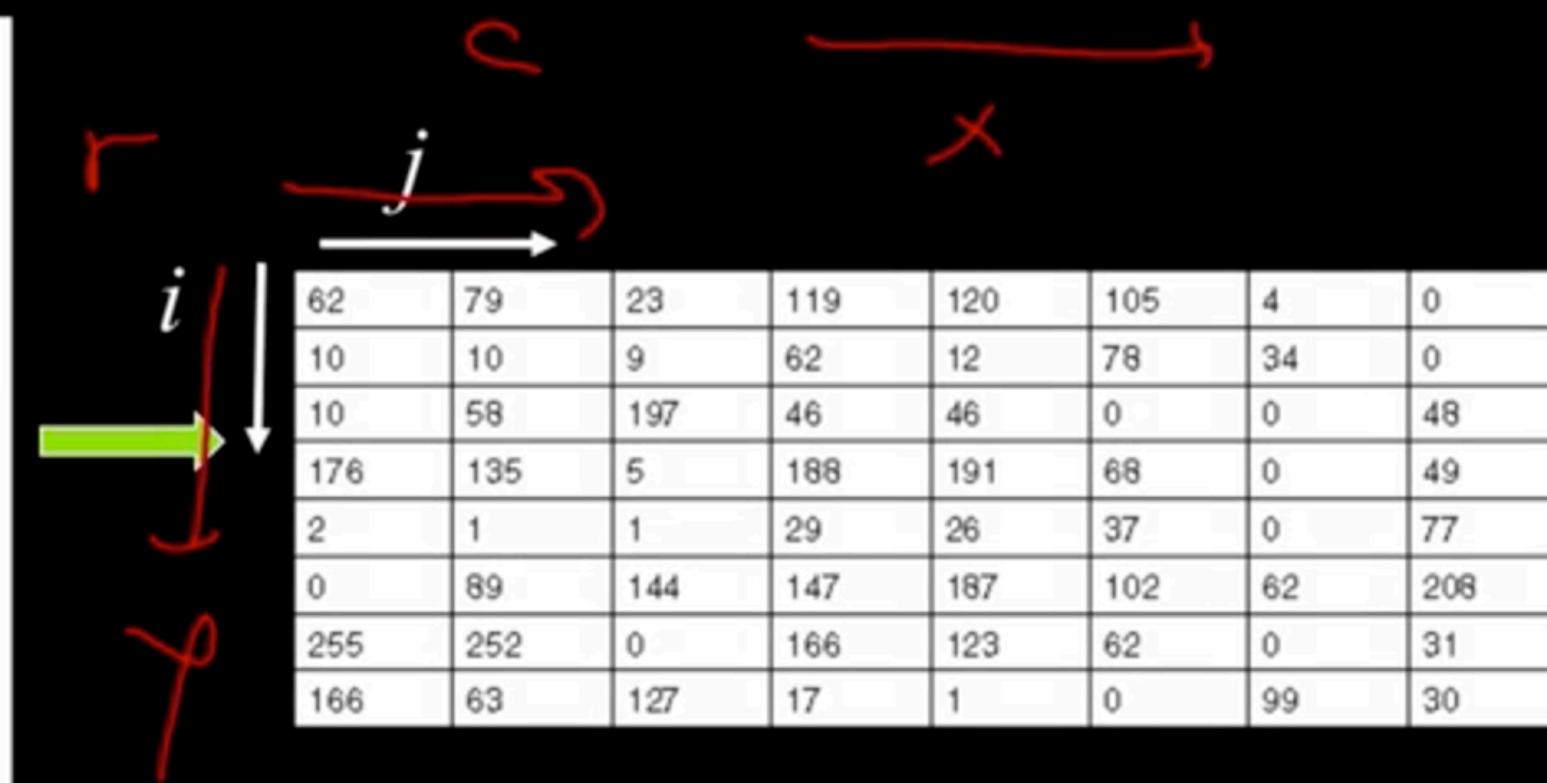
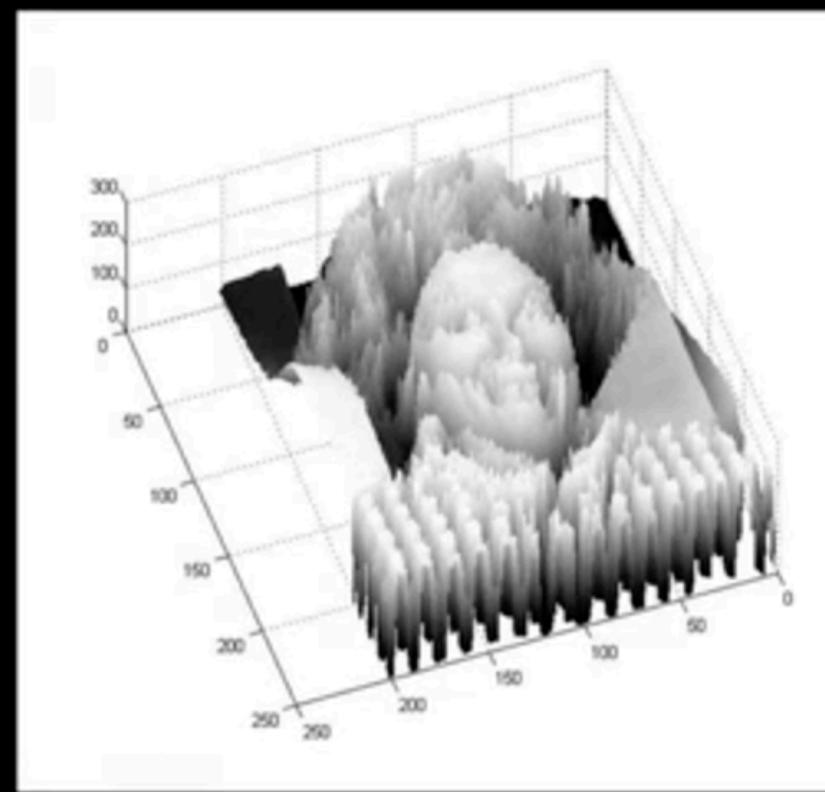
62	79	23	119	120	105	4	0
10	10	9	62	12	78	34	0
10	58	197	46	46	0	0	48
176	135	5	188	191	68	0	49
2	1	1	29	26	37	0	77
0	89	144	147	187	102	62	208
255	252	0	166	123	62	0	31
166	63	127	17	1	0	99	30

2D

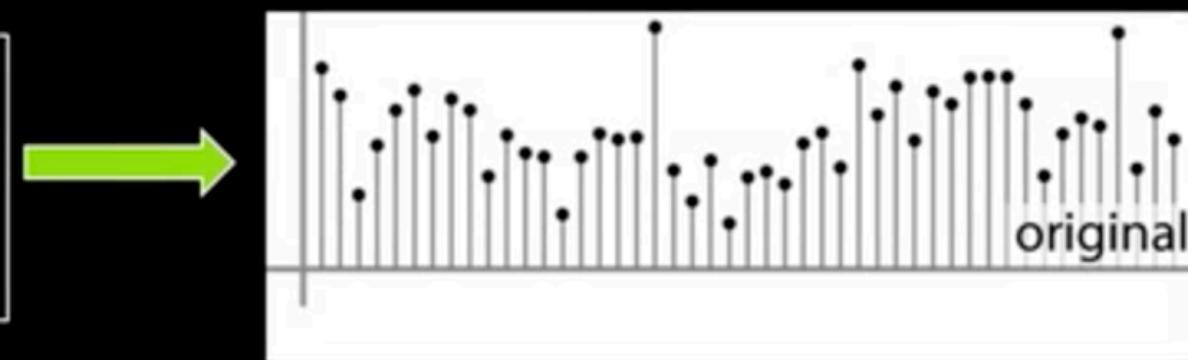
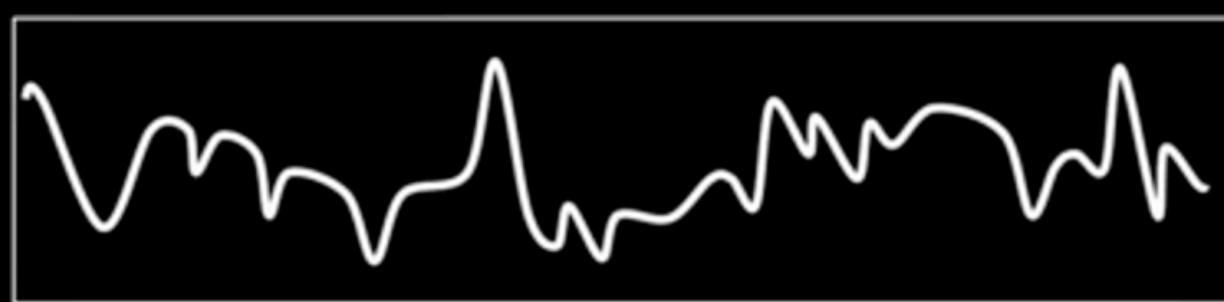
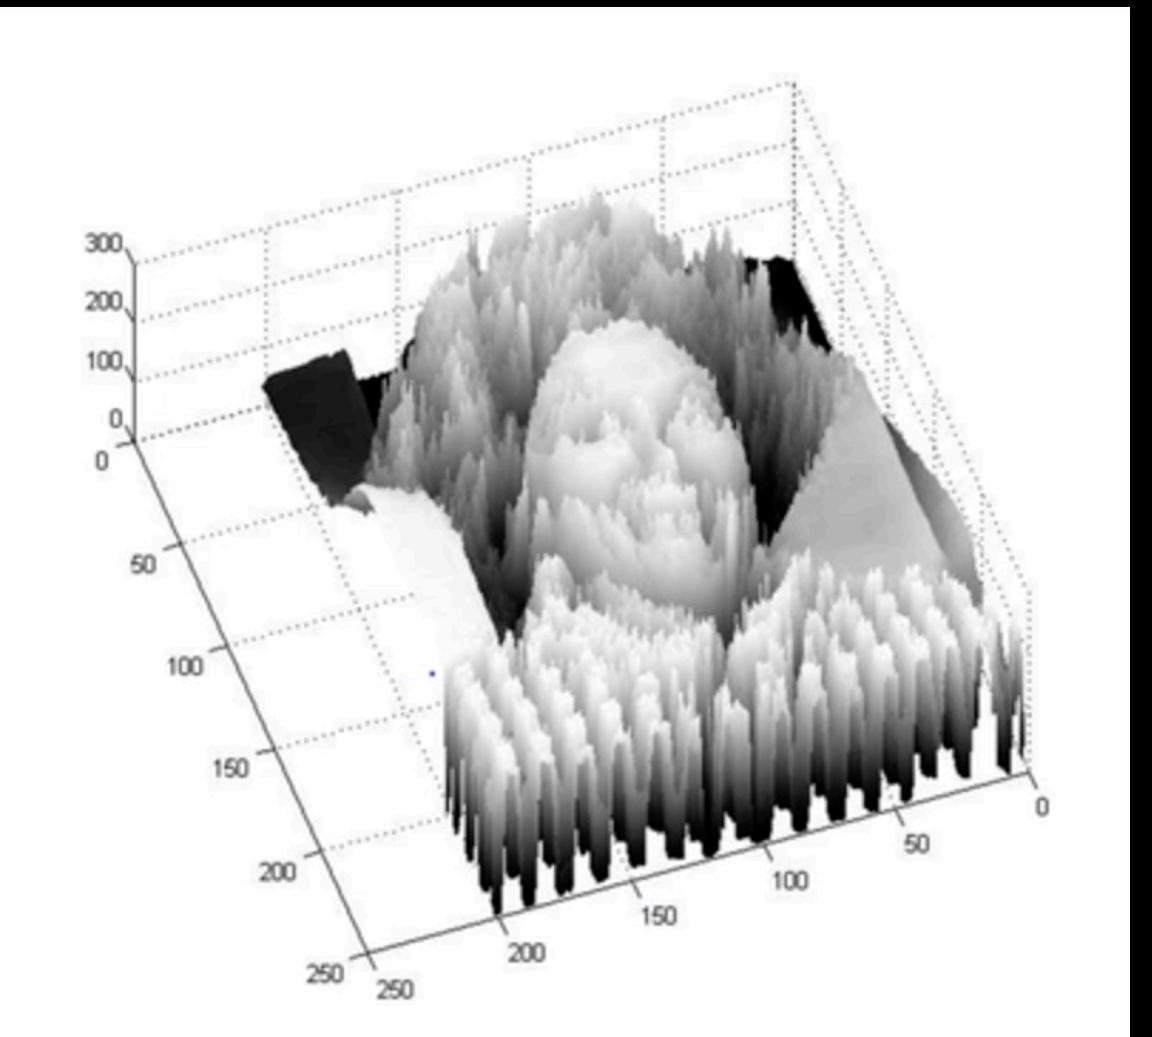
A color image is just three functions “stacked” together. We can write this as a “vector-valued” function:

$$f(x, y) = \begin{bmatrix} r(x, y) \\ g(x, y) \\ b(x, y) \end{bmatrix}$$

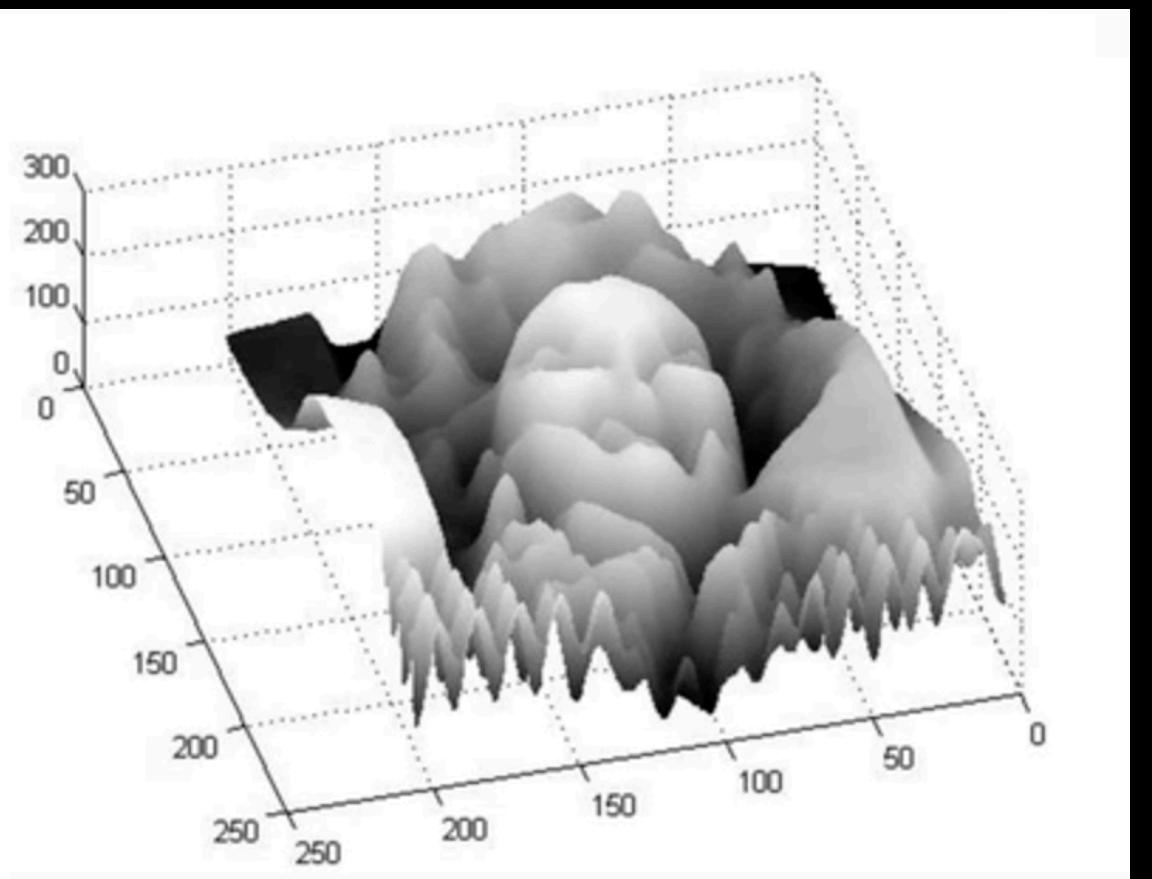
Images as functions



2D

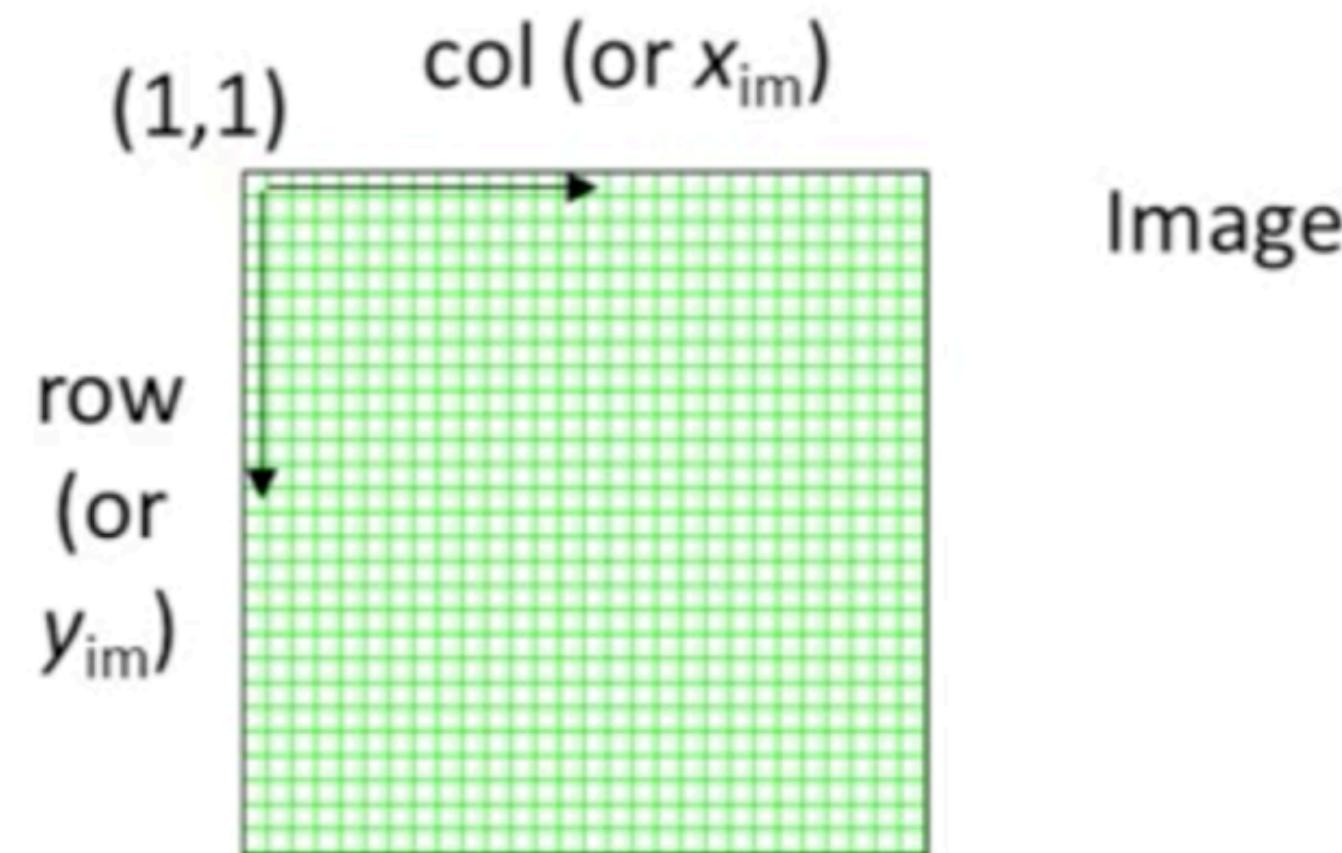


1D



Images as functions

- Image can be represented by an xy plane
- Sampling partitions the plane into a grid of pixels, whose indices are integers
- We can also think of it as an $M \times N$ matrix of numbers
- We can index an element either by (row, col) or (x, y)
- We will use the convention
 - Top left pixel is $(1, 1)$
 - x index (or column) increases to the right
 - y index (or row) increases down



Example

- An image has 8 bits per pixel
 - If unsigned, range of values is?
 - If signed (ie., two's complement), range is?
 - Number of bytes in a 3872×2592 pixel image?

Example (2)

- An image has 8 bits per pixel

- If unsigned, range of values is?

$$0..(2^8-1) = 0..(2^8-1) = 0..255$$

- If signed (ie., two's complement), range is?

$$-(2^{8-1}) .. + (2^{8-1}-1) \rightarrow -128 .. +127$$

- Number of bytes in a 3872 x 2592 pixel image?

$$3872 \times 2592 = 10,036,224$$

10 MB

N bits per pixel
allows 2^N values



Image size, image resolution and spatial resolution

Image size

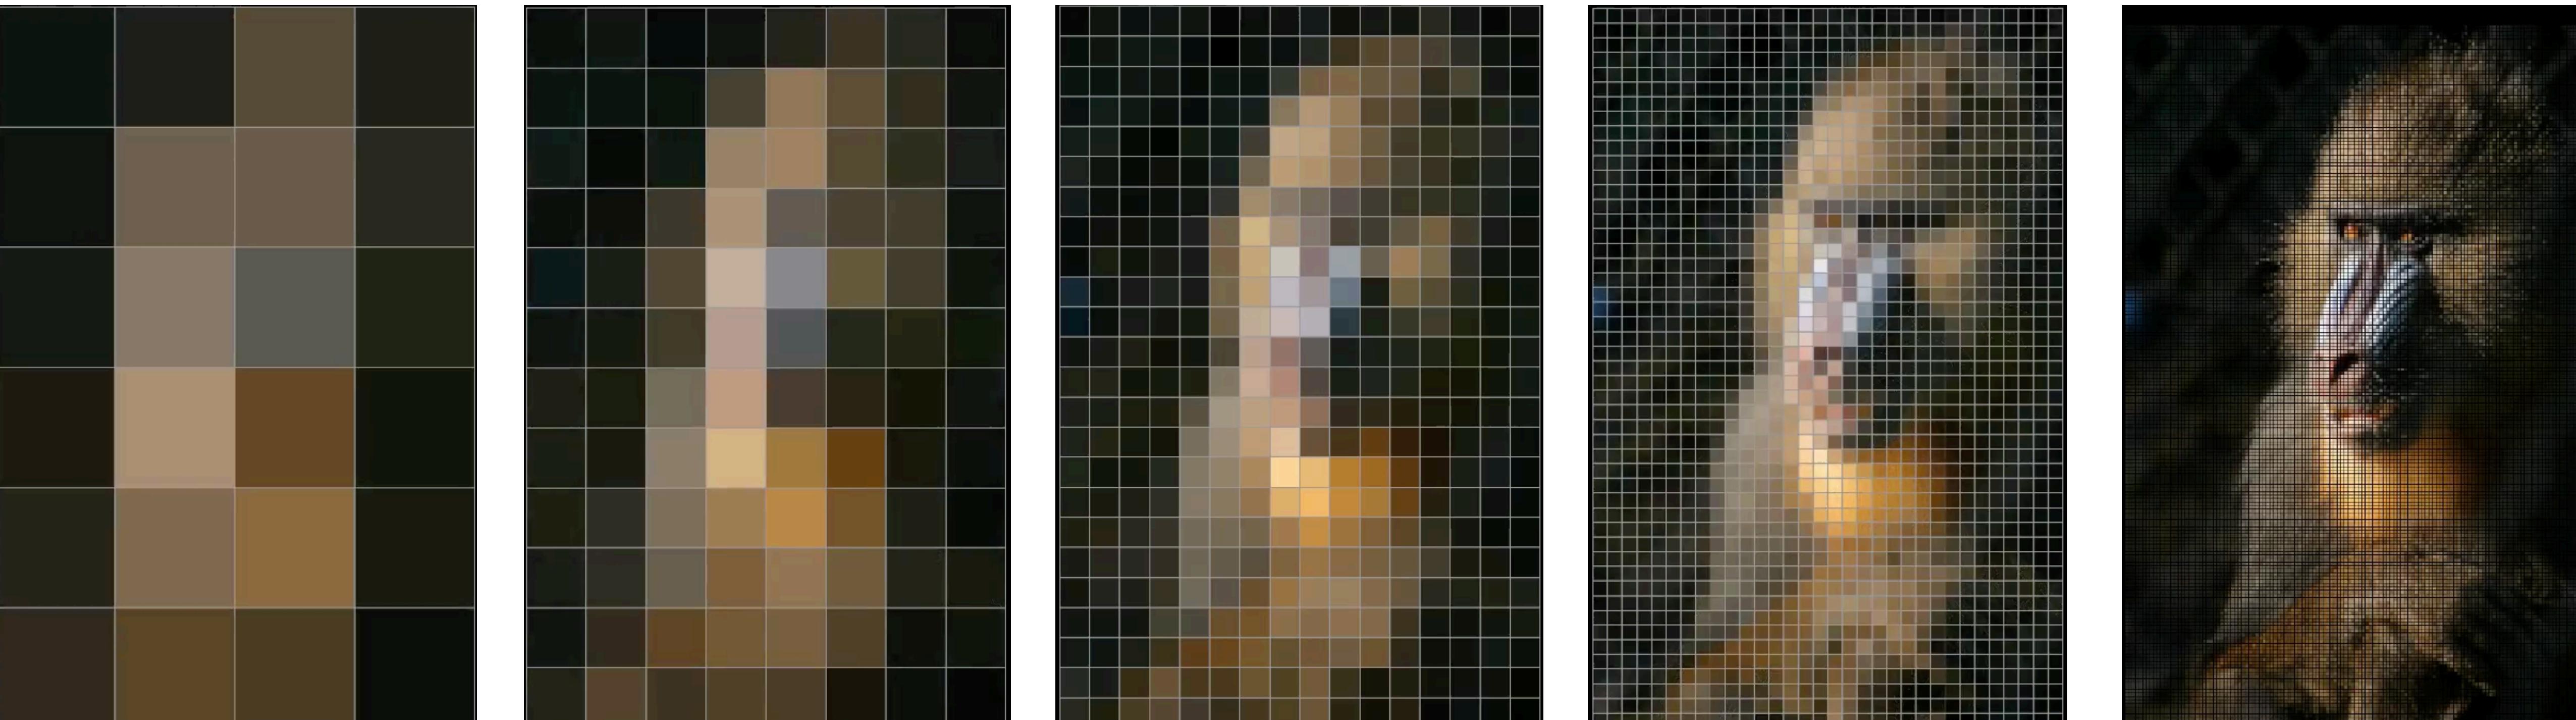
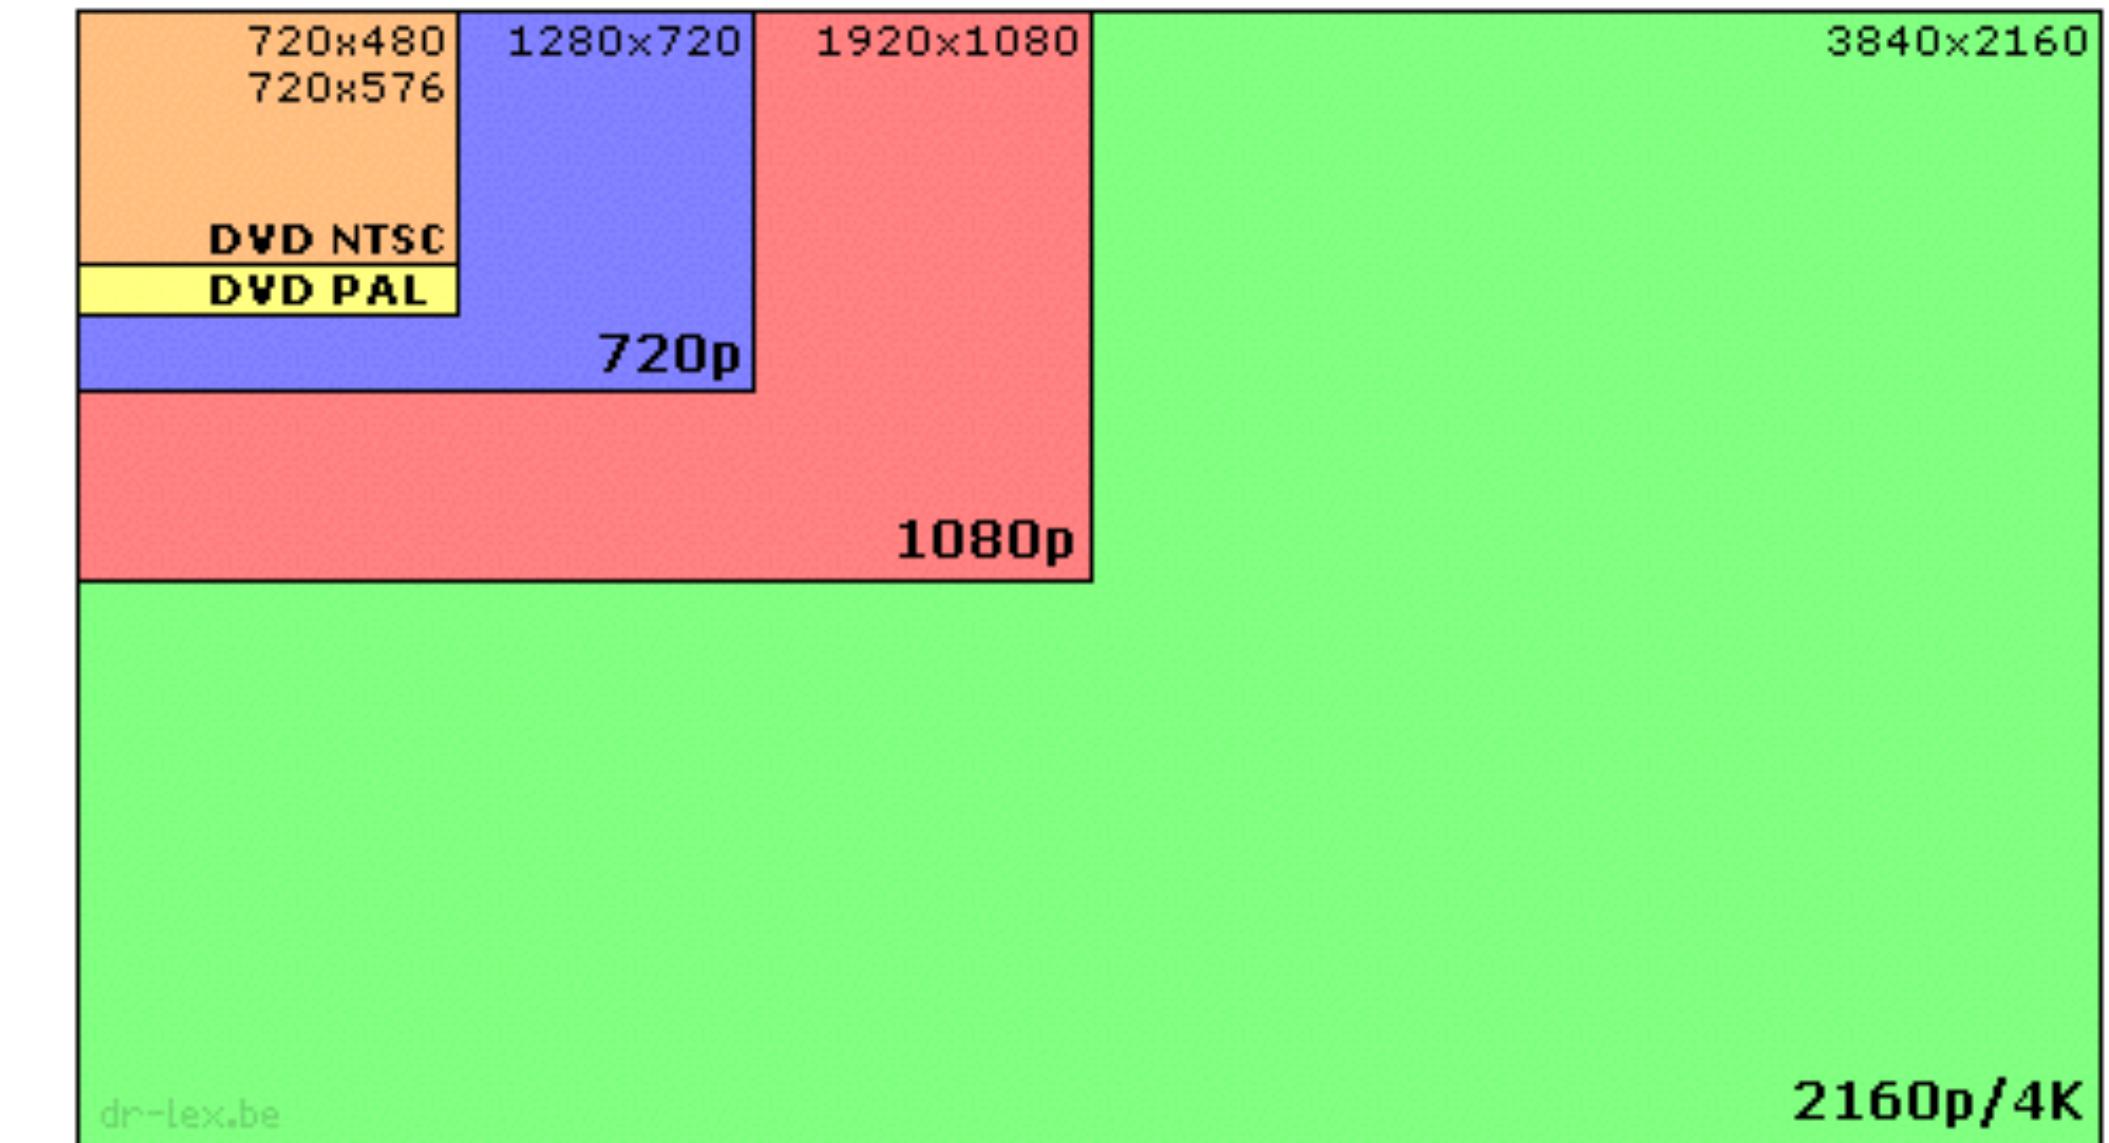
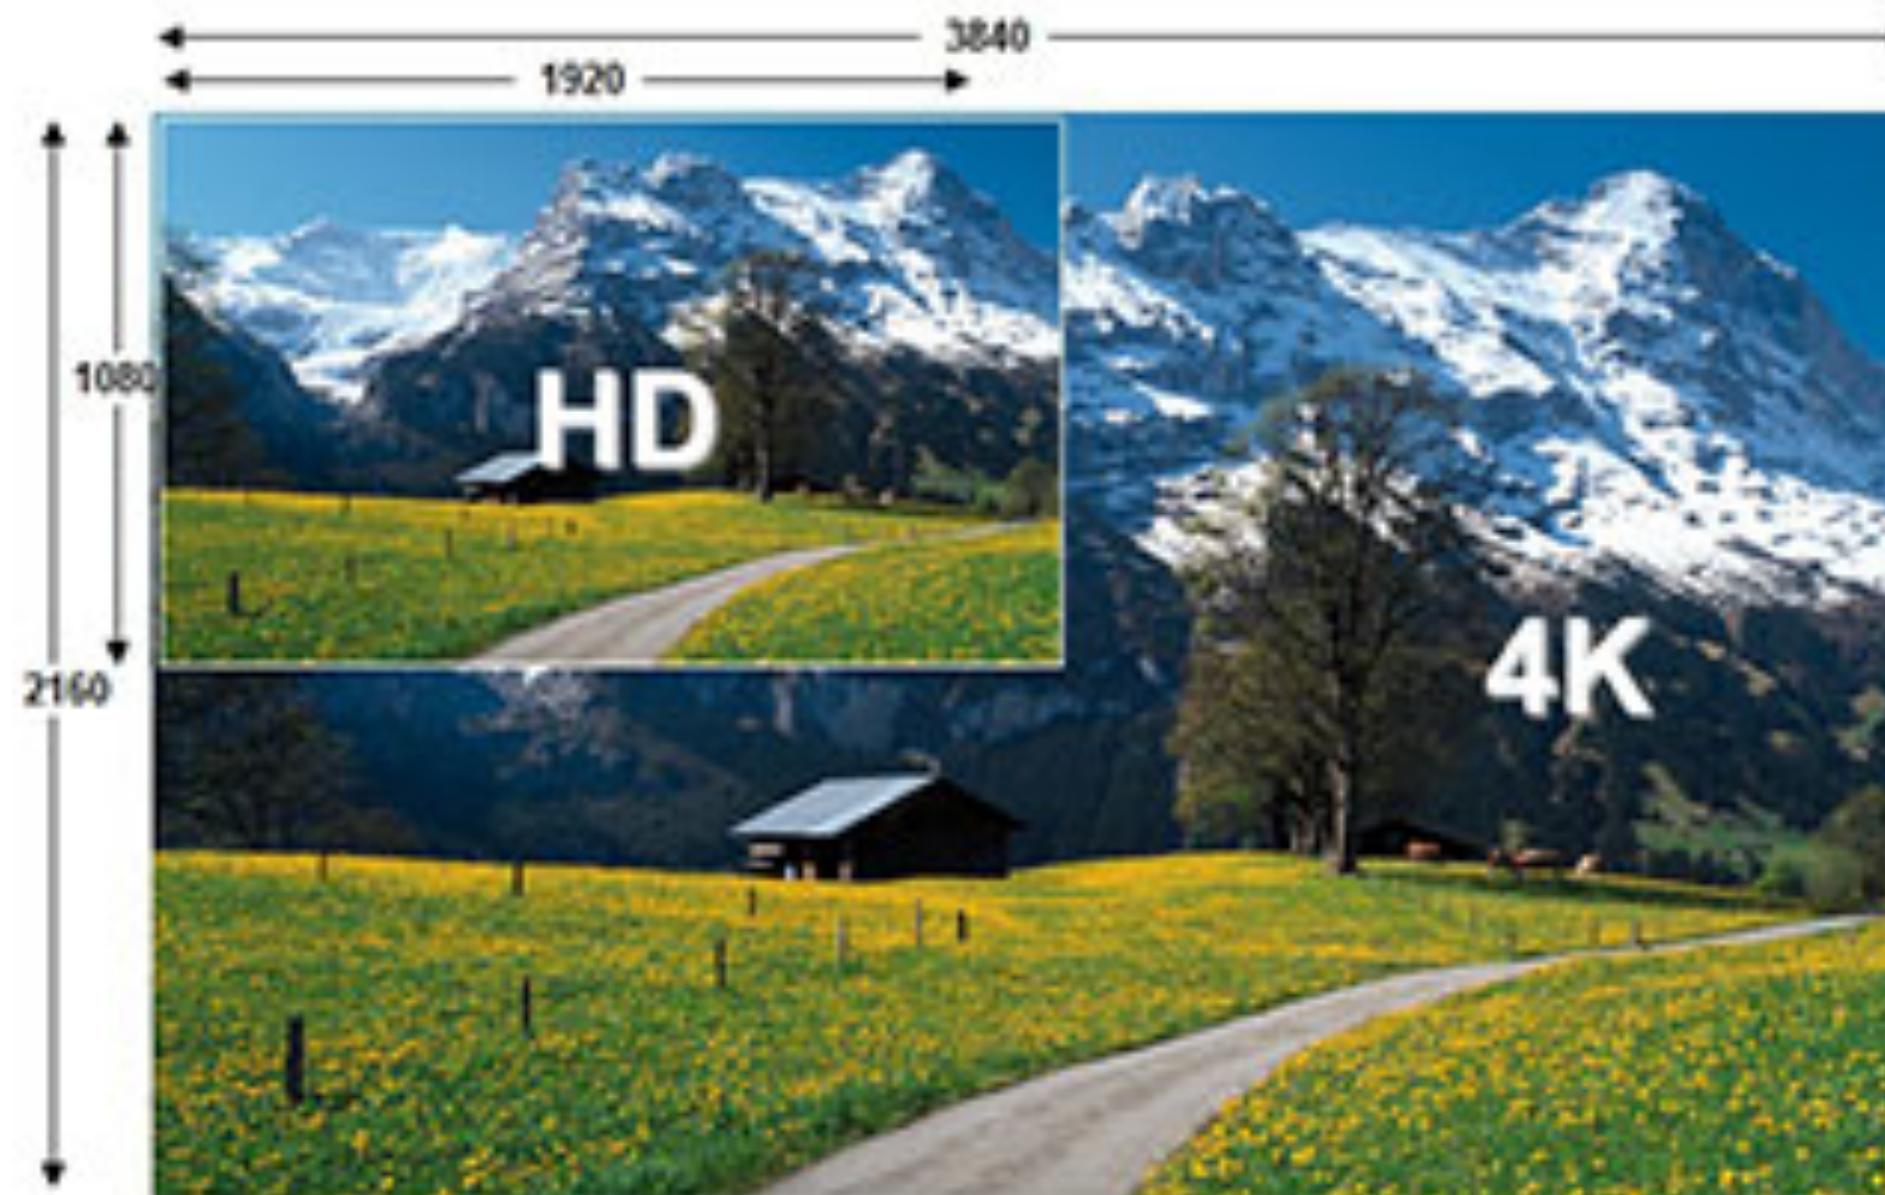


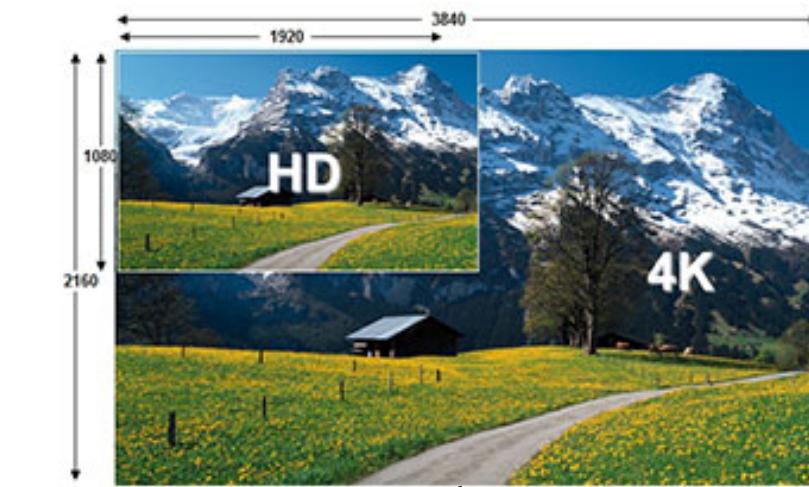
Image Size



pixels (width x height)

Image Resolution

Image



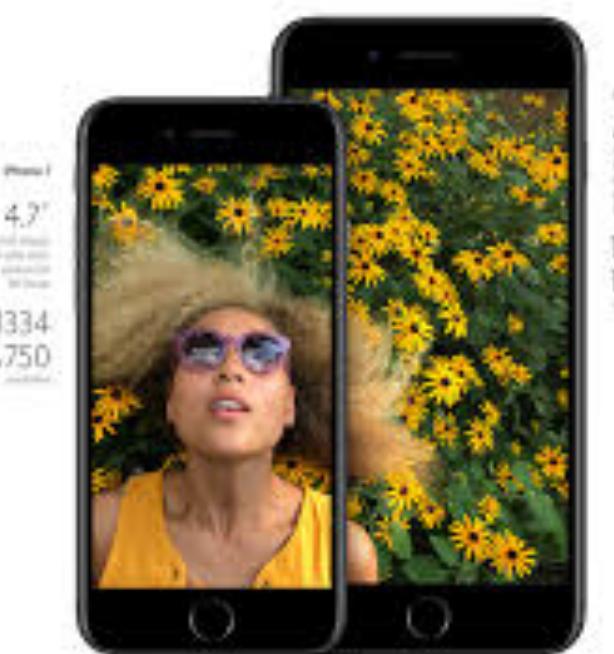
Display



TV



Laptop



Phone

DPI

Spatial Resolution

- To resolve an object, we must use a wavelength equal to the size of the object or smaller

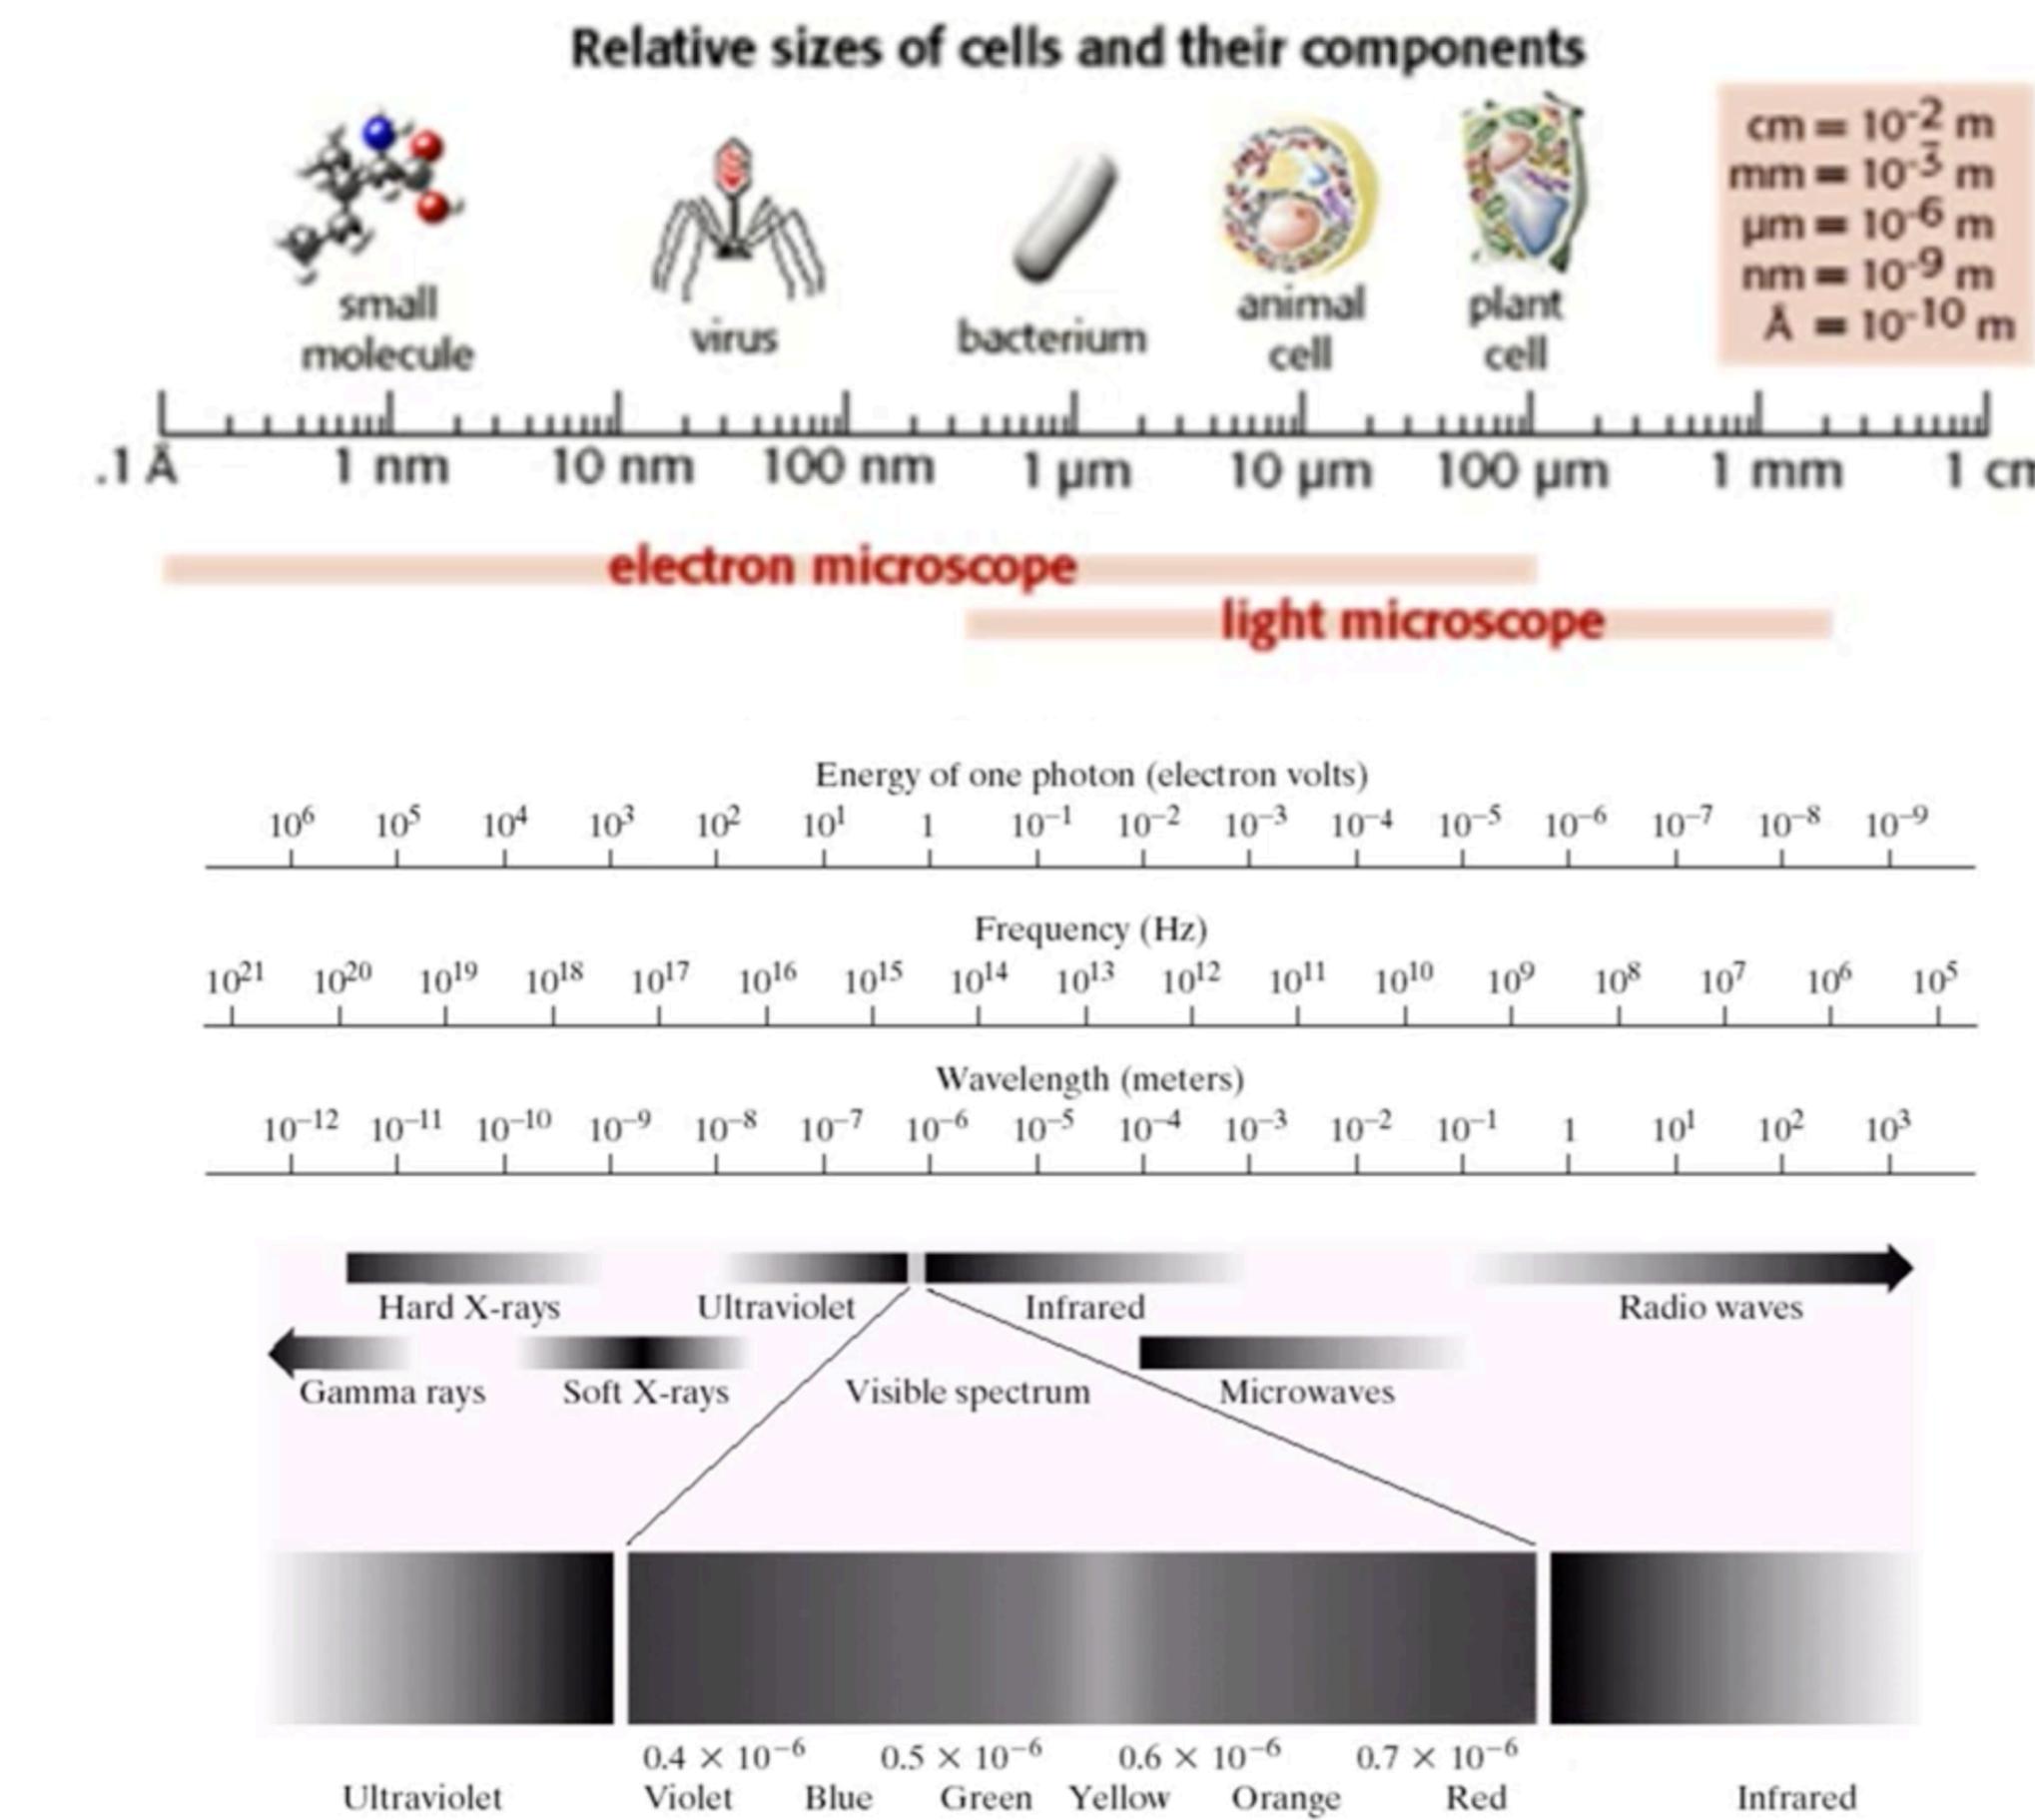
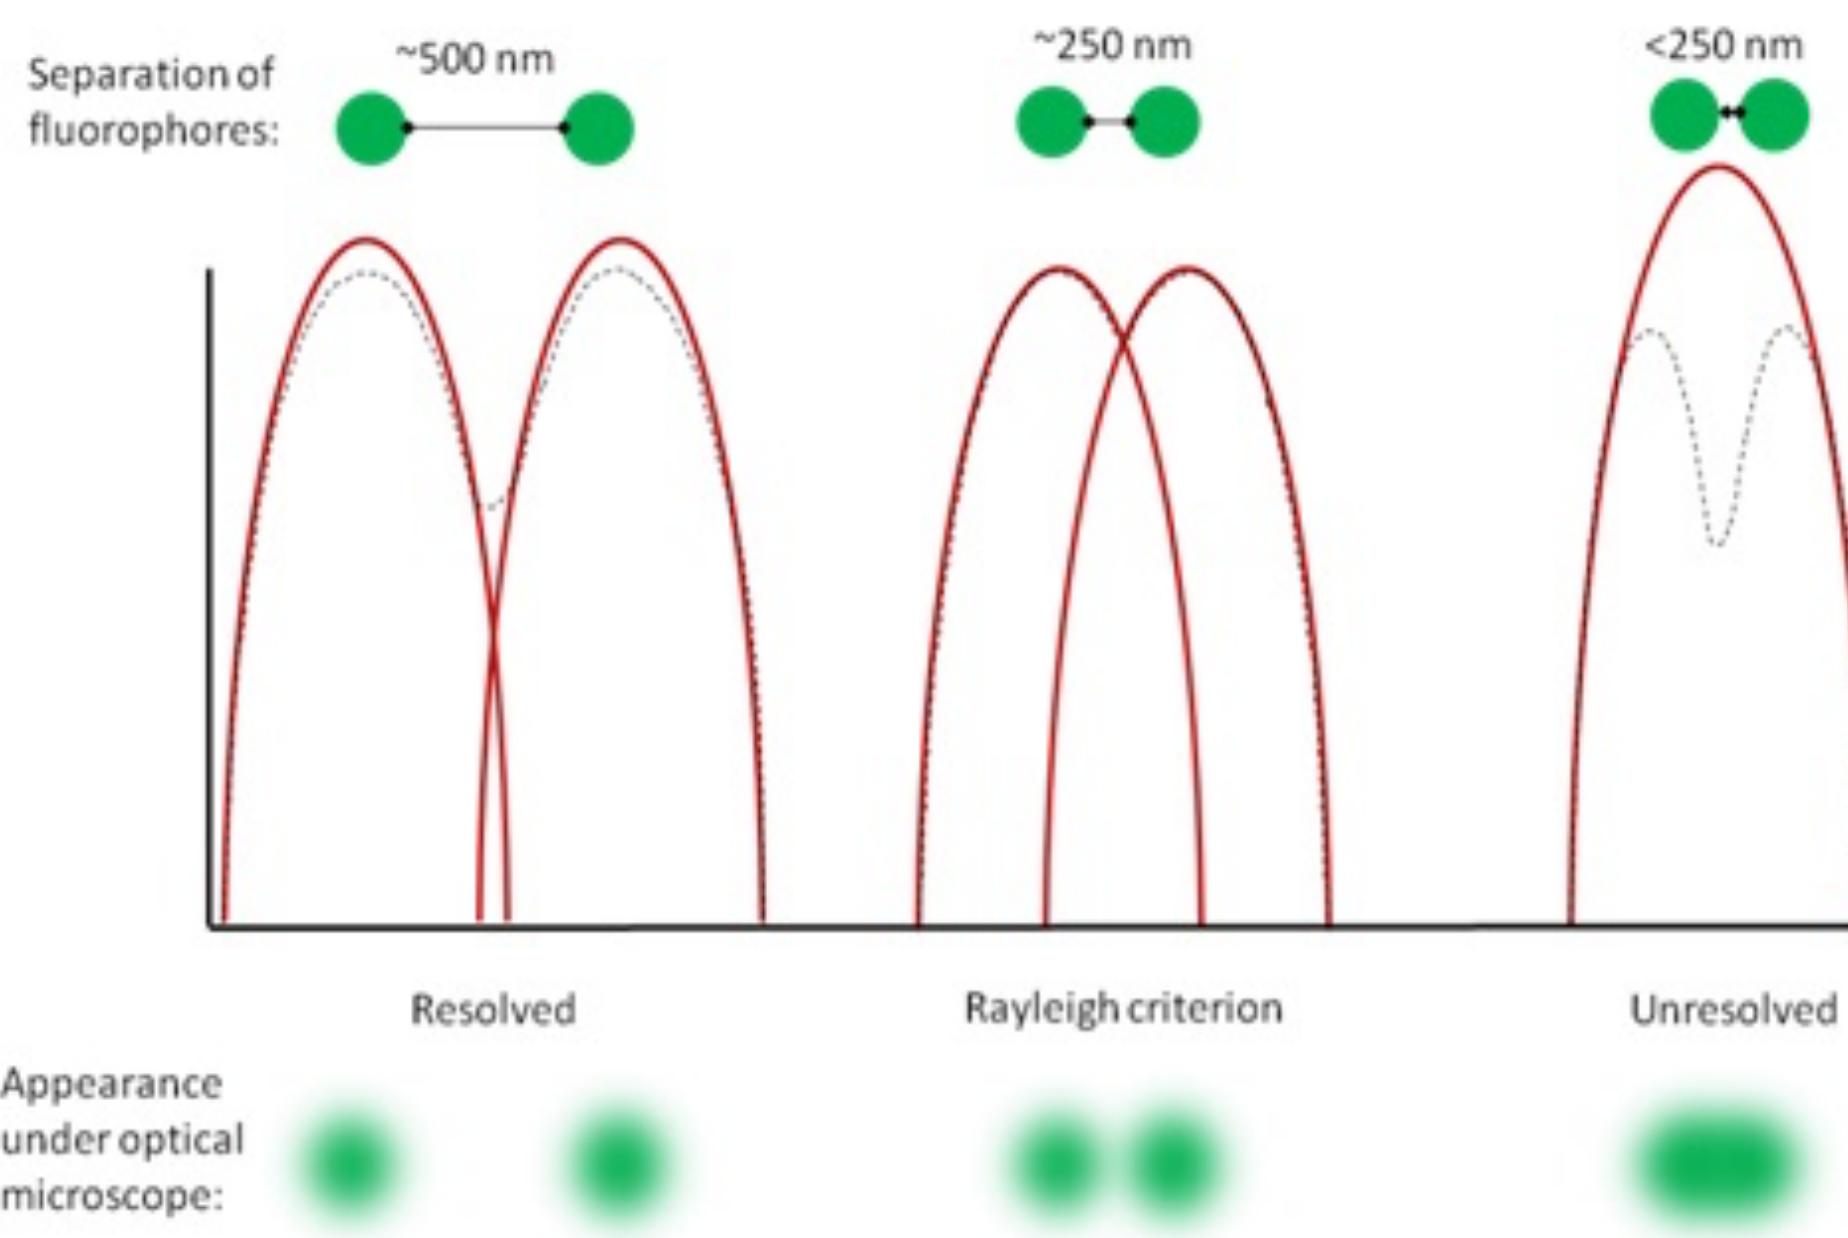
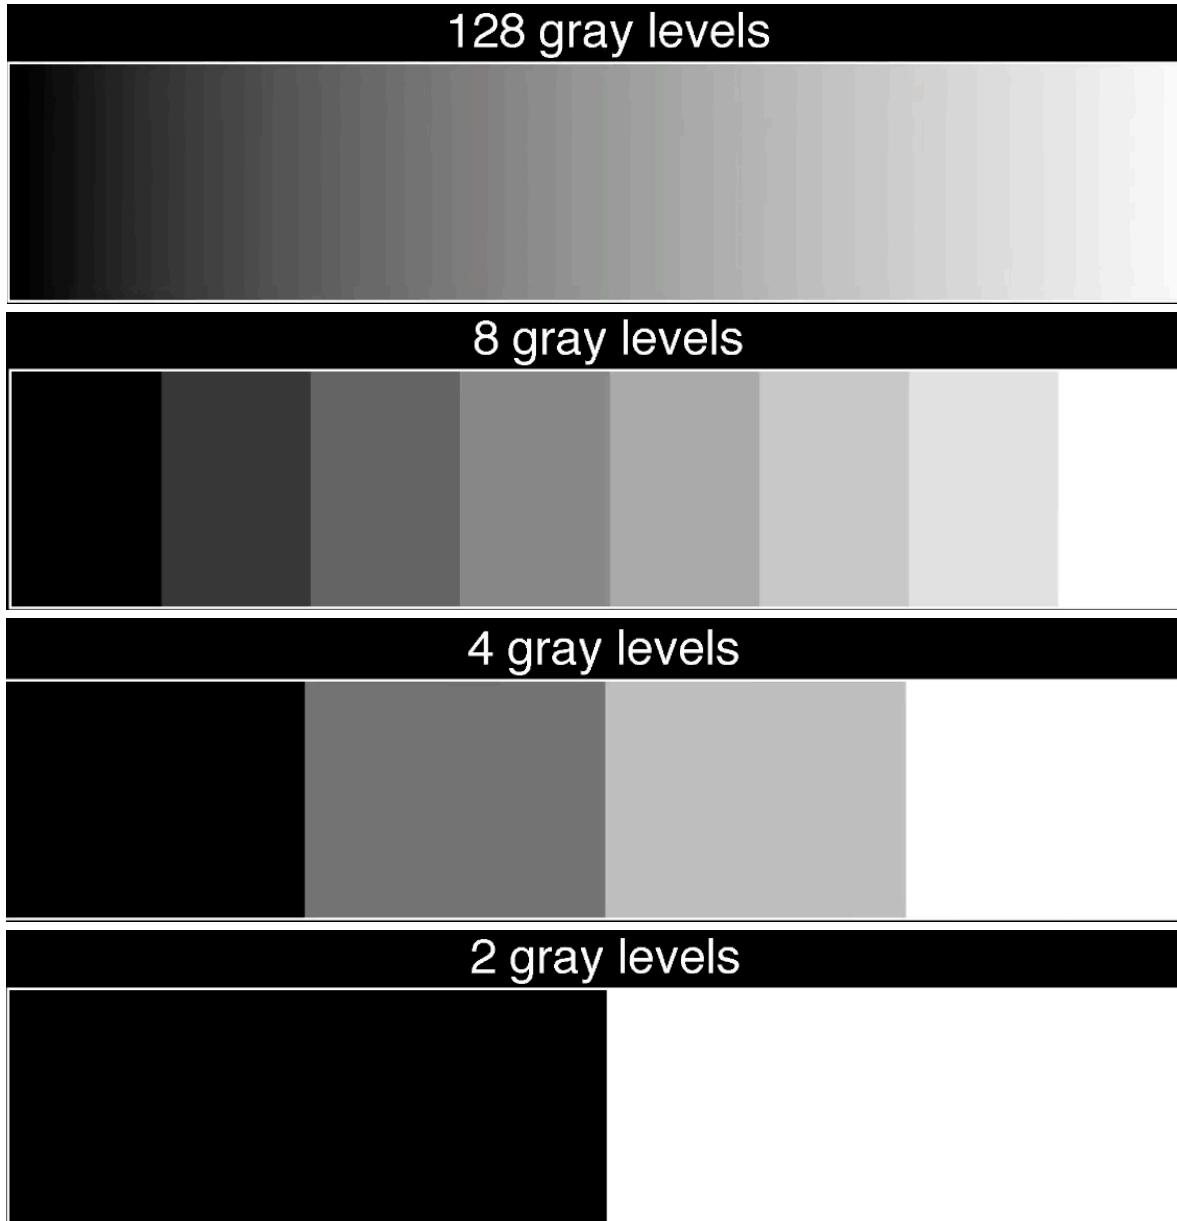


FIGURE 2.10 The electromagnetic spectrum. The visible spectrum is shown zoomed to facilitate explanation, but note that the visible spectrum is a rather narrow portion of the EM spectrum.

Color levels and Intensity Resolution

128 gray levels



128 gray levels



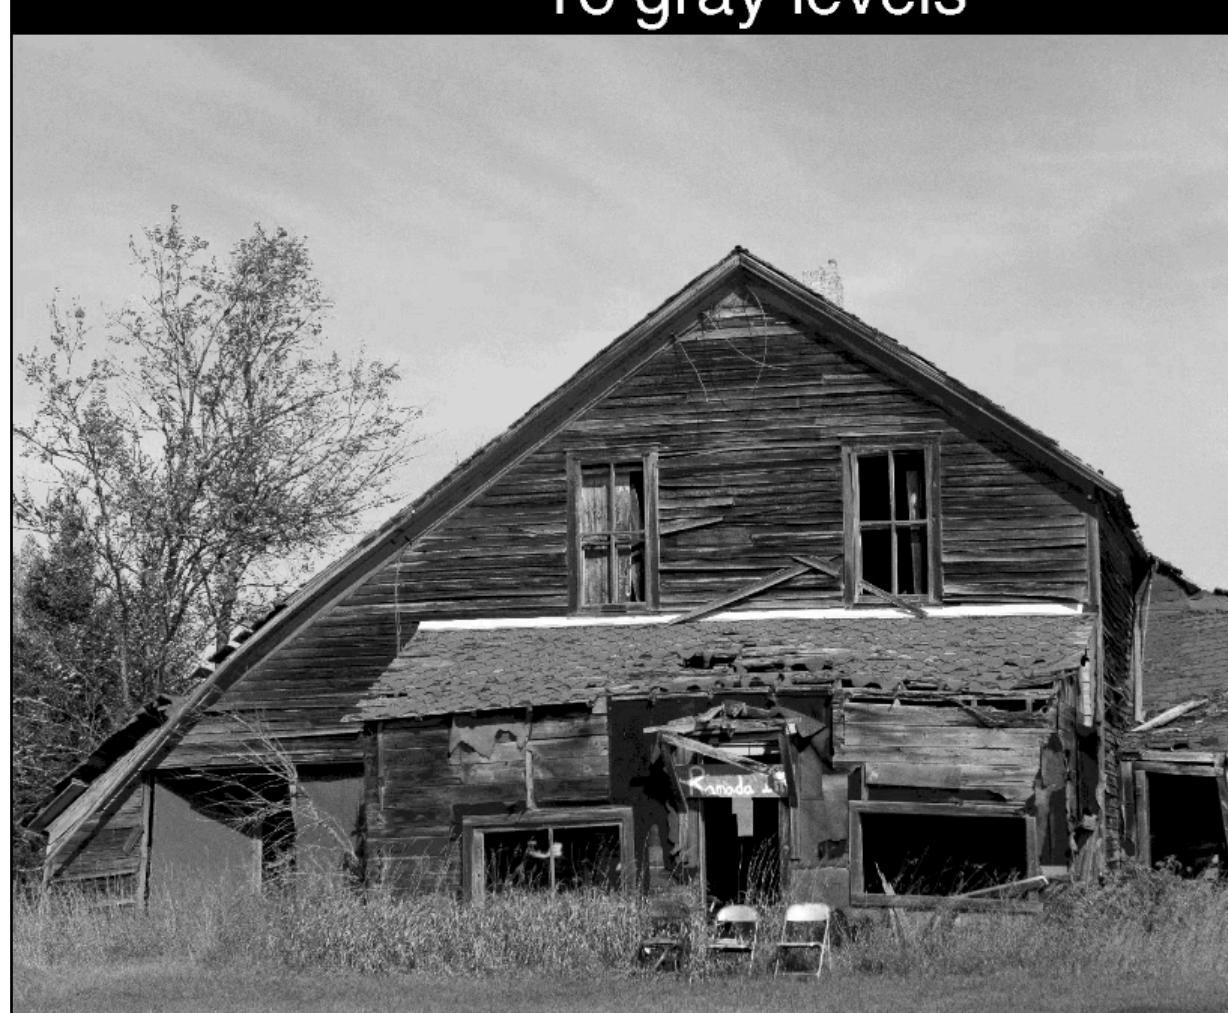
64 gray levels



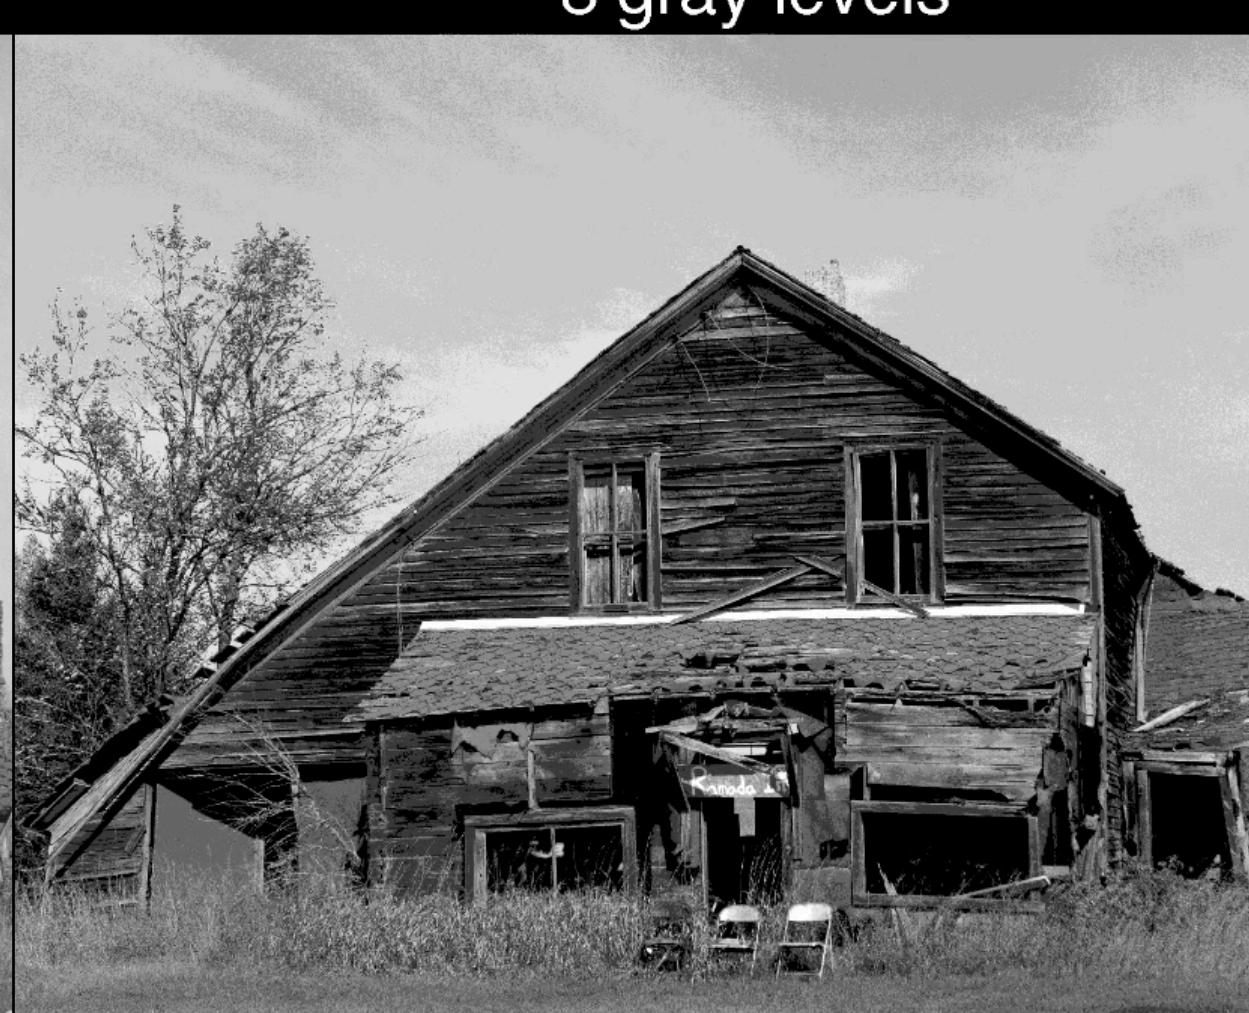
32 gray levels



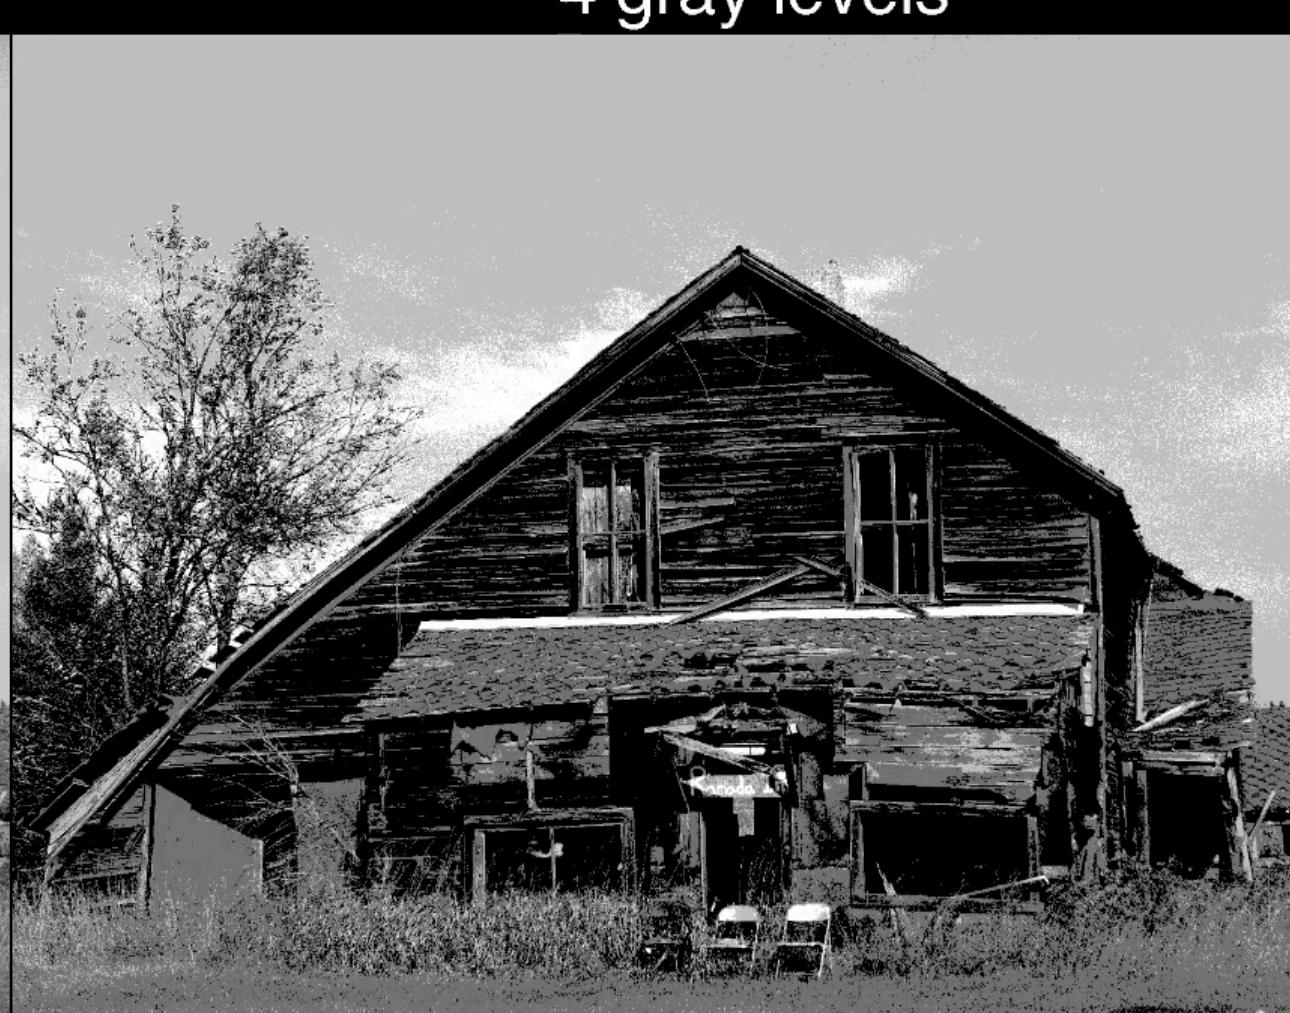
16 gray levels



8 gray levels



4 gray levels



2 gray levels



Summary

- image size (# pixels / pixel size) influences image resolution
- # color / gray levels influences intensity resolution
- Spatial / Intensity resolution intended to quantify the smallest details that can be perceived or recognized in a digital image.
- Spatial resolution = smallest change in position that can be recognized / *smallest distance between physical points that the imaging system can resolve.*
- Intensity resolution = smallest change in brightness that can be recognized.

$$f(x,y) = \begin{bmatrix} r(x,y) \\ g(x,y) \\ b(x,y) \end{bmatrix}$$

Graylevel vs. Color Images

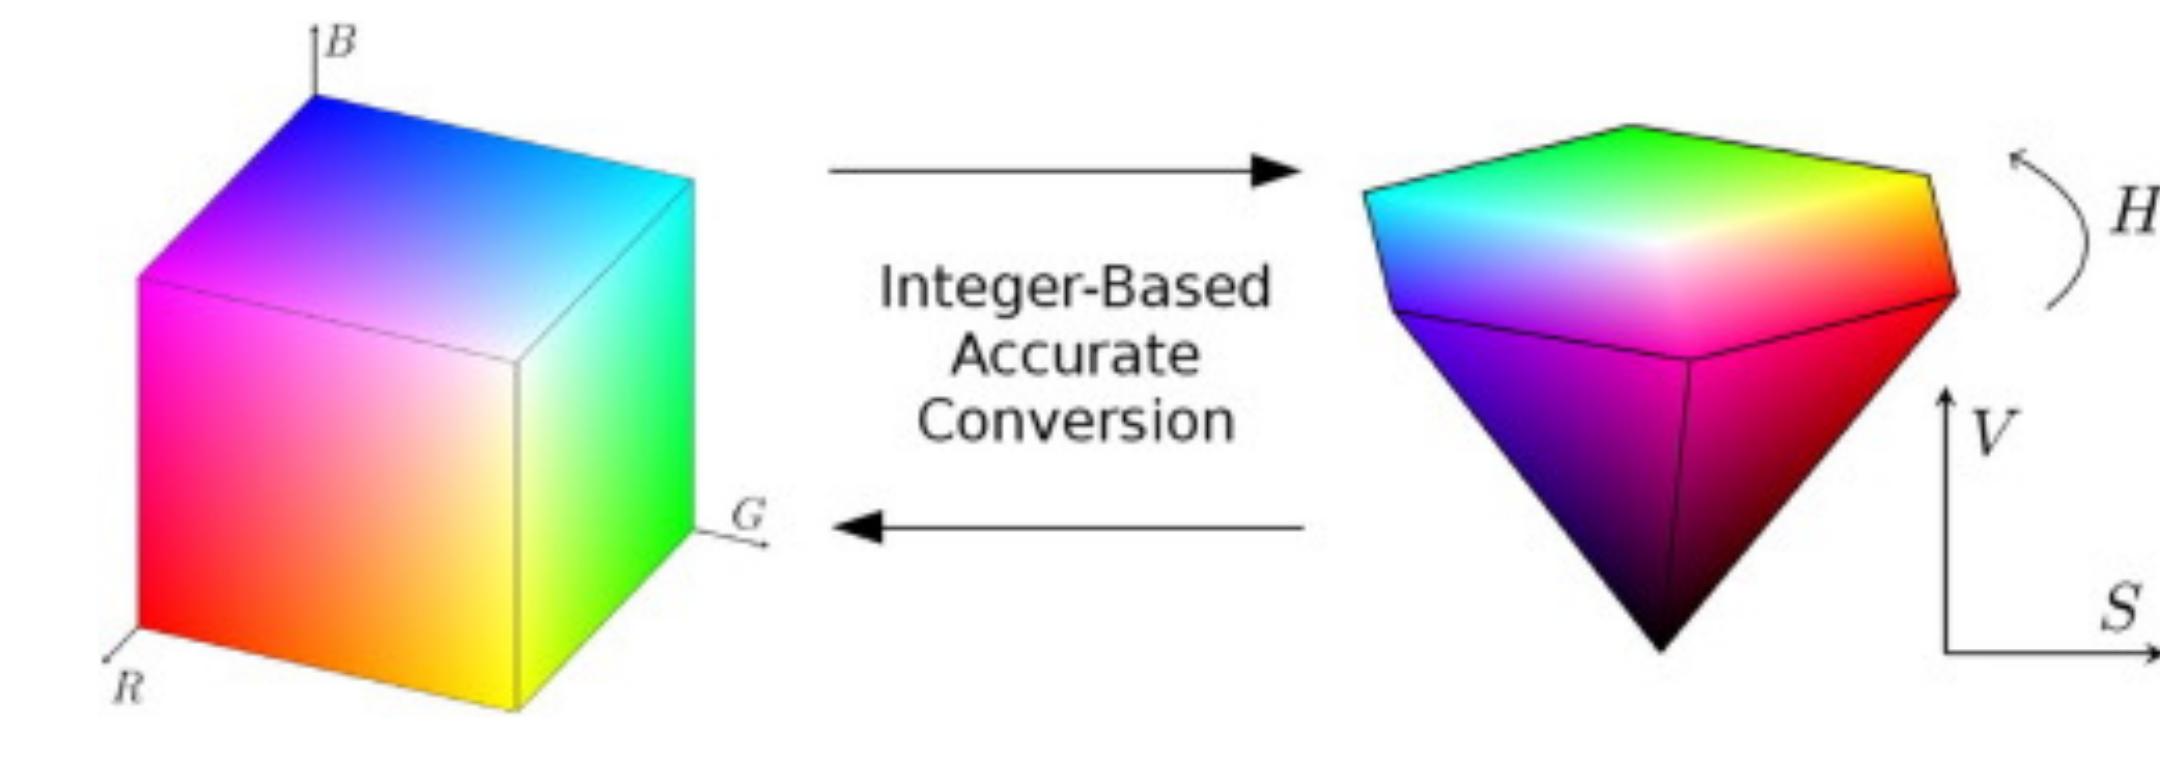
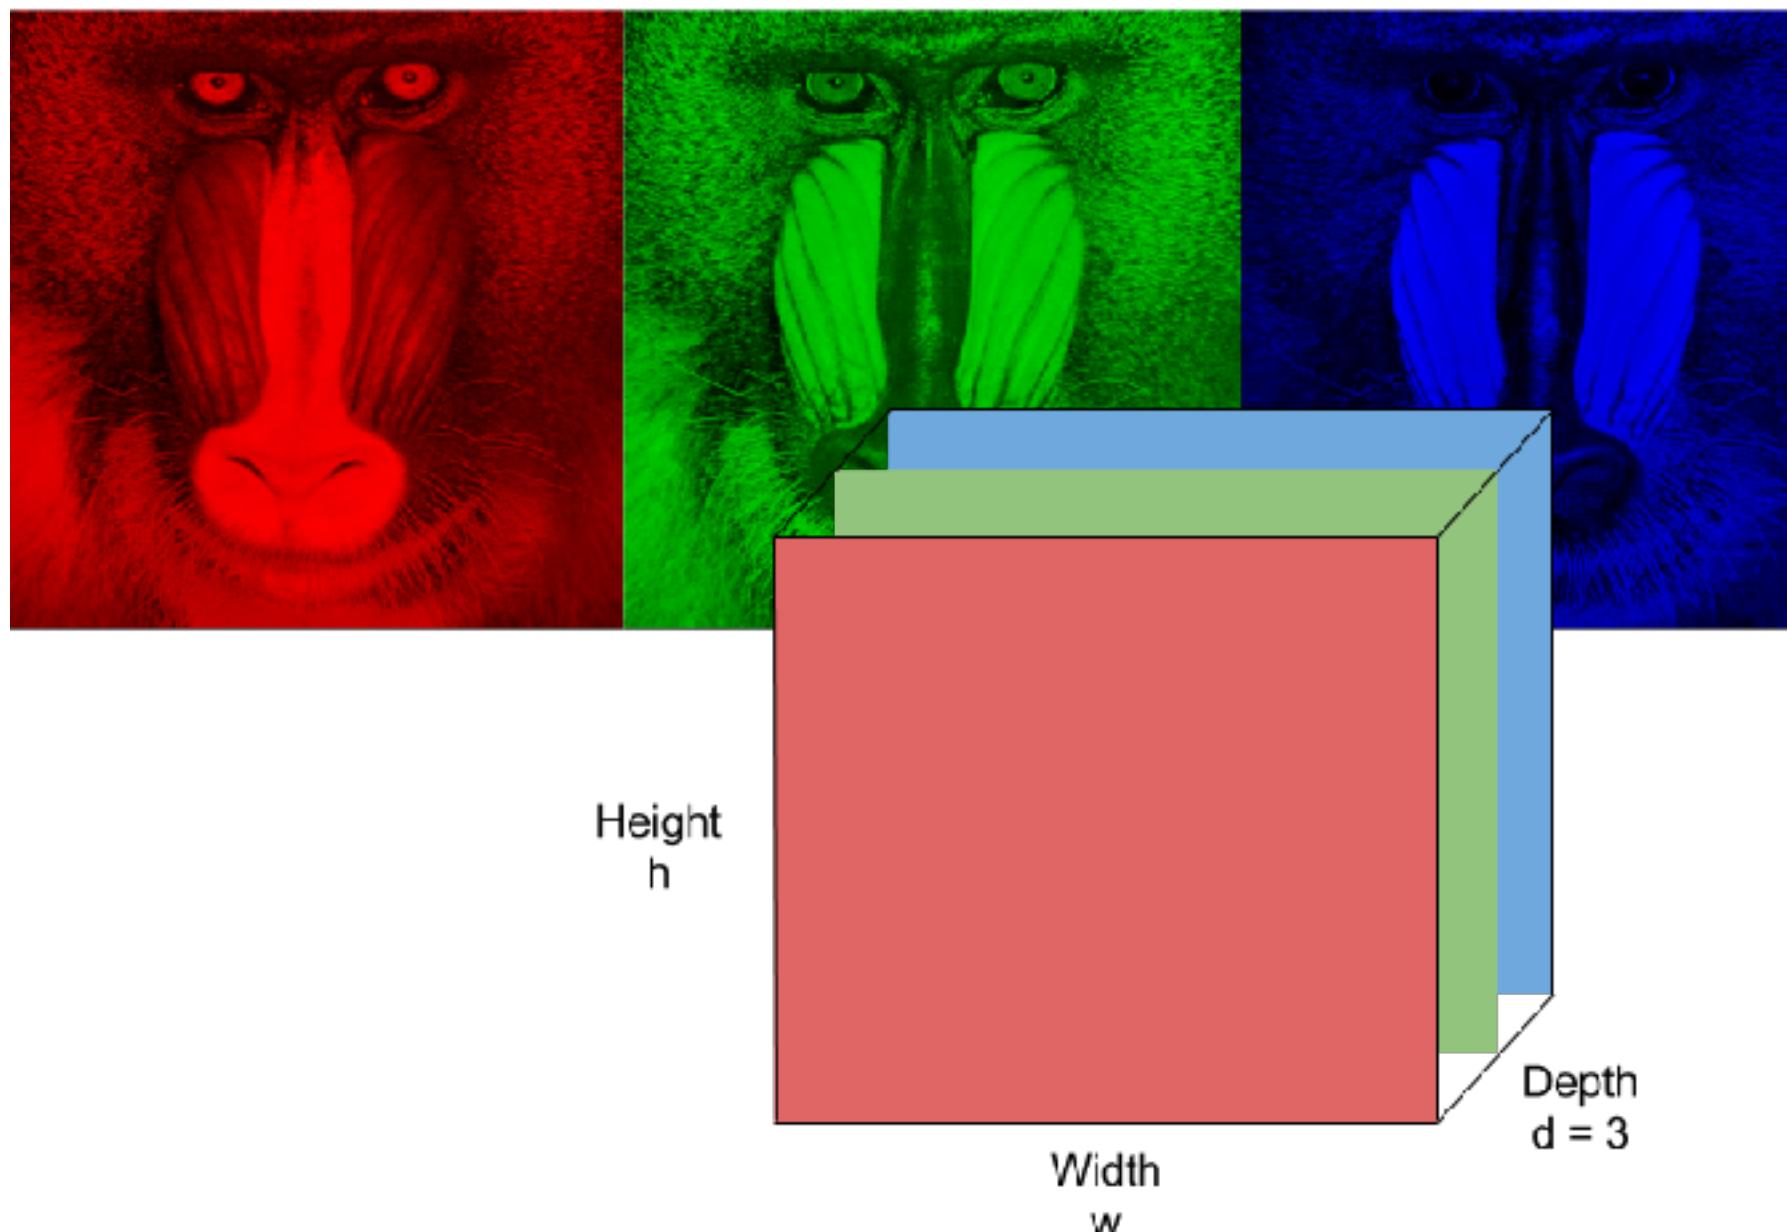
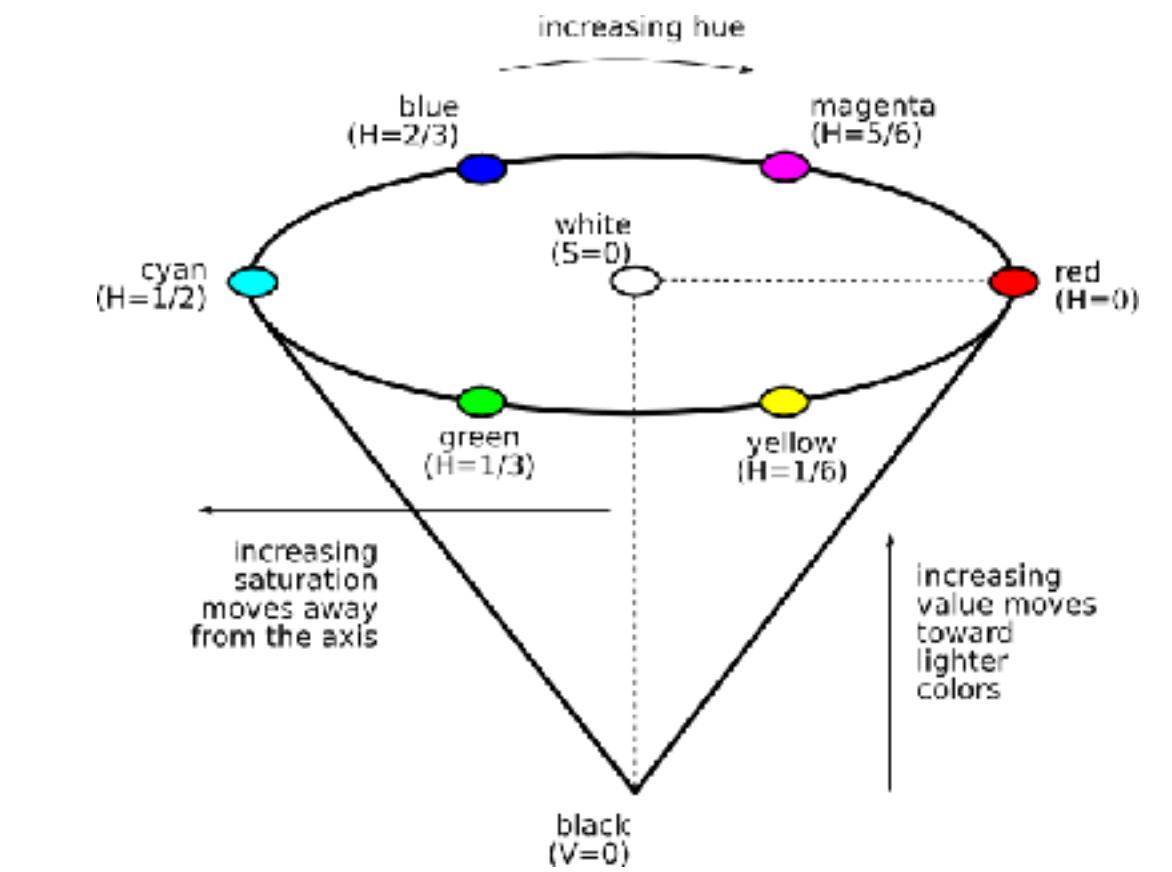
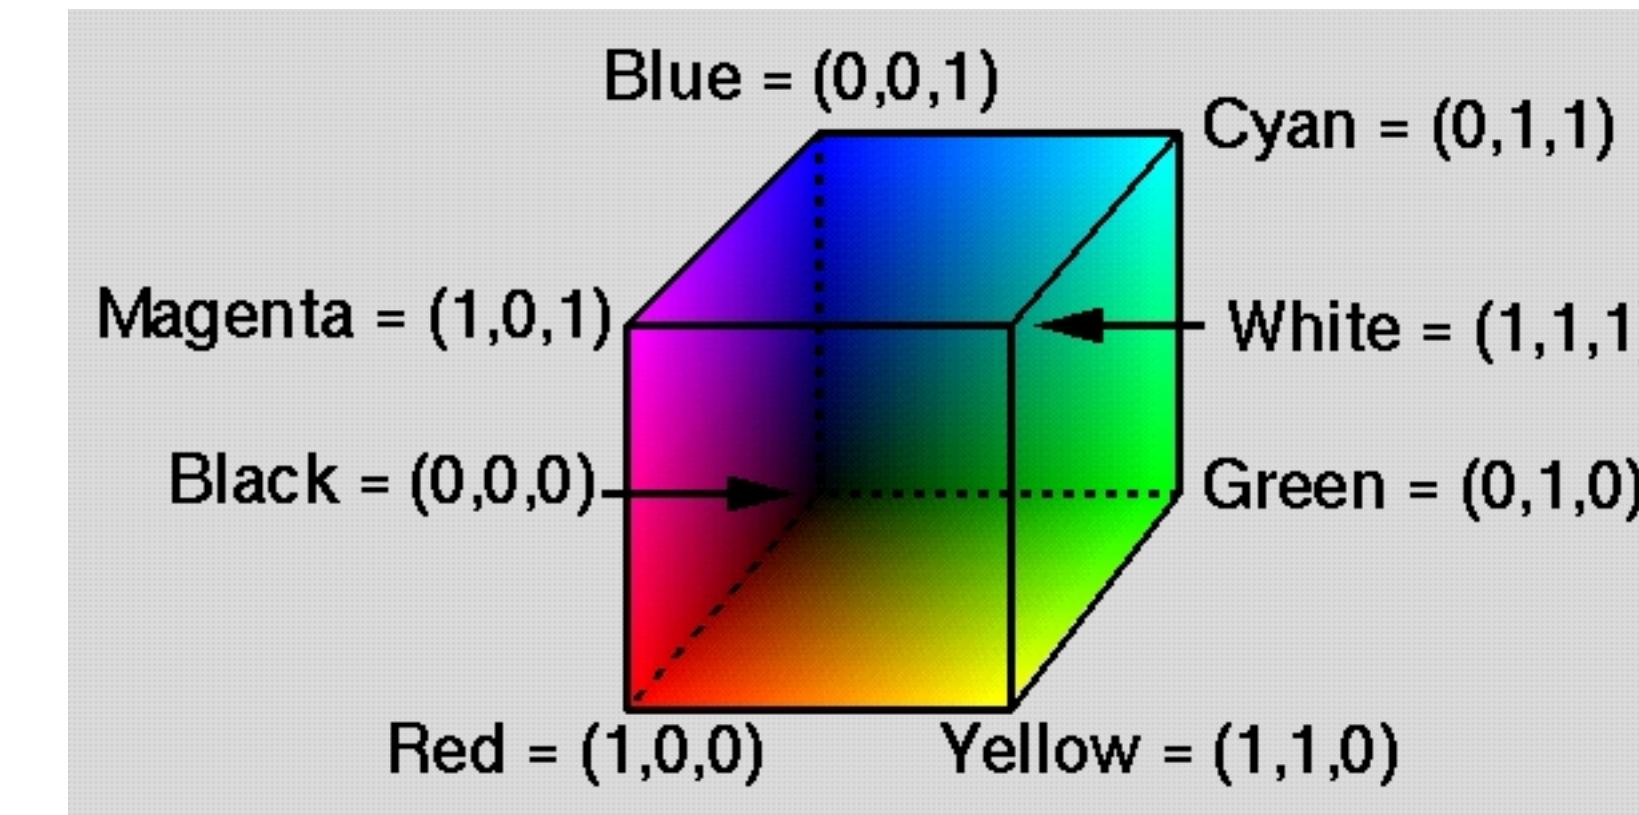
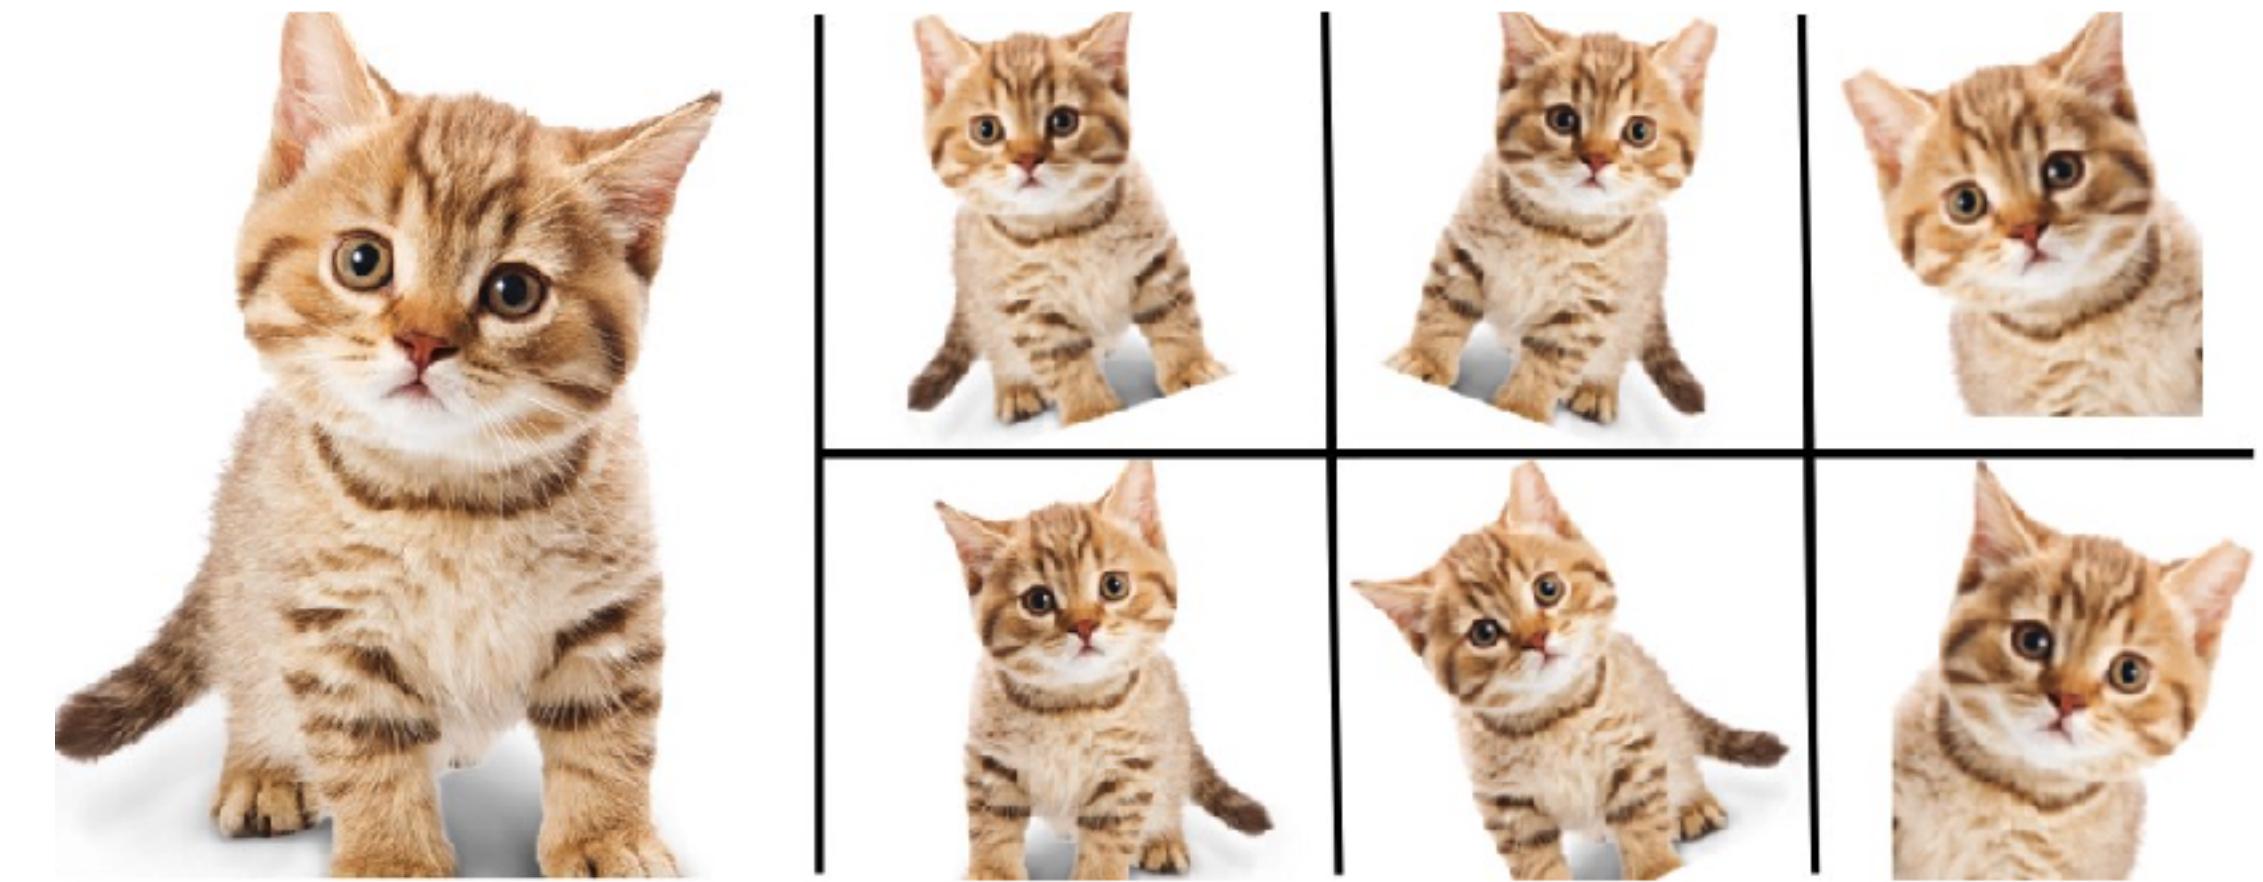
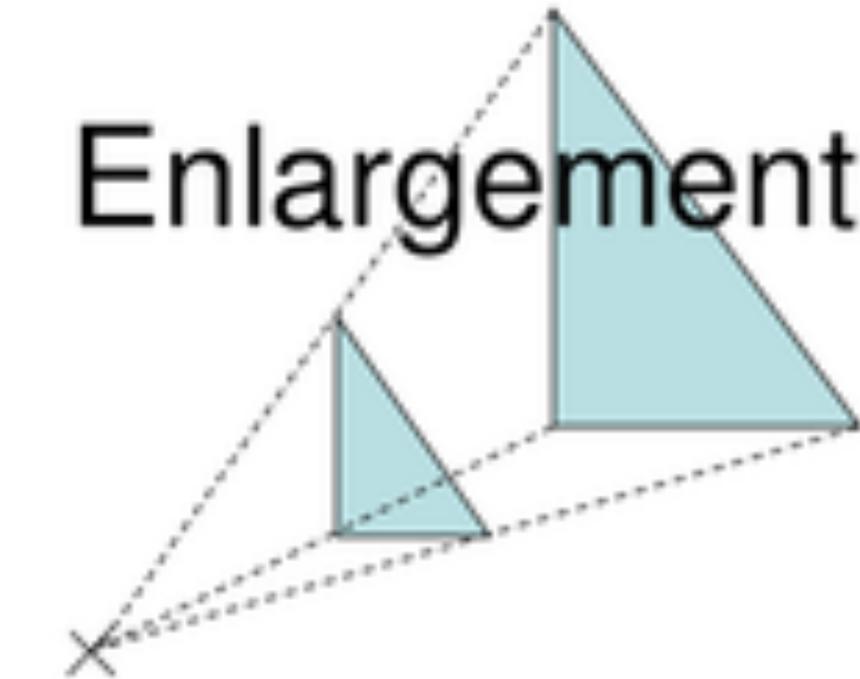
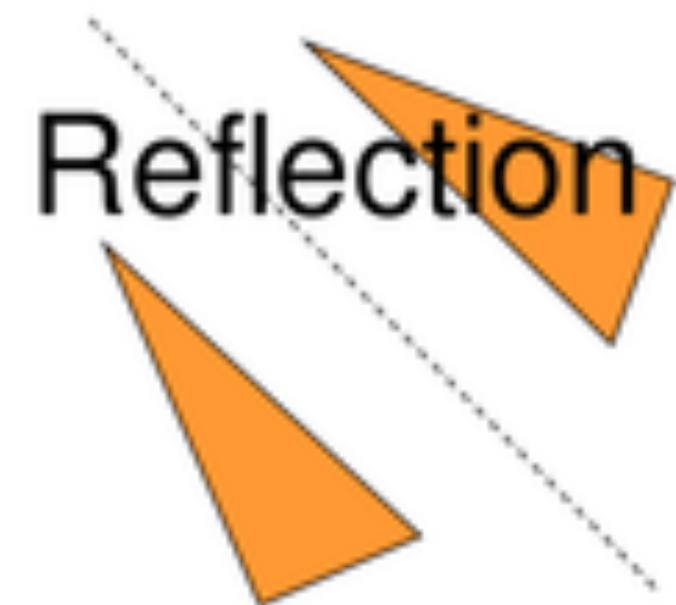
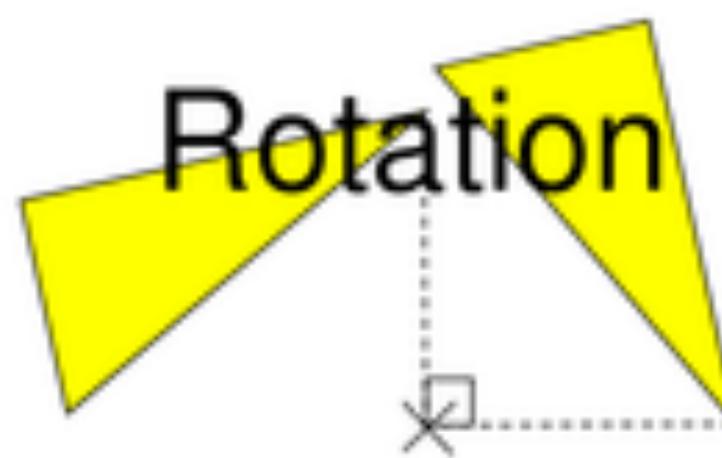
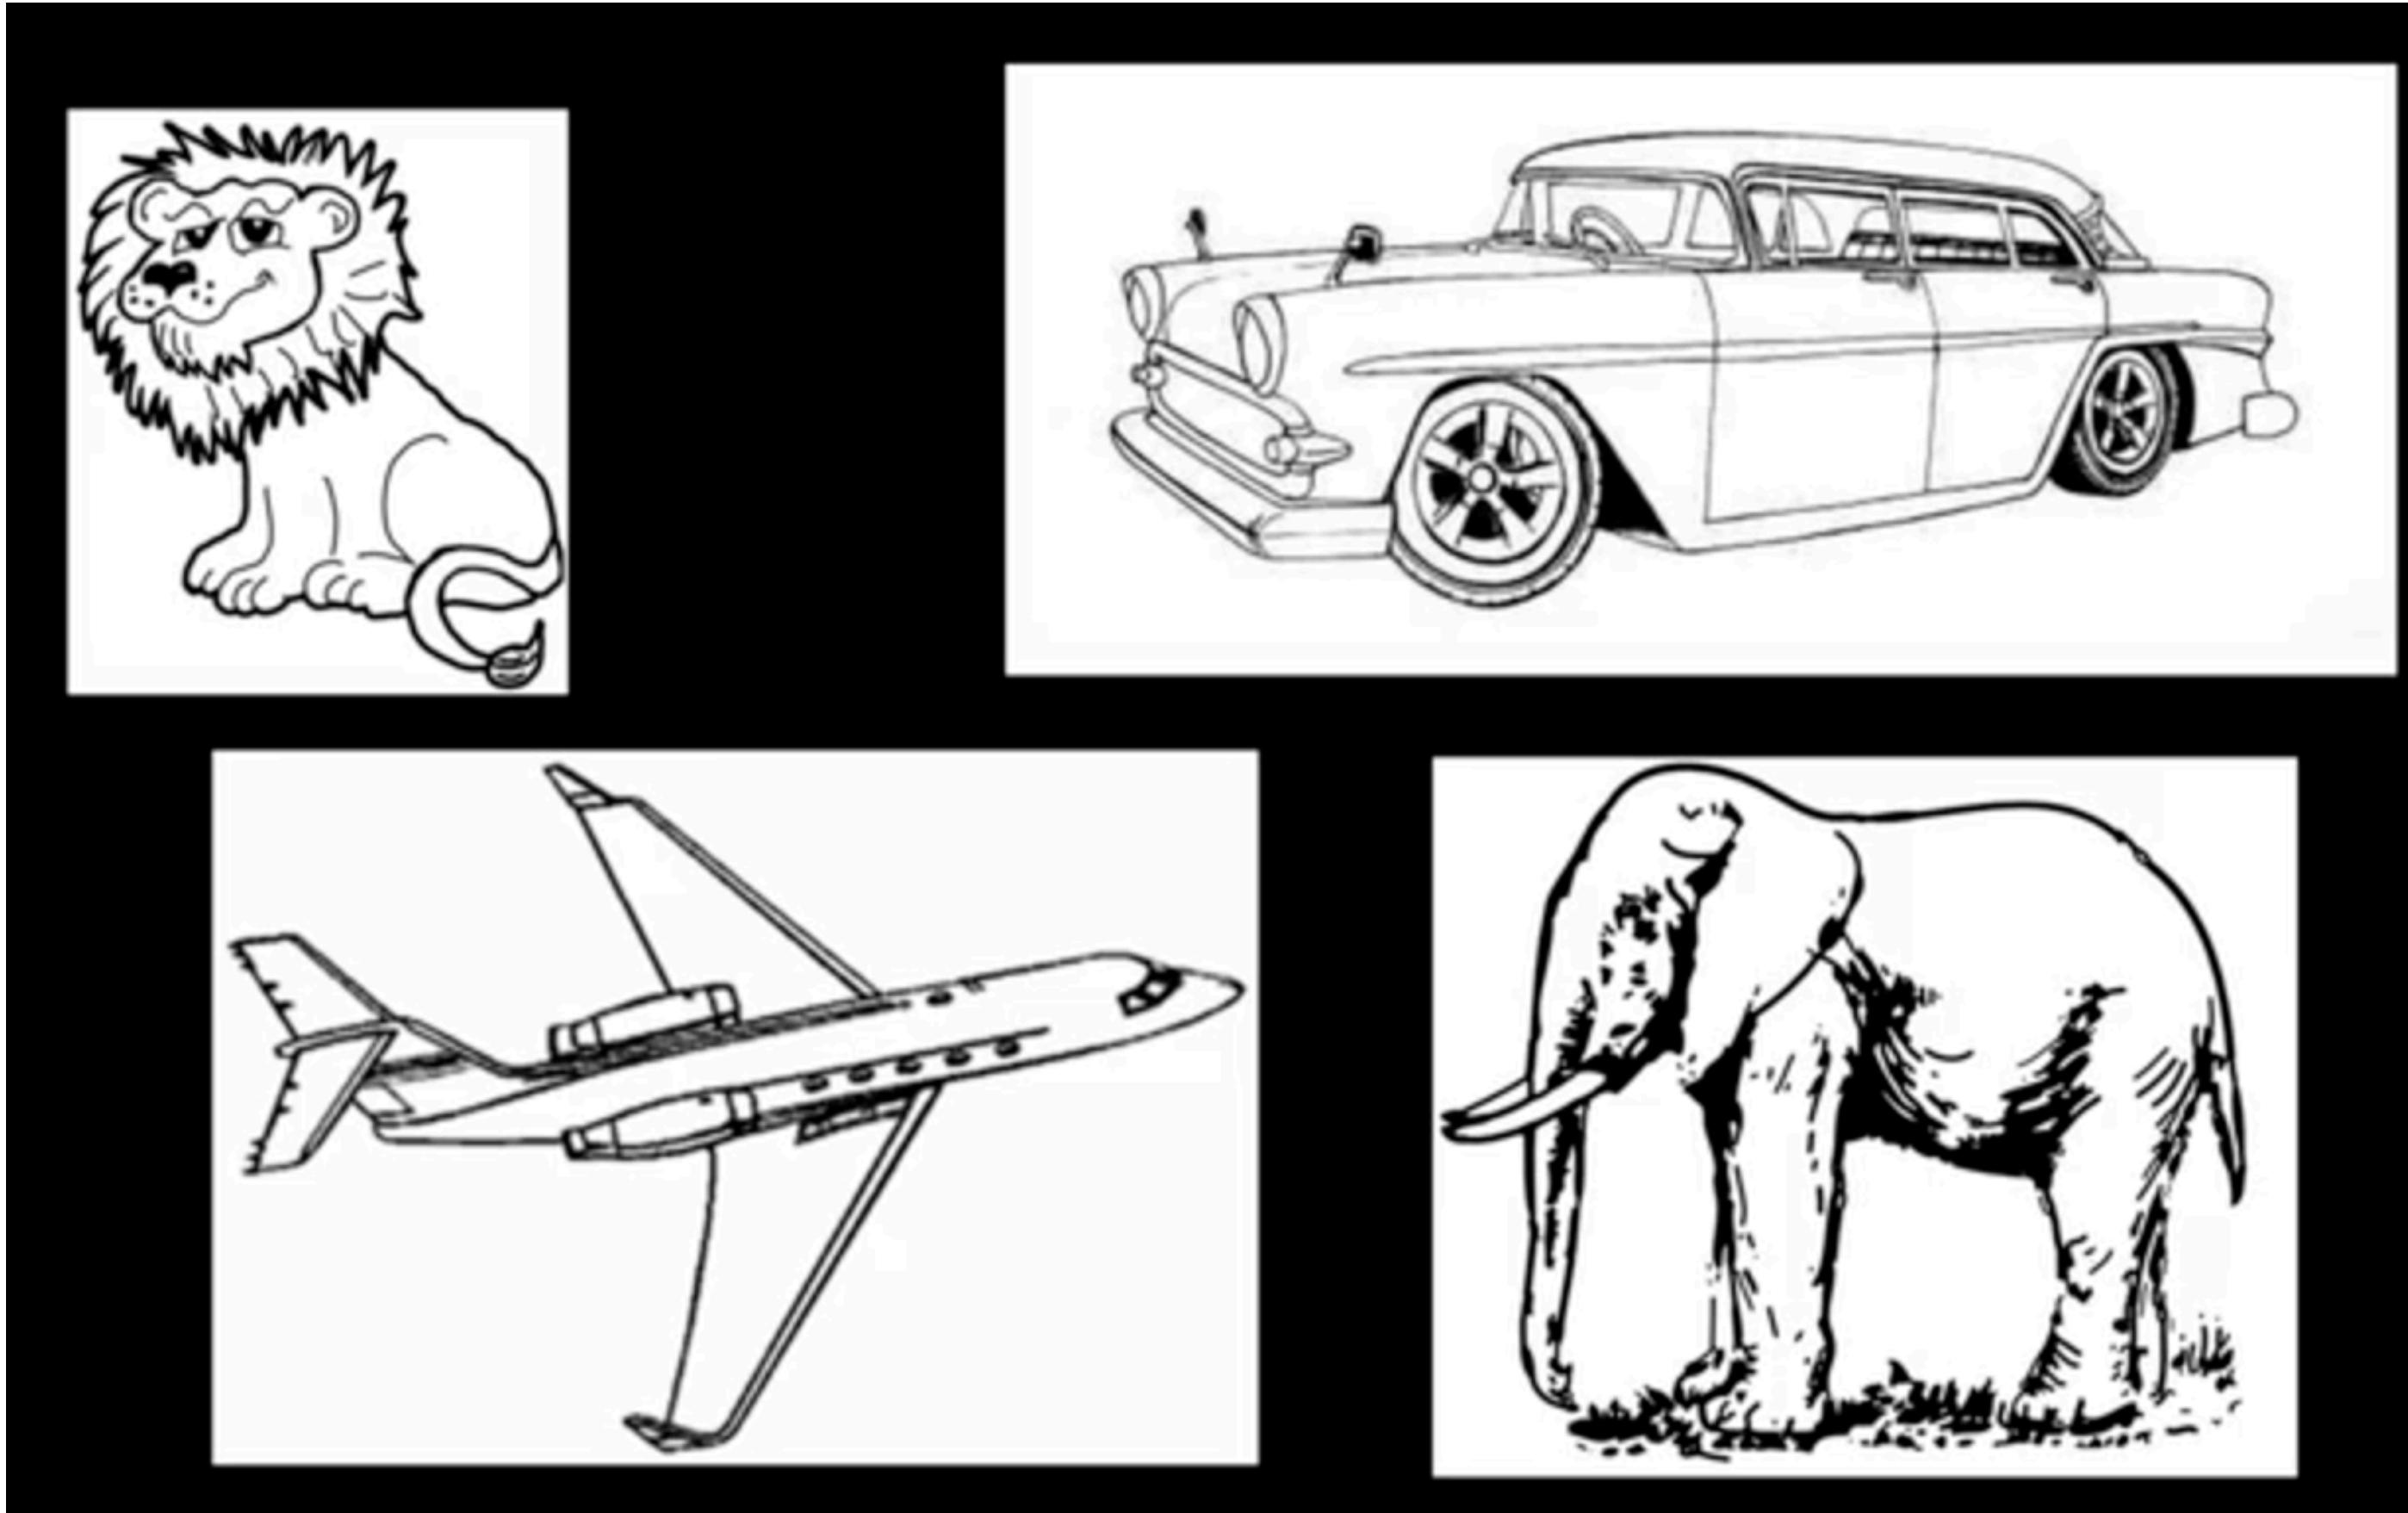


Image Transformations



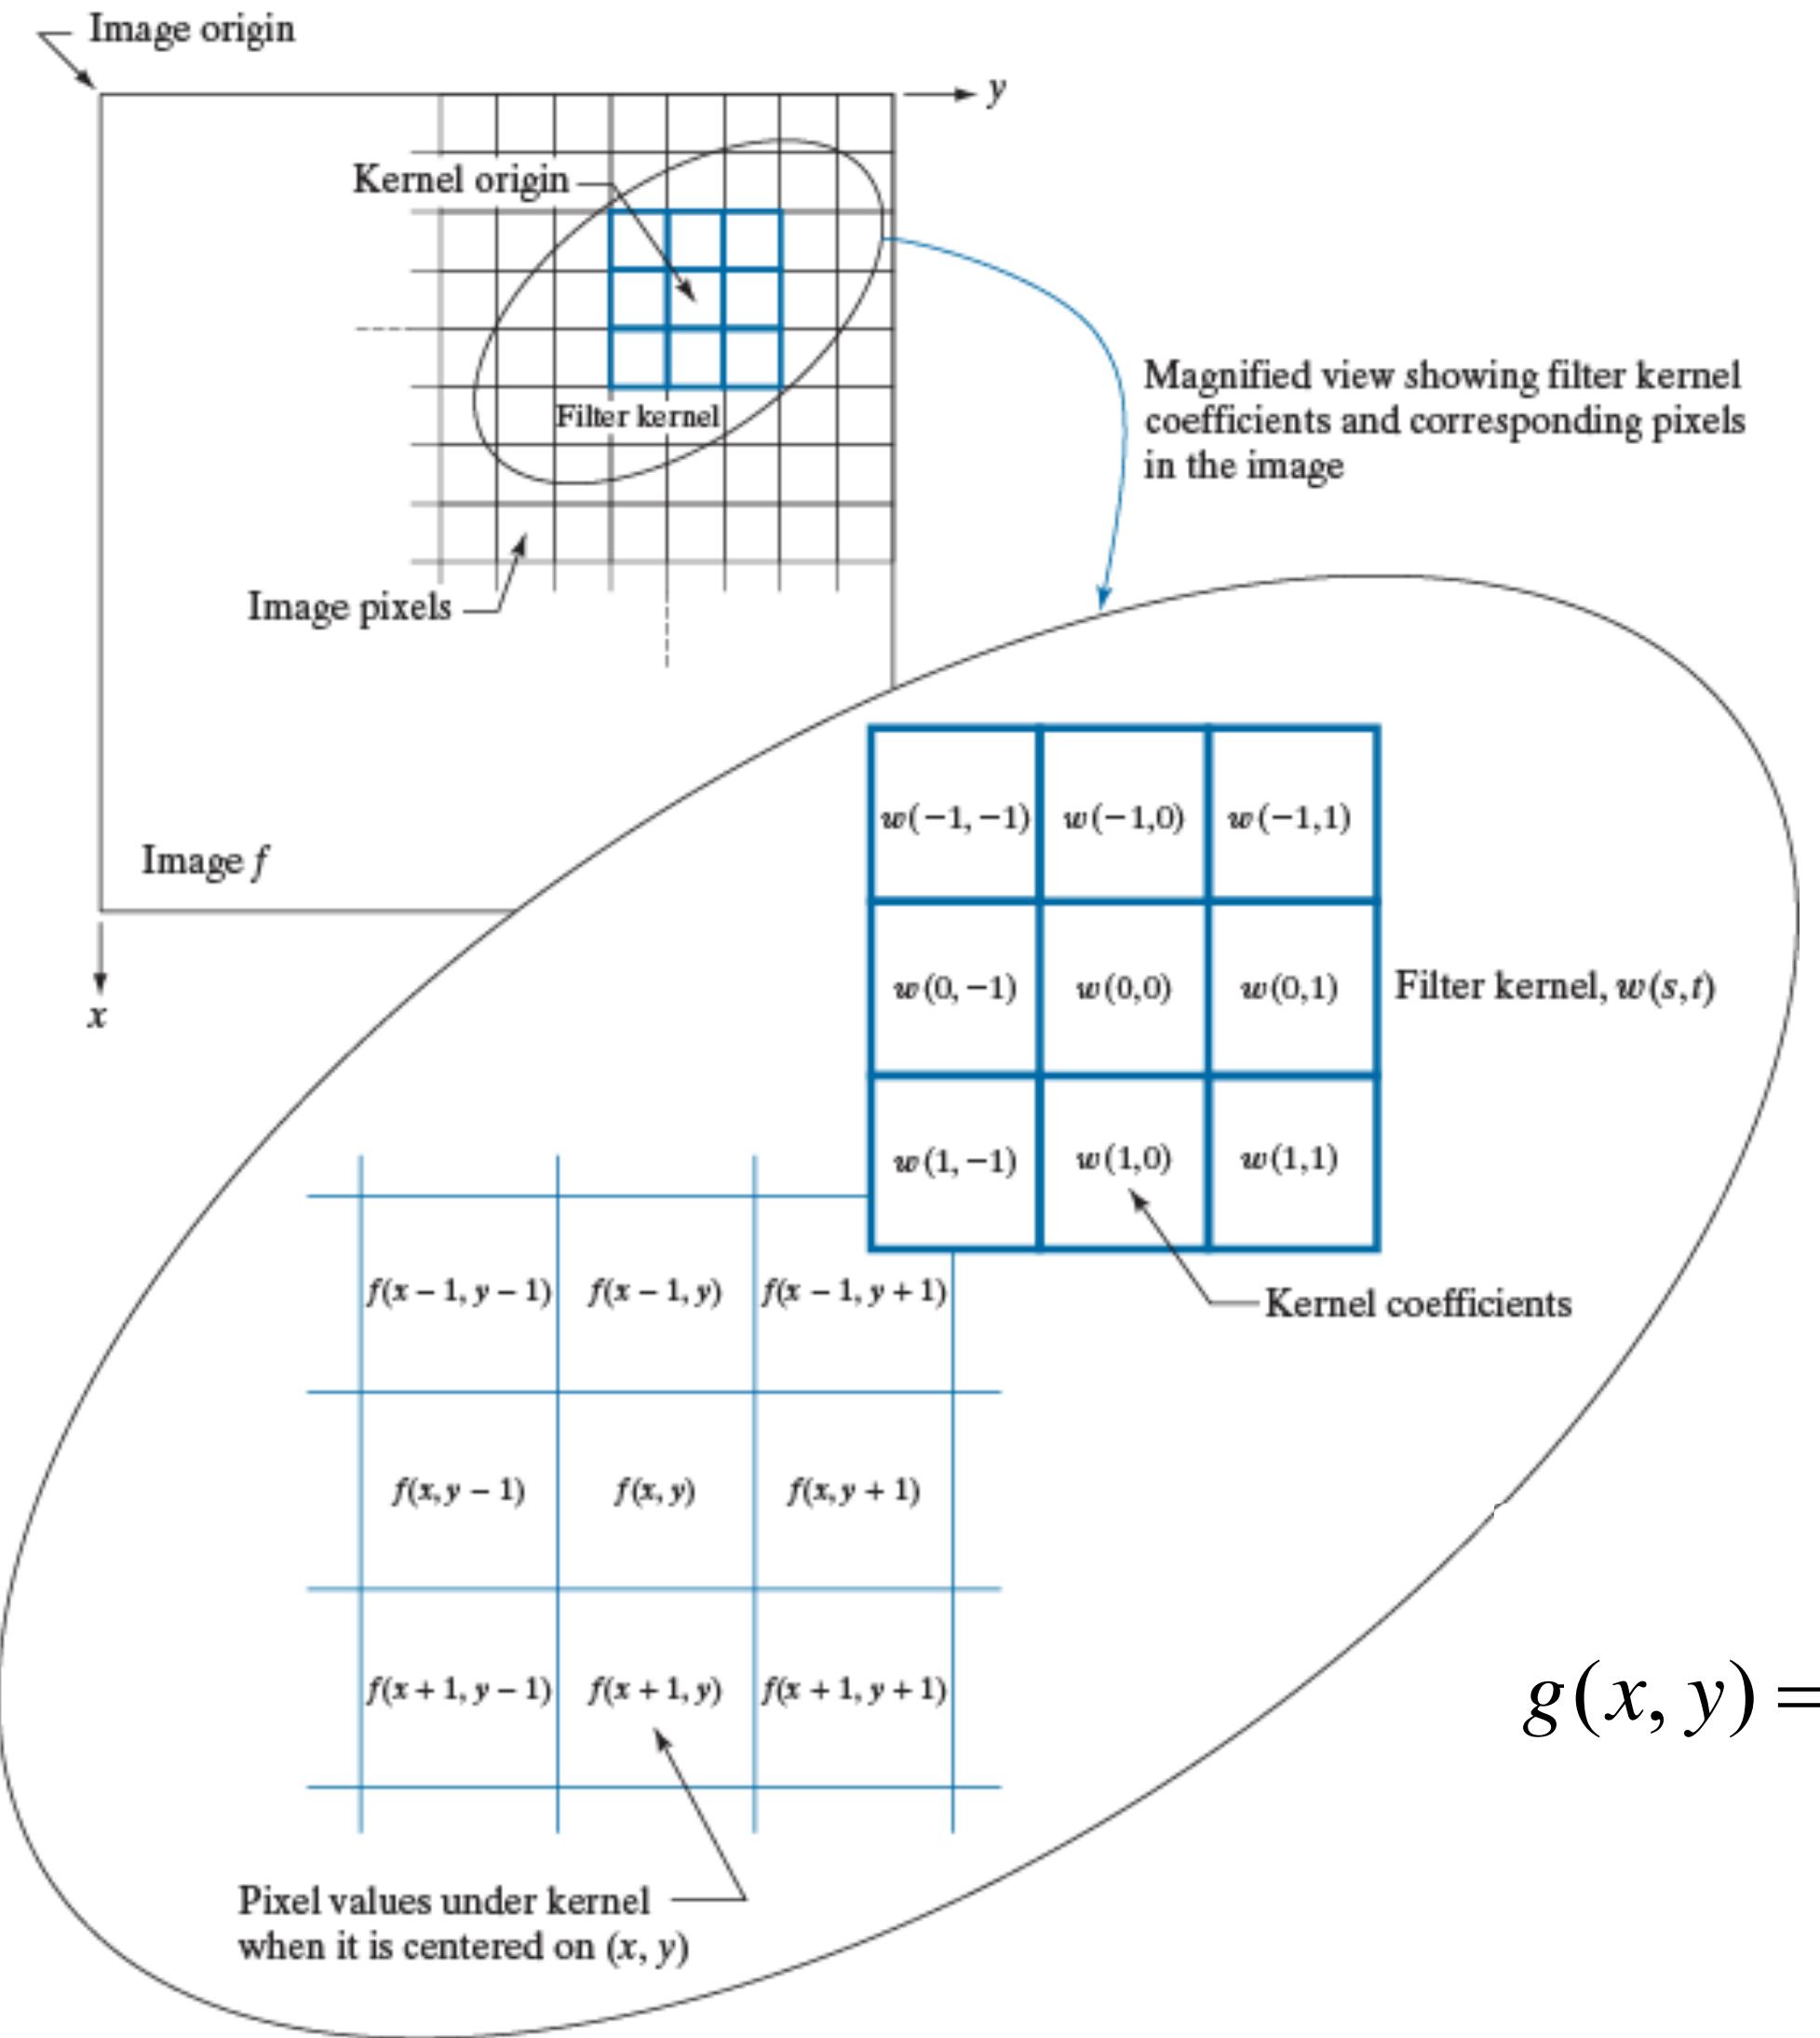
Enlarge your Dataset

Spatial filtering++



«Reduced images» - Edges seem to be important

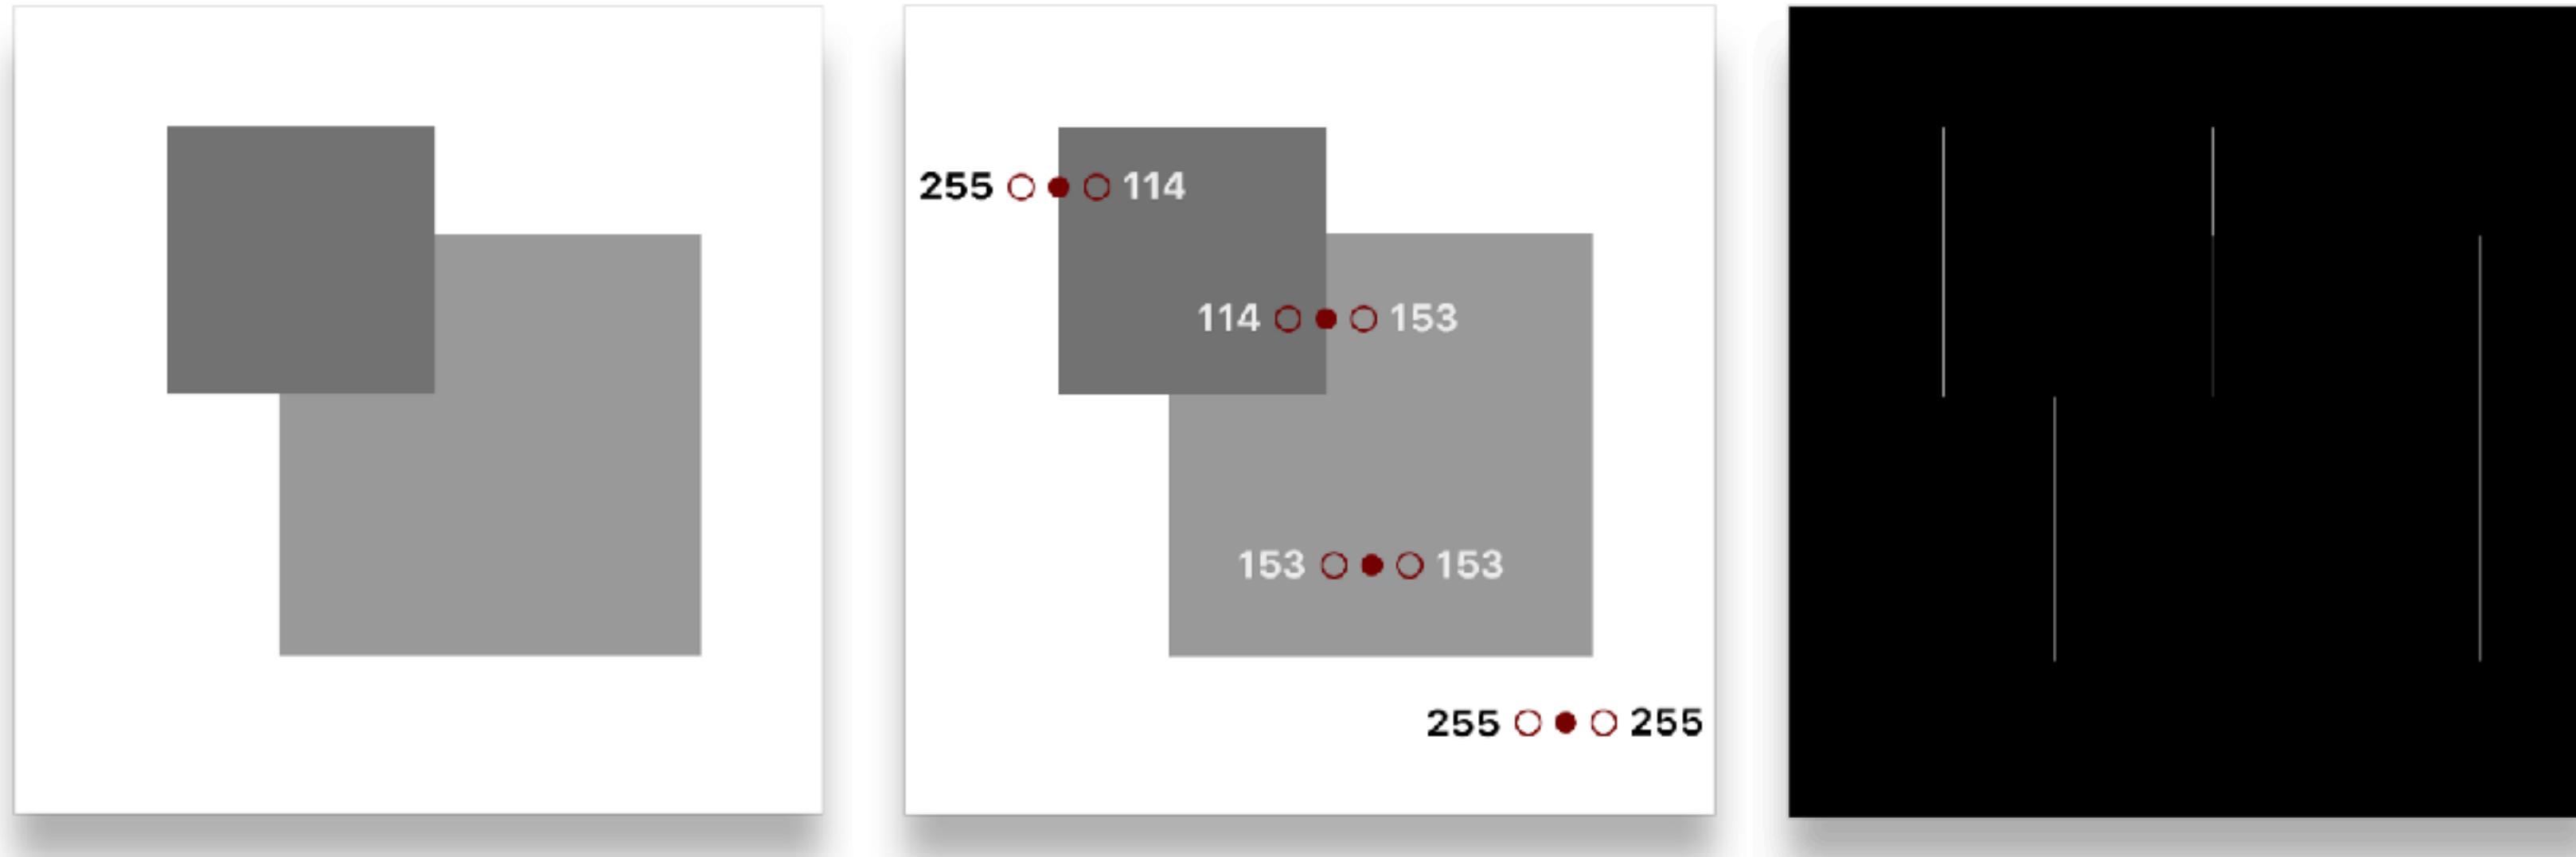
Spatial filtering++



- Neighborhood processing / filtering:
 - Smoothing
 - **Sharpening**

$$g(x, y) = \sum_{s=-m/2}^{m/2} \sum_{t=-n/2}^{n/2} w(s, t) f(x+s, y+t)$$

Sharpening - Finite Differences



```
kernel = [1 0 -1]
```

```
kernel = [  
    [0 0 0]  
    [1 0 -1]  
    [0 0 0]  
]
```

255	255	255	114	114	114	114
x 0	x 0	x 0	114	114	114	114
255	255	255	114	114	114	114
x 1	x 0	x -1	114	114	114	114
255	255	255	114	114	114	114
x 0	x 0	x 0	114	114	114	114
255	255	255	114	114	114	114
255	255	255	255	255	255	255
255	255	255	255	255	255	255
255	255	255	255	255	255	255

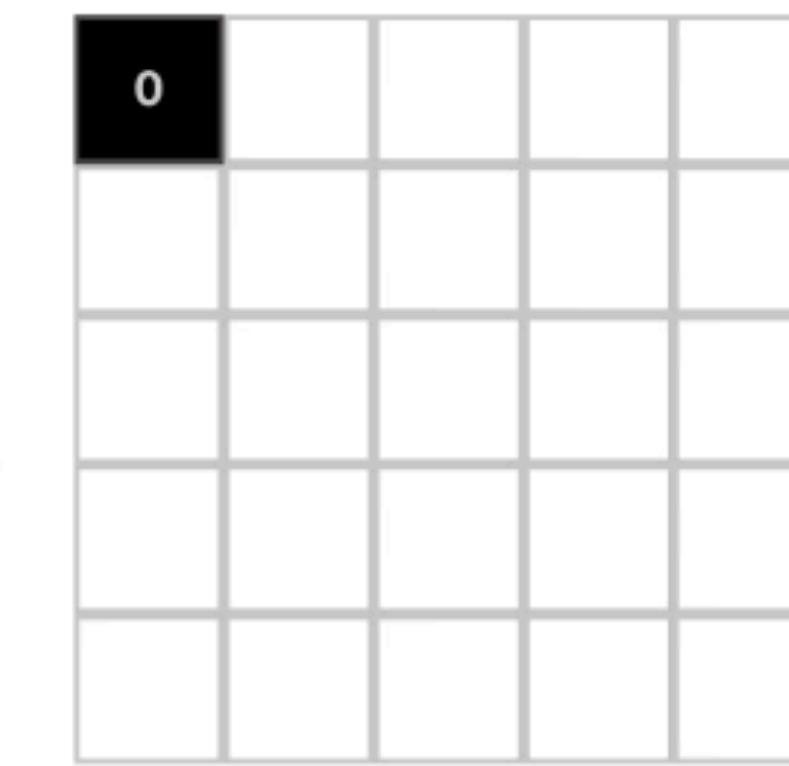


Image gradient

The gradient of an image:

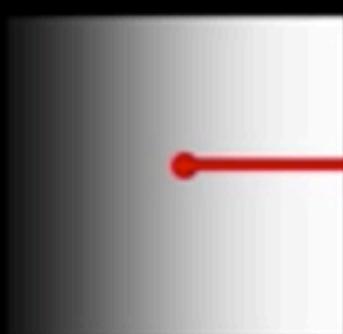
$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

The gradient direction is given by:

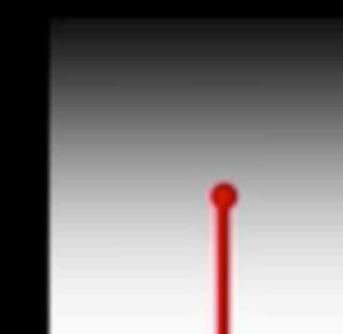
$$\theta = \tan^{-1}\left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x}\right)$$

The *amount of change* is given by
the gradient magnitude:

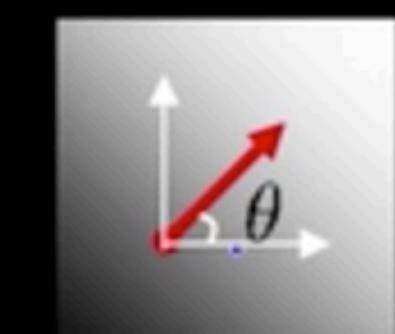
$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$



$$\nabla f = \left[\frac{\partial f}{\partial x}, 0 \right]$$

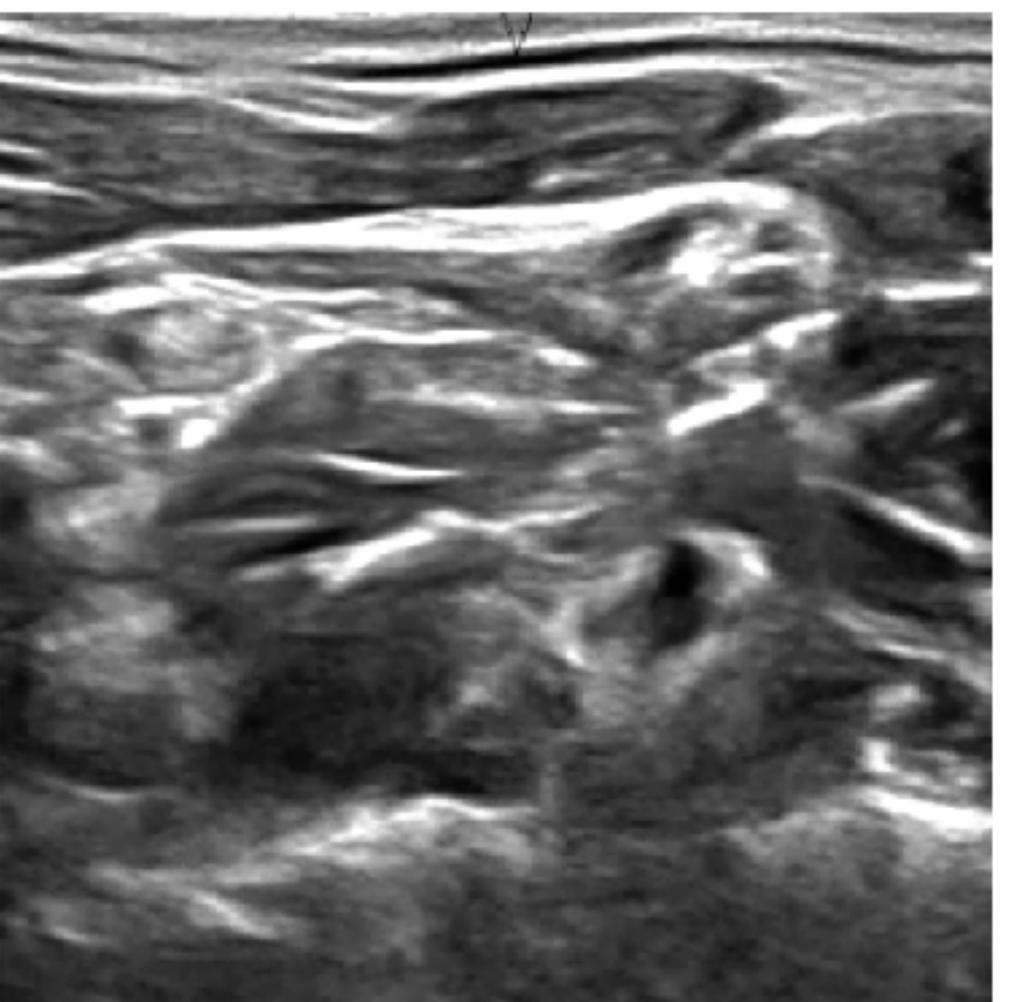
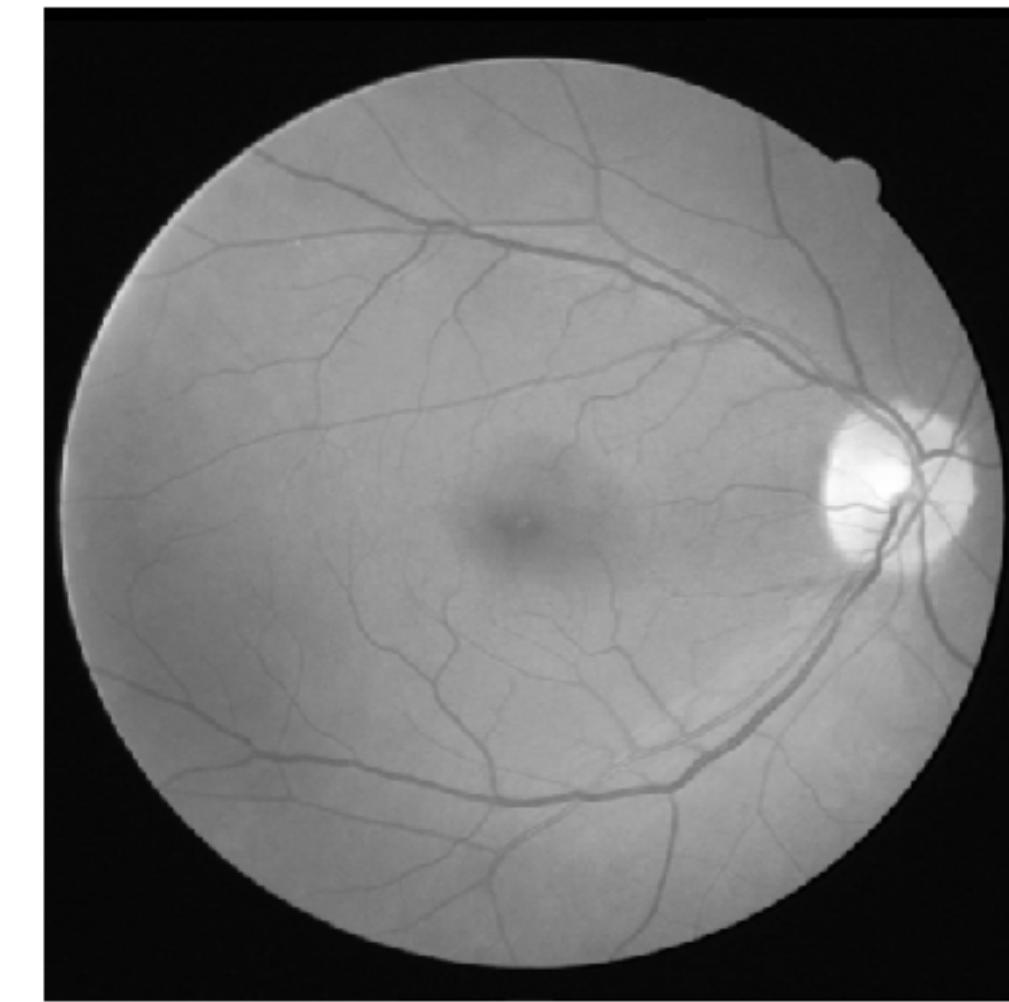
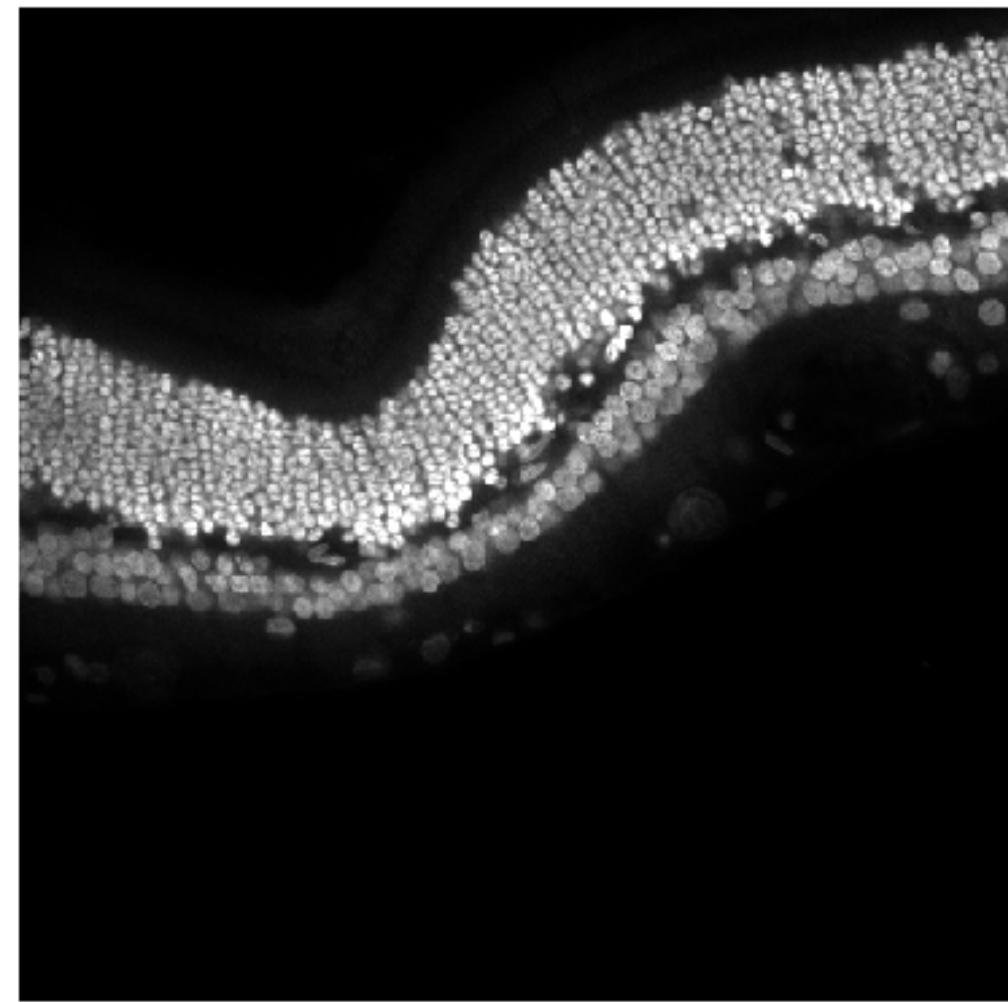
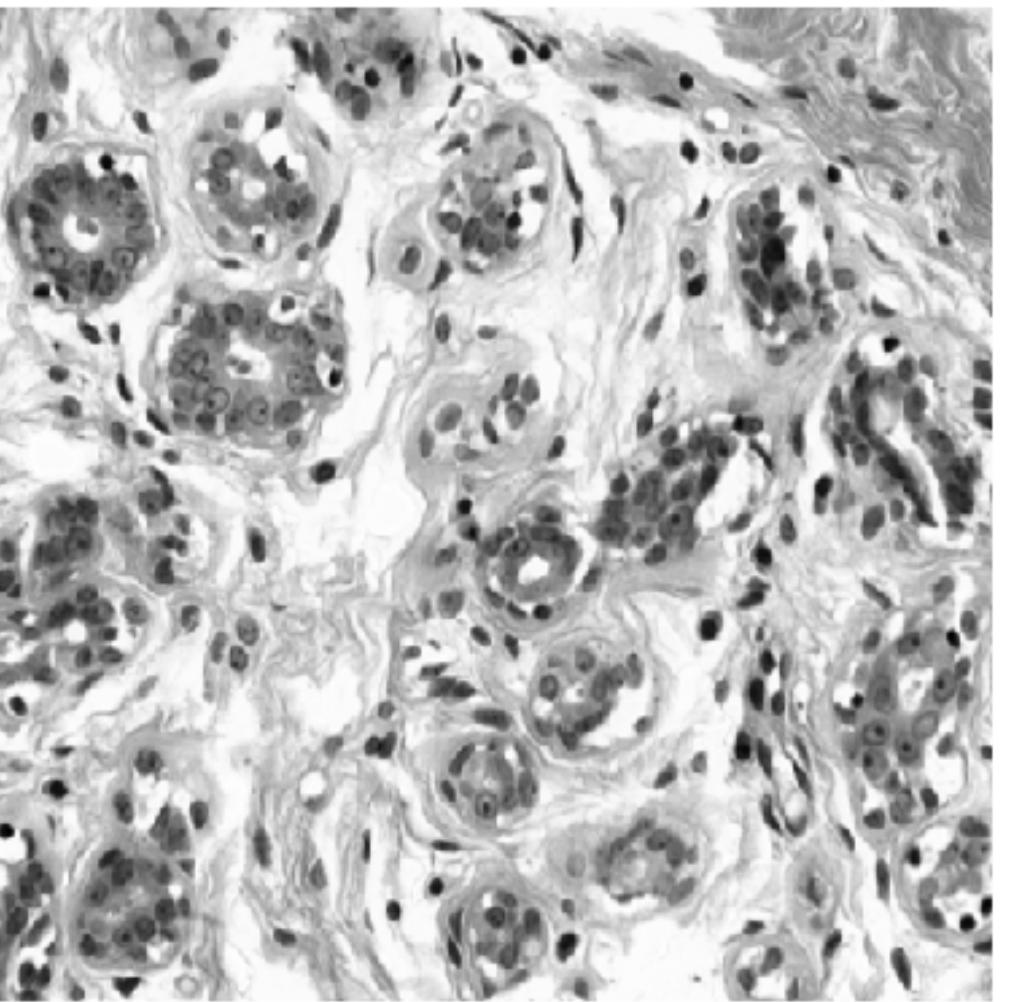


$$\nabla f = \left[0, \frac{\partial f}{\partial y} \right]$$

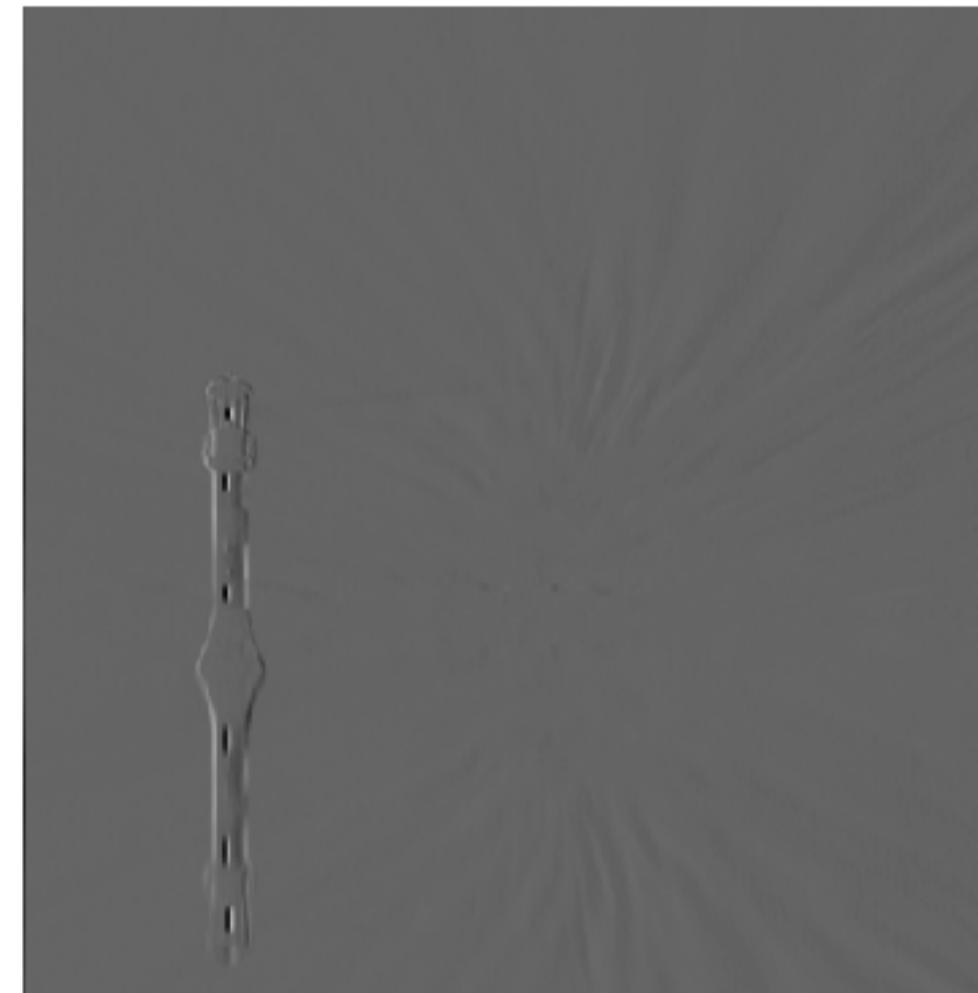
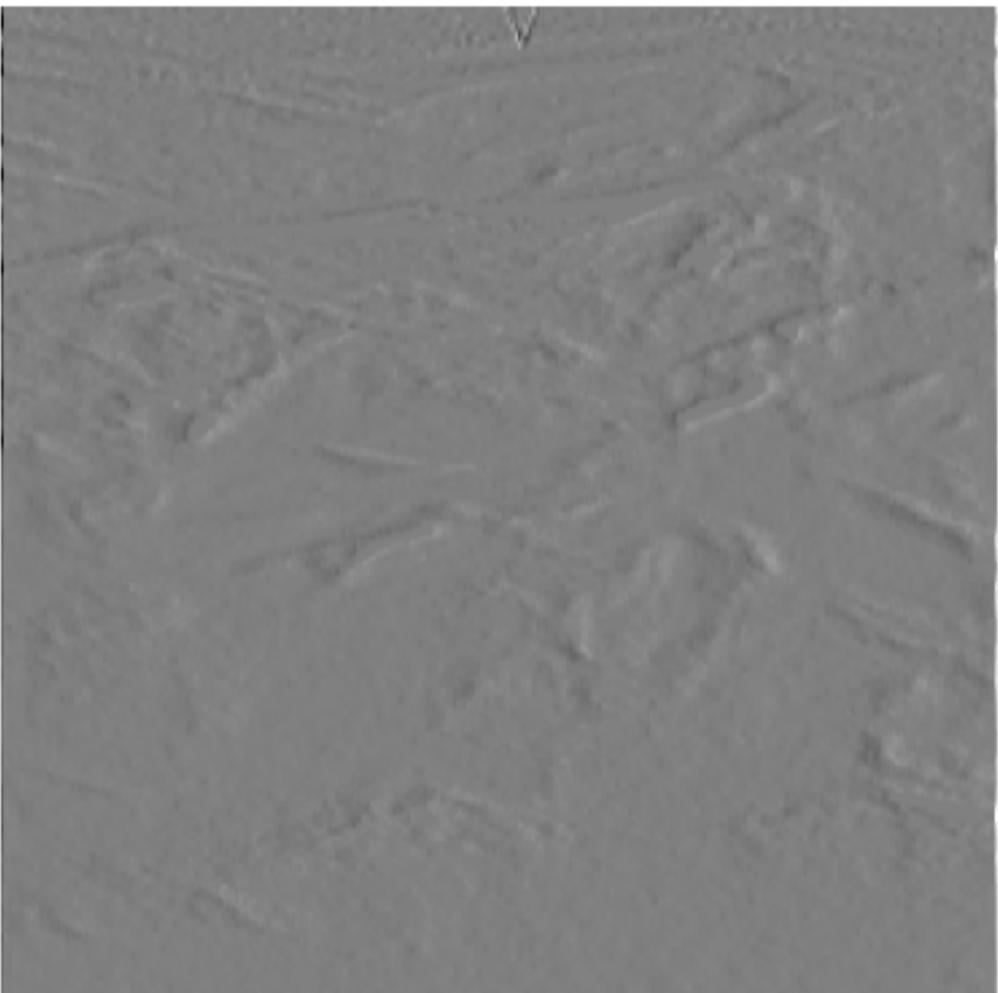
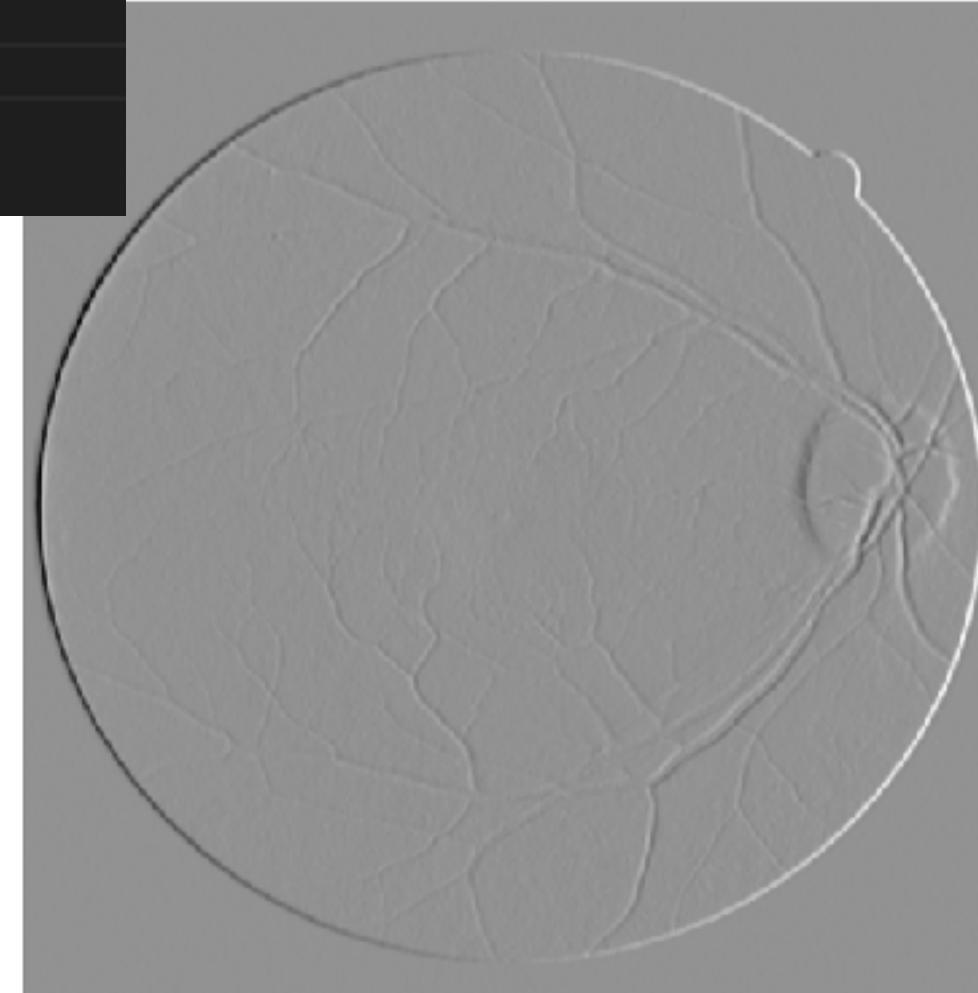
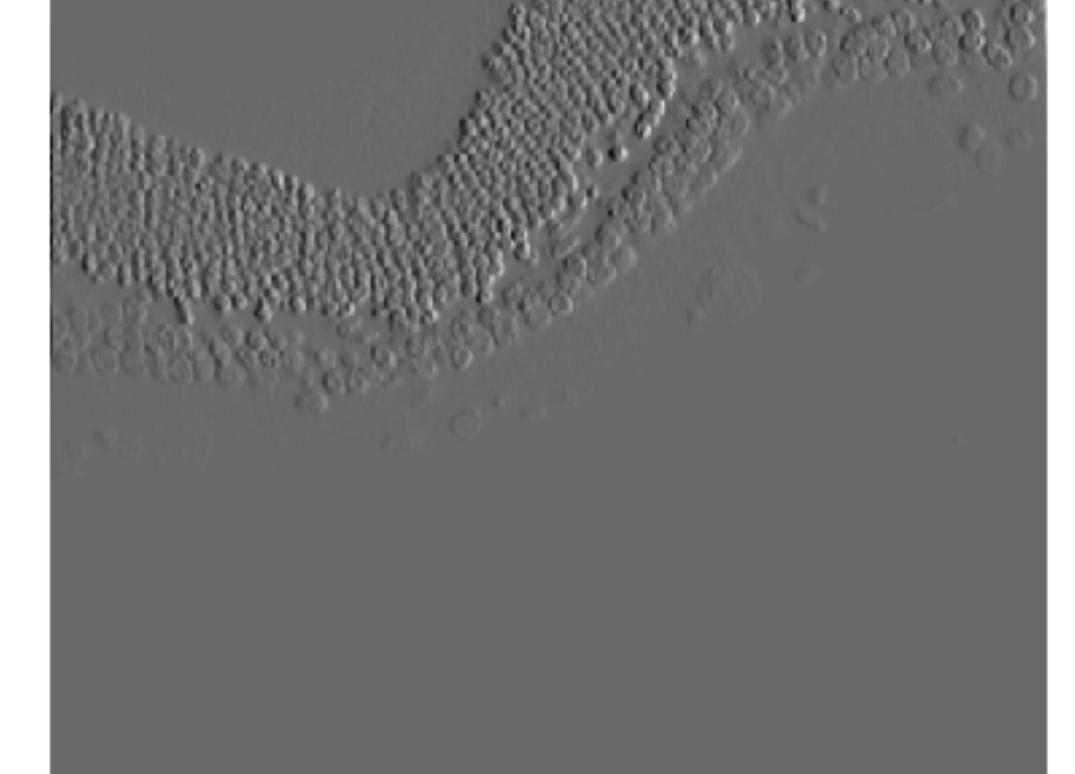
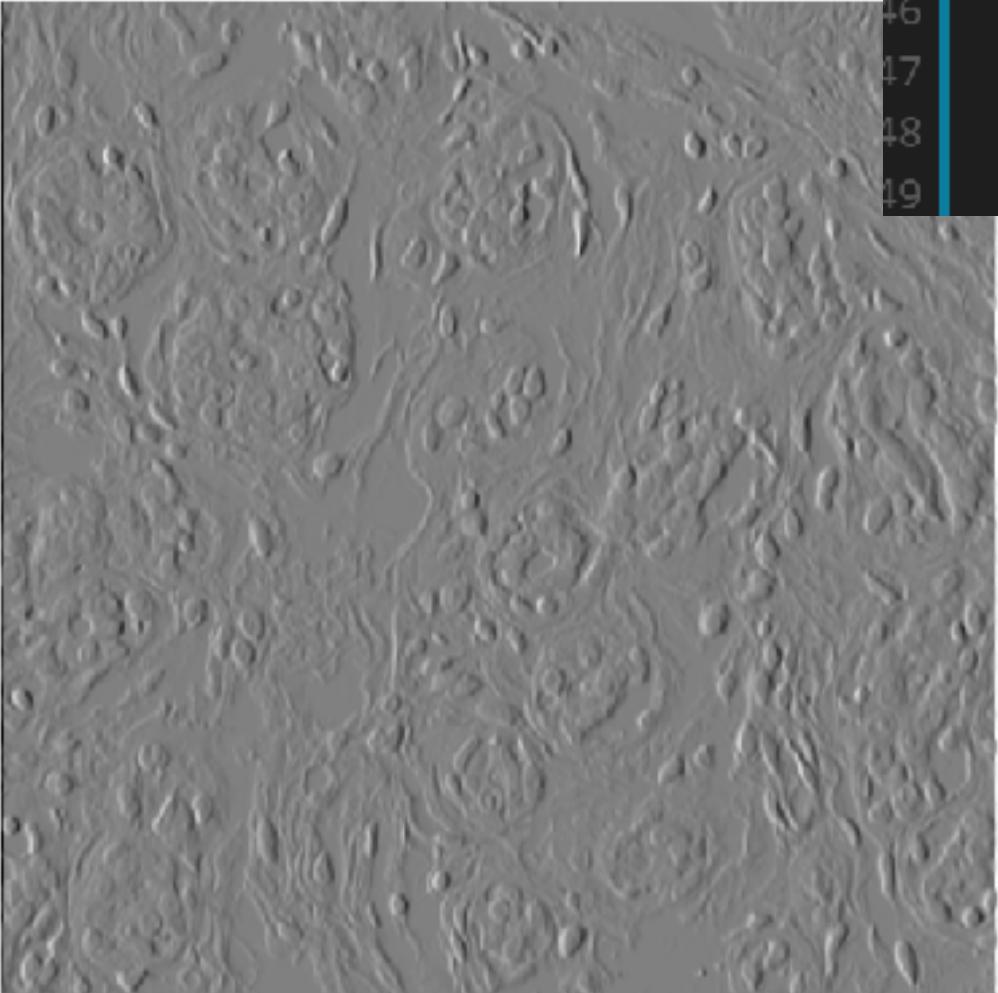


$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

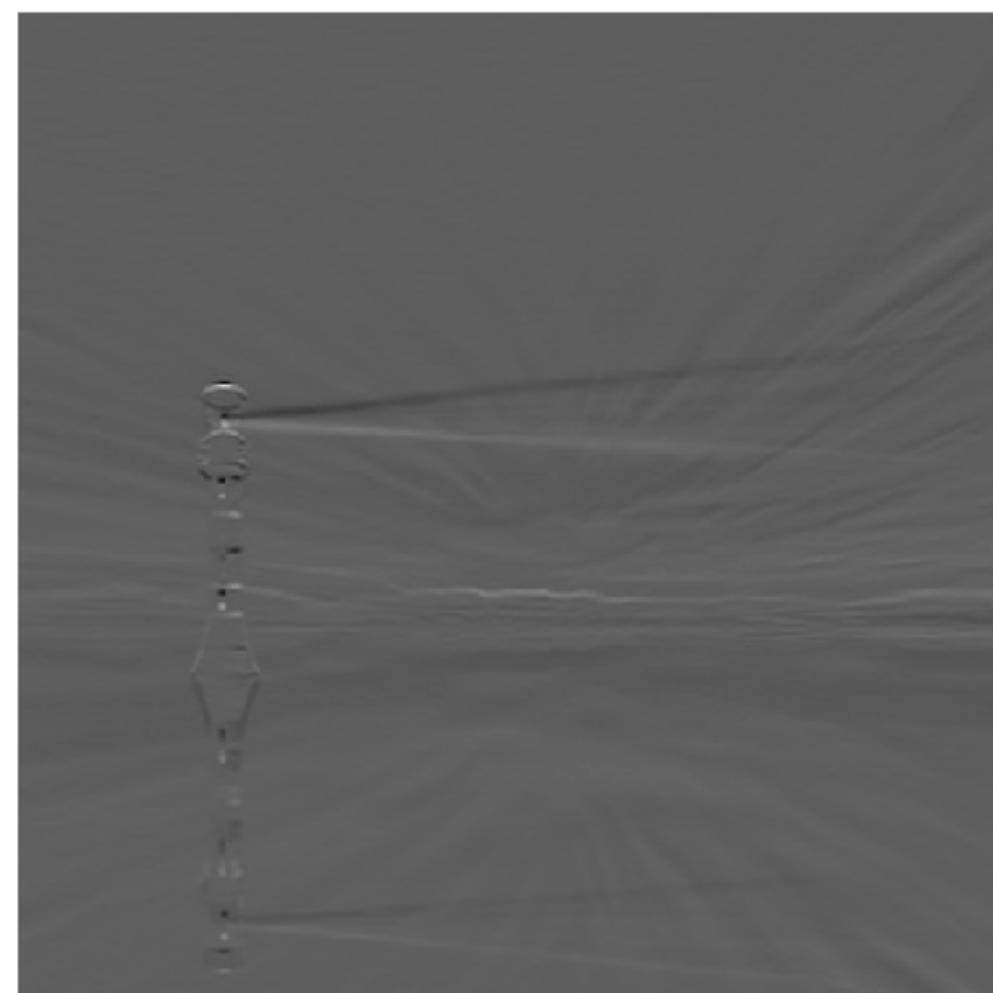
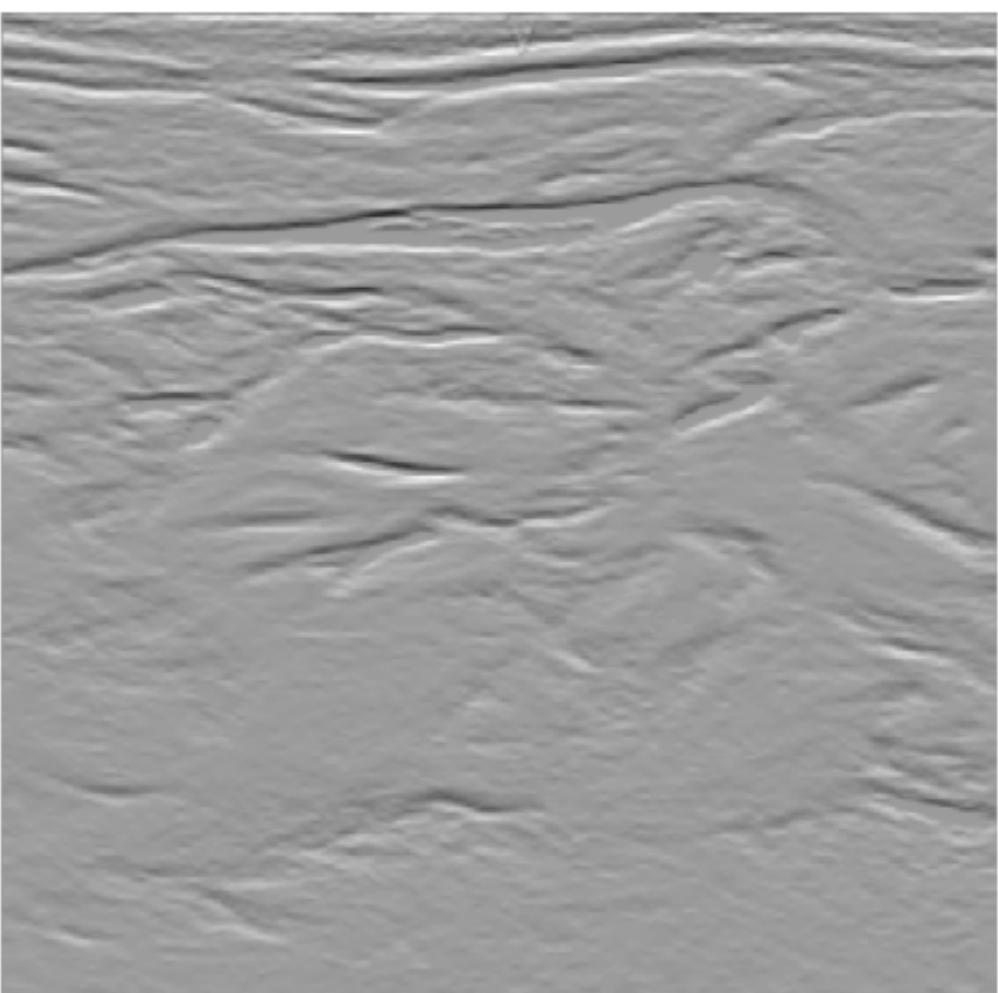
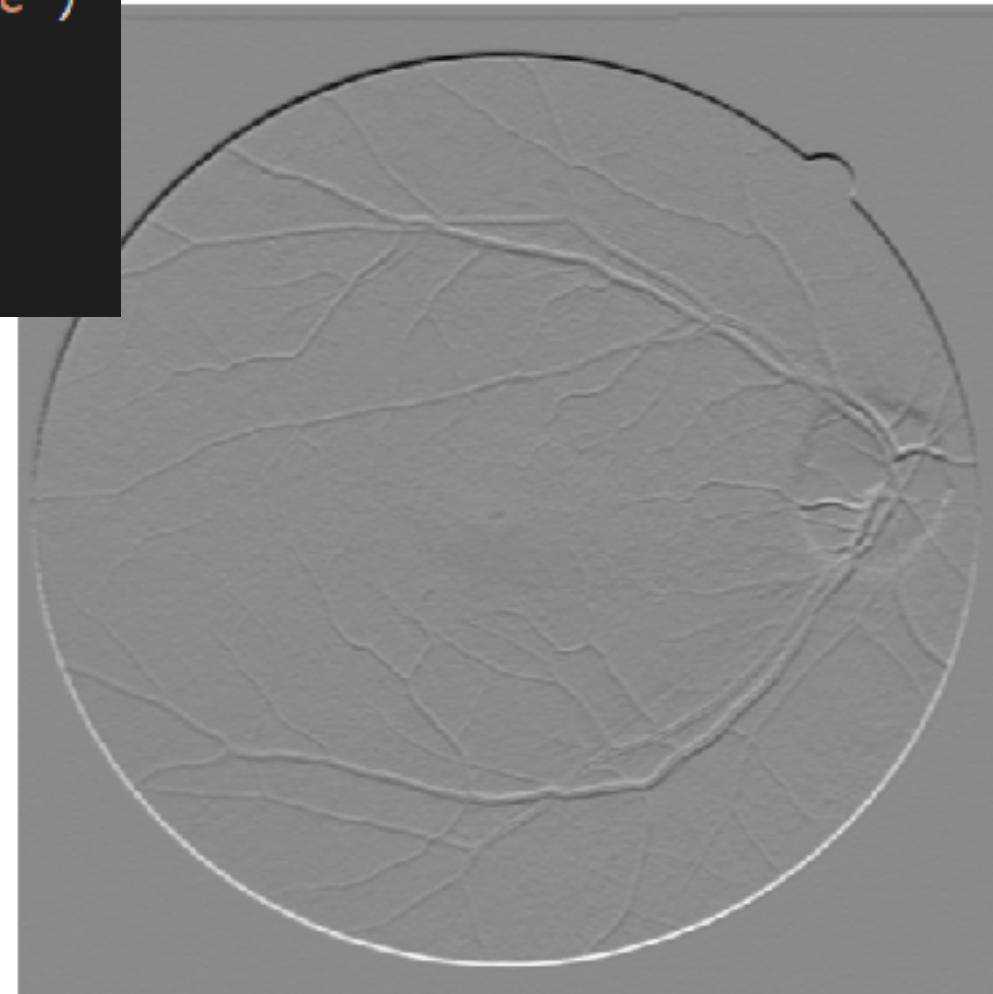
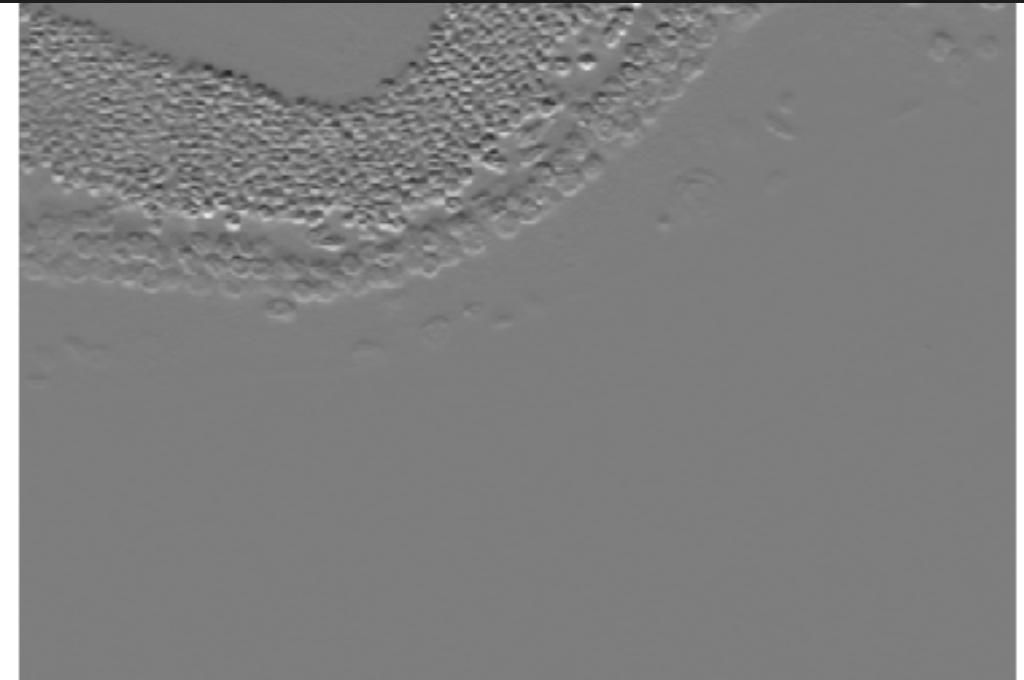
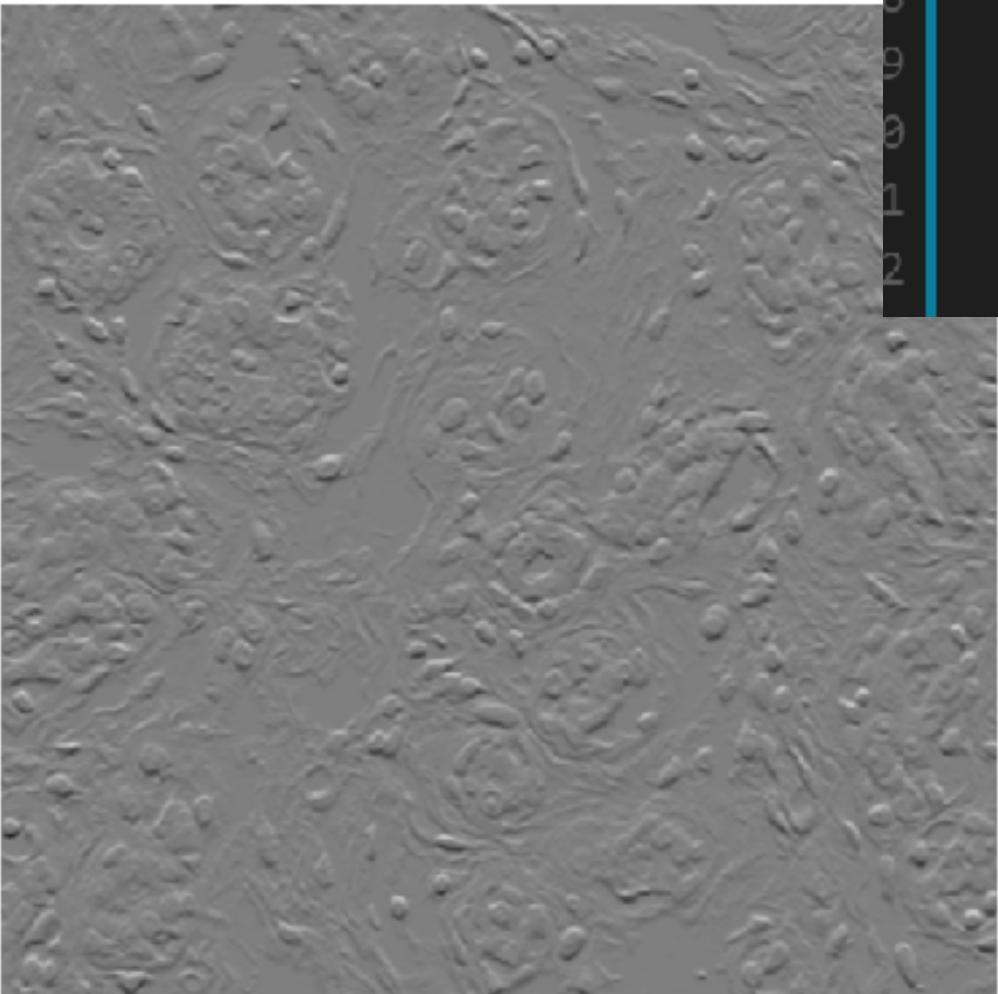
Sharpening...



```
36 # 2. Edge Detection (Horizontal)
37 kernel = np.array([
38     [-1,0,1],
39     [-1,0,1],
40     [-1,0,1]
41 ])
42 ]
43
44 for x in range(len(one)):
45     temp = convolve2d(one[x,:,:],kernel,mode='same')
46     plt.axis('off')
47     plt.imshow(temp,cmap='gray')
48     plt.savefig(str(x)+'.png',bbox_inches='tight')
```



```
# 3. Edge Detection (Vertical)
1 kernel = np.array([
2     [-1,-1,-1],
3     [0,0,0],
4     [1,1,1]
5 ])
6
7 for x in range(len(one)):
8     temp = convolve2d(one[:, :, :], kernel, mode='same')
9     plt.axis('off')
10    plt.imshow(temp, cmap='gray')
11    plt.savefig(str(x)+'.png', bbox_inches='tight')
```



$$\mathbf{G} = \sqrt{\mathbf{G}_x^2 + \mathbf{G}_y^2}$$

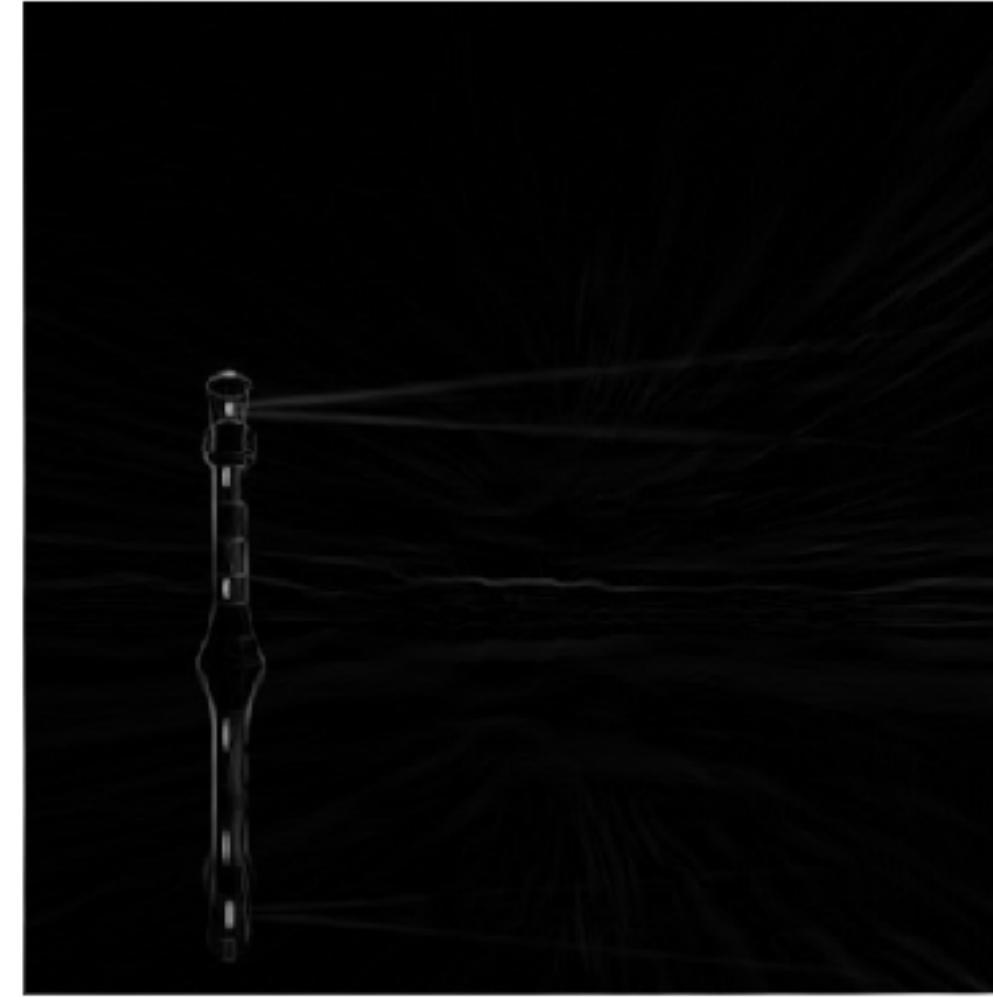
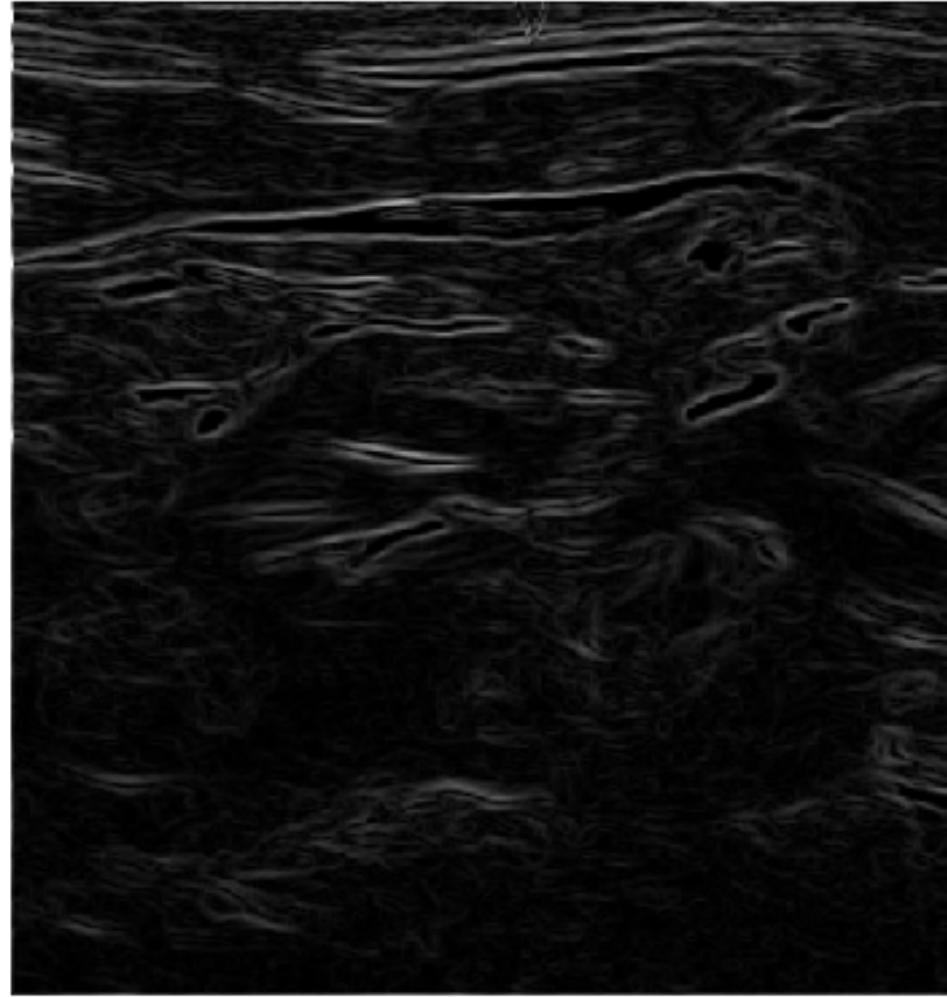


```
# 4. Gradient Magnitude
kernerl1 = np.array([
    [-1,-1,-1],
    [0,0,0],
    [1,1,1]
])
kernerl2 = np.array([
    [-1,0,1],
    [-1,0,1],
    [-1,0,1]
])

for x in range(len(one)):
    temp1 = convolve2d(one[x,:,:],kernerl1,mode='same')
    temp2 = convolve2d(one[x,:,:],kernerl2,mode='same')

    temp3 = np.sqrt(temp1**2 + temp2**2)

    plt.axis('off')
    plt.imshow(temp3,cmap='gray')
    plt.savefig(str(x)+'.png',bbox_inches='tight')
```



$$\Theta = \text{atan}\left(\frac{\mathbf{G}_y}{\mathbf{G}_x}\right)$$

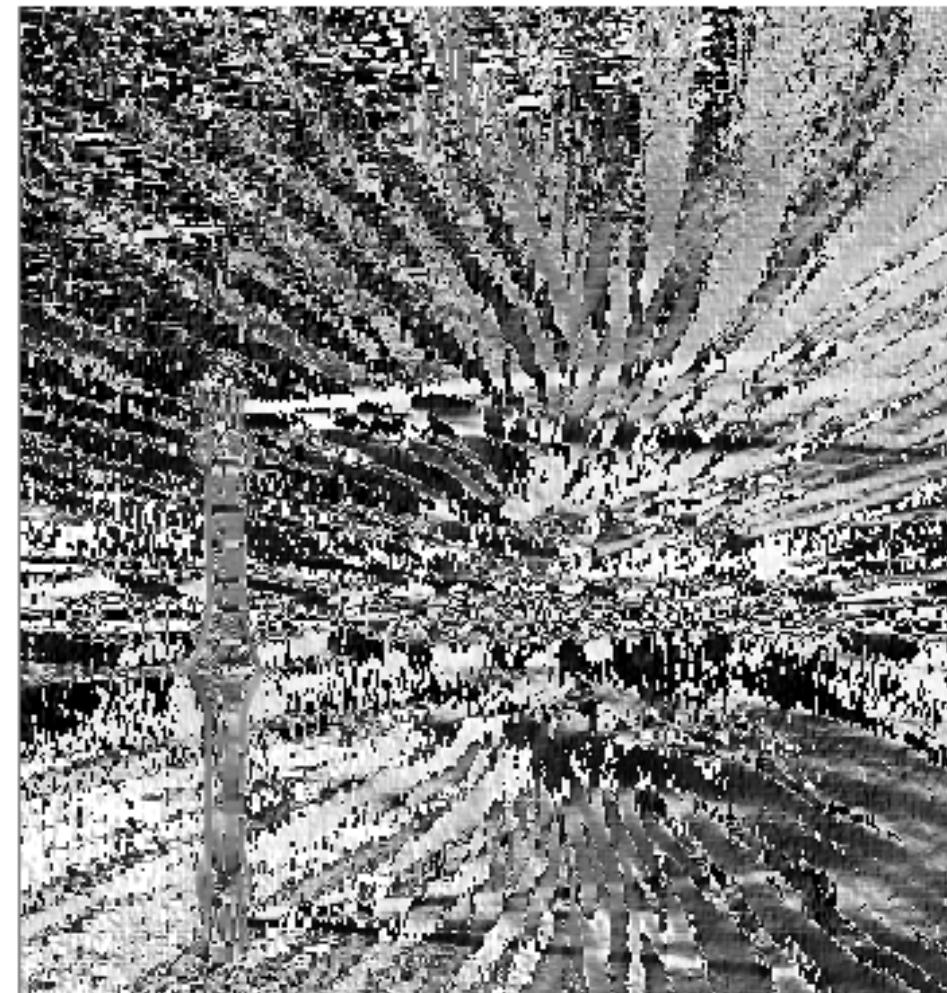
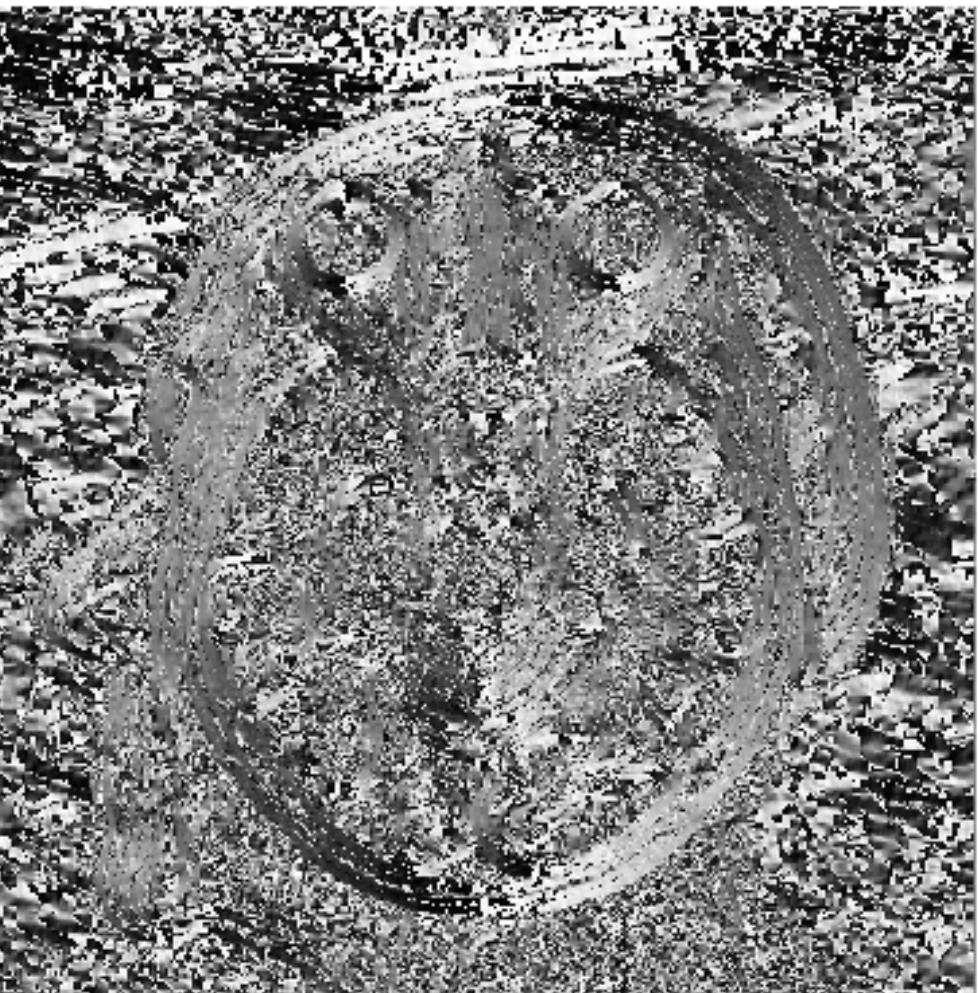
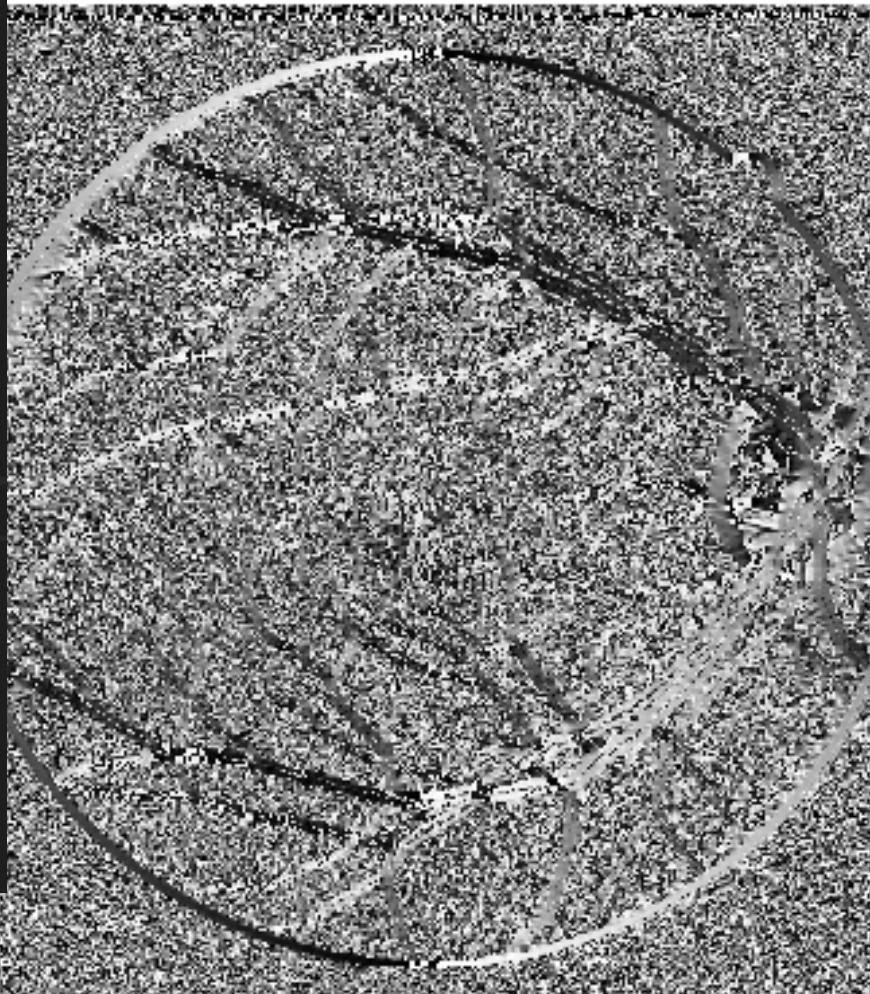


```
# 5. Gradient Direction
kernerl1 = np.array([
    [-1,-1,-1],
    [0,0,0],
    [1,1,1]
])
kernerl2 = np.array([
    [-1,0,1],
    [-1,0,1],
    [-1,0,1]
])

for x in range(len(one)):
    temp1 = convolve2d(one[x,:,:],kernerl1,mode='same')
    temp2 = convolve2d(one[x,:,:],kernerl2,mode='same')

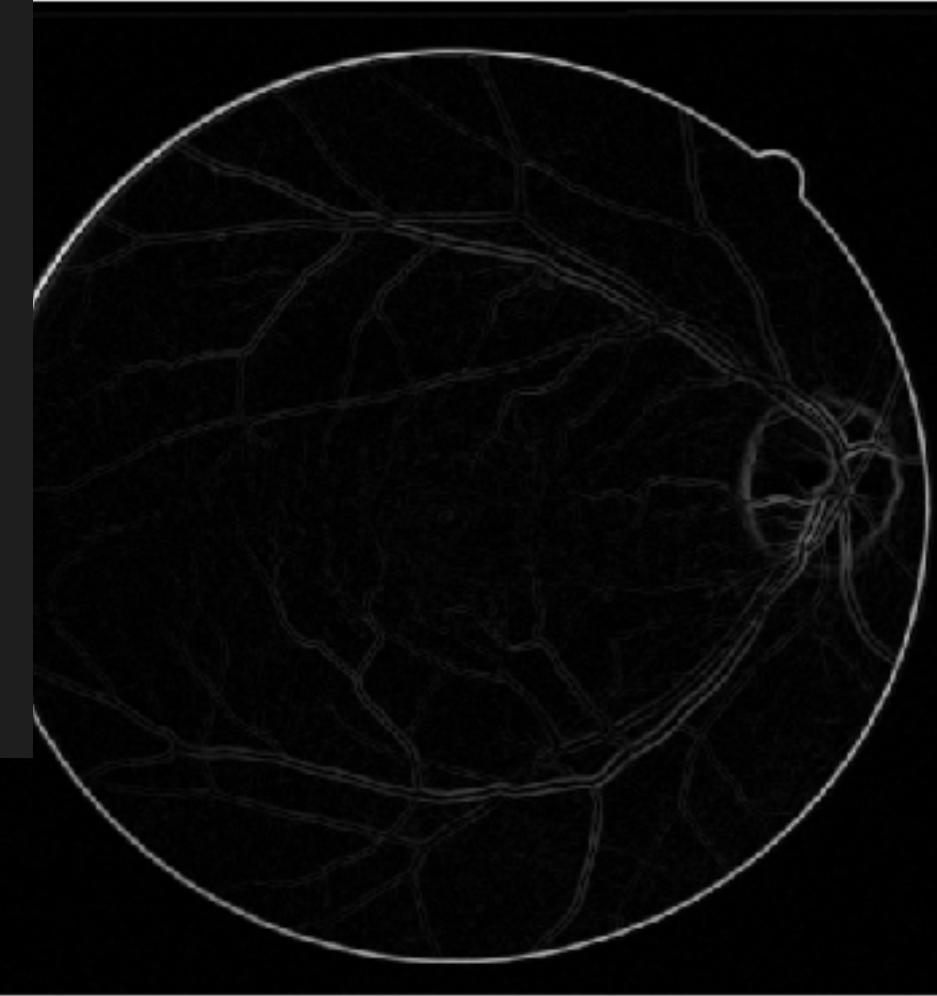
    temp3 = np.arctan(temp1/temp2)

    plt.axis('off')
    plt.imshow(temp3,cmap='gray')
    plt.savefig(str(x)+'.png',bbox_inches='tight')
```

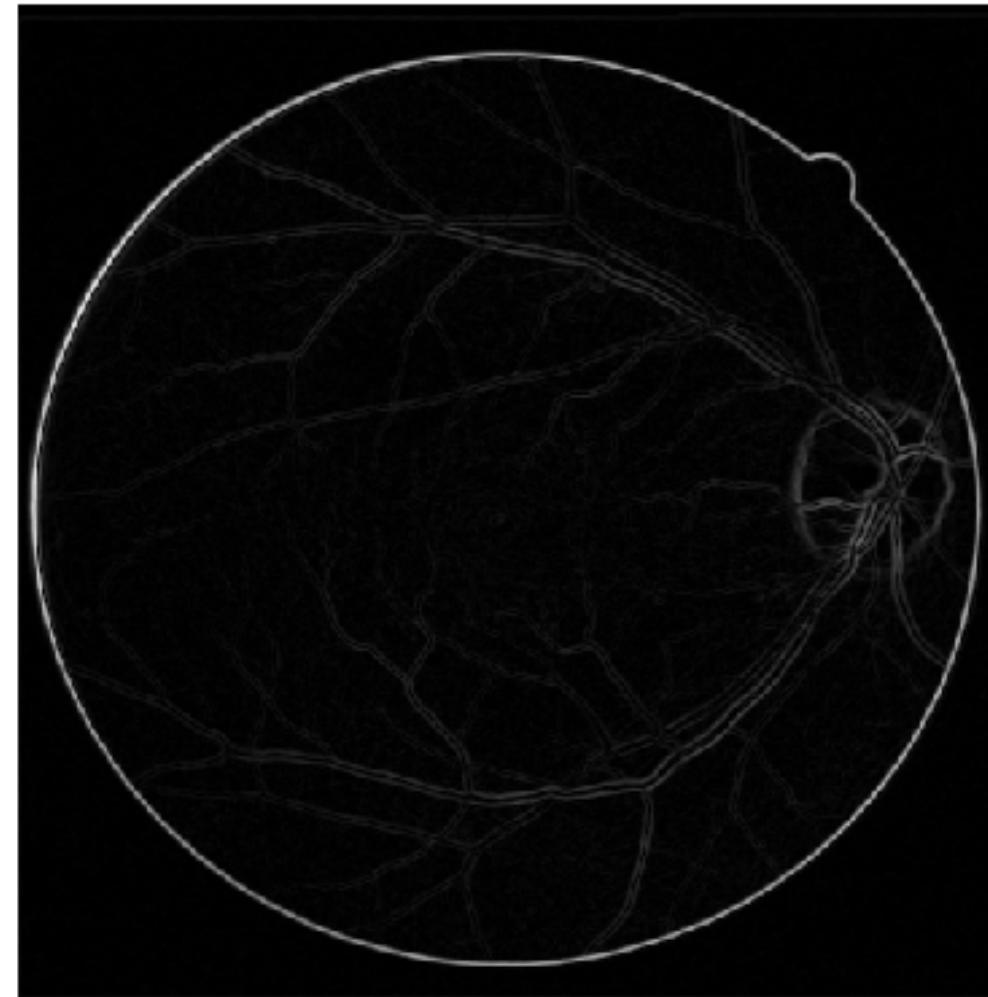
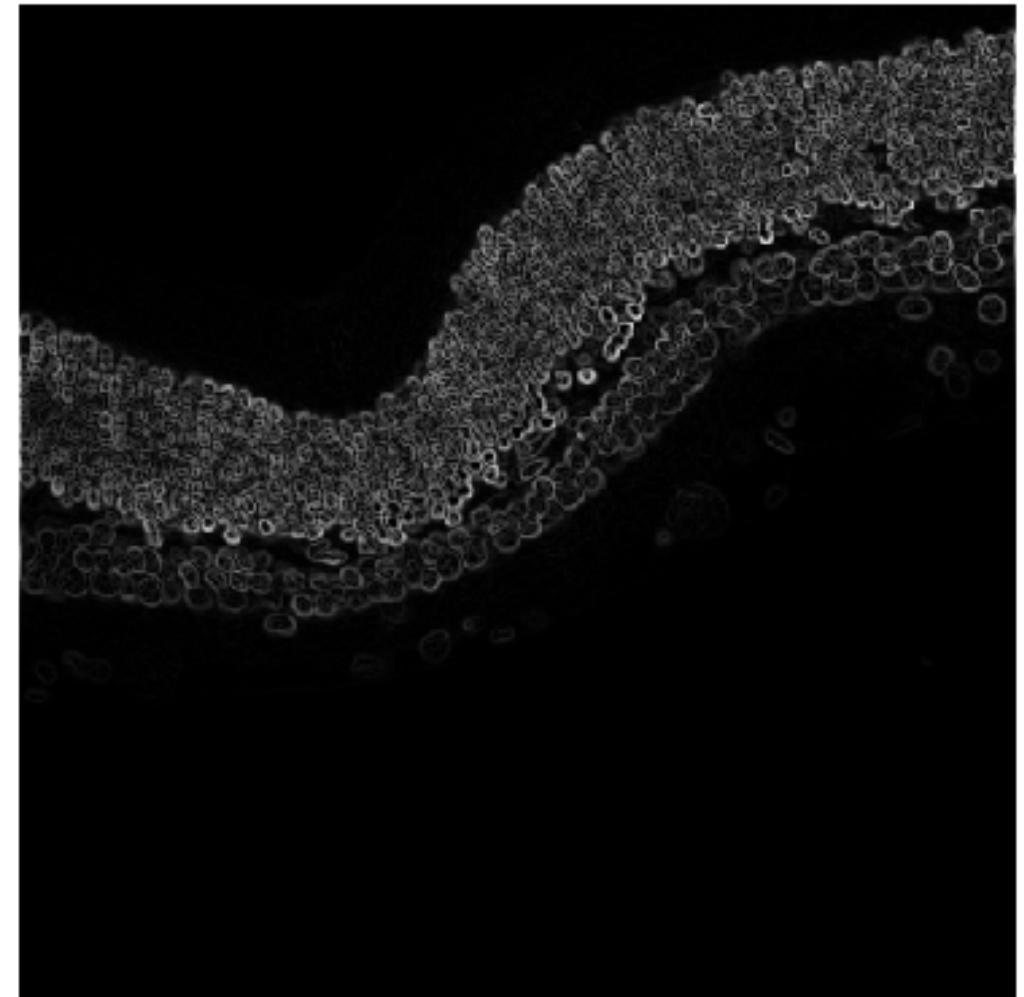
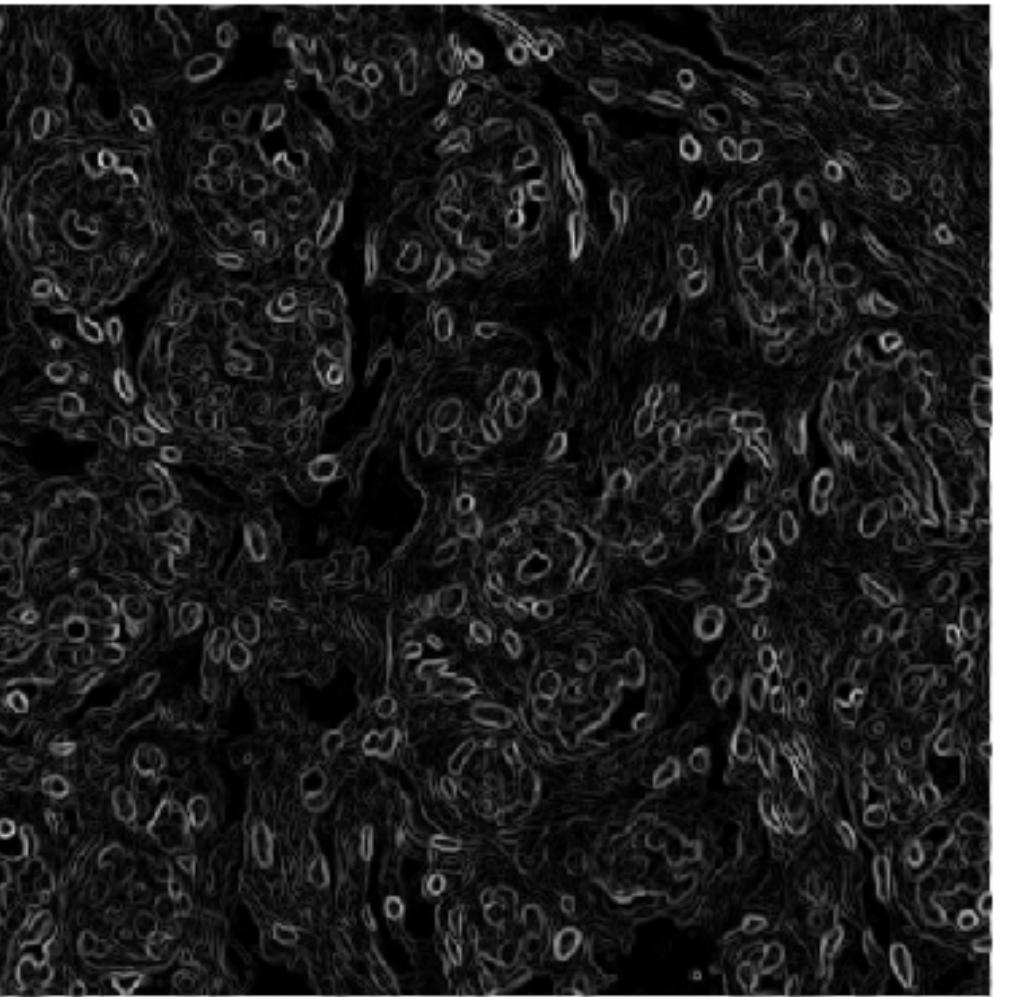


```
9 # 6. Sobel Gradient Magnitude
0 kernerl1 = np.array([
1     [1,2,1],
2     [0,0,0],
3     [-1,-2,-1]
4 ])
5 kernerl2 = np.array([
6     [1,0,-1],
7     [2,0,-2],
8     [1,0,-1]
9 ])
10
11 for x in range(len(one)):
12     temp1 = convolve2d(one[x,:,:],kernerl1,mode='same')
13     temp2 = convolve2d(one[x,:,:],kernerl2,mode='same')
14
15     temp3 = np.sqrt(temp1**2 + temp2**2)
16
17     plt.axis('off')
18     plt.imshow(temp3,cmap='gray')
19     plt.savefig(str(x)+'.png',bbox_inches='tight')
```

Sobel



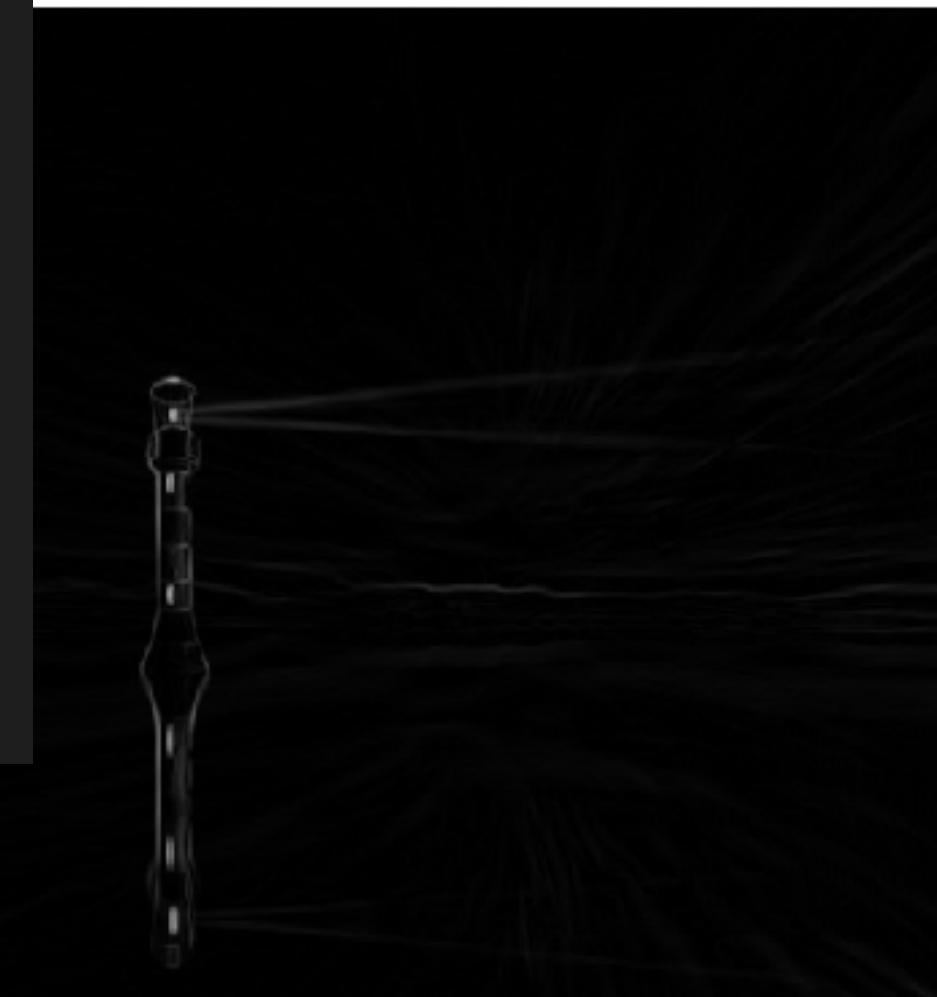
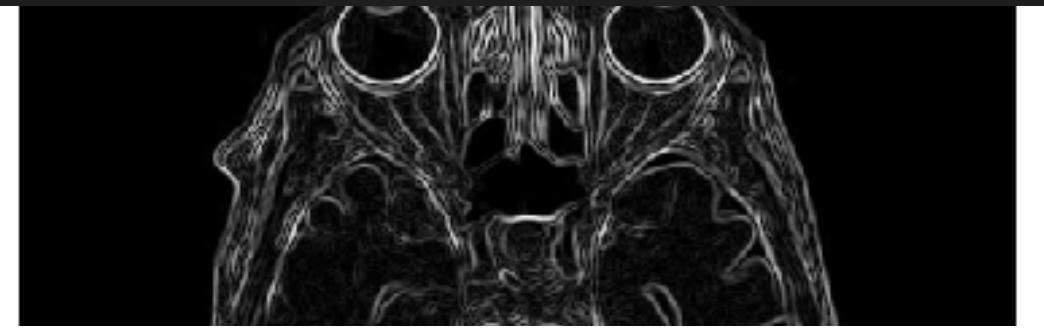
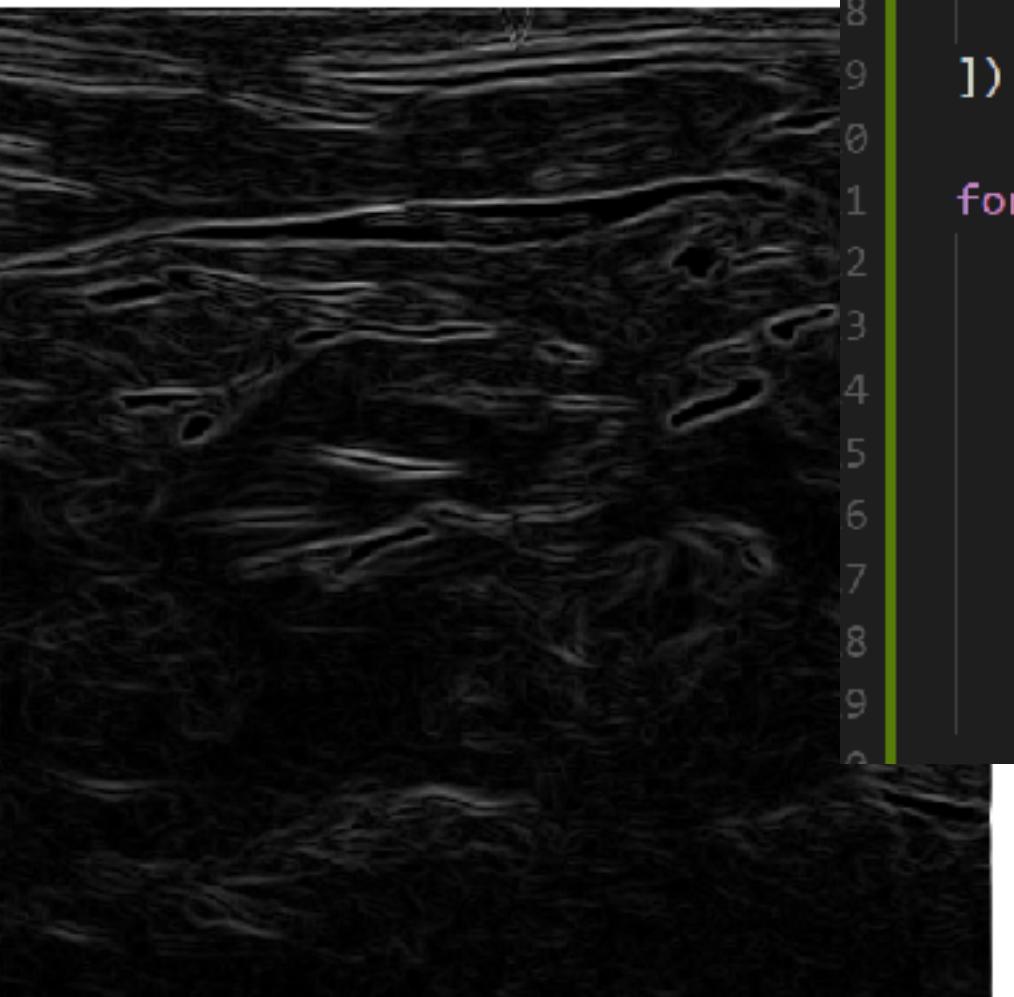
«Ordinary»
Prewitt



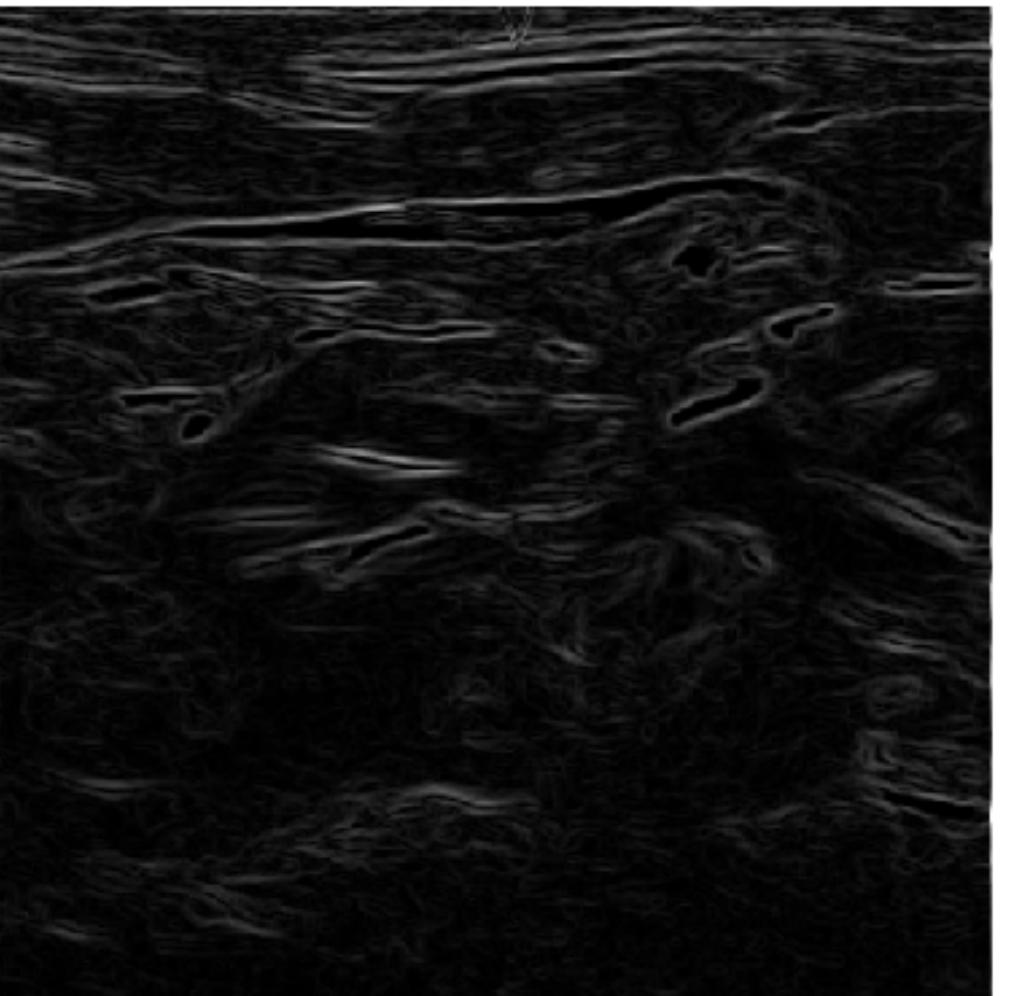
Sobel

```
# 6. Sobel Gradient Magnitude
0 kernerl1 = np.array([
1     [1,2,1],
2     [0,0,0],
3     [-1,-2,-1]
4 ])
5 kernerl2 = np.array([
6     [1,0,-1],
7     [2,0,-2],
8     [1,0,-1]
9 ])
0
10 for x in range(len(one)):
11     temp1 = convolve2d(one[x,:,:],kernerl1,mode='same')
12     temp2 = convolve2d(one[x,:,:],kernerl2,mode='same')
13
14     temp3 = np.sqrt(temp1**2 + temp2**2)
15
16     plt.axis('off')
17     plt.imshow(temp3,cmap='gray')
18     plt.savefig(str(x)+'.png',bbox_inches='tight')
```

Sobel



«Ordinary»
Prewitt

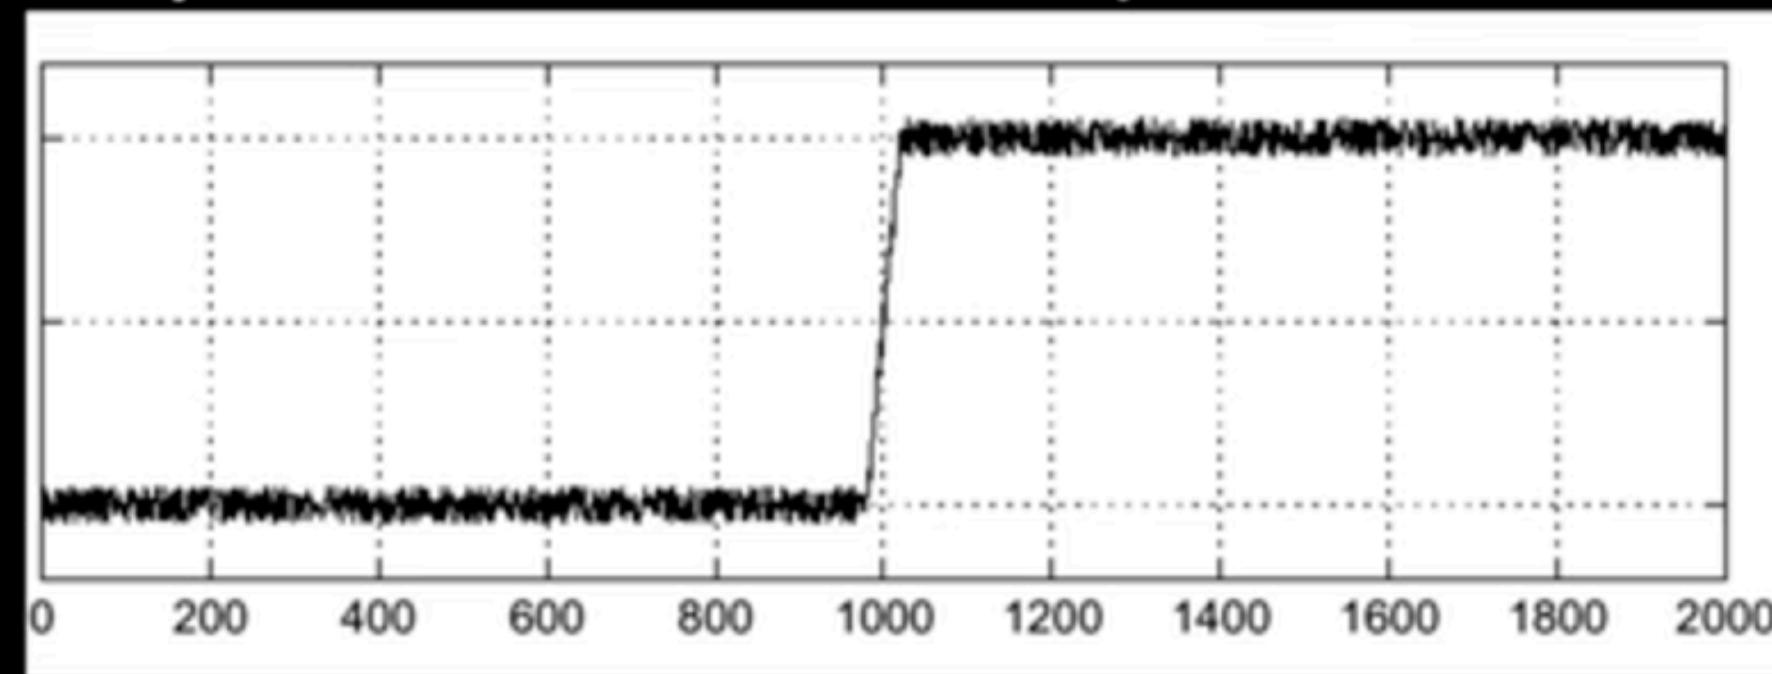


Sobel

Noise + derivative filter = NotGood

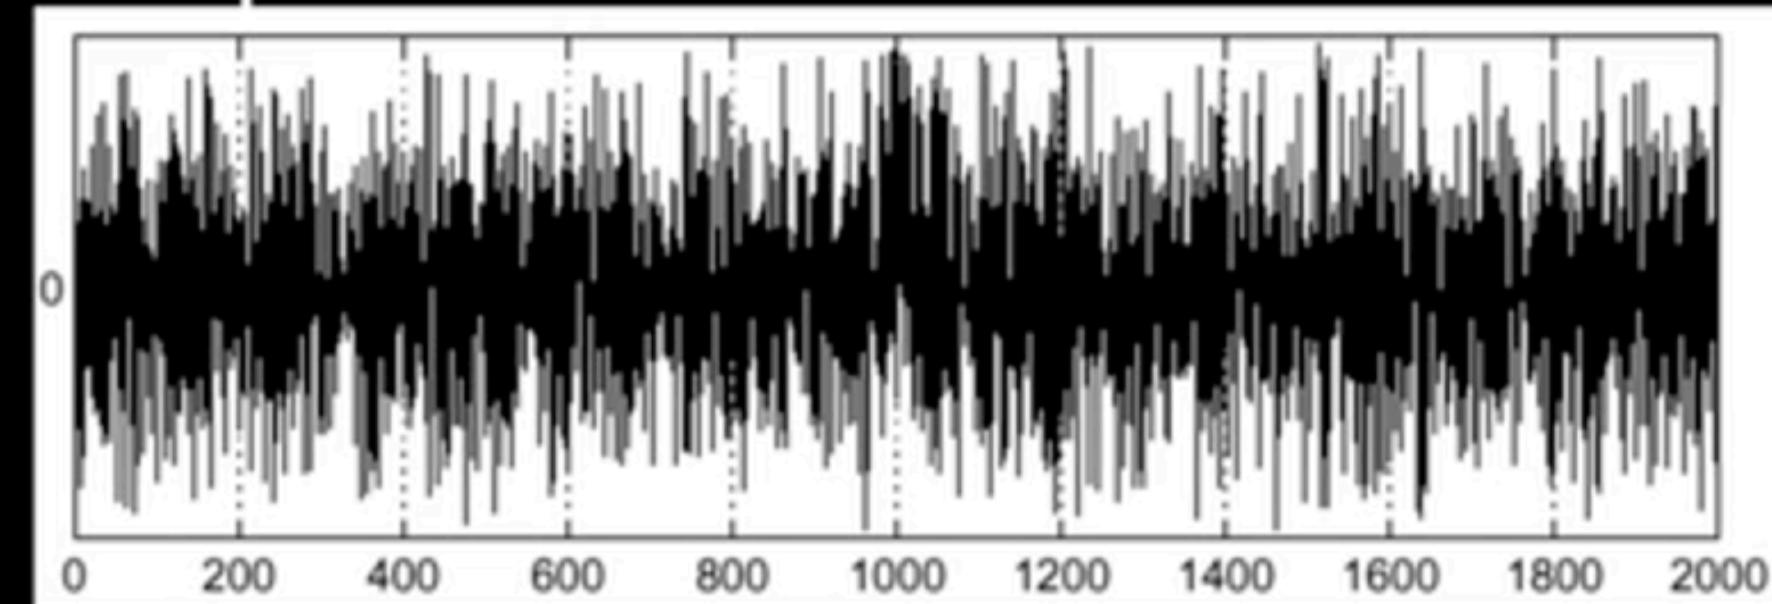
Consider a single row or column of the image
(plotting intensity as a function of x)

$$f(x)$$



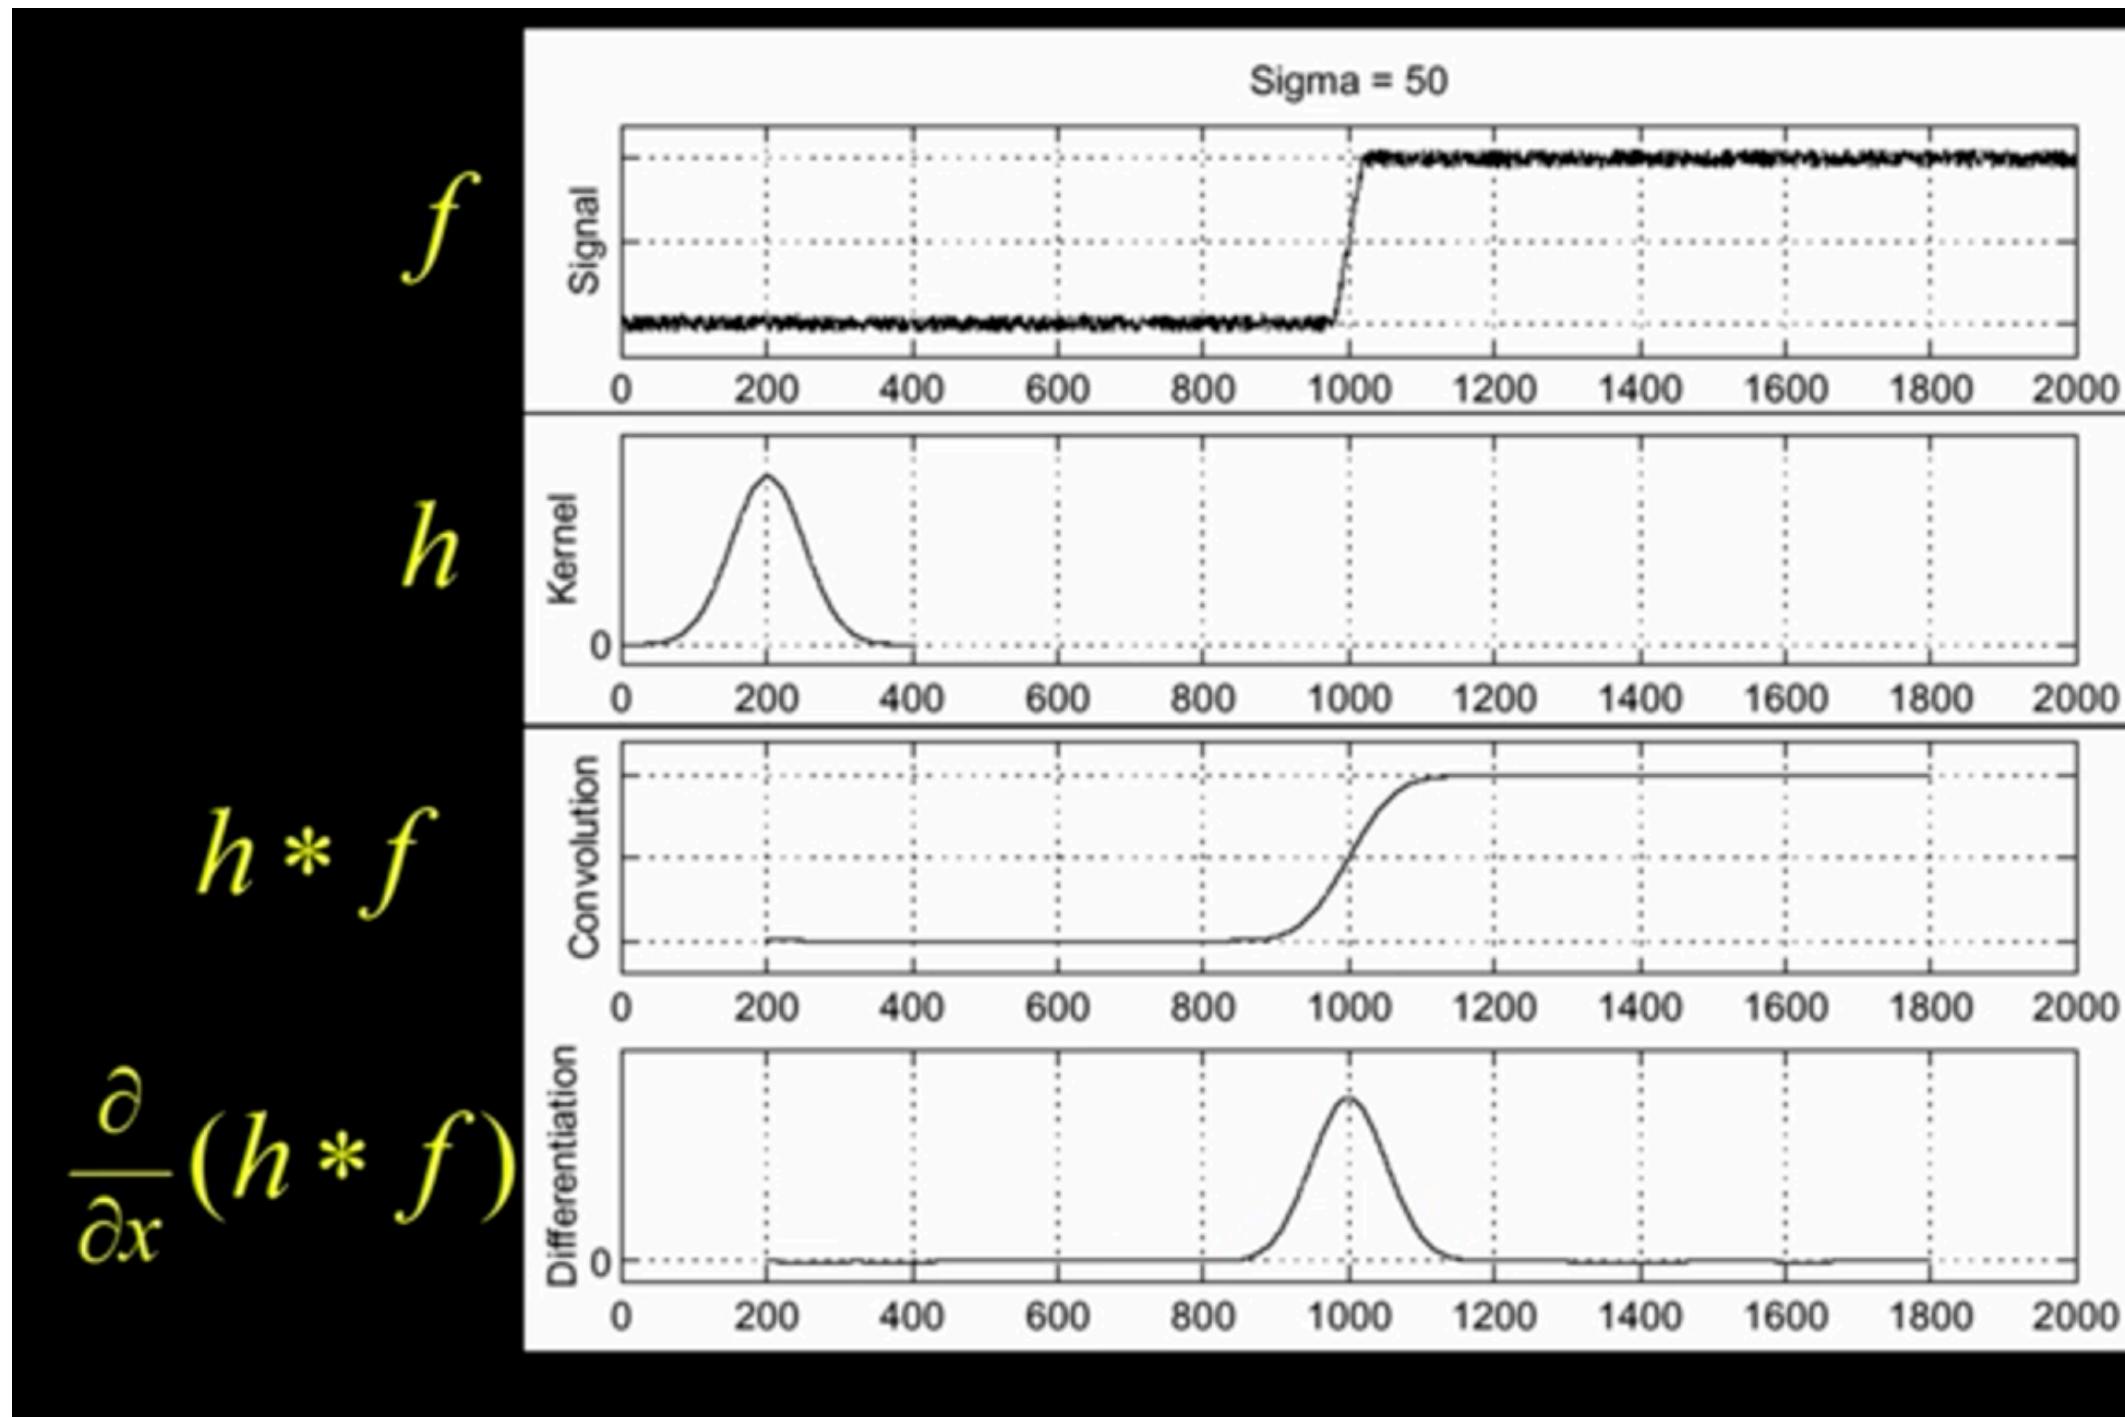
Apply derivative operator....

$$\frac{d}{dx} f(x)$$

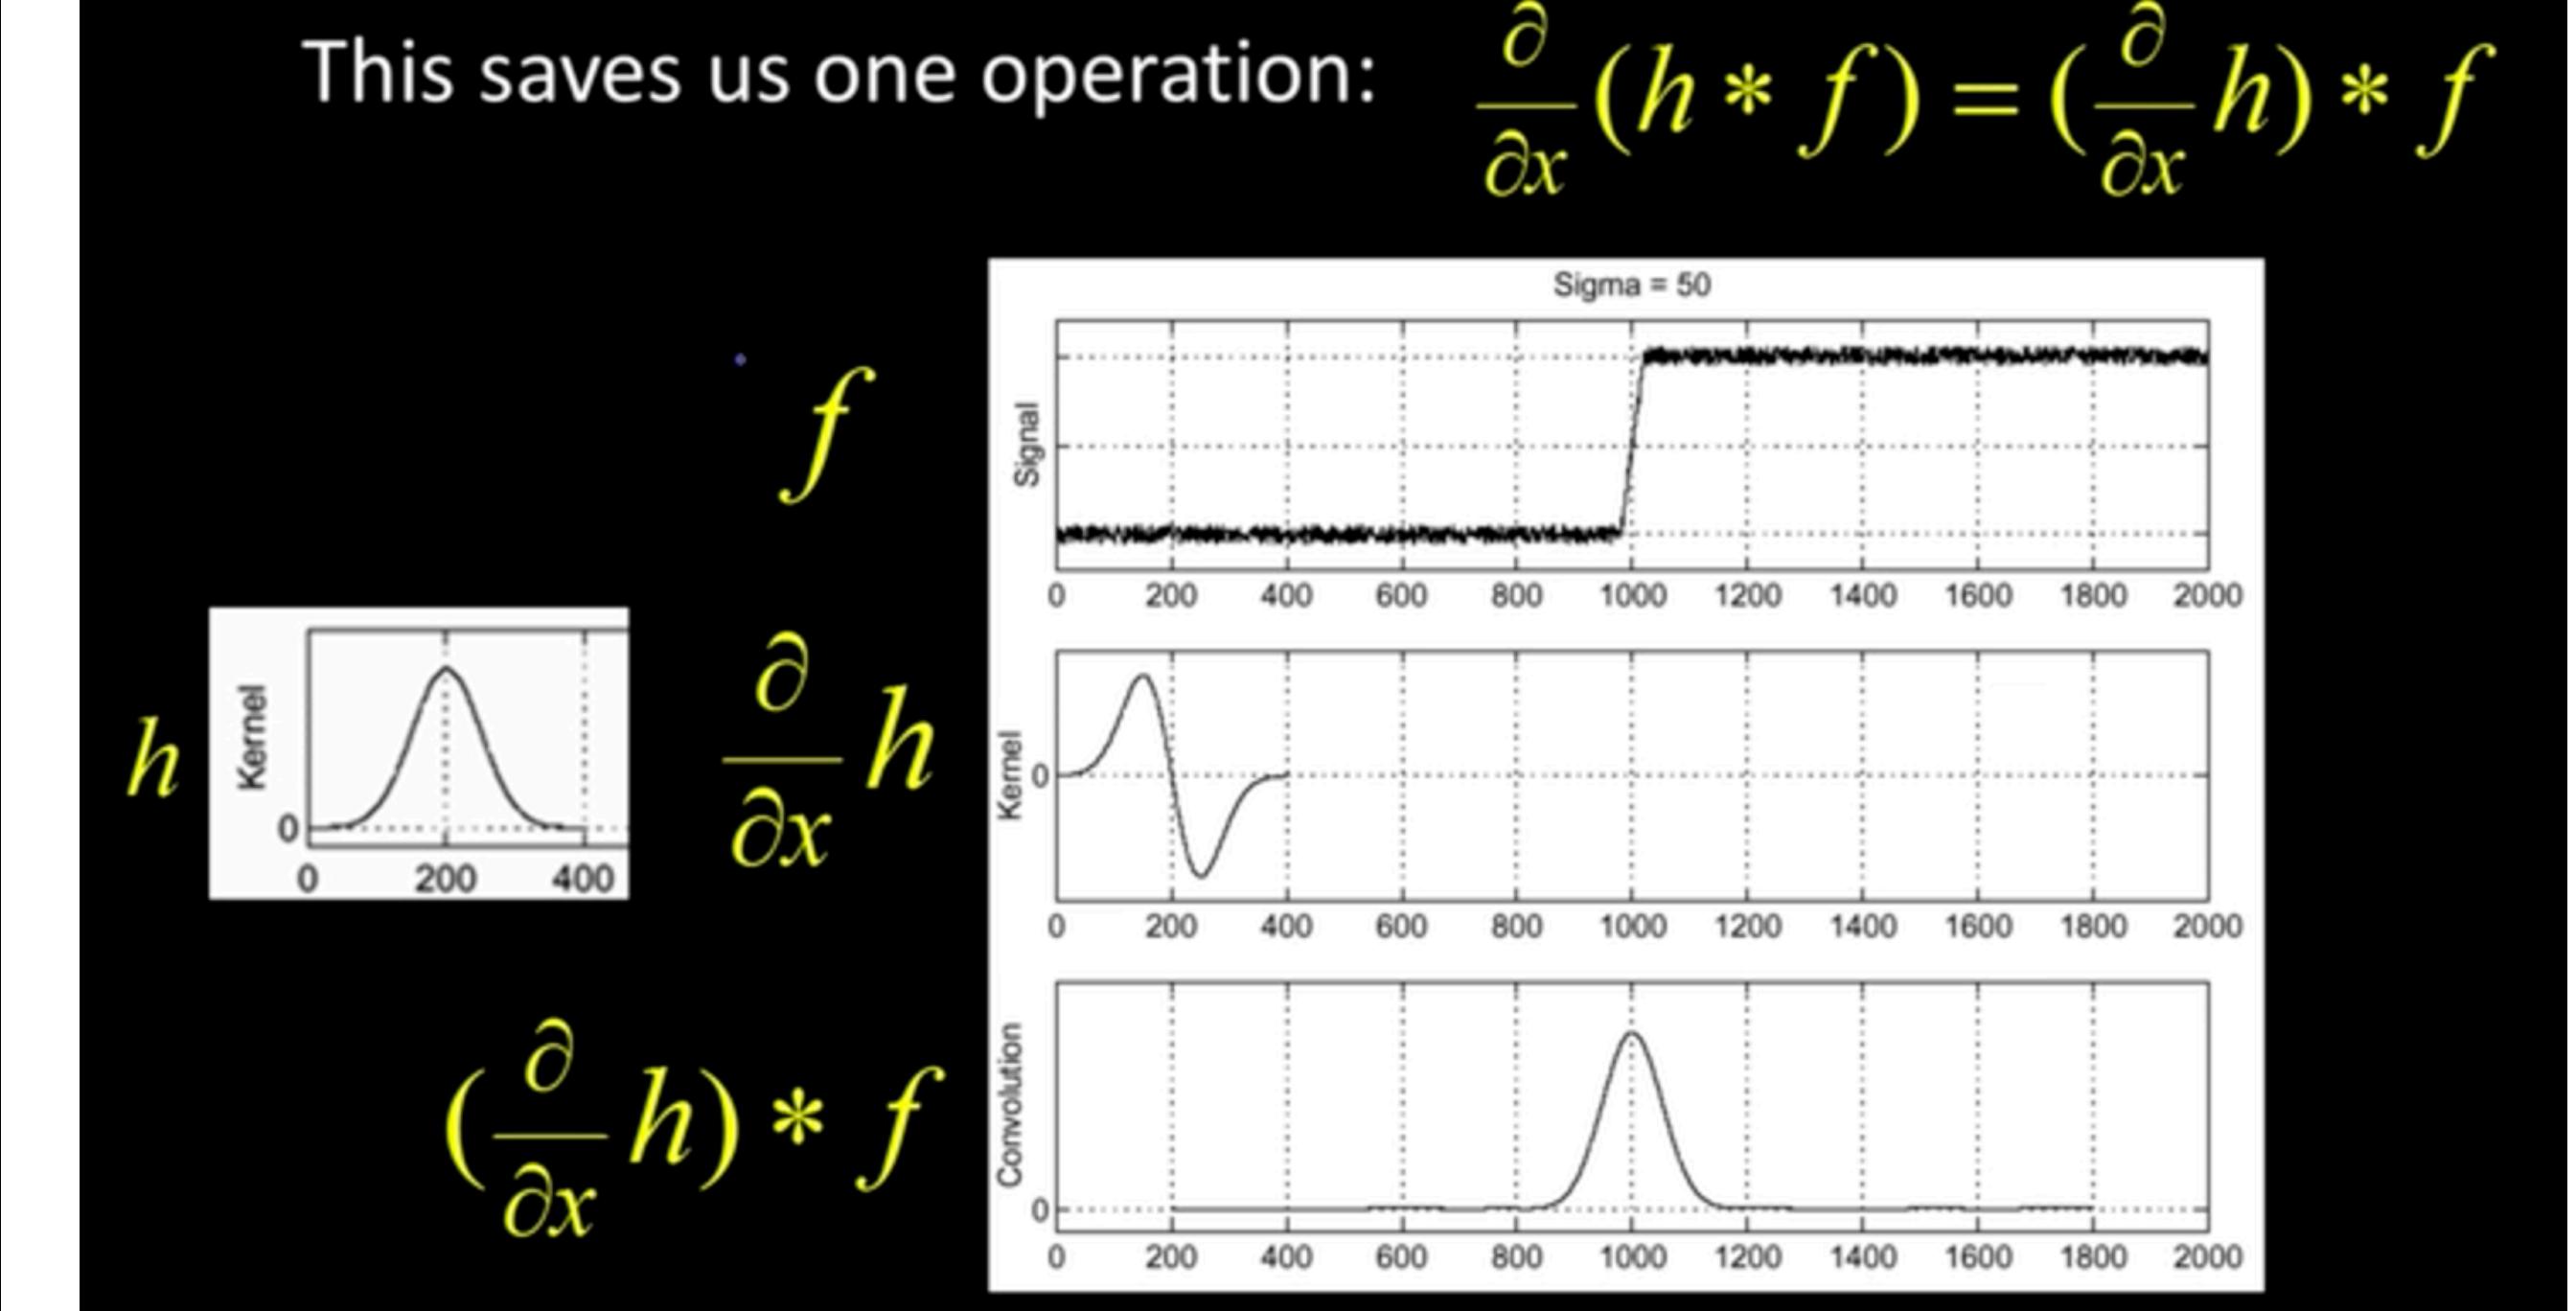


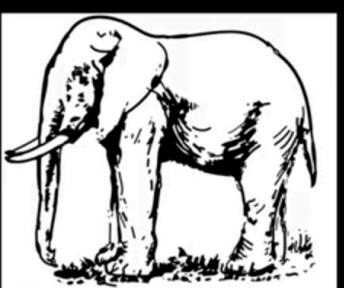
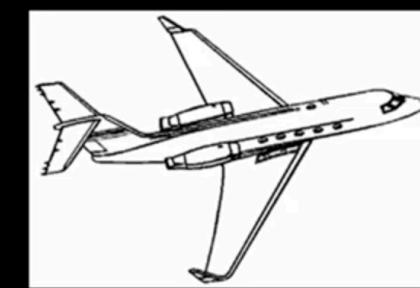
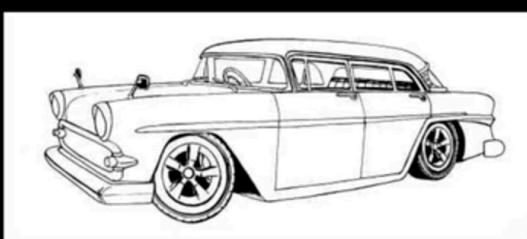
*Uh, where's
the edge?*

Smooth first / or combine

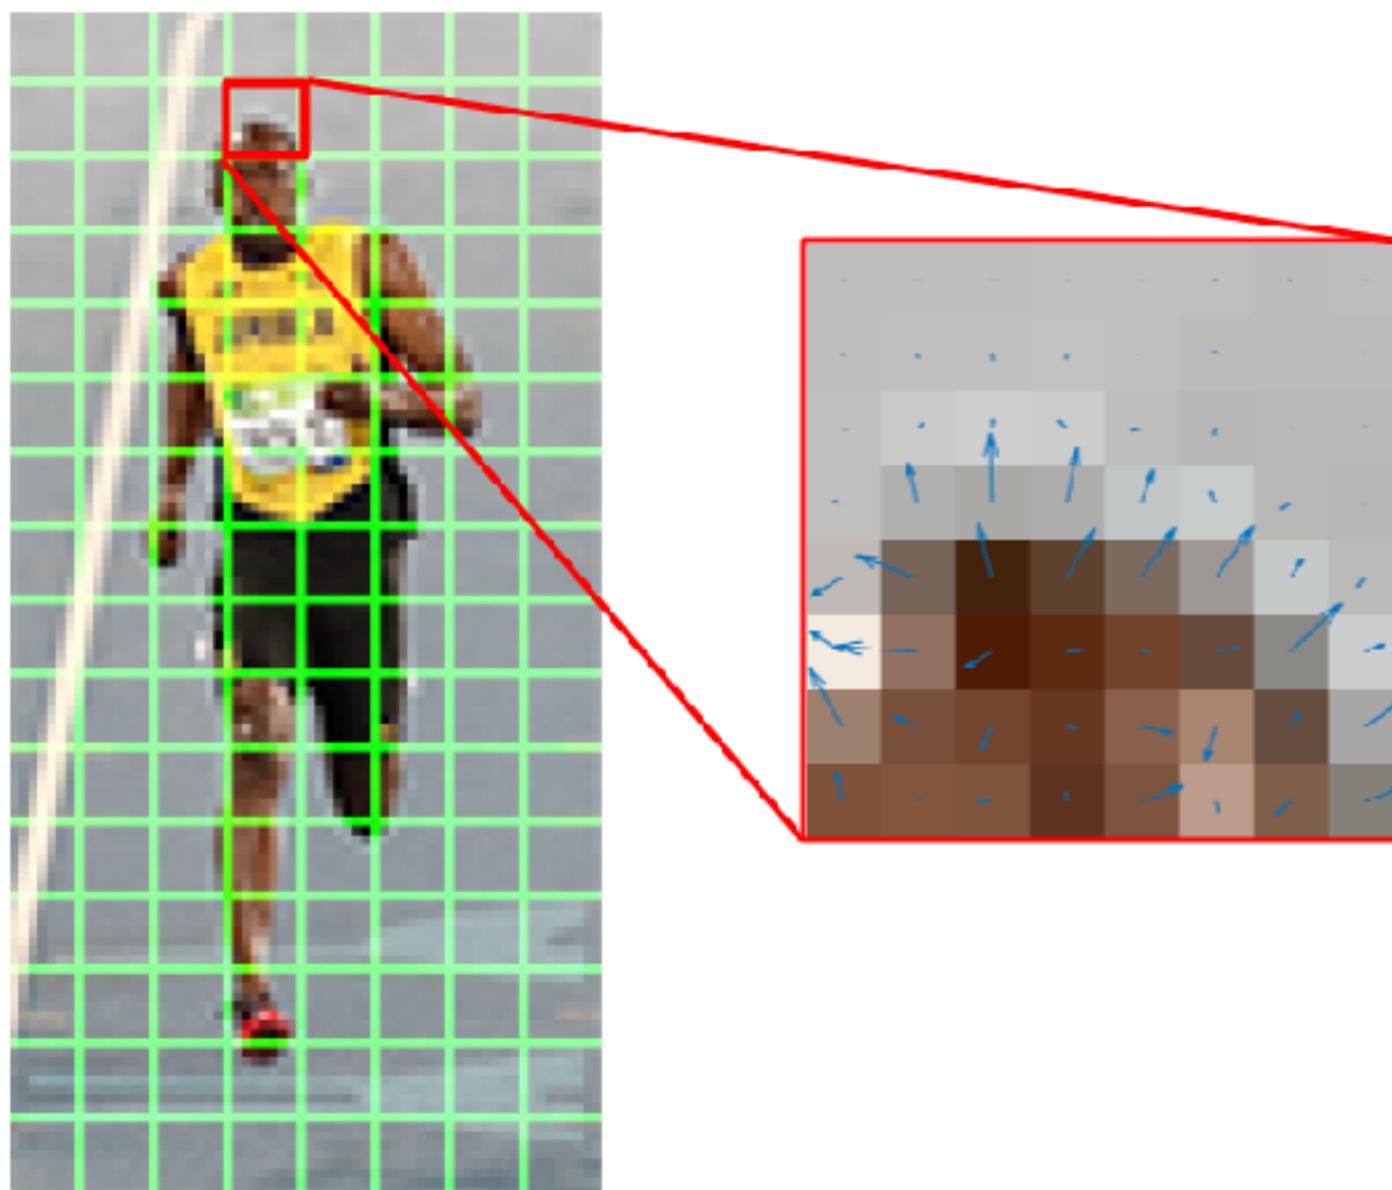


This saves us one operation: $\frac{\partial}{\partial x}(h * f) = (\frac{\partial}{\partial x} h) * f$





SotA before DL



Gradient Magnitude & Direction

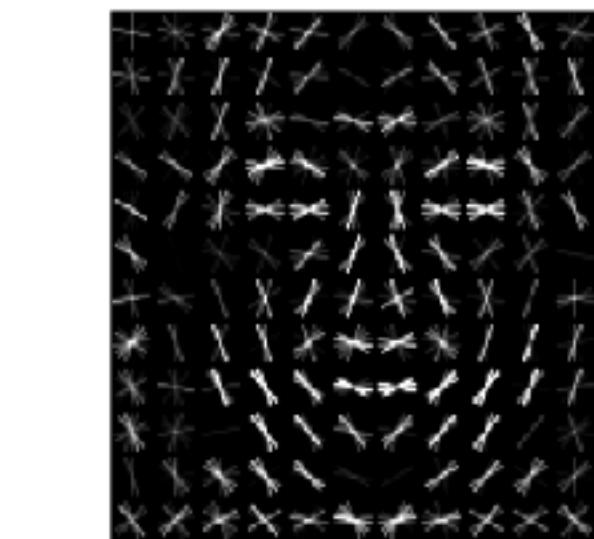
2	3	4	4	3	4	2
5	11	17	13	7	9	3
11	21	23	27	22	17	4
23	99	165	135	85	32	26
91	155	133	136	144	152	57
98	196	76	38	26	60	170
165	60	60	27	77	85	43

Gradient Magnitude

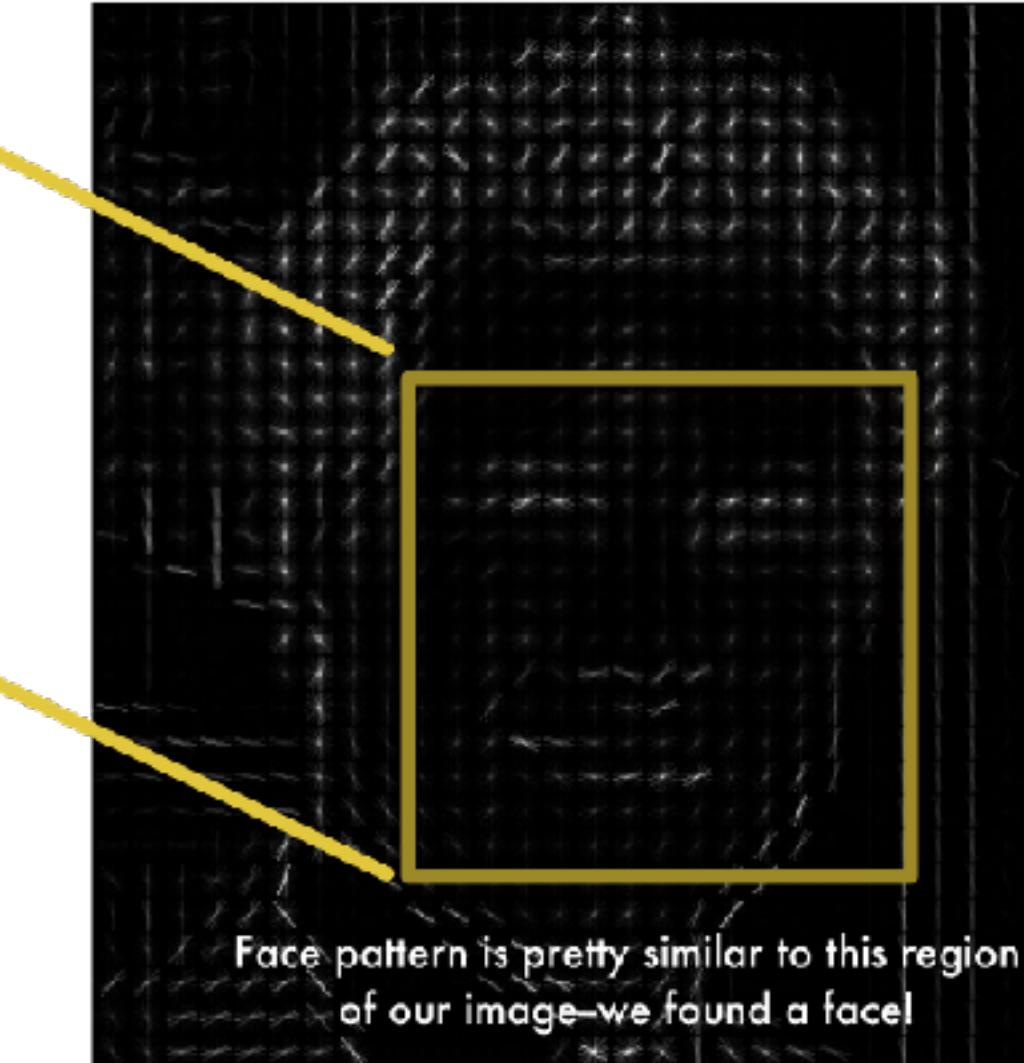
80	36	5	10	0	64	90	1
37	9	9	179	78	27	169	1
87	136	173	39	102	163	152	1
76	13	1	168	159	22	125	1
120	70	14	150	145	144	145	1
58	86	119	98	100	101	133	1
30	65	157	75	78	165	145	1
11	170	91	4	110	17	133	1

Gradient Direct

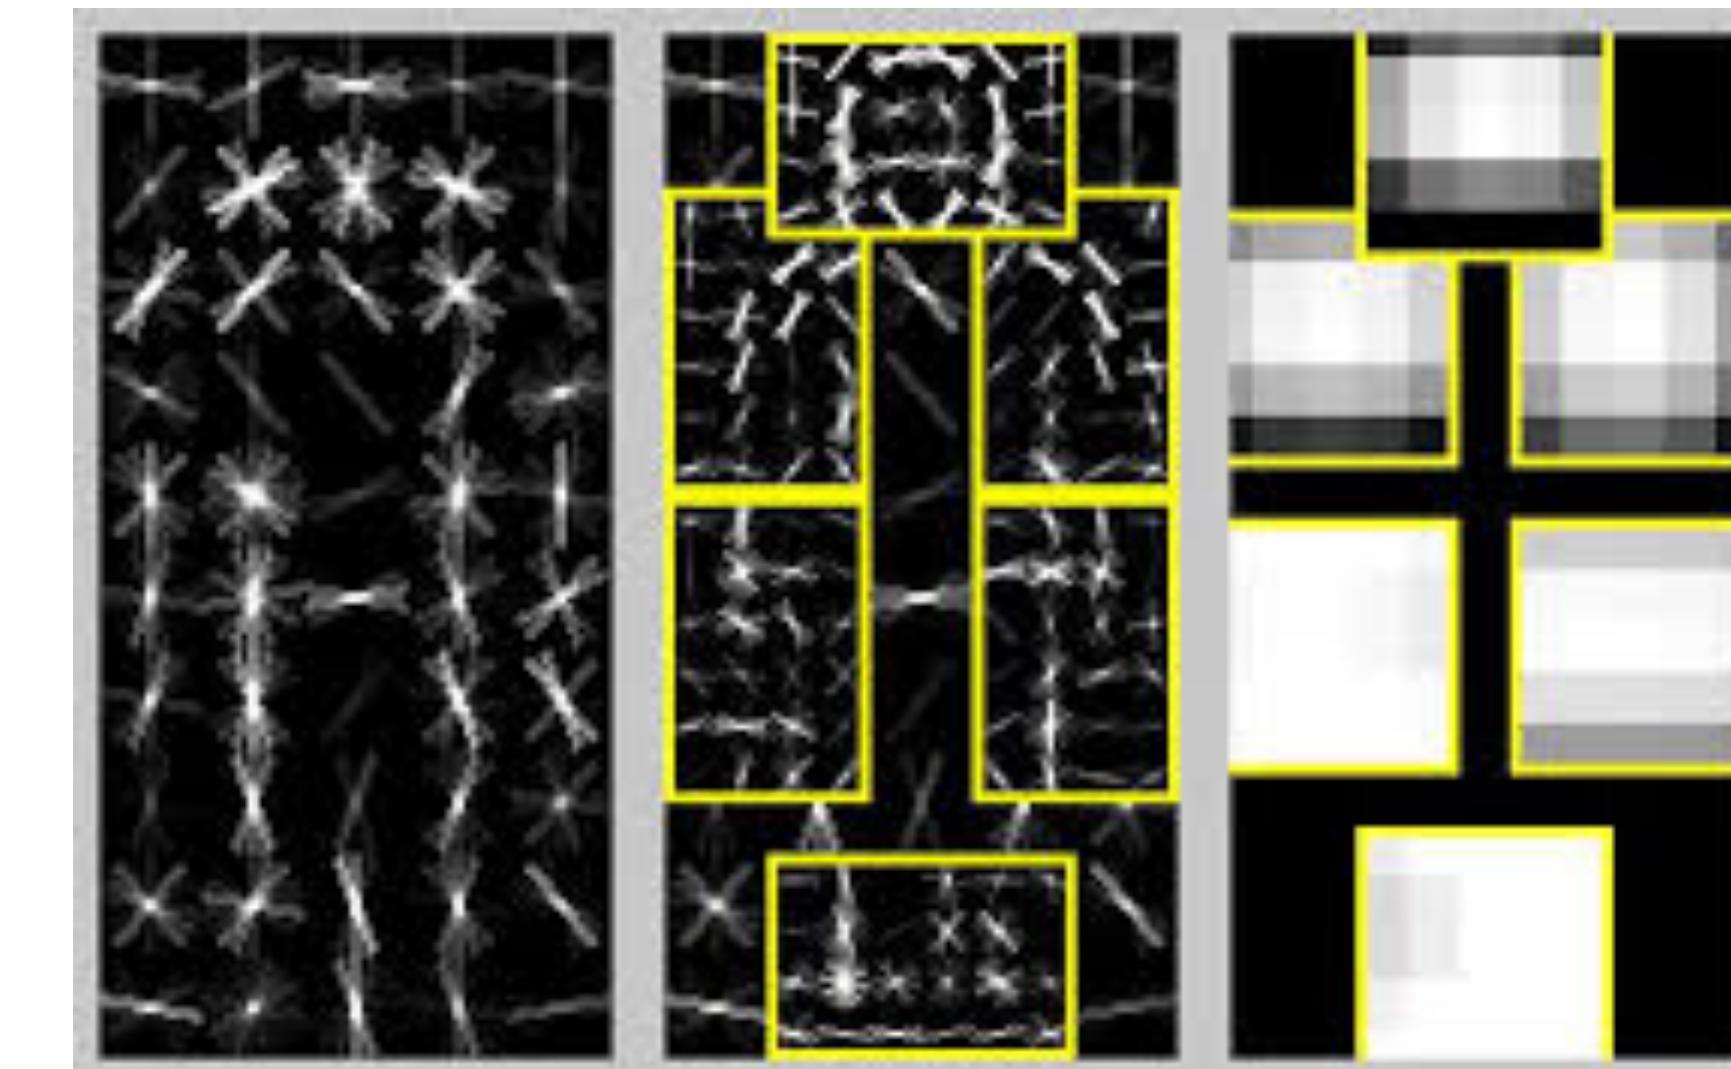
HOG face pattern generation from lots of face images



HOG version of our image



HOGL



DPM

Feature engineering vs. learning

